

# Image and Video Understanding

2VO 710.095 WS

Christoph Feichtenhofer, Axel Pinz

Further reading: Szeliski, Richard. *Computer Vision: Algorithms and Applications*. Springer, 2010, Chapter 3, Section 3.1-3.5

Slide credits:

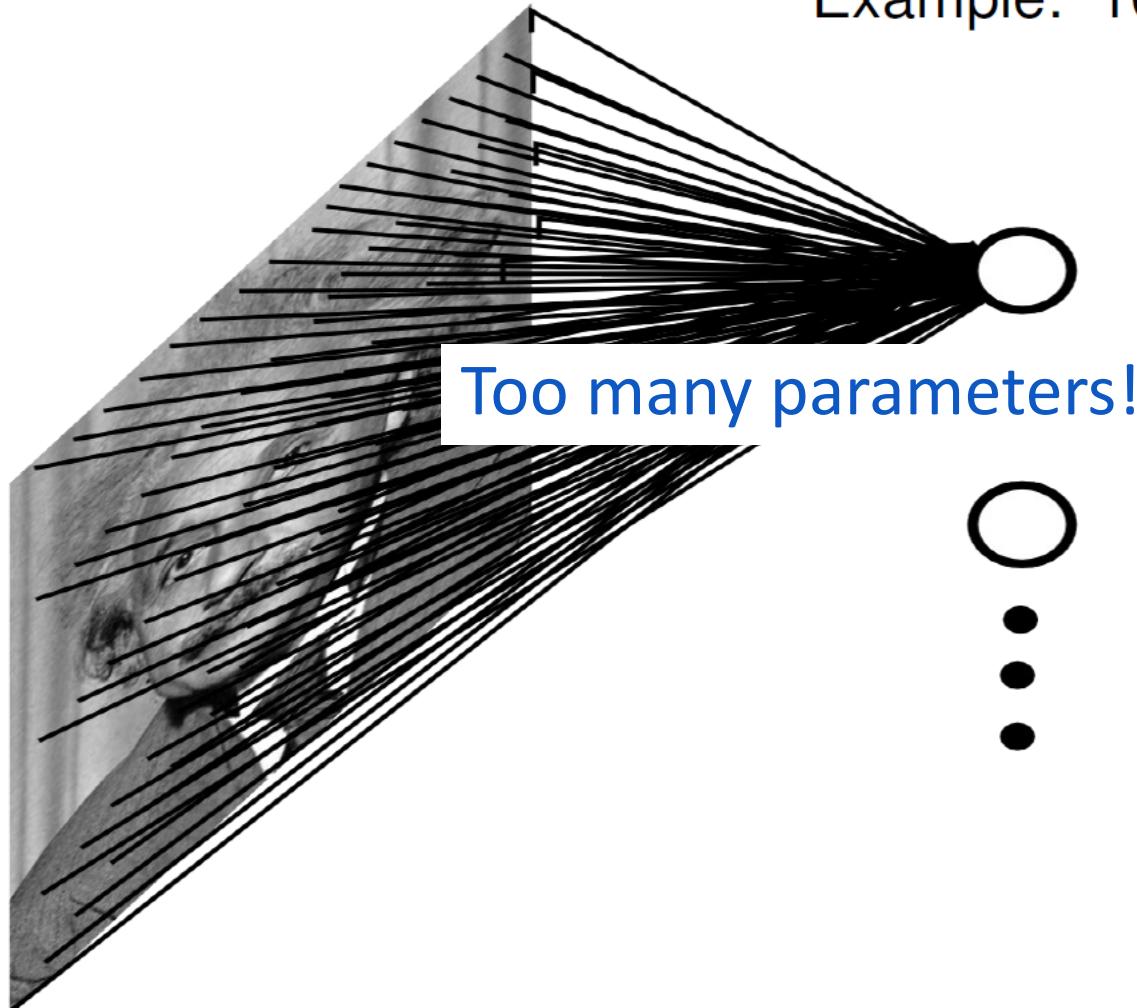
Many thanks to all the great computer vision researchers on which this presentation relies on.

# Outline

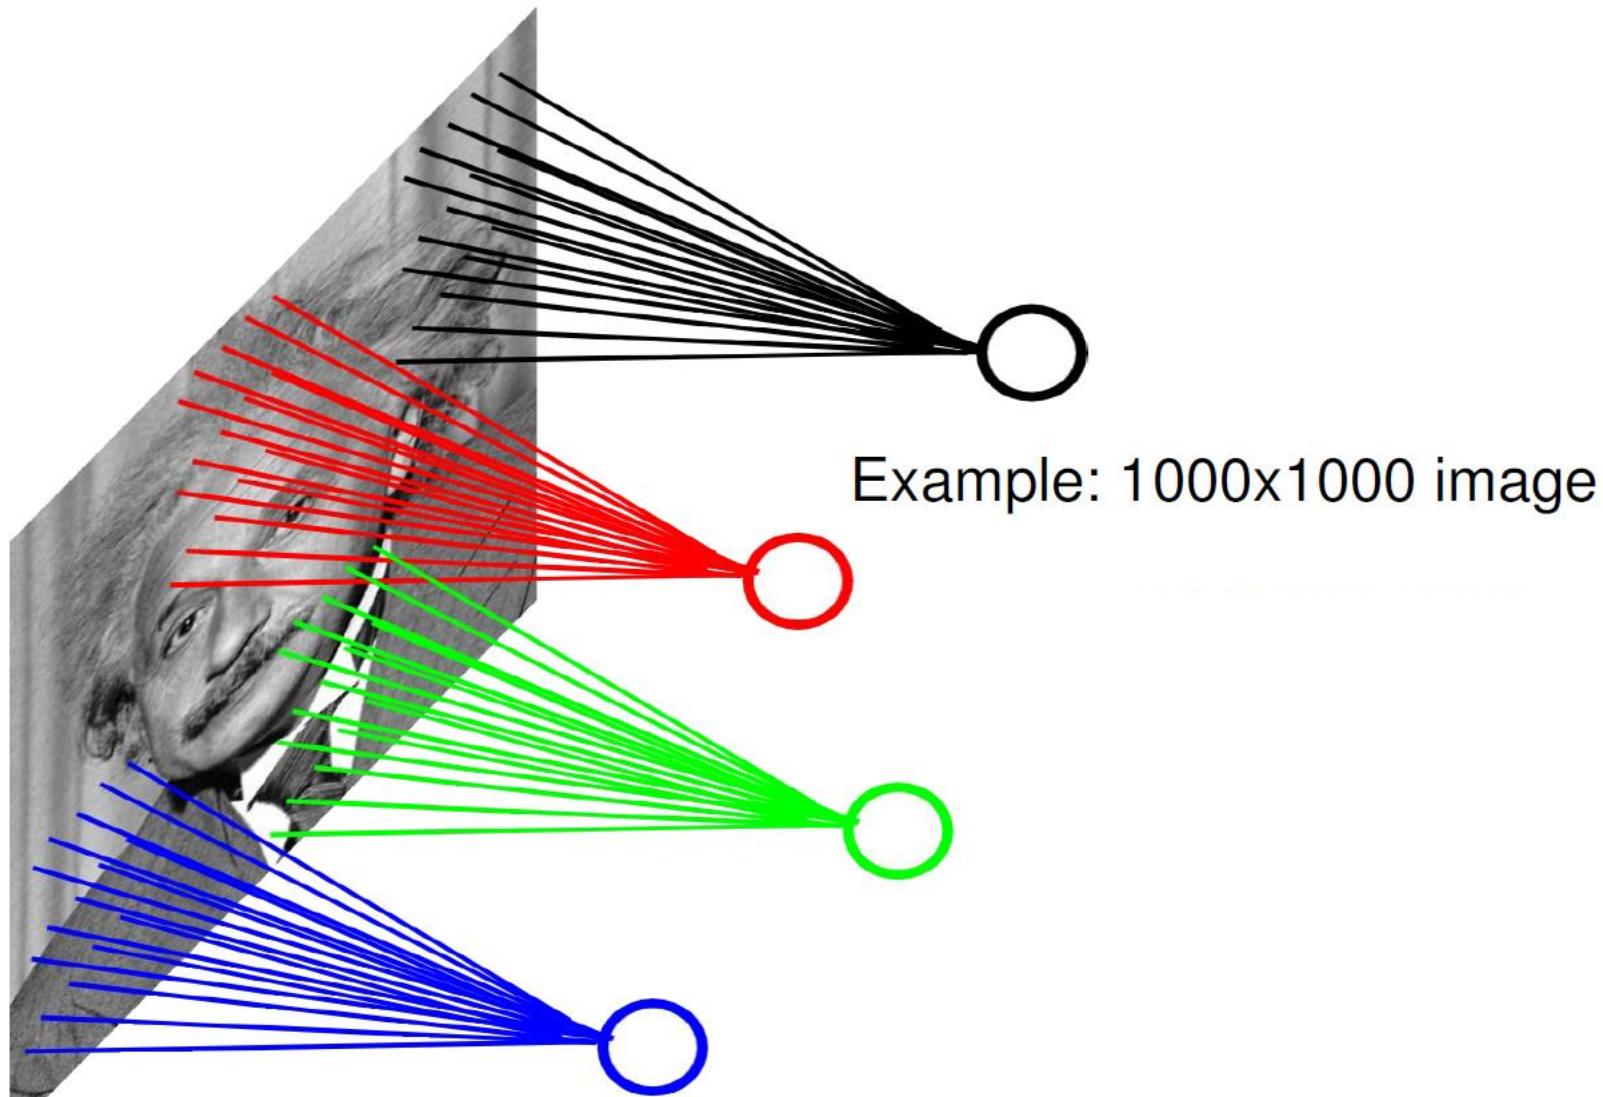
- Linear filtering and the importance of **convolution**
  - Apply a filtermask to the local neighborhood at each pixel in the image
  - The filtermask defines how to combine values from neighbors.
  - Can be used for
    - Extract **intermediate representations** to abstract images by higher-level “**features**”, for further processing (i.e., preserve the useful information only and discard redundancy)
    - Image **modification**, e.g., to reduce noise, resize, increase contrast, etc.
    - Match **template** images (e.g. by correlating two image patches)
- Image filtering in the frequency domain
  - Provides a nice way to illustrate the effect of linear filtering
    - Filtering is a way to modify the frequencies of images
  - Efficient signal filtering is possible in that domain
  - The frequency domain offers an alternative way to understanding and manipulating the image.

Motivation: Images as a composition of local parts  
“Pixel-based” representation

Example: 1000x1000 image



Motivation: Images as a composition of local parts  
“Patch-based” representation



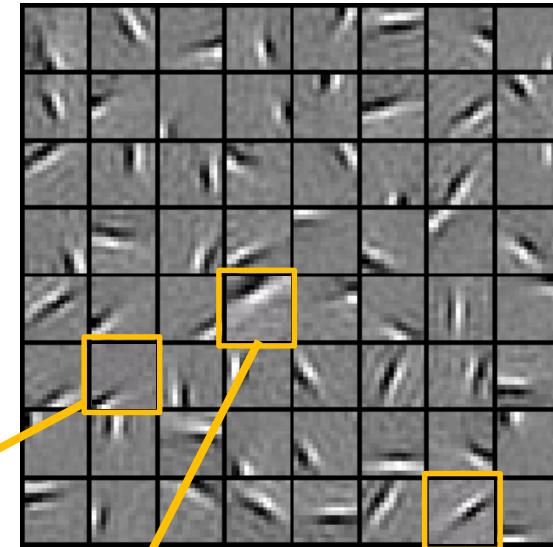
# Motivation: Images as a composition of local parts

## Sparse coding example

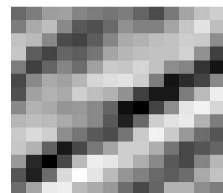
### Natural Images



### Learned bases ( $\phi_1, \dots, \phi_{64}$ ): "Edges"



### Test example



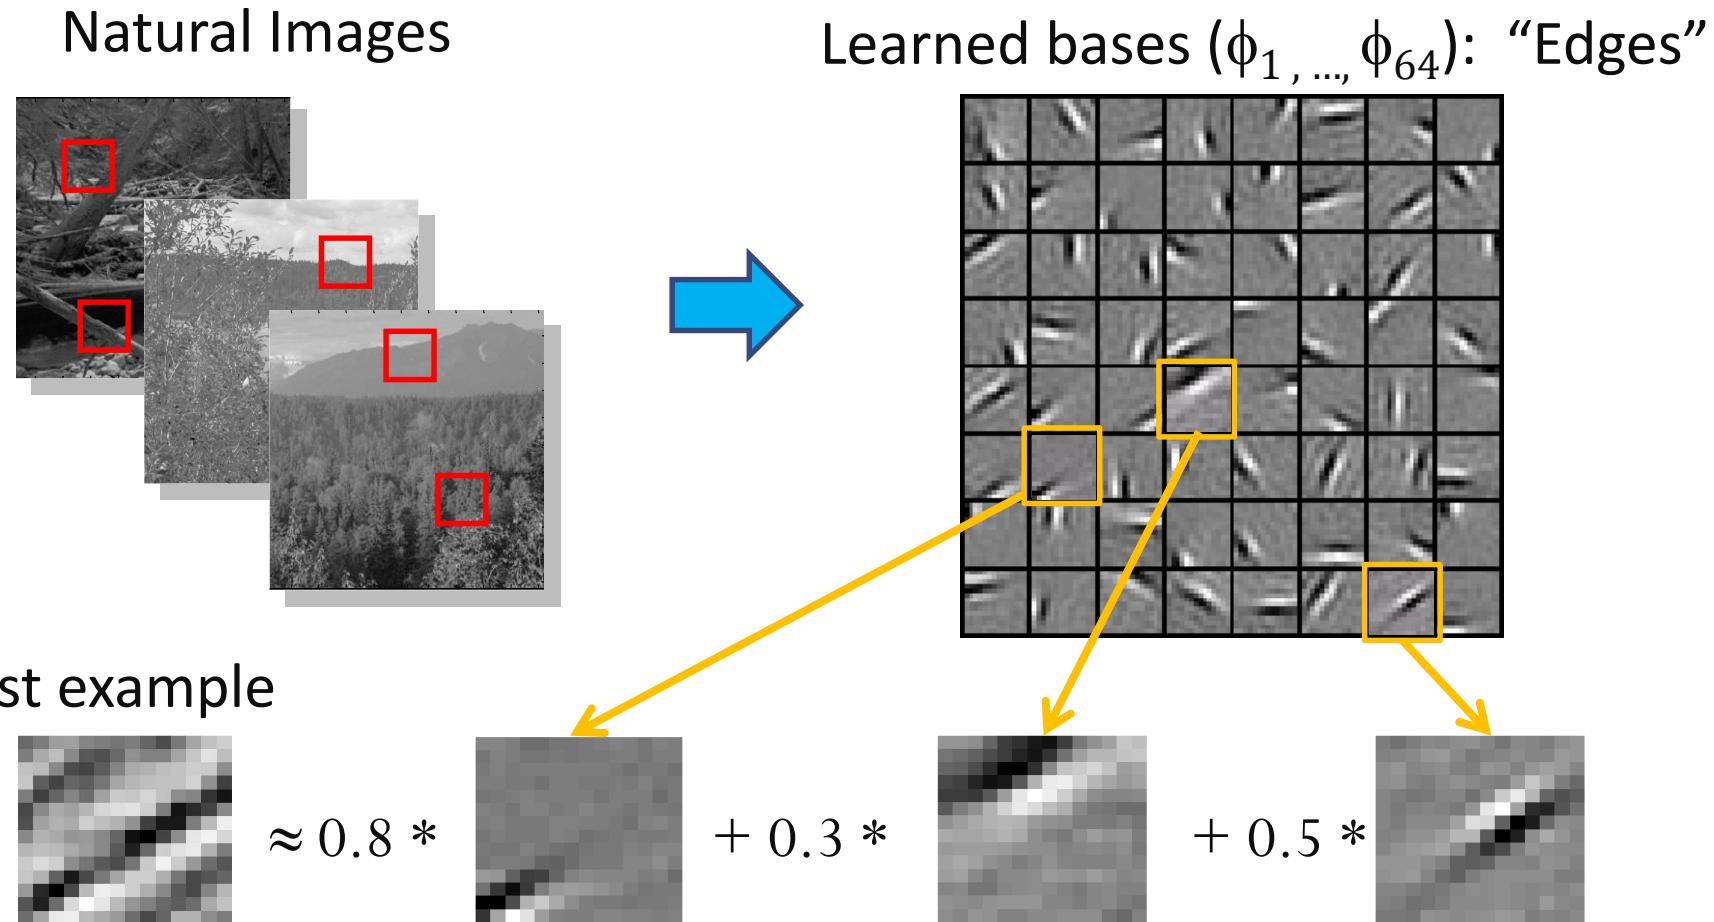
$$x \approx 0.8 * \phi_{36} + 0.3 * \phi_{42} + 0.5 * \phi_{63}$$

$$[a_1, \dots, a_{64}] = [0, 0, \dots, 0, \mathbf{0.8}, 0, \dots, 0, \mathbf{0.3}, 0, \dots, 0, \mathbf{0.5}, 0]$$

(feature representation)

# Motivation: Images as a composition of local parts

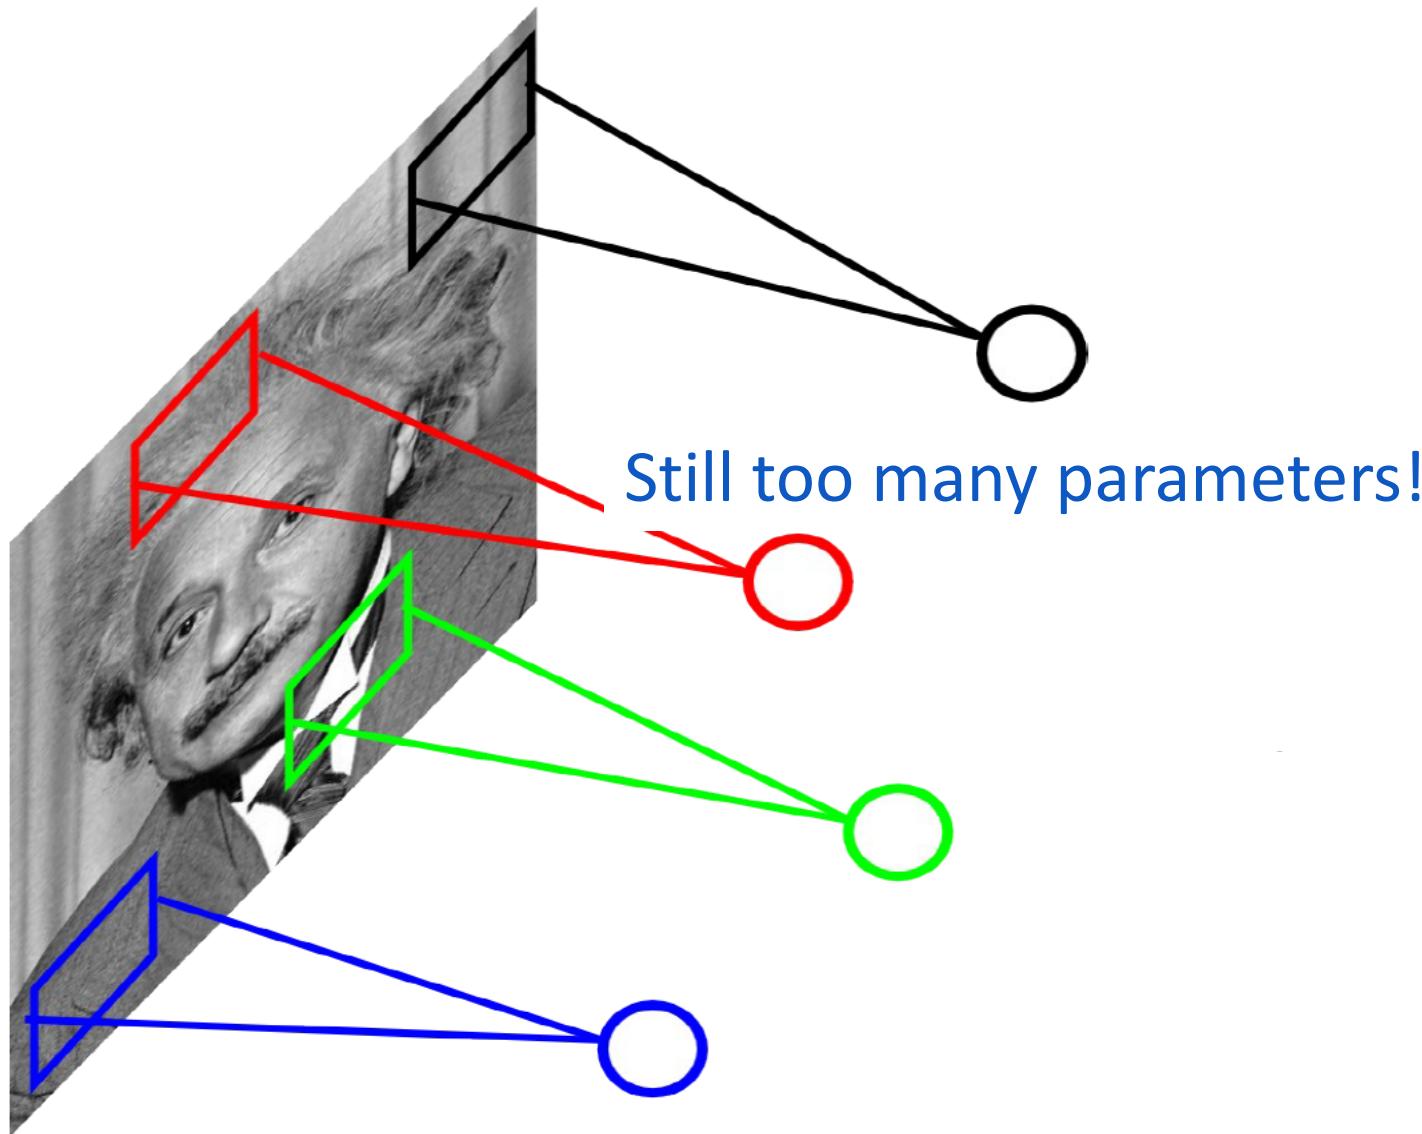
## Sparse coding example



- Method “invents” edge detection
- Automatically learns to represent an image in terms of the edges that appear in it
- Gives a more succinct, higher-level representation than the raw pixels
- Quantitatively similar to primary visual cortex (area V1) in brain

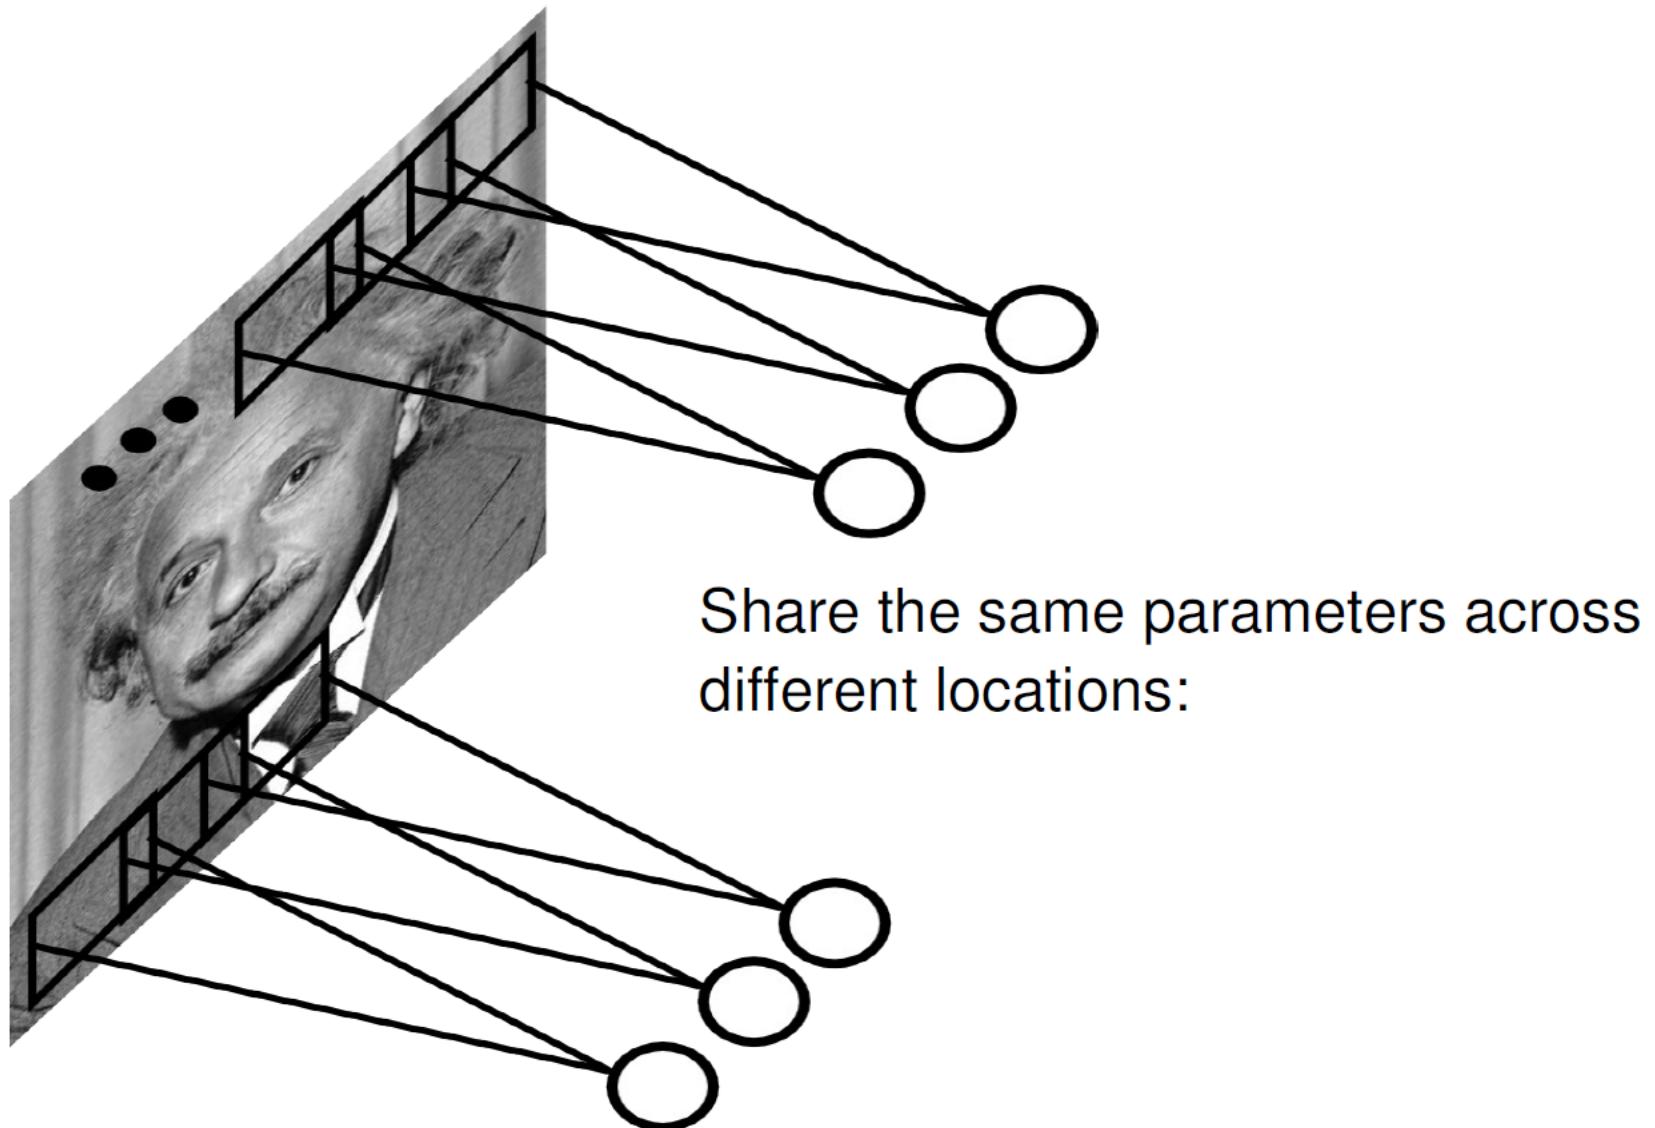
# Motivation: Images as a composition of local parts

## “Patch-based” representation



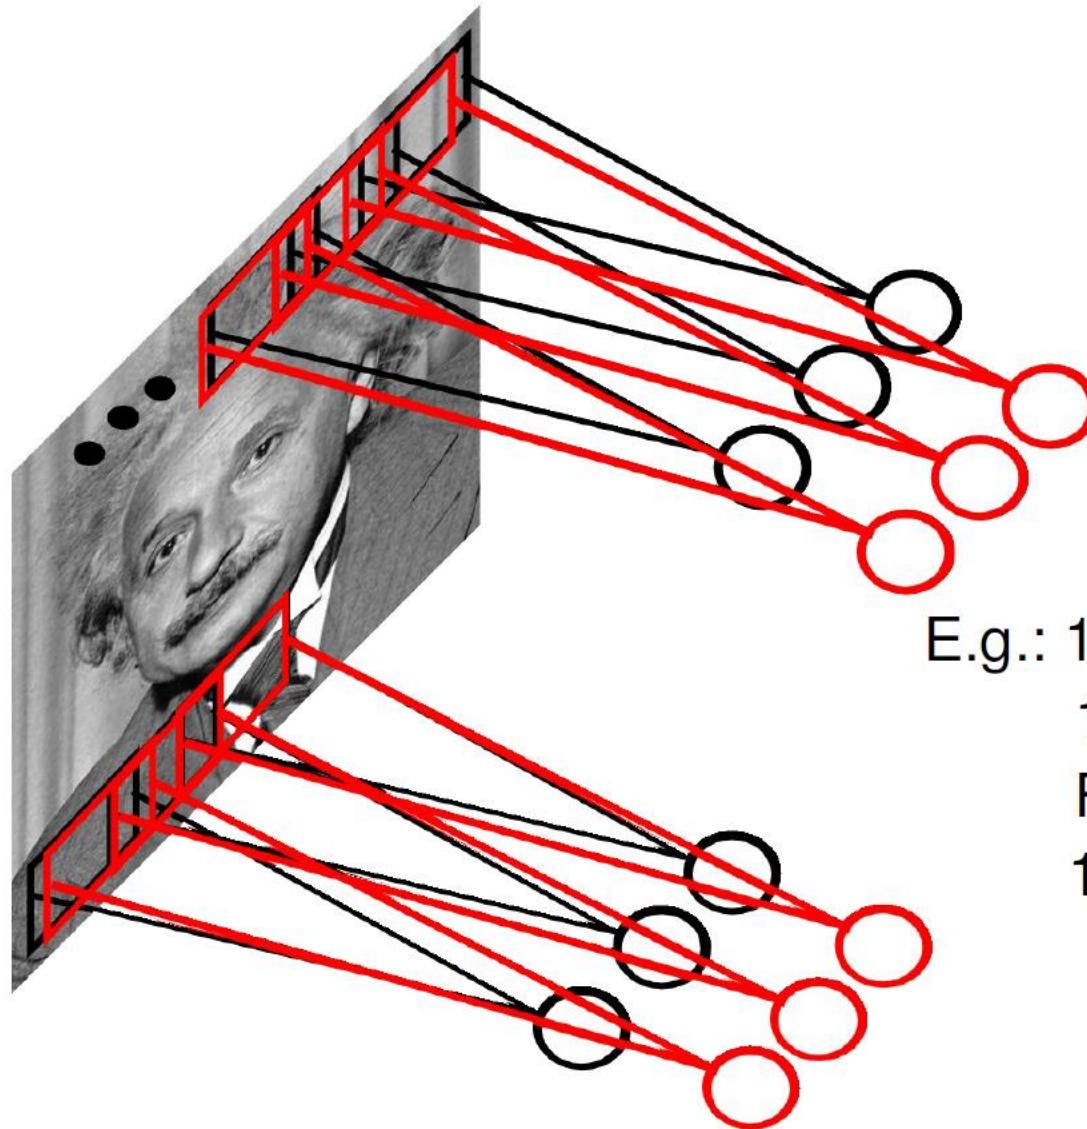
# Motivation: Images as a composition of local parts

## Convolution example



# Motivation: Images as a composition of local parts

## Convolution example



# Motivation: Images as a composition of local parts

## Convolution example

- Why convolution?

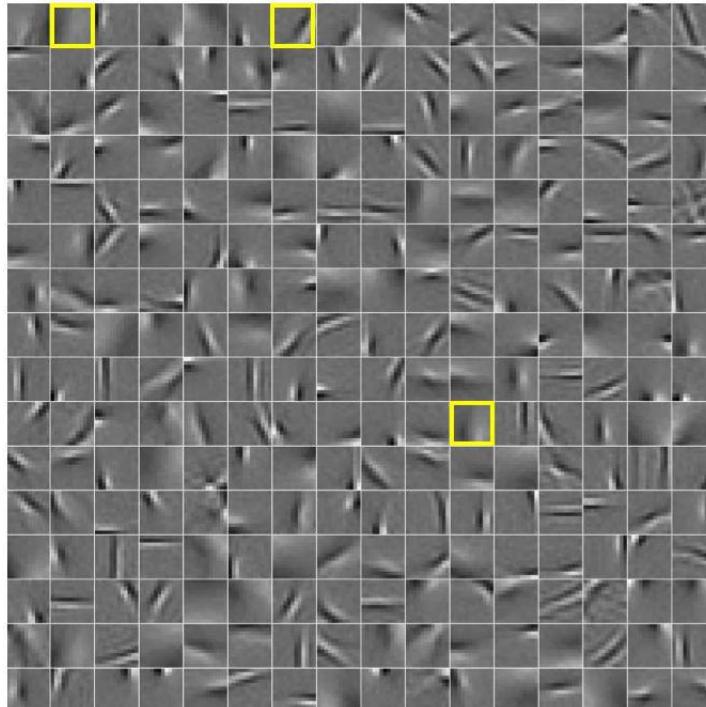
- Statistics of images look similar at different locations
- Dependencies are very local
- Filtering is an operation with translation *equivariance*



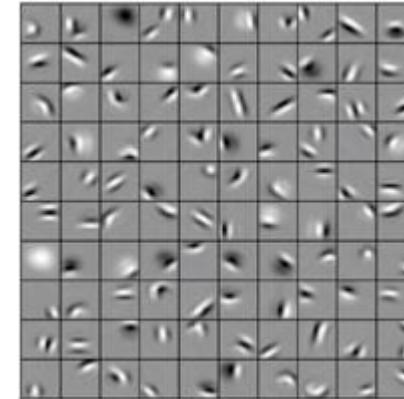
# Motivation: Images as a composition of local parts

## Filtering example

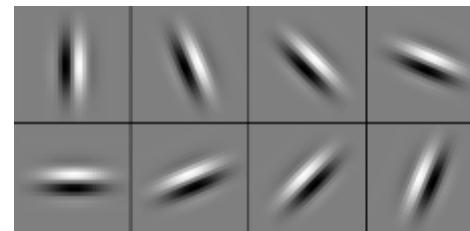
- Why translation *equivariance*?
  - Input translation leads to a translation of features
    - Fewer filters needed: no translated replications
    - But still need to cover orientation/frequency



Patch-based



Patch-based



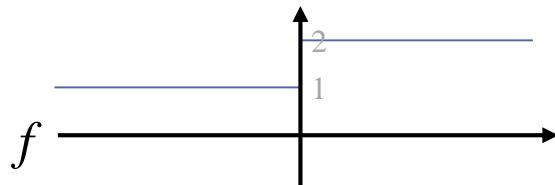
Convolutional

## Linear filtering

Further reading: Szeliski, Richard. *Computer Vision: Algorithms and Applications*. Springer, 2010, Chapter 3, Section 3.2

# Basics: Smoothing via local averaging

## Graphic depiction



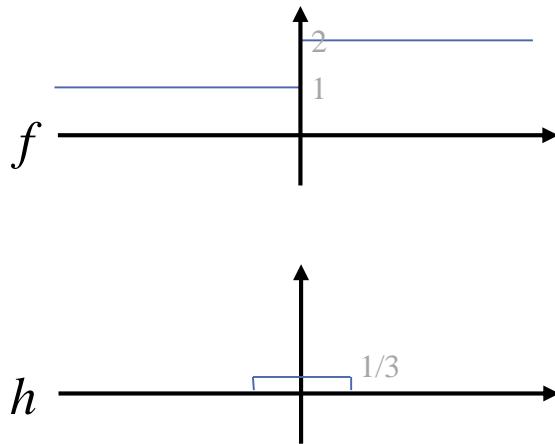
## Formalization

- We begin by considering a function:

$$f(x)$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Formalization

- We begin by considering a function:

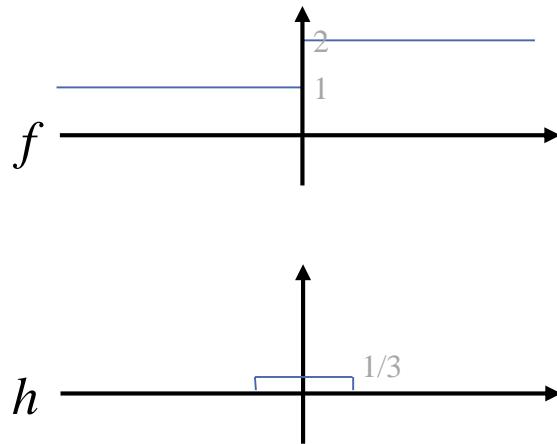
$$f(x)$$

- And we multiply it with values of another function:

$$f(x)h$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Formalization

- We begin by considering a function:

$$f(x)$$

- And we multiply it with values of another function:

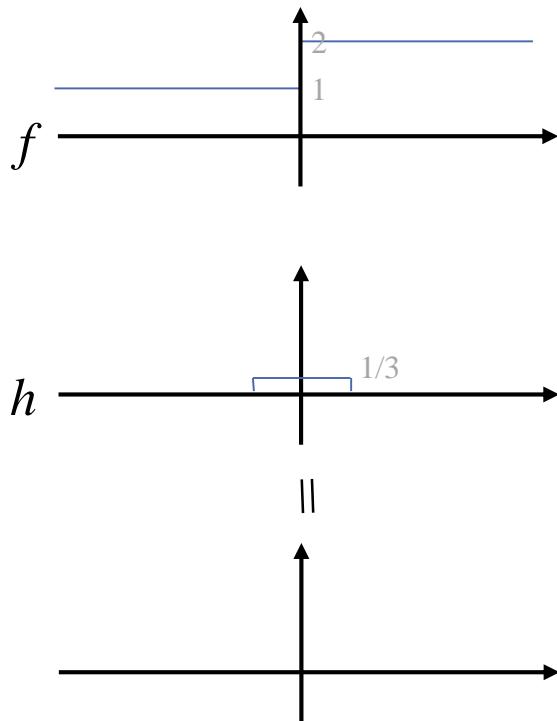
$$f(x)h$$

- But we do this at various offsets:

$$f(x-s)h(s)$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Formalization

- We begin by considering a function:

$$f(x)$$

- And we multiply it with values of another function:

$$f(x)h$$

- But we do this at various offsets:

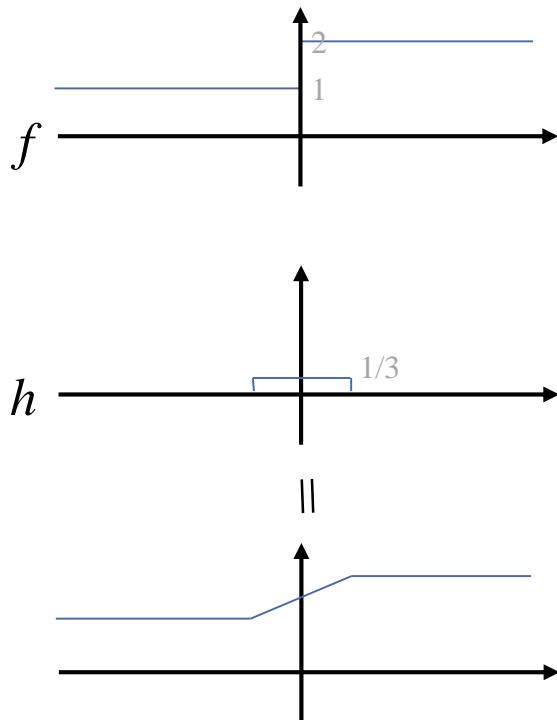
$$f(x-s)h(s)$$

- and multiply by infinitesimal support elements:

$$f(x-s)h(s)ds$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Formalization

- We begin by considering a function:

$$f(x)$$

- And we multiply it with values of another function:

$$f(x)h$$

- But we do this at various offsets:

$$f(x-s)h(s)$$

- and multiply by infinitesimal support elements:

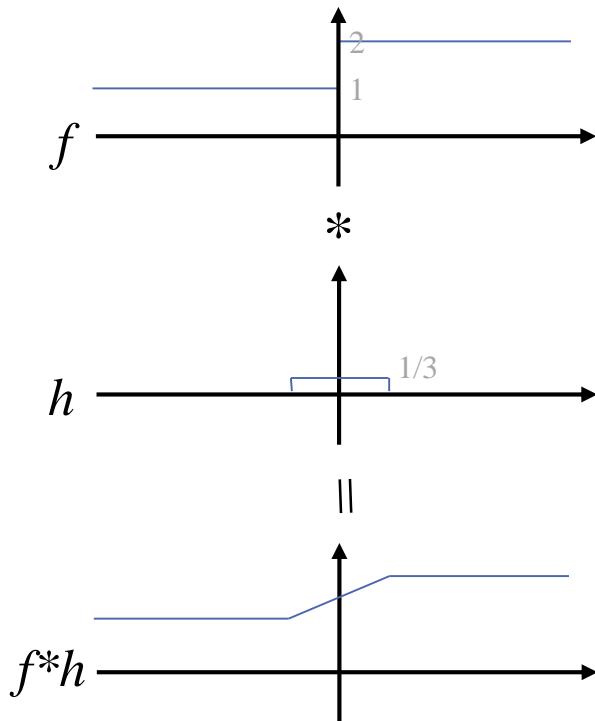
$$f(x-s)h(s)ds$$

- Finally, we sum up (integrate):

$$\int f(x-s)h(s)ds$$

# Basics: Smoothing via local averaging

## Graphic depiction



- We call this operation  $*$  a *convolution*

$$f*h = \int f(x-s)h(s)ds$$

## Formalization

- We begin by considering a function:

$$f(x)$$

- And we multiply it with values of another function:

$$f(x)h$$

- But we do this at various offsets:

$$f(x-s)h(s)$$

- and multiply by infinitesimal support elements:

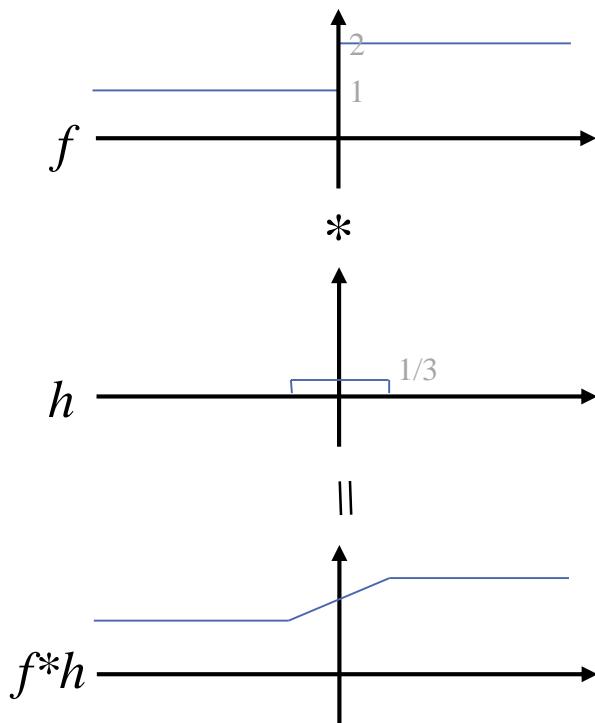
$$f(x-s)h(s)ds$$

- Finally, we sum up (integrate):

$$\int f(x-s)h(s)ds$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Numerical calculation

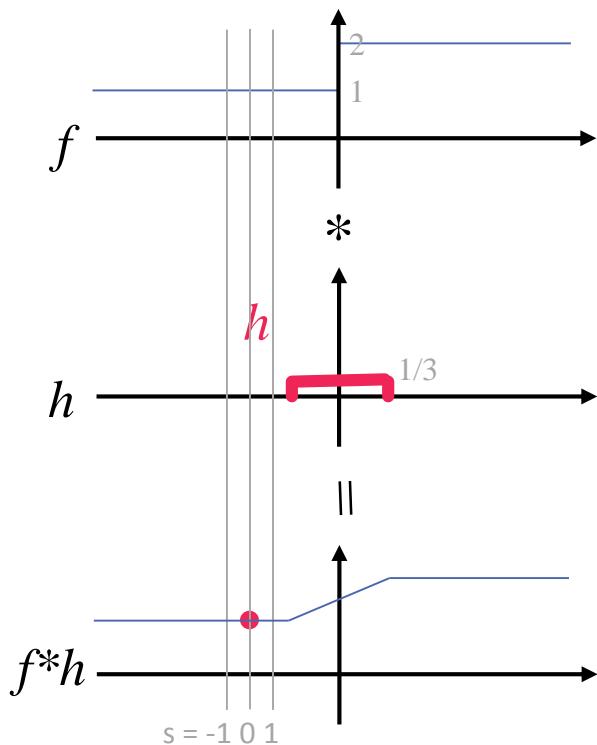
$x$	$s$	$f(x-s)h(s)$	$+$
$-2$	$-1$		
	$0$		
	$1$		
$-1$	$-1$		
	$0$		
	$1$		
$0$	$-1$		
	$0$		
	$1$		
$1$	$-1$		
	$0$		
	$1$		
$2$	$-1$		
	$0$		
	$1$		

## Formalization

$$\int f(x-s)h(s)ds$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Numerical calculation

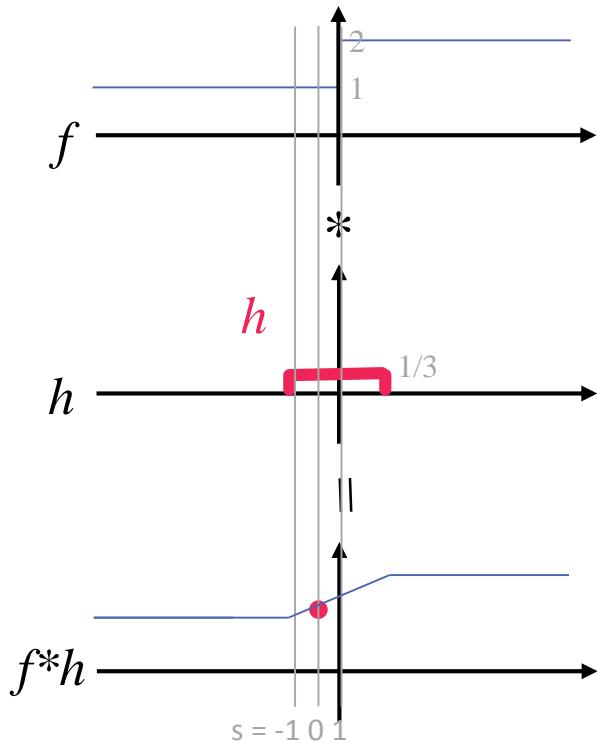
$x$	$s$	$f(x-s)h(s)$	$+$
$-2$	$-1$	$(1)(1/3)$	
	$0$	$(1)(1/3)$	
	$1$	$(1)(1/3)$	$1$
$-1$	$-1$		
	$0$		
	$1$		
$0$	$-1$		
	$0$		
	$1$		
$1$	$-1$		
	$0$		
	$1$		
$2$	$-1$		
	$0$		
	$1$		

## Formalization

$$\int f(x-s)h(s)ds$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Numerical calculation

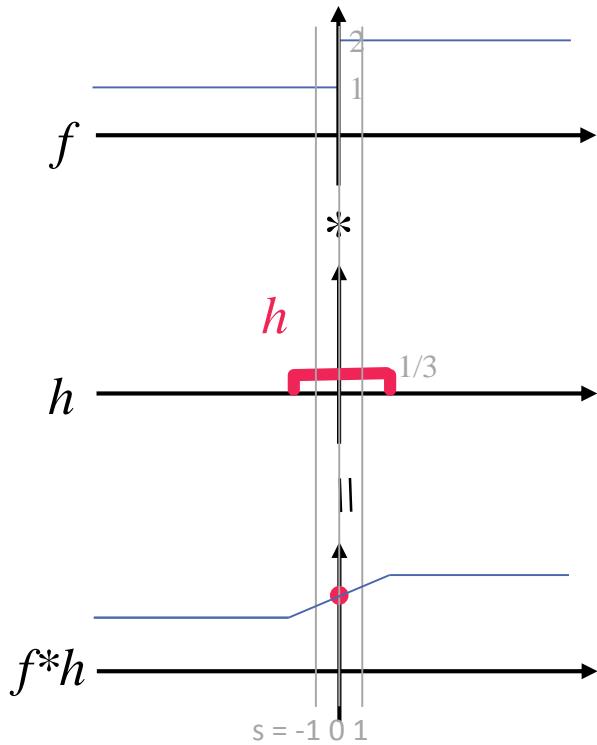
$x$	$s$	$f(x-s)h(s)$	$+$
-2	-1	$(1)(1/3)$	
	0	$(1)(1/3)$	
	1	$(1)(1/3)$	1
-1	-1	$(3/2)(1/3)$	
	0	$(1)(1/3)$	
	1	$(1)(1/3)$	7.0/6.0
0	-1		
	0		
	1		
1	-1		
	0		
	1		
2	-1		
	0		
	1		

## Formalization

$$\int f(x-s)h(s)ds$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Numerical calculation

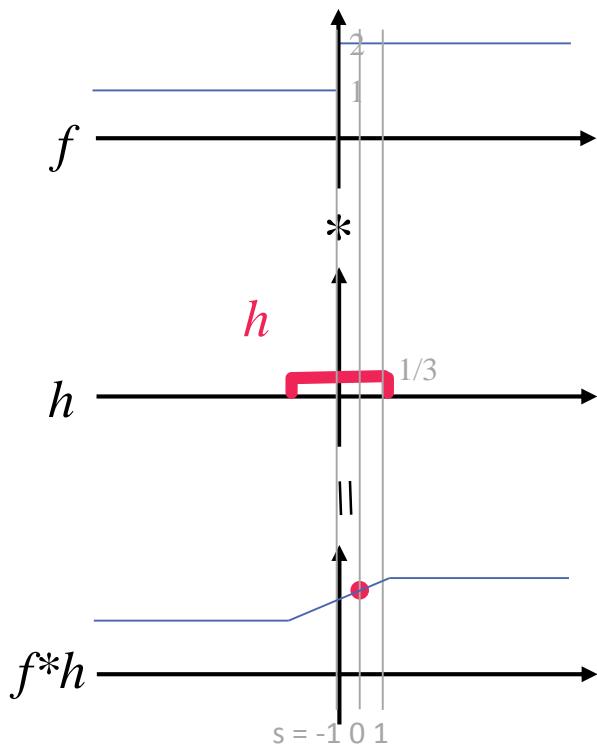
x	s	$f(x-s)h(s)$	+
-2	-1	(1)(1/3)	
	0	(1)(1/3)	
	1	(1)(1/3)	1
-1	-1	(3/2)(1/3)	
	0	(1)(1/3)	
	1	(1)(1/3)	7.0/6.0
0	-1	(2)(1/3)	
	0	(3/2)(1/3)	
	1	(1)(1/3)	3.0/2.0
1	-1		
	0		
	1		
2	-1		
	0		
	1		

## Formalization

$$\int f(x-s)h(s)ds$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Numerical calculation

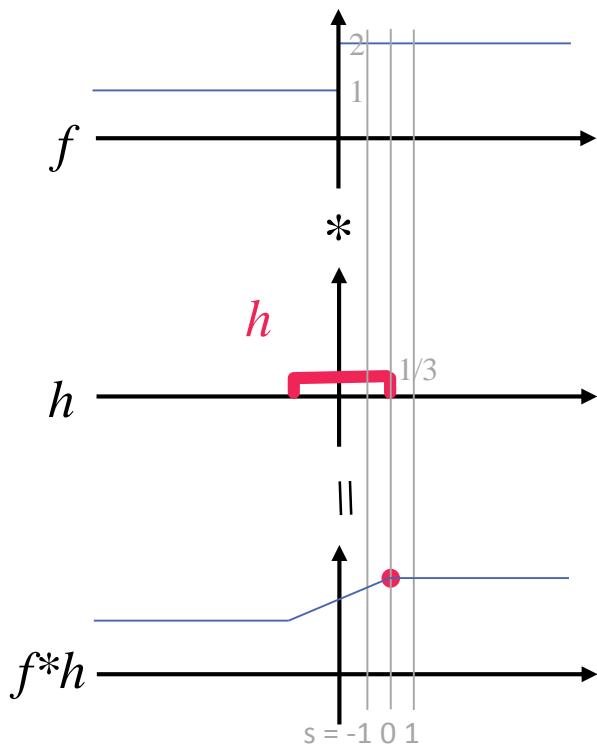
$x$	$s$	$f(x-s)h(s)$	$+$
-2	-1	$(1)(1/3)$	1
	0	$(1)(1/3)$	
	1	$(1)(1/3)$	
-1	-1	$(3/2)(1/3)$	7.0/6.0
	0	$(1)(1/3)$	
	1	$(1)(1/3)$	
0	-1	$(2)(1/3)$	3.0/2.0
	0	$(3/2)(1/3)$	
	1	$(1)(1/3)$	
1	-1	$(2)(1/3)$	11.0/6.0
	0	$(2)(1/3)$	
	1	$(3/2)(1/3)$	
2	-1		
	0		
	1		

## Formalization

$$\int f(x-s)h(s)ds$$

# Basics: Smoothing via local averaging

## Graphic depiction



## Numerical calculation

x	s	$f(x-s)h(s)$	+
-2	-1	$(1)(1/3)$	1
	0	$(1)(1/3)$	
	1	$(1)(1/3)$	
-1	-1	$(3/2)(1/3)$	7.0/6.0
	0	$(1)(1/3)$	
	1	$(1)(1/3)$	
0	-1	$(2)(1/3)$	3.0/2.0
	0	$(3/2)(1/3)$	
	1	$(1)(1/3)$	
1	-1	$(2)(1/3)$	11.0/6.0
	0	$(2)(1/3)$	
	1	$(2)(1/3)$	
2	-1	$(2)(1/3)$	2
	0	$(2)(1/3)$	
	1	$(2)(1/3)$	

## Formalization

$$\int f(x-s)h(s)ds$$

# Basics: 2D Convolution

## Definition

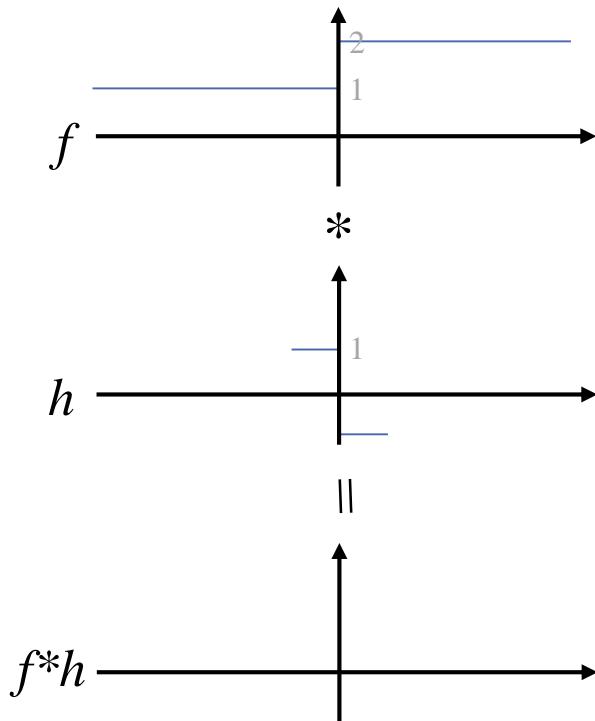
- Consider a system that, given  $f(x, y)$  as input, produces output

$$g(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x - \xi, y - \eta) h(\xi, \eta) d\xi d\eta$$

- We say that  $g$  is the **convolution** of  $f$  and  $h$ , written as  $g = f * h$ .

## Basics: Another convolution example

Graphic depiction



Numerical calculation

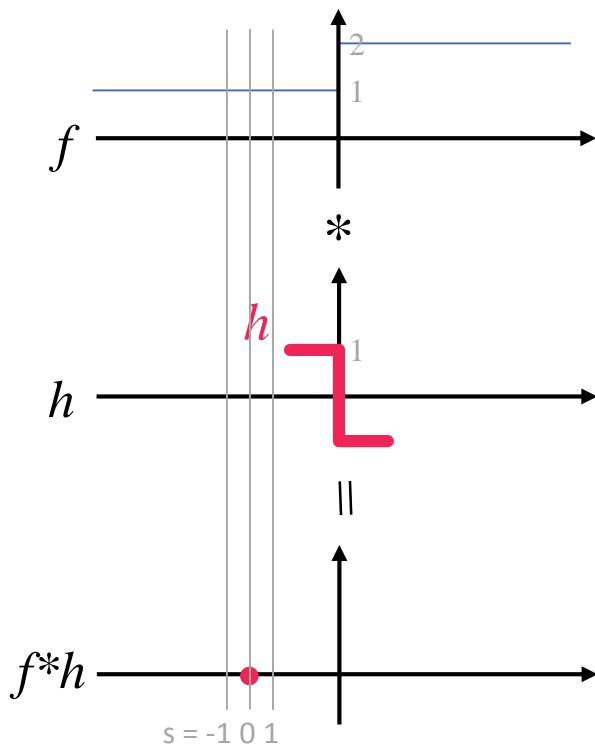
x	s	$f(x-s)h(s)$	+
-2	-1		
	0		
	1		
-1	-1		
	0		
	1		
0	-1		
	0		
	1		
1	-1		
	0		
	1		
2	-1		
	0		
	1		

Formalization

$$\int_{-\infty}^{\infty} f(x-s)h(s)ds$$

## Basics: Another convolution example

### Graphic depiction



### Numerical calculation

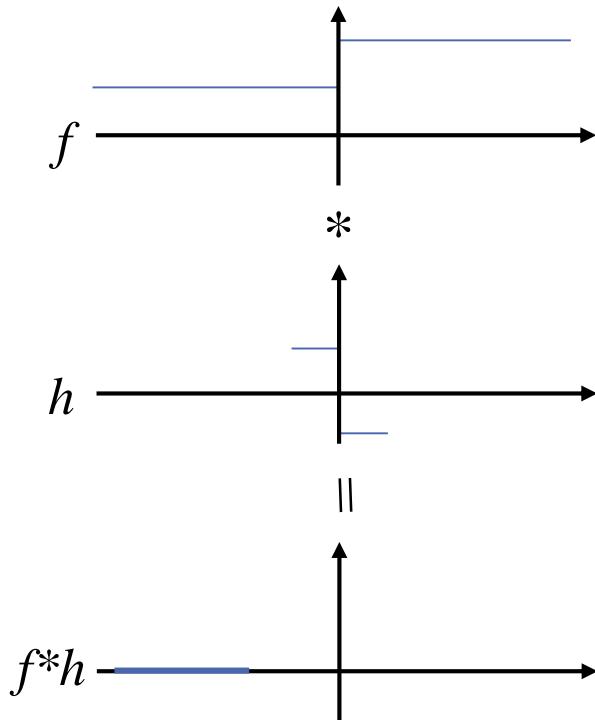
$x$	$s$	$f(x-s)h(s)$	$+$
-2	-1	(1)(1)	
	0	(1)(0)	
	1	(1)(-1)	0
-1	-1		
	0		
	1		
0	-1		
	0		
	1		
1	-1		
	0		
	1		
2	-1		
	0		
	1		

### Formalization

$$\int_{-\infty}^{\infty} f(x-s)h(s)ds$$

## Basics: Another convolution example

### Graphic depiction



### Numerical calculation

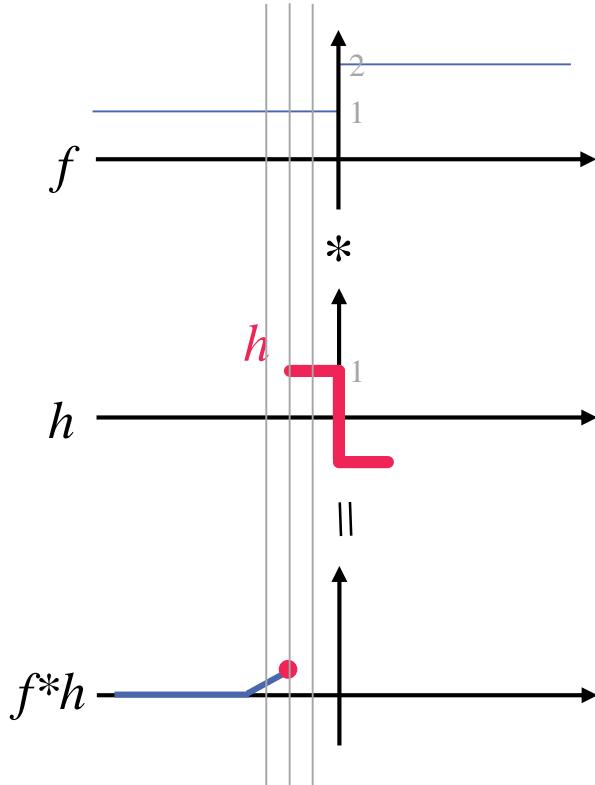
x	s	$f(x-s)h(s)$	+
-2	-1	(1)(1)	
	0	(1)(0)	
	1	(1)(-1)	0
-1	-1		
	0		
	1		
0	-1		
	0		
	1		
1	-1		
	0		
	1		
2	-1		
	0		
	1		

### Formalization

$$\int_{-\infty}^{\infty} f(x-s)h(s)ds$$

# Basics: Another convolution example

## Graphic depiction



## Numerical calculation

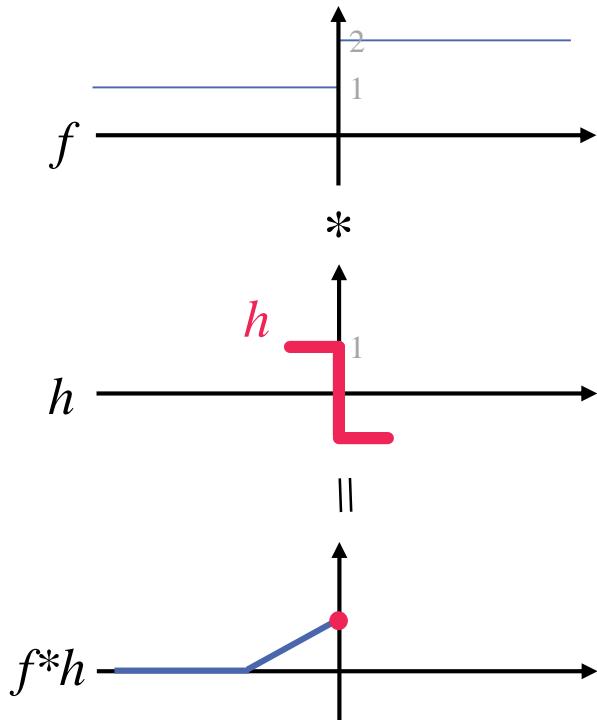
$x$	$s$	$f(x-s)h(s)$	$+$
-2	-1	(1)(1)	
	0	(1)(0)	
	1	(1)(-1)	0
-1	-1	(3/2)(1)	
	0	(1)(0)	
	1	(1)(-1)	1/2
0	-1		
	0		
	1		
1	-1		
	0		
	1		
2	-1		
	0		
	1		

## Formalization

$$\int_{-\infty}^{\infty} f(x-s)h(s)ds$$

# Basics: Another convolution example

## Graphic depiction



## Numerical calculation

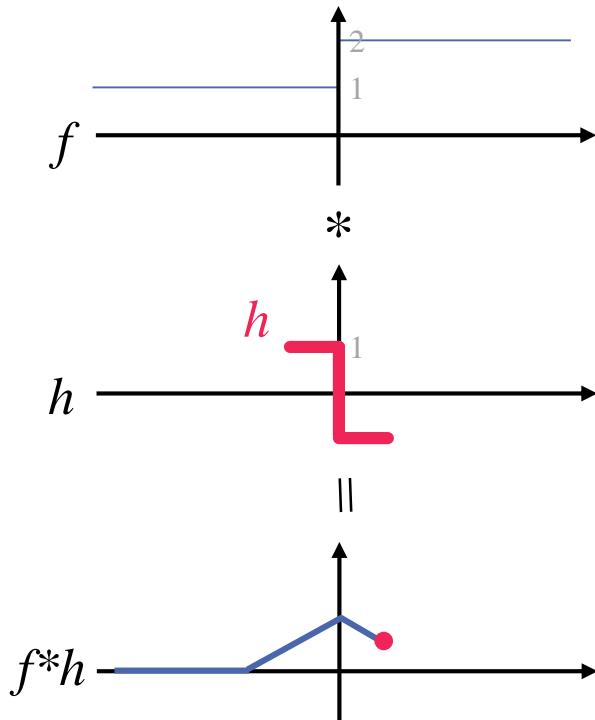
x	s	$f(x-s)h(s)$	+
-2	-1	(1)(1)	
	0	(1)(0)	
	1	(1)(-1)	0
-1	-1	(3/2)(1)	
	0	(1)(0)	
	1	(1)(-1)	1/2
0	-1	(2)(1)	
	0	(3/2)(0)	
	1	(1)(-1)	1
1	-1		
	0		
	1		
2	-1		
	0		
	1		

## Formalization

$$\int_{-\infty}^{\infty} f(x-s)h(s)ds$$

# Basics: Another convolution example

## Graphic depiction



## Numerical calculation

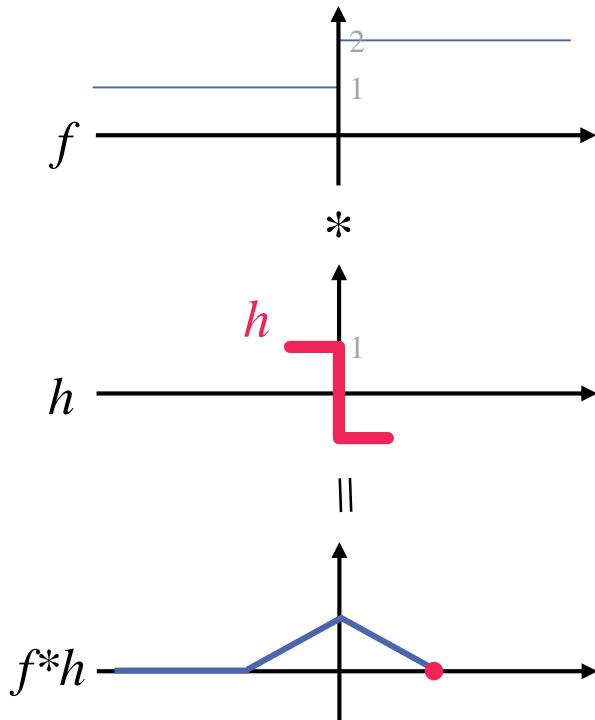
x	s	$f(x-s)h(s)$	+
-2	-1	(1)(1)	
	0	(1)(0)	
	1	(1)(-1)	0
-1	-1	(3/2)(1)	
	0	(1)(0)	
	1	(1)(-1)	1/2
0	-1	(2)(1)	
	0	(3/2)(0)	
	1	(1)(-1)	1
1	-1	(2)(1)	
	0	(2)(0)	
	1	(3/2)(-1)	1/2
2	-1		
	0		
	1		

## Formalization

$$\int_{-\infty}^{\infty} f(x-s)h(s)ds$$

# Basics: Another convolution example

## Graphic depiction



## Numerical calculation

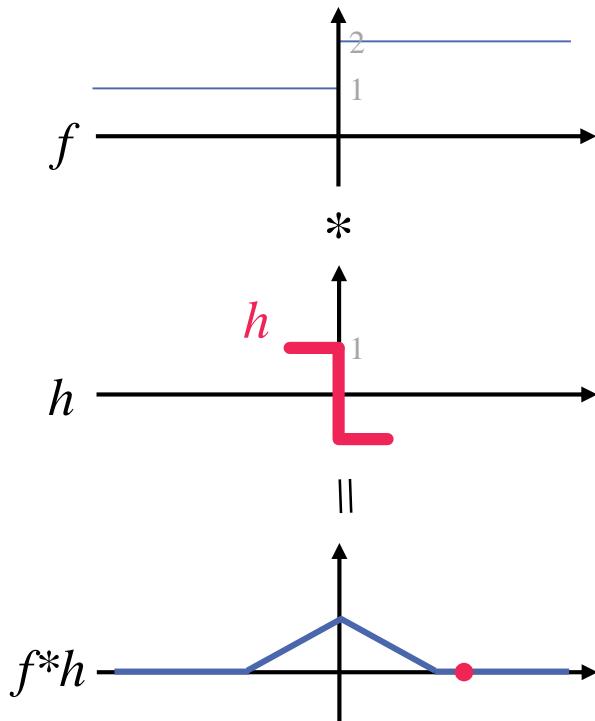
x	s	$f(x-s)h(s)$	+
-2	-1	(1)(1)	
	0	(1)(0)	
	1	(1)(-1)	0
-1	-1	(3/2)(1)	
	0	(1)(0)	
	1	(1)(-1)	1/2
0	-1	(2)(1)	
	0	(3/2)(0)	
	1	(1)(-1)	1
1	-1	(2)(1)	
	0	(2)(0)	
	1	(3/2)(-1)	1/2
2	-1	(2)(1)	
	0	(2)(0)	
	1	(2)(-1)	0

## Formalization

$$\int_{-\infty}^{\infty} f(x-s)h(s)ds$$

# Basics: Another convolution example

## Graphic depiction



## Numerical calculation

$x$	$s$	$f(x-s)h(s)$	$+$
-2	-1	(1)(1)	
	0	(1)(0)	
	1	(1)(-1)	0
-1	-1	(3/2)(1)	
	0	(1)(0)	
	1	(1)(-1)	1/2
0	-1	(2)(1)	
	0	(3/2)(0)	
	1	(1)(-1)	1
1	-1	(2)(1)	
	0	(2)(0)	
	1	(3/2)(-1)	1/2
2	-1	(2)(1)	
	0	(2)(0)	
	1	(2)(-1)	0

## Formalization

$$\int_{-\infty}^{\infty} f(x-s)h(s)ds$$

# Basics: Convolution

## Definition

- Consider a system that, given  $f(x,y)$  as input, produces output

$$g(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x - \xi, y - \eta) h(\xi, \eta) d\xi d\eta$$

- We say that  $g$  is the **convolution** of  $f$  and  $h$ , written as  $g = f * h$ .

## Convolution is linear

- Applying the system to  $(a f1(x,y) + b f2(x,y))$  yields  $(a g1(x,y) + b g2(x,y))$ .
- Follows from rule for integrating the product of a constant and a function
- and the rule for integrating the sum of two functions.

$$\int [a\alpha(\xi) + b\beta(\xi)] d\xi = a \int \alpha(\xi) d\xi + b \int \beta(\xi) d\xi$$

# Basics: Convolution

## Definition

- Consider a system that, given  $f(x,y)$  as input, produces output

$$g(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x - \xi, y - \eta) h(\xi, \eta) d\xi d\eta$$

- We say that  $g$  is the **convolution** of  $f$  and  $h$ , written as  $g = f * h$ .

## Convolution is linear

- Applying the system to  $(a f1(x,y) + b f2(x,y))$  yields  $(a g1(x,y) + b g2(x,y))$ .
- Follows from rule for integrating the product of a constant and a function
- and the rule for integrating the sum of two functions.

## Convolution is shift invariant

- Applying the system to  $f(x-a, y-b)$  yields  $g(x-a, y-b)$ .
- Follows from the convolution integral being independent of  $(x, y)$
- So a change of variables  $(x, y) \rightarrow (x-a, y-b) = (x', y')$  just shifts the result.
- Another way to think of shift invariance is that the operation (e.g.  $*$ ) “behaves the same everywhere”
- For images: The value of the output depends on the pattern in the image neighborhood, not the position of the neighborhood

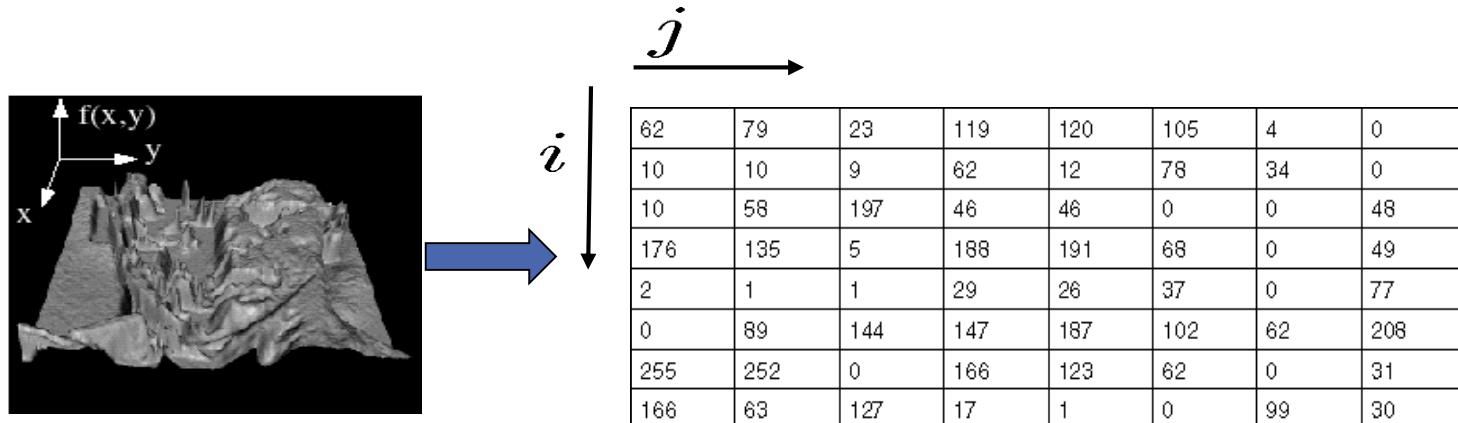
## Images as functions

- We can think of an image as a function,  $f$ , from  $\mathbf{R}^2$  to  $\mathbf{R}$ :
  - $f(x, y)$  gives the intensity at position  $(x, y)$
  - Realistically, we expect the image only to be defined over a rectangle, with a finite range:
    - $f: [a,b] \times [c,d] \rightarrow [0, 1.0]$
- A color image is just three functions pasted together. We can write this as a “vector-valued” function:

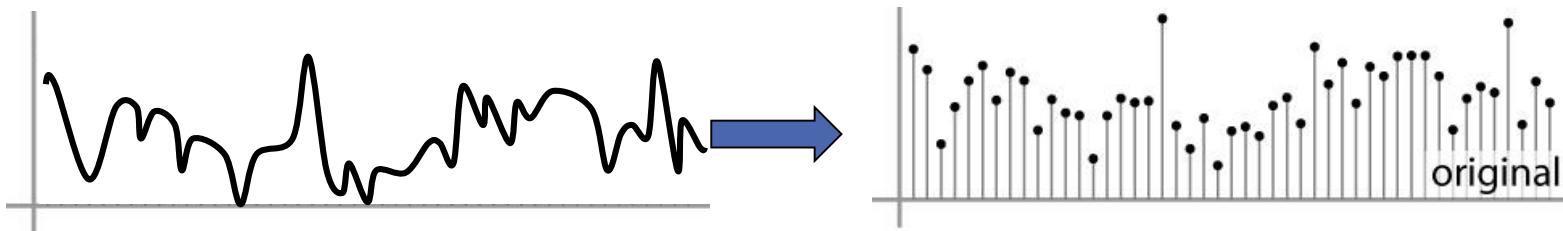
$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

# Digital images

- In computer vision we operate on **digital (discrete)** images:
  - **Sample** the 2D space on a regular grid
  - **Quantize** each sample (round to nearest integer)
- Image thus represented as a matrix of integer values.



2D



1D

## Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some function

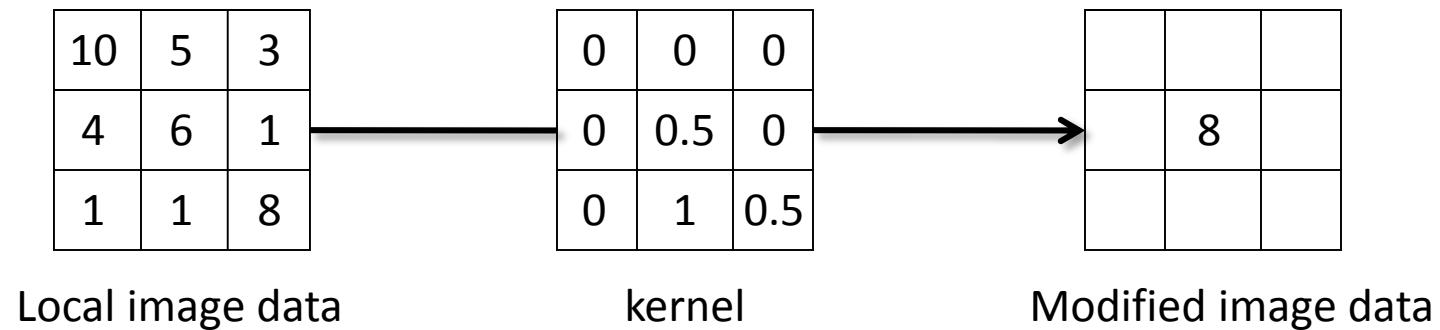


	7	

Modified image data

## Linear filtering

- One simple version: linear filtering (cross-correlation, convolution)
  - Replace each pixel by a linear combination (a weighted sum) of its neighbors
- The prescription for the linear combination is called the “kernel” (or “mask”, “filter”)



## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$


## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0						

## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$g =$

			0	10						

## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0	10	20				

## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$


## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$


## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0	10	20	30	30	30	

## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

	0	10	20	30	30	30			

## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0	10	20	30	30	30	

## Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

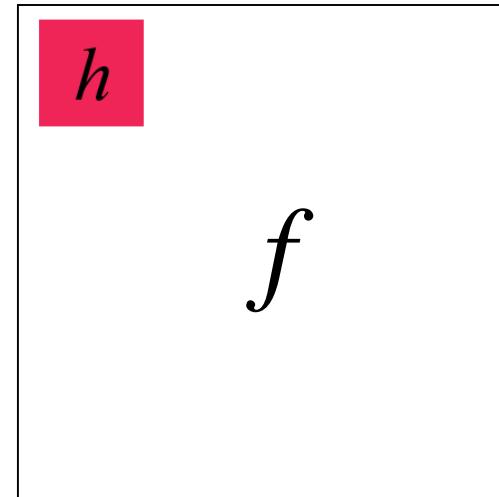
$g =$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

## Convolution vs. correlation

Let  $f$  be the image and  $g$  be the kernel, the cross-correlation operation  $\otimes$  is defined as

$$g(x, y) = \sum_{k, l} f(x + k, y + l)h(k, l)$$



**Note:** We have defined convolution as  $g=f*h$

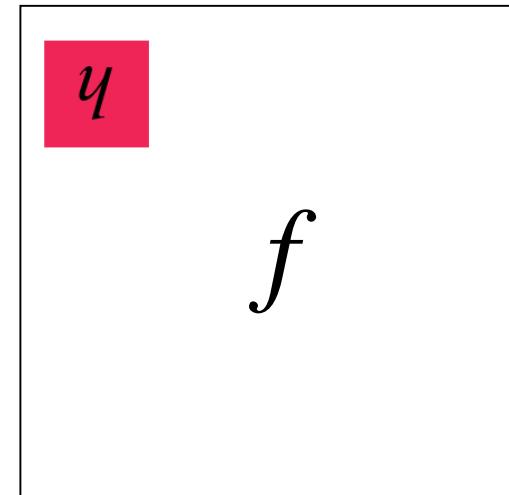
$$g(x, y) = \sum_{k, l} f(x - k, y - l)h(k, l)$$

Filter kernel is “flipped” in both dimensions (bottom to top, right to left)



Then cross-correlation is applied

For a symmetric kernel, how will the outputs differ?



If the input is an impulse signal, how will the outputs differ?

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$


$$G[x, y]$$

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G = H \otimes F$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0

$$G[x, y]$$

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G = H \otimes F$$

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

$$G[x, y]$$

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G = H \otimes F$$

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	i	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

$$G[x, y]$$

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G = H \otimes F$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	i	h			

$$G[x, y]$$

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G = H \otimes F$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	i	h	g	0	0
0	0	f	e			

$$G[x, y]$$

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

$$F[x, y]$$

$$G = H \otimes F$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	i	h	g	0	0
0	0	f	e	d	0	0
0	0	c	b	a		

$$G[x, y]$$

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	i	h	g	0	0
0	0	f	e	d	0	0
0	0	c	b	a	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$$F[x, y]$$

$$G = H \otimes F$$

$$G[x, y]$$

## Filtering an impulse signal

What is the result of filtering the impulse signal (image)  $F$  with the arbitrary kernel  $H$ ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0



a	b	c
d	e	f
g	h	i

$$H[u, v]$$

		i	h	g		
		f	e	d		
		c	b	a		

$$F[x, y]$$

$$G = H \otimes F$$

$$G[x, y]$$

Filter output is reversed.

# Convolution

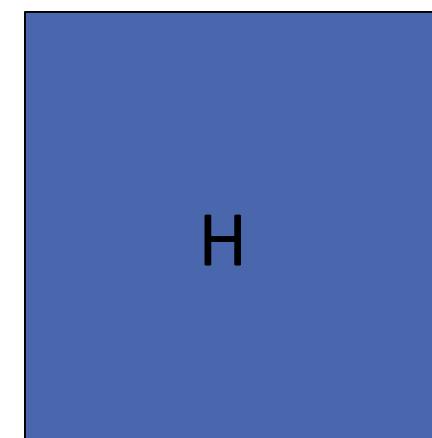
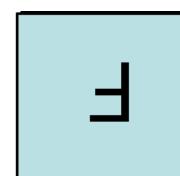
- Convolution:
  - Flip the filter in both dimensions (bottom to top, right to left)
  - Then apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$



*Notation for  
convolution  
operator*



H

## Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

## Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)

## Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?

## Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted left  
By 1 pixel

## Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)

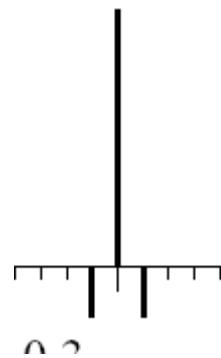
## Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$


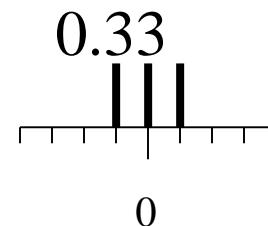
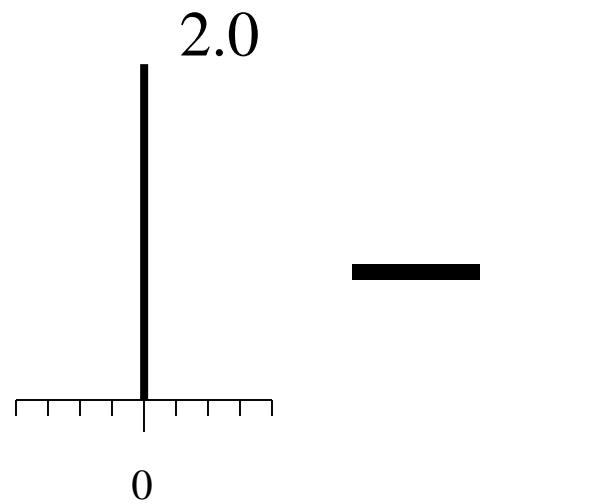
### Sharpening filter

- Accentuates differences with local average

# Sharpening

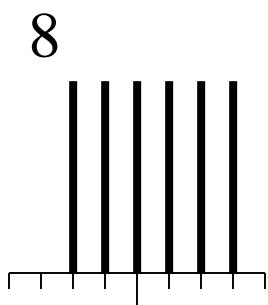


original

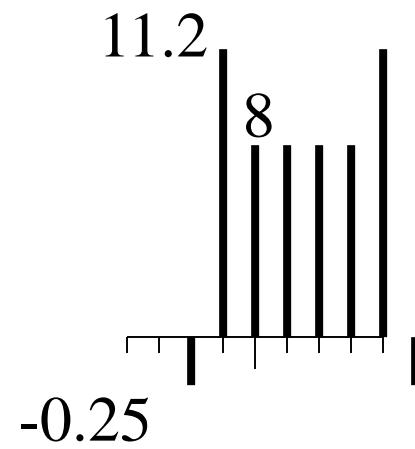
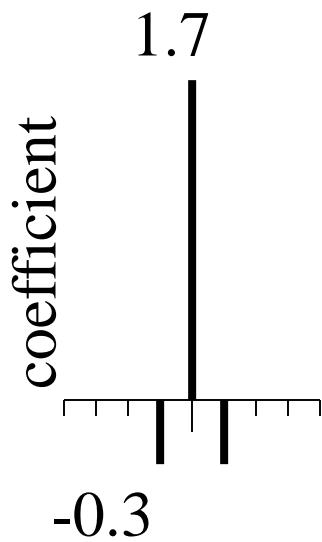


Sharpened  
original

## Sharpening example

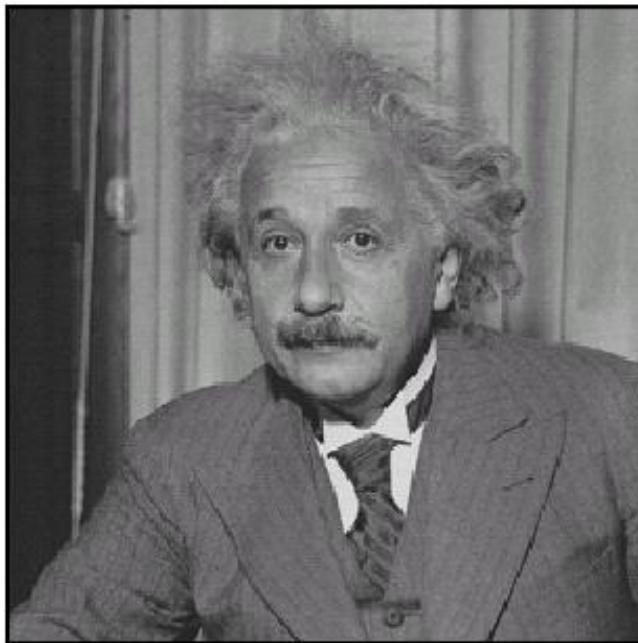


original

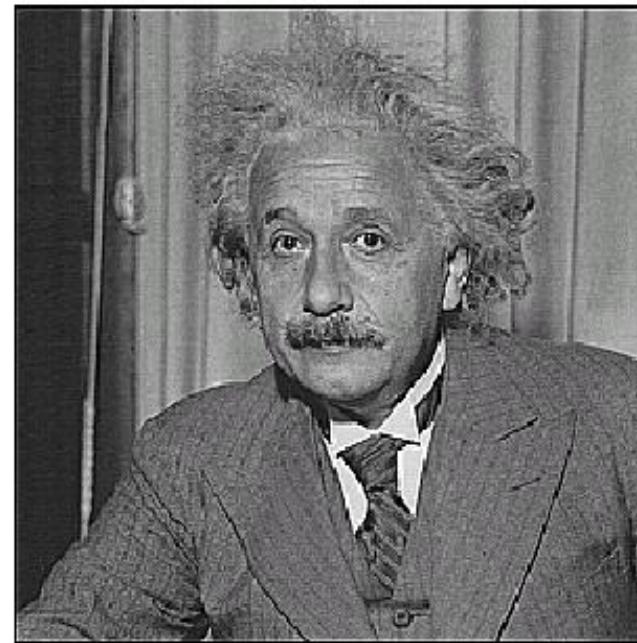


Sharpened  
(differences are  
accentuated; constant  
areas are left untouched).

# Sharpening

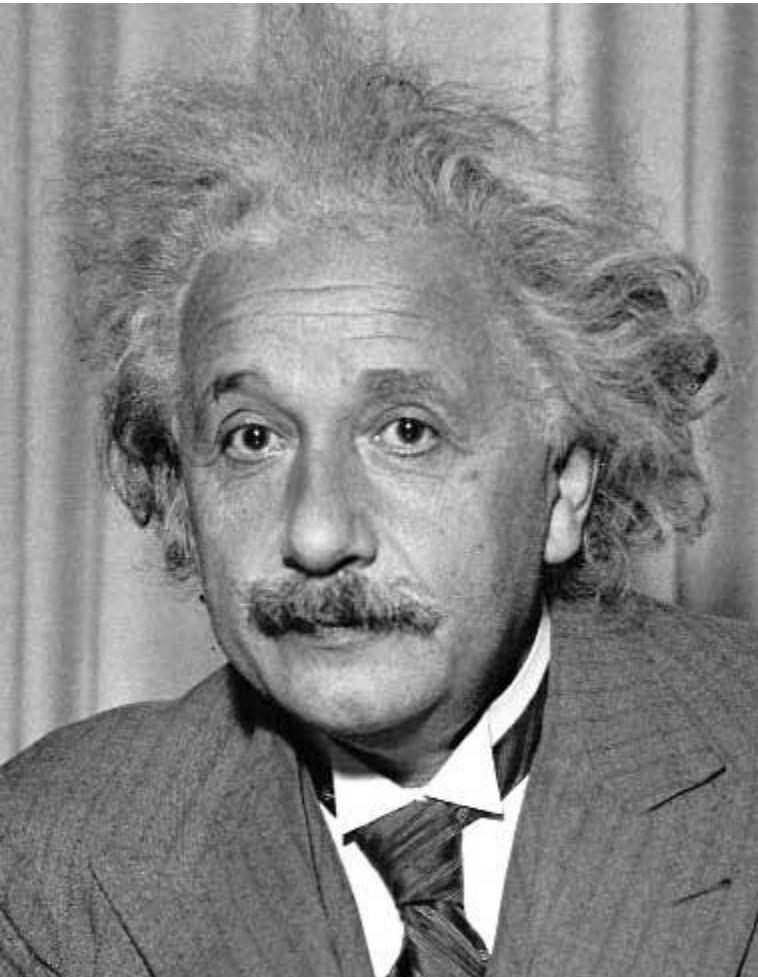


**before**



**after**

## Other filters



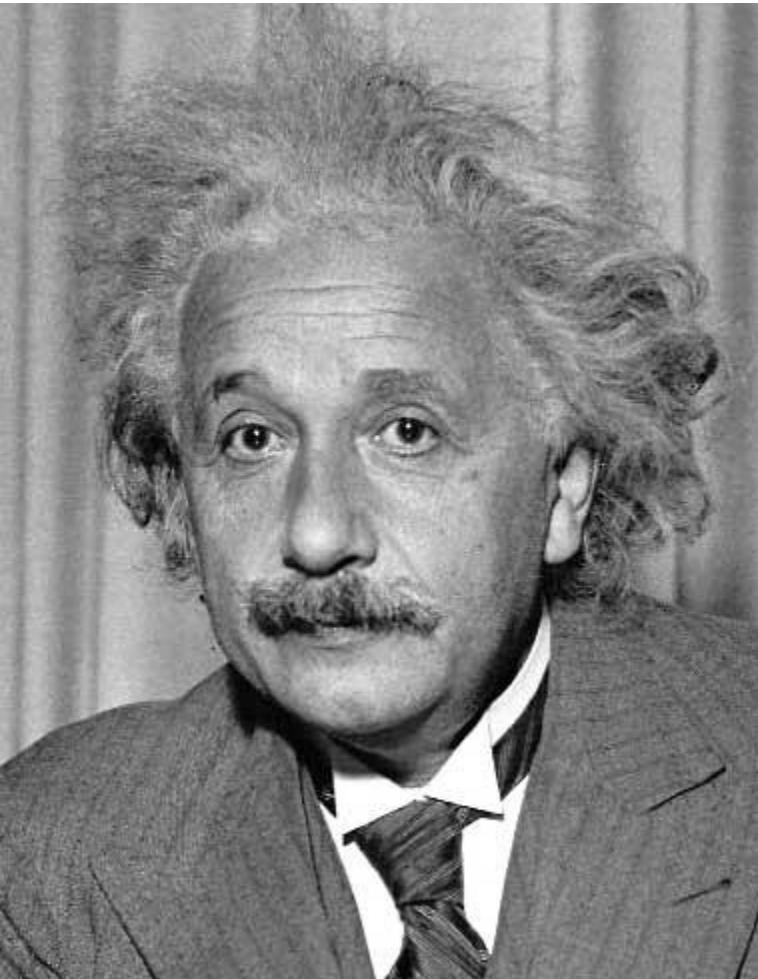
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge  
(absolute value)

## Other filters



1	2	1
0	0	0
-1	-2	-1

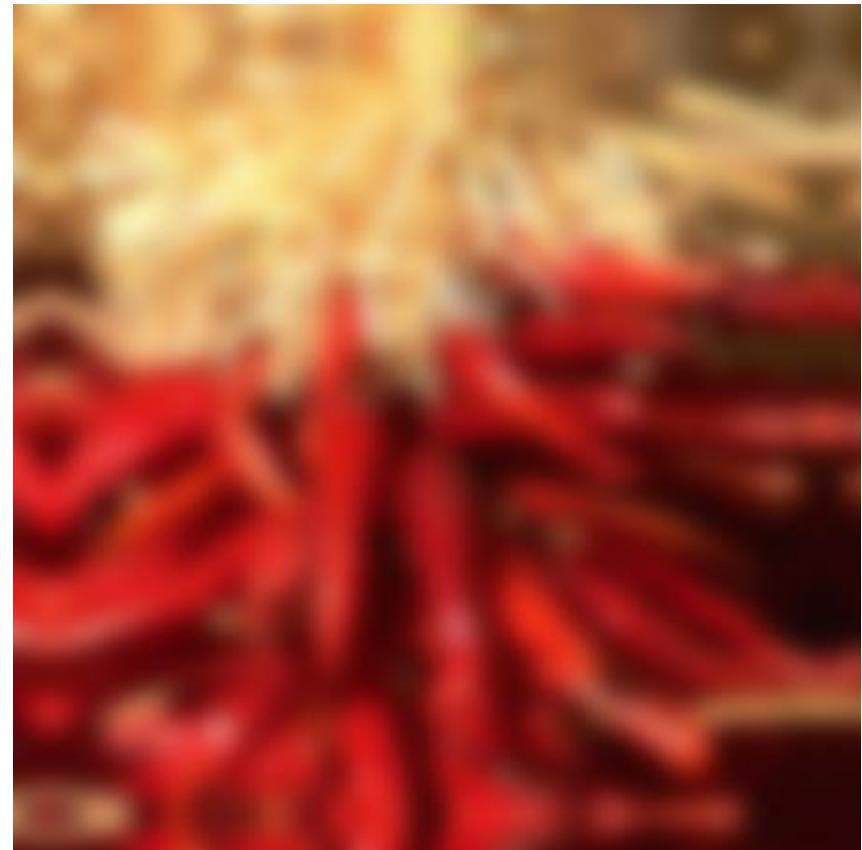
Sobel



Horizontal Edge  
(absolute value)

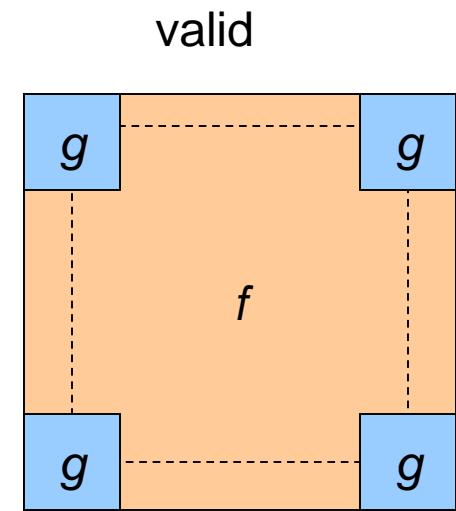
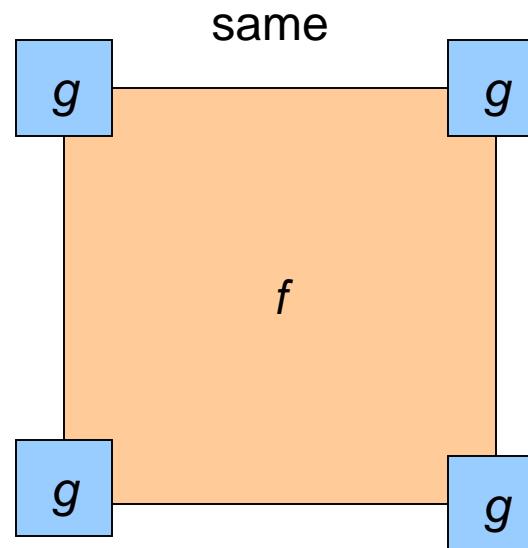
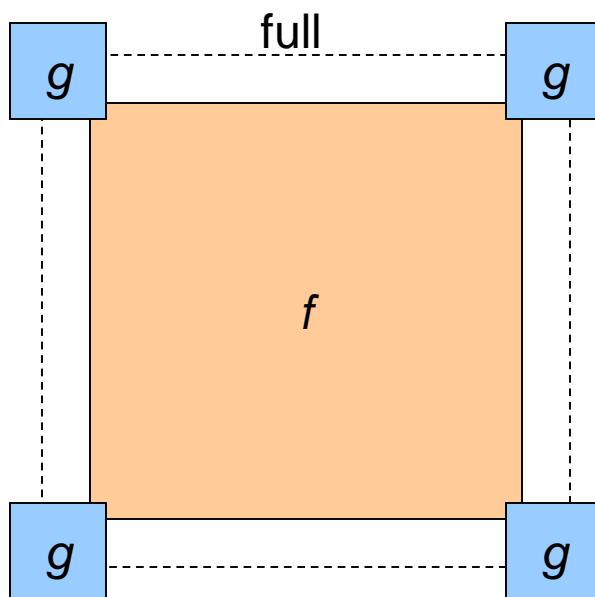
## Border treatment

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge

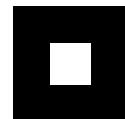


## Border treatment

- What is the size of the output?
  - *shape* = ‘full’: output size is sum of sizes of *f* and *g*
  - *shape* = ‘same’: output size is same as *f*
  - *shape* = ‘valid’: output size is difference of sizes of *f* and *g*



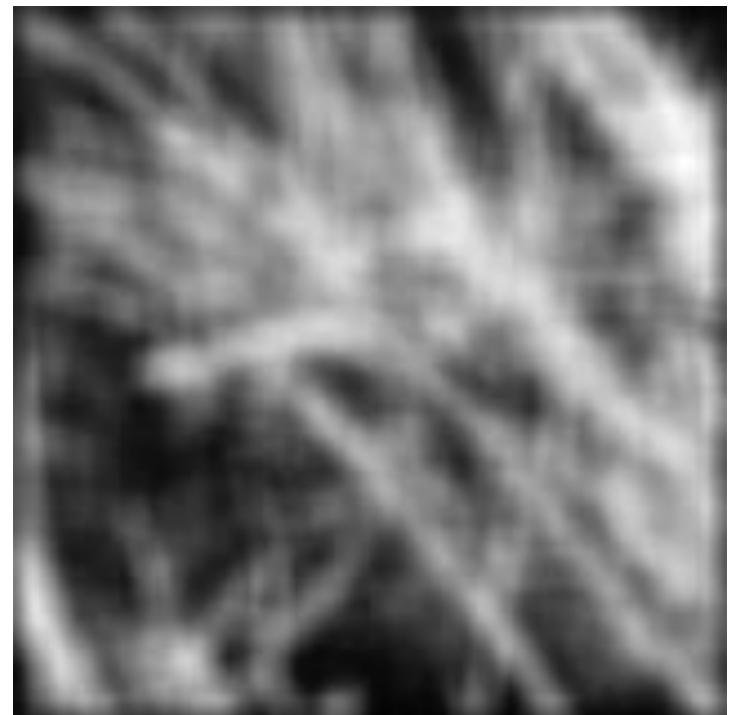
## Smoothing by averaging



depicts box filter:  
white = high value, black = low value



original



filtered

## Important filter: Gaussian

- What if we want nearest neighboring pixels to have the most influence on the output?

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

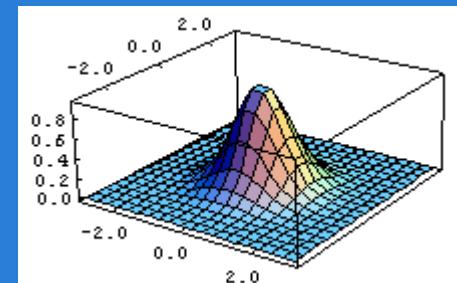
$$F[x, y]$$

$$\frac{1}{16} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

$H[u, v]$

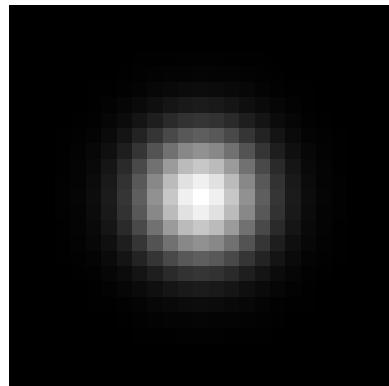
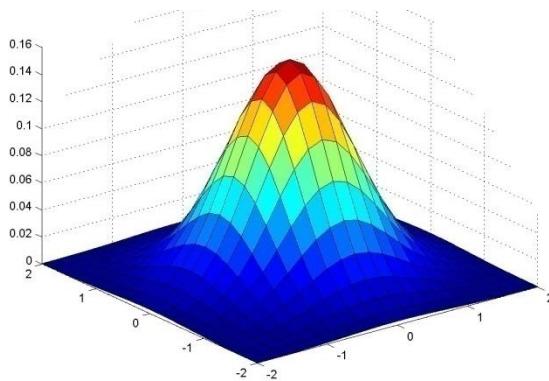
This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



## Important filter: Gaussian

- Weight contributions of neighboring pixels by nearness



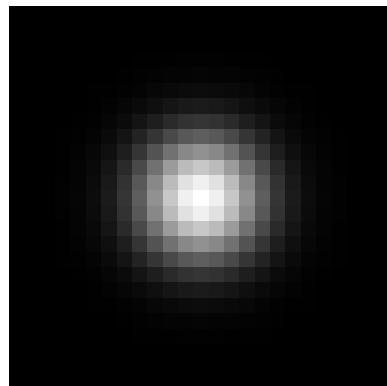
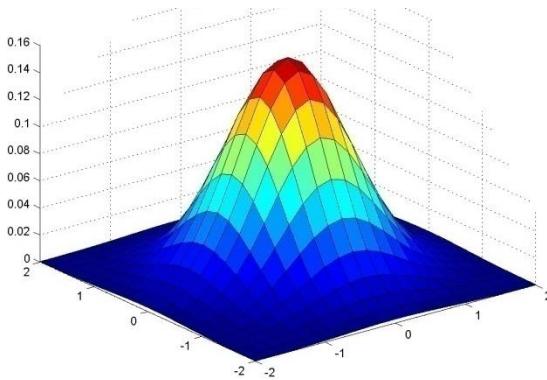
0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

## Important filter: Gaussian

- **Gaussian filters have some interesting properties**



$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

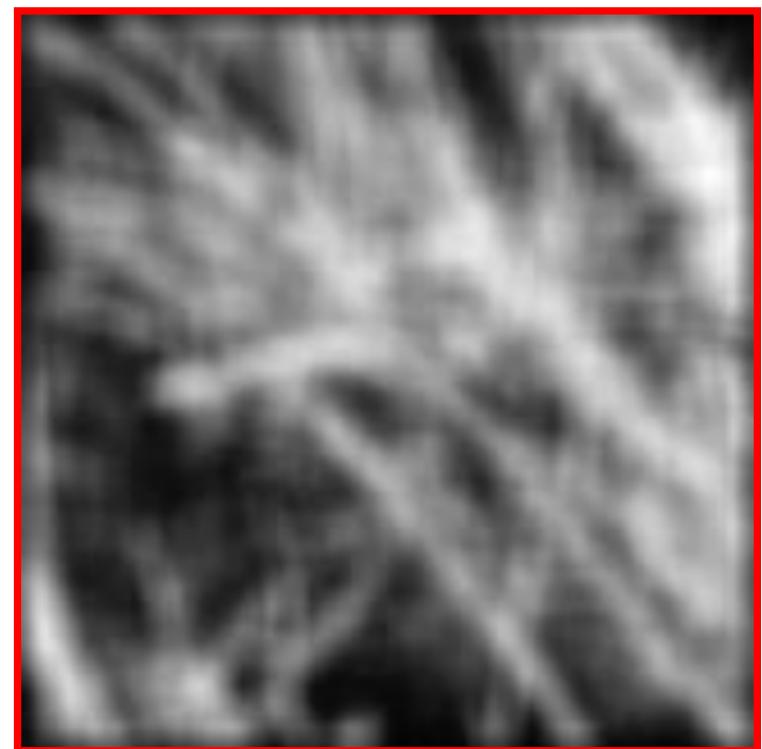
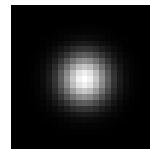
0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

- Common in many natural models
- Smooth and symmetric function → it has an infinite number of derivatives
- Fourier Transform of Gaussian is Gaussian (see later)
- Convolution of a Gaussian with itself is a Gaussian
- Gaussian is separable (e.g. 2D convolution can be performed by two 1-D convolutions)
- There is evidence that the human visual system performs Gaussian filtering

$5 \times 5, \sigma = 1$

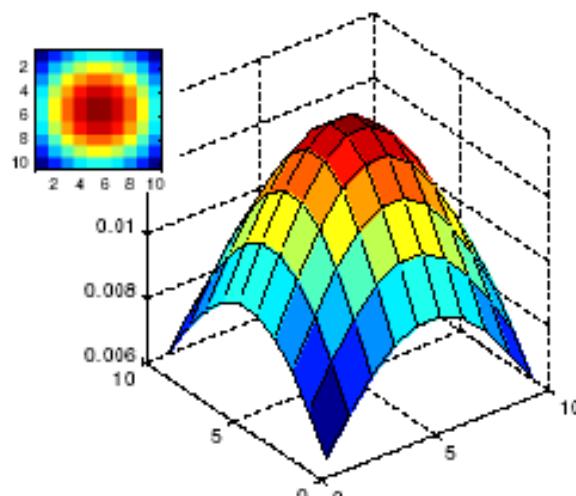
## Smoothing with a Gaussian

- Remove “high-frequency” components from the image (low-pass filter)  
Images become more smooth

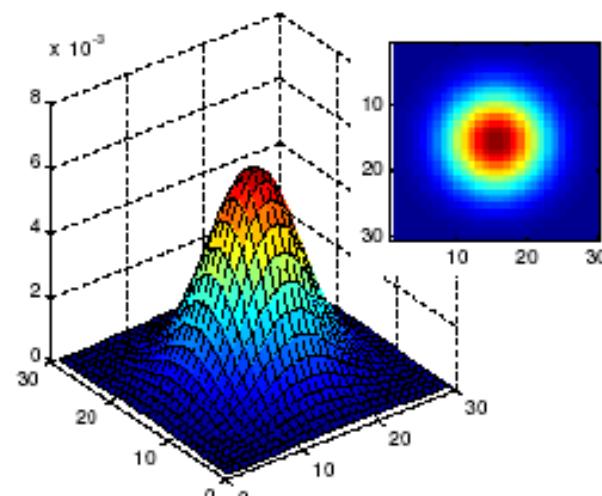


## Gaussian filters

- What parameters matter here?
- **Size of kernel or mask**
  - Note, Gaussian function has infinite support, but discrete filters use finite kernels



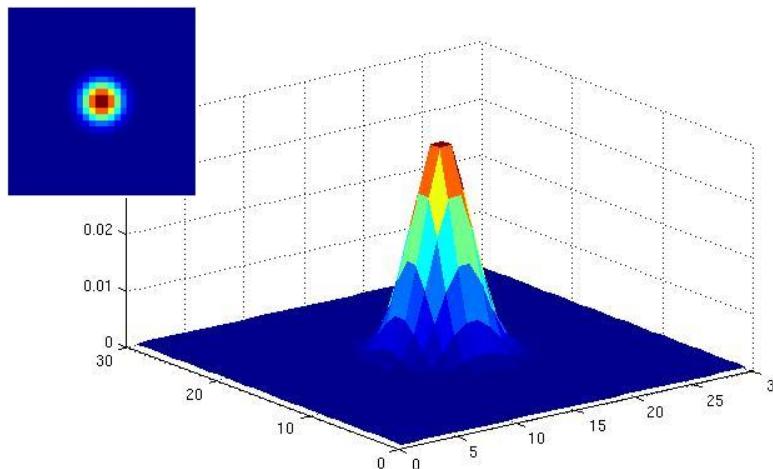
$\sigma = 5$  with 10  
x 10 kernel



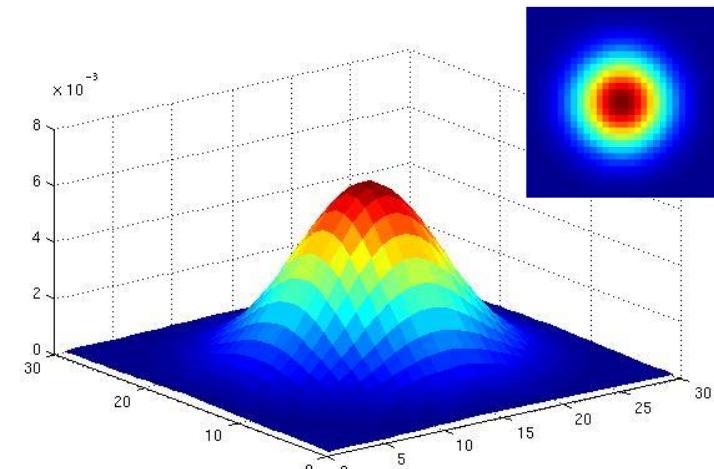
$\sigma = 5$  with 30  
x 30 kernel

## Gaussian filters

- What parameters matter here?
- **Variance of Gaussian:** determines extent of smoothing



$\sigma = 2$  with 30  
x 30 kernel



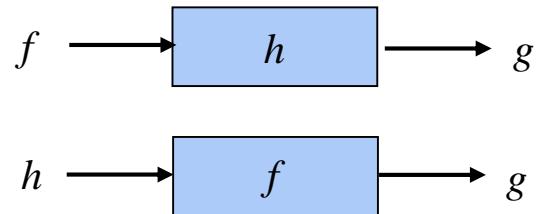
$\sigma = 5$  with 30  
x 30 kernel

- Rule of thumb: filter width of about  $6\sigma$

# Basics: More facts about convolution

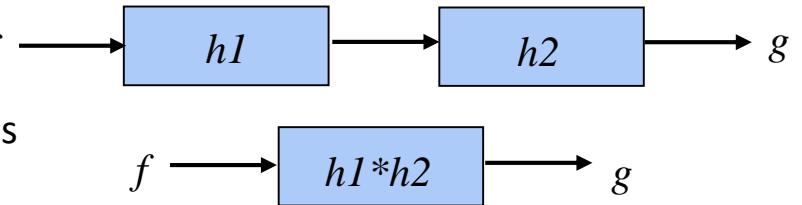
## Convolution is commutative

- That is  $f * h = h * f$
- Interchange of  $h$  and  $f$  possible
- Order does not care



## Convolution is associative

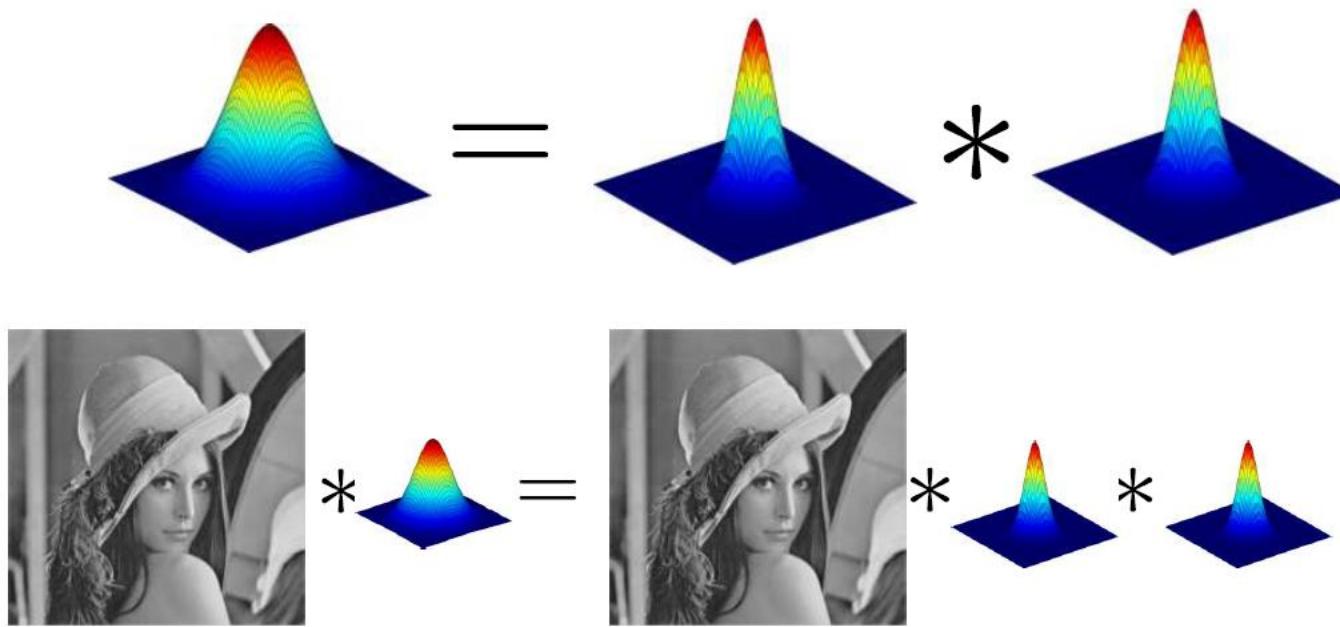
- That is  $(f * h1) * h2 = f * (h1 * h2)$
- Can be exploited for efficient implementations



# Gaussian filters

**Note: Convolution is associative:**  $(f*g)*h = f*(g*h)$

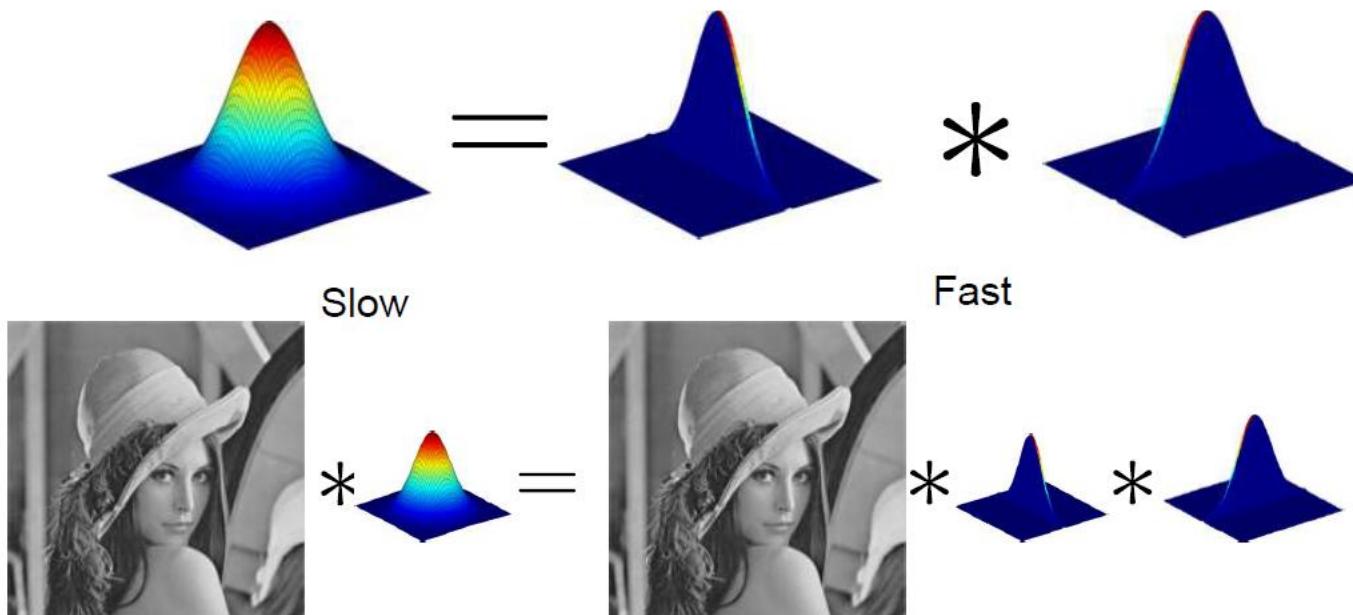
- Can be exploited for multi-scale processing and efficiency:
  - Convolving two times with Gaussian kernel of width  $\sigma$  is same as convolving once with kernel of width  $\sigma\sqrt{2}$
  - Efficiency: multiple smoothing with small-width kernel delivers same result as larger-width kernel



# Gaussian filters: Separability

**Note: Convolution is associative:**  $(f*g)*h = f*(g*h)$

- Can be further exploited for efficiency:
  - Separable kernel factors into product of two 1D Gaussians
  - Efficiency: multiple smoothing with 1D filter delivers same result as with high dimensional filter



## Separability of the Gaussian filter for 2D

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

## Separability: Numeric example for 2D

**Note: Convolution is associative:**  $(f*g)*h = f*(g*h)$

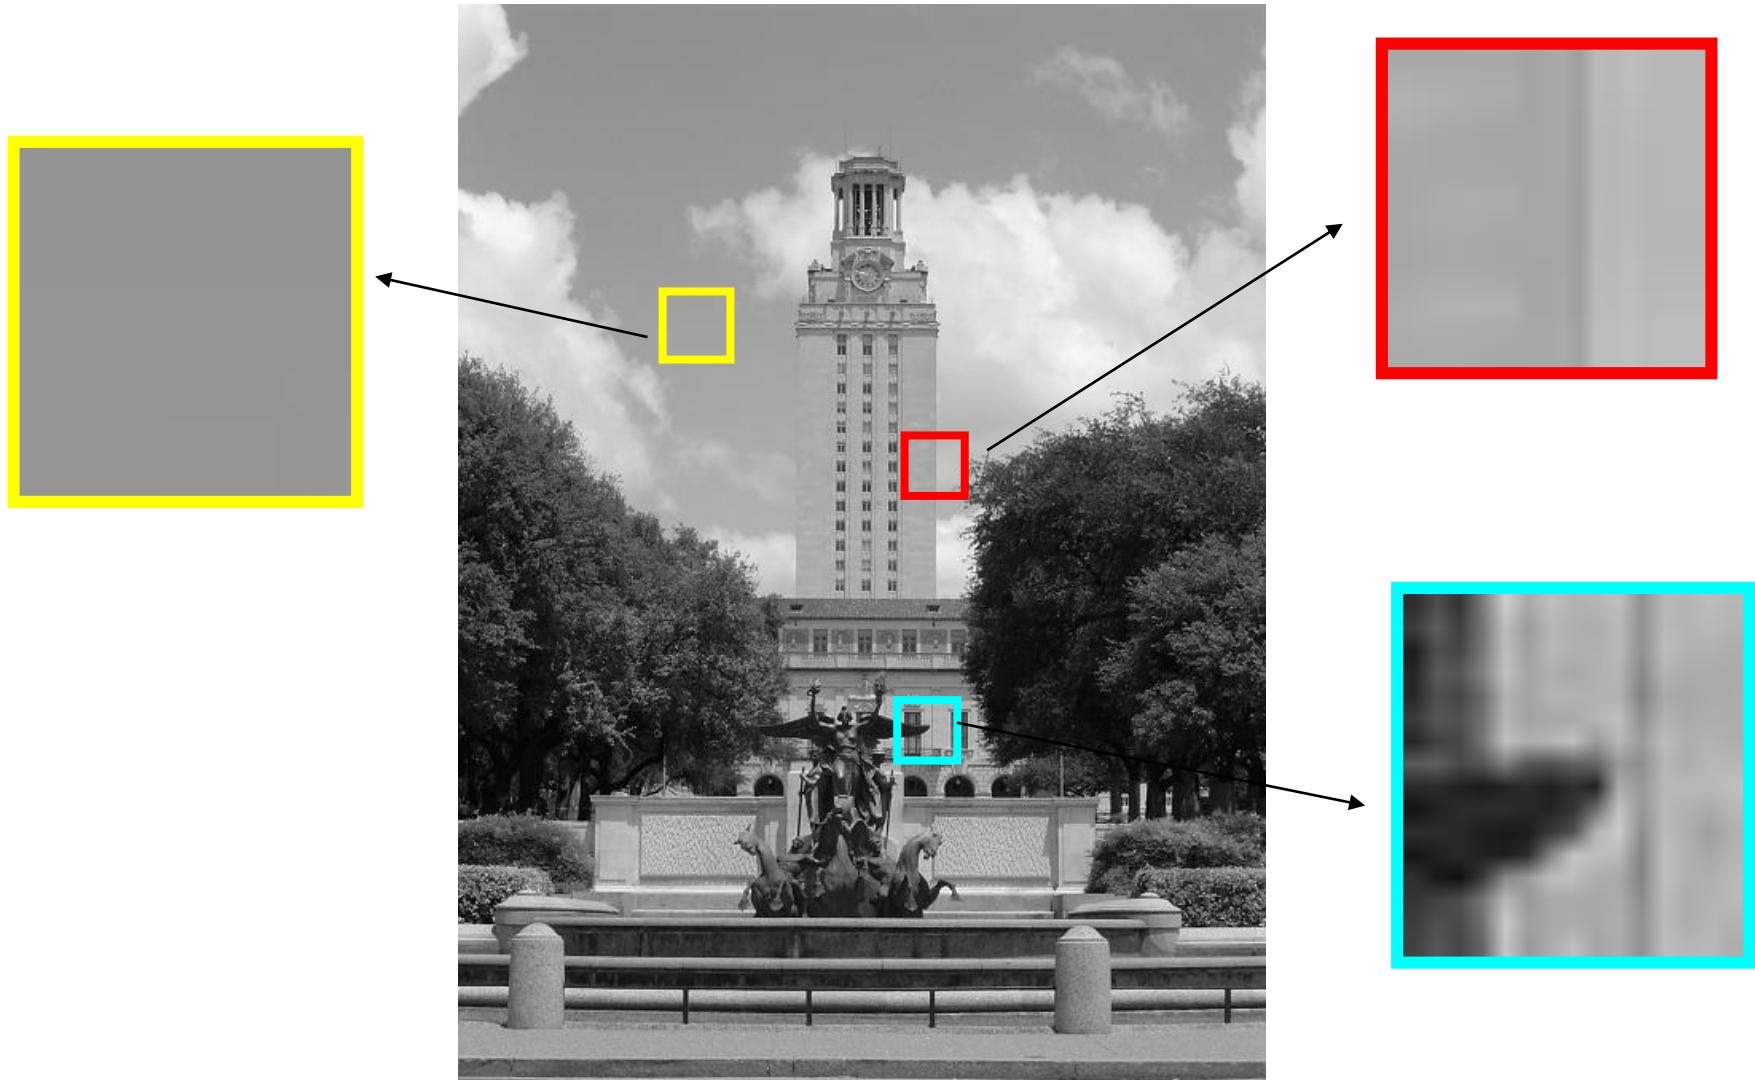
- 2D filters are *separable* if they can be expressed as the outer product of two vectors. For example:

$h$				
$g$	1 2 1	2 3 3 3 5 5 4 4 6		
$f$	1 2 1	11 18 18		
1 2 1	$\times$ 1 2 1 =	1 2 1 2 4 2 1 2 1	2 3 3 3 5 5 4 4 6	$= 2 + 6 + 3 = 11$ $= 6 + 20 + 10 = 36$ $= 4 + 8 + 6 = 18$ <hr style="width: 100px; margin: 10px 0;"/> <span style="float: right;">65</span>

For MN image, PQ filter: 2D takes  $MNPQ$  add/times,  
while 1D takes  $MN(P + Q)$

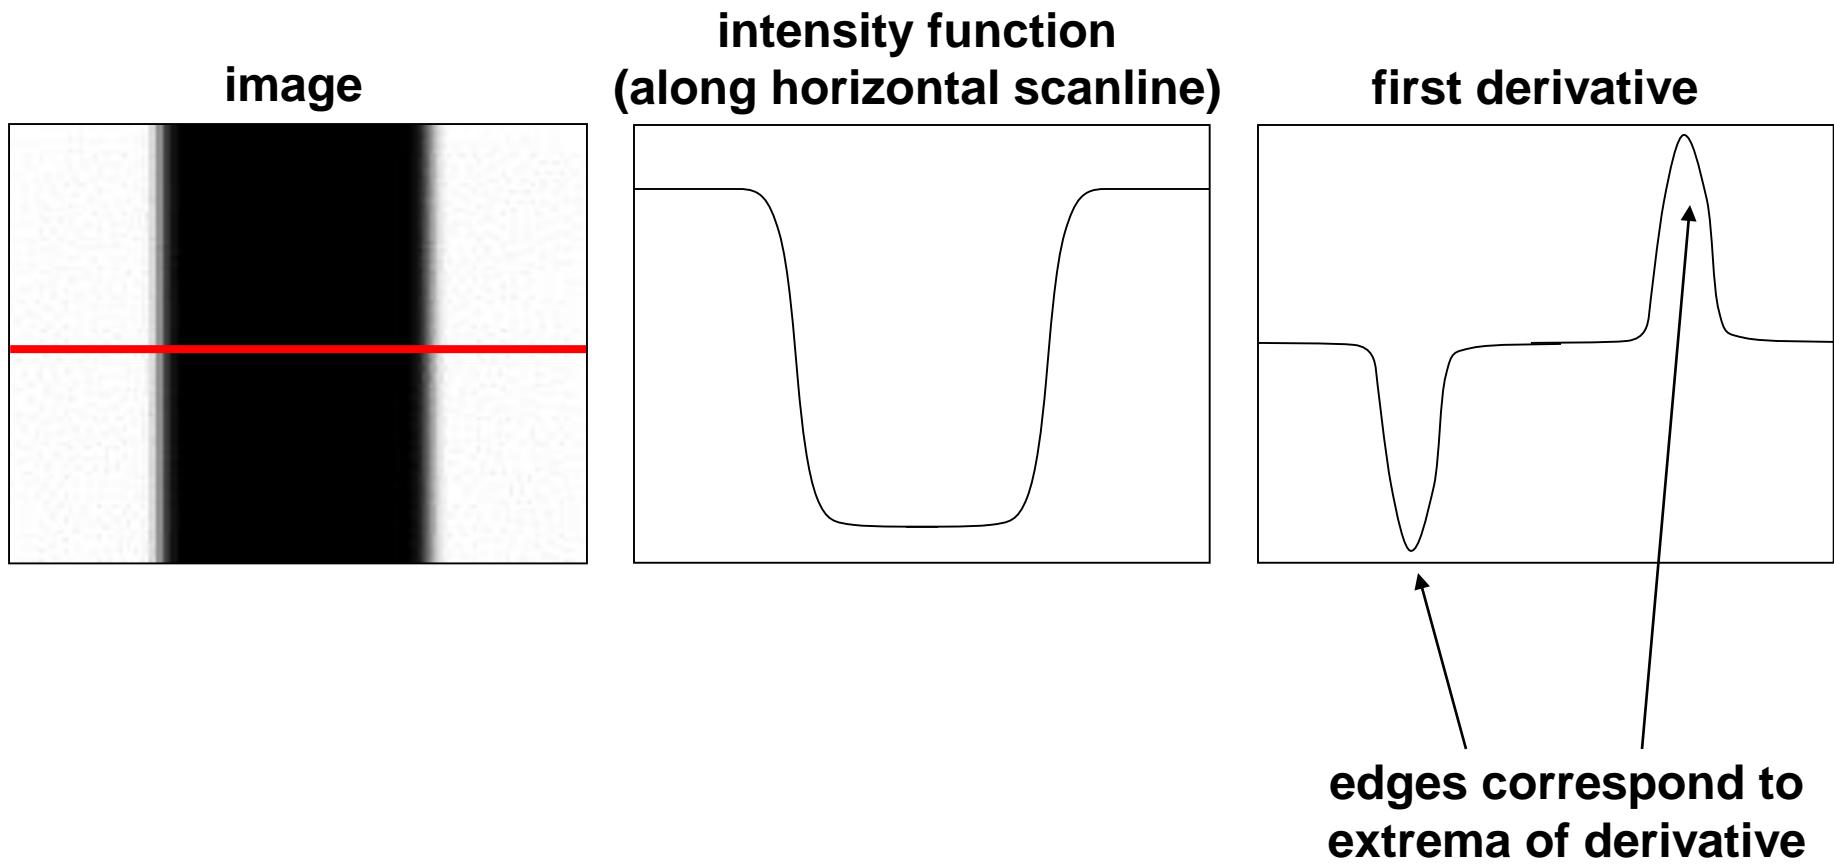
## More important filters: Edges/gradients and invariance

- Motivation: We want to represent distinctive parts of an image



## Derivatives and edges

An edge is a place of rapid change in the image intensity function.

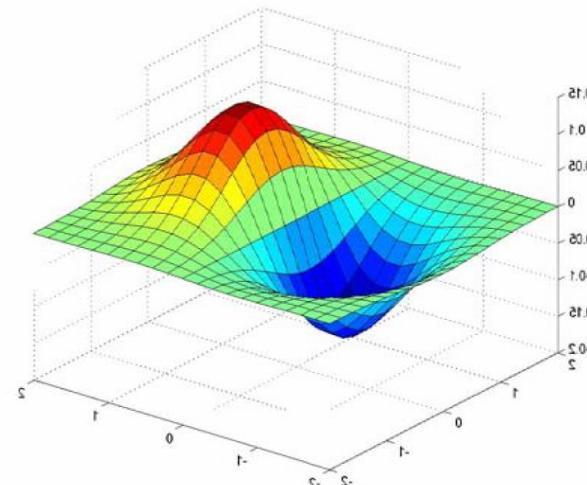


## More important filters: Derivative of Gaussian

$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$

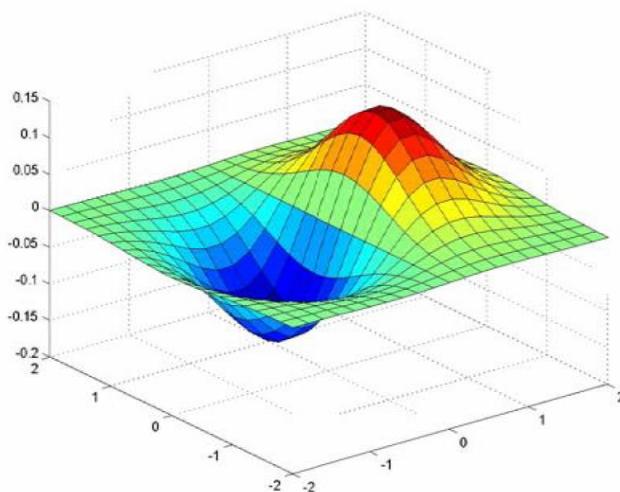
$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$

$$\begin{bmatrix} & & & & \\ & & & & \\ & & & & \\ & & & & \\ & & & & \end{bmatrix} \otimes \begin{bmatrix} 1 & -1 \end{bmatrix}$$

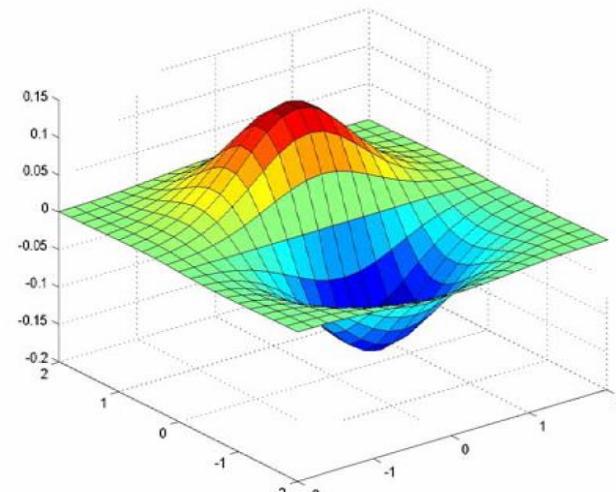


Credit: K. Grauman

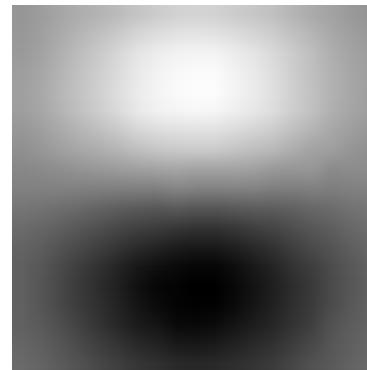
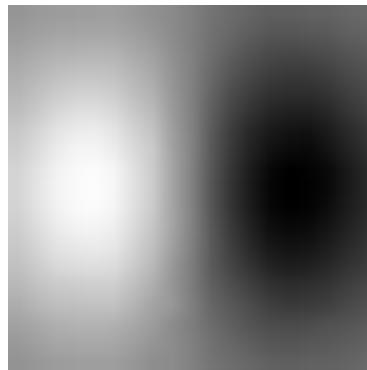
## More important filters: Derivative of Gaussian



x-direction



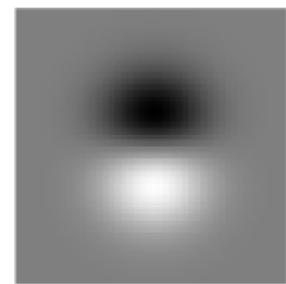
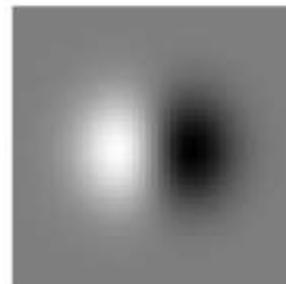
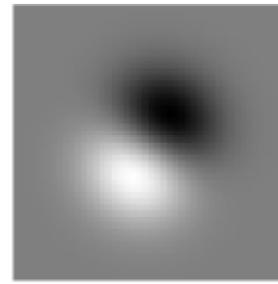
y-direction



## Gaussian filters: Steerability

Distributive property:  $f * (g + h) = f * g + f * h$

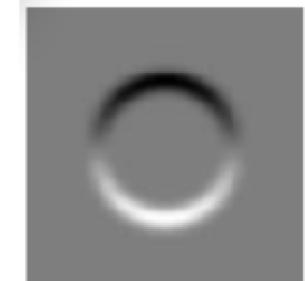
Steerable filter:  $g_\theta(x, y) = \cos(\theta)g_0(x, y) + \sin(\theta)g_{\pi/2}(x, y)$



$I$

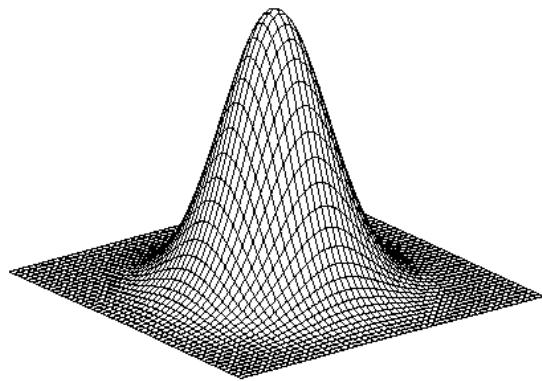


$$I * g_\theta = \cos(\theta)(I * g_0) + \sin(\theta)(I * g_{\pi/2})$$



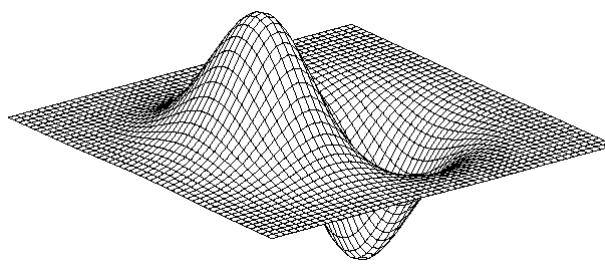
W. Freeman and E. Adelson, 'The design and use of steerable filters', PAMI, 1991

## More important filters: Derivative(s) of Gaussian



Gaussian

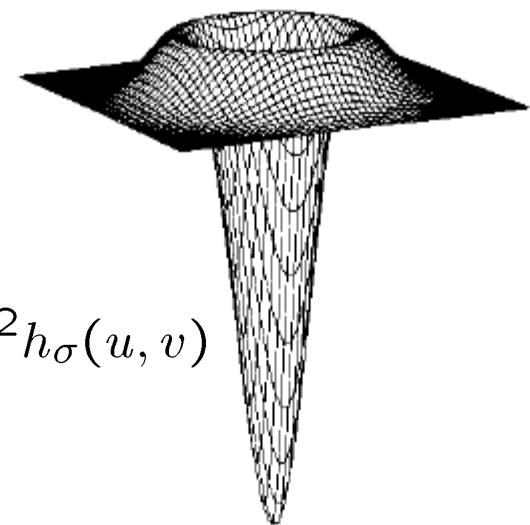
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



$$\nabla^2 h_\sigma(u, v)$$

- $\nabla^2$  is the Laplacian operator:

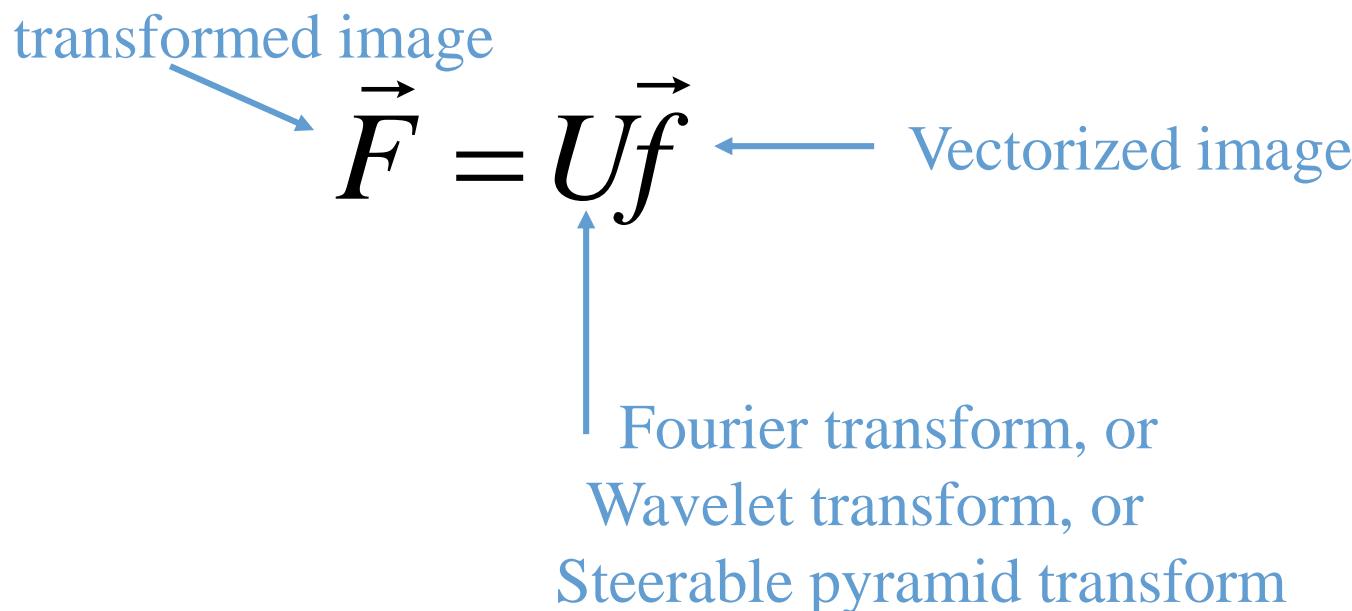
$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

## The Fourier transform

Further reading: Szeliski, Richard. Computer Vision: Algorithms and Applications. Springer, 2010, Chapter 3, Section 3.4

## Linear image transformations

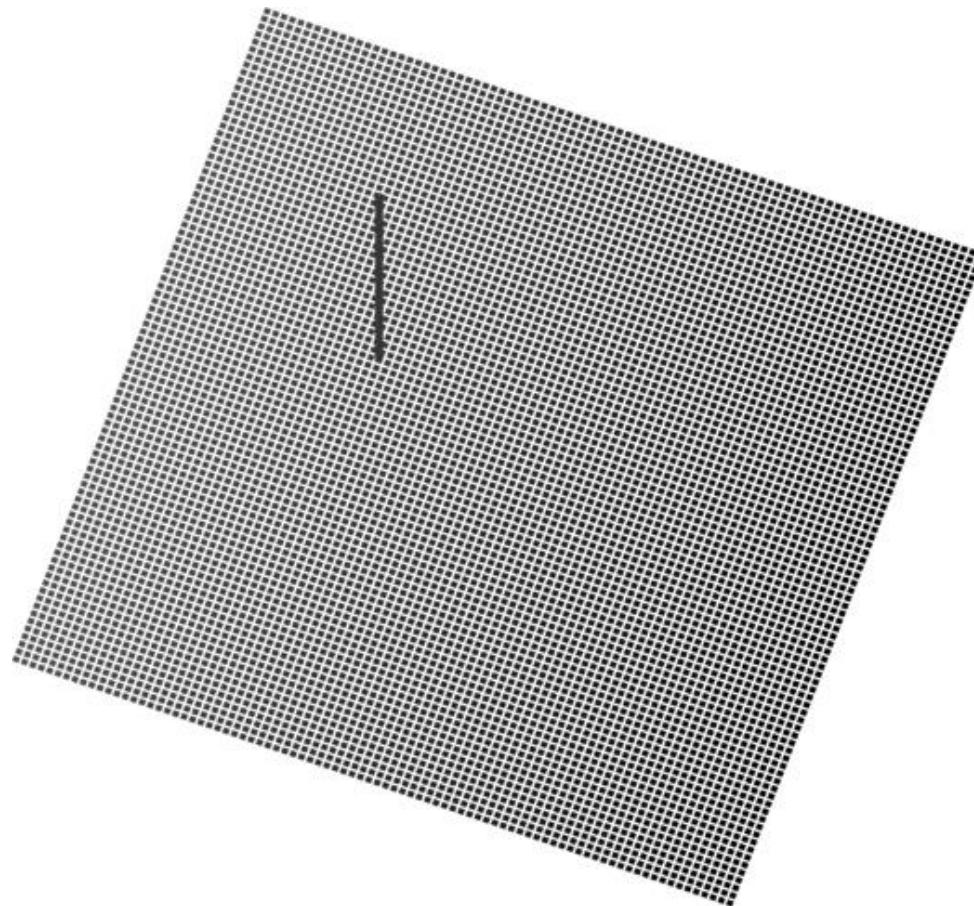
- In analyzing images, it's often useful to make a change of basis.



Kronecker delta

Canonical basis for 2D signals

$$d_k[n] = \begin{cases} 1, & n = k \\ 0, & \text{otherwise} \end{cases}$$

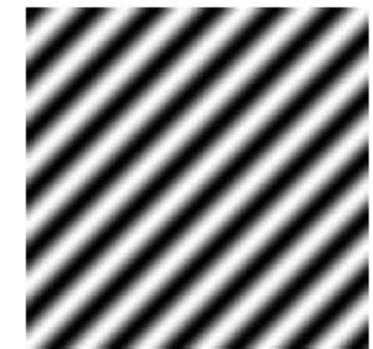


## Fourier Transform = Change of Basis



$$F(\omega_1) \cdot e^{j\omega_1 x}$$

+

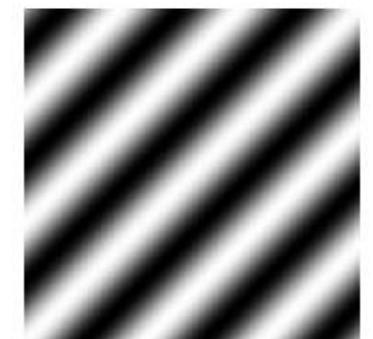


$$= F(\omega_2) \cdot e^{j\omega_2 x}$$

+



$$F(\omega_K) \cdot e^{j\omega_K x}$$



Credit: I. Kokkinos

## Self-inverting transforms

Same basis functions are used for the inverse transform

$$\vec{f} = U^{-1} \vec{F}$$

$$= U^+ \vec{F}$$



U transpose and complex conjugate

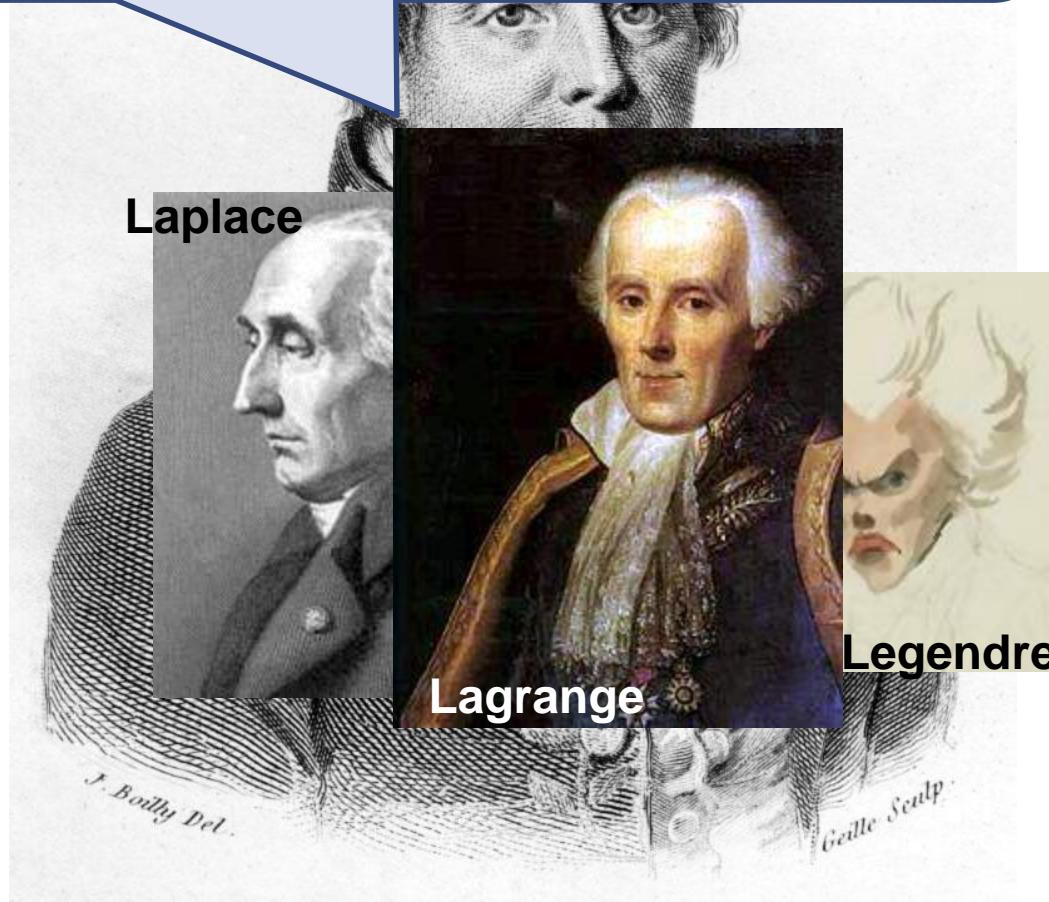
# Jean Baptiste Joseph Fourier (1768-1830)

had crazy idea (1807):

*Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.*

*...the manner in which the author arrives at these equations is not exempt of difficulties and...his analysis to integrate them still leaves something to be desired on the score of generality and even rigour.*

- Don't believe it?
  - Neither did Lagrange, Laplace, Poisson and others
  - Not translated into English until 1878!
- But it's true!
  - called Fourier Series



## Fourier Transform

Our building block:

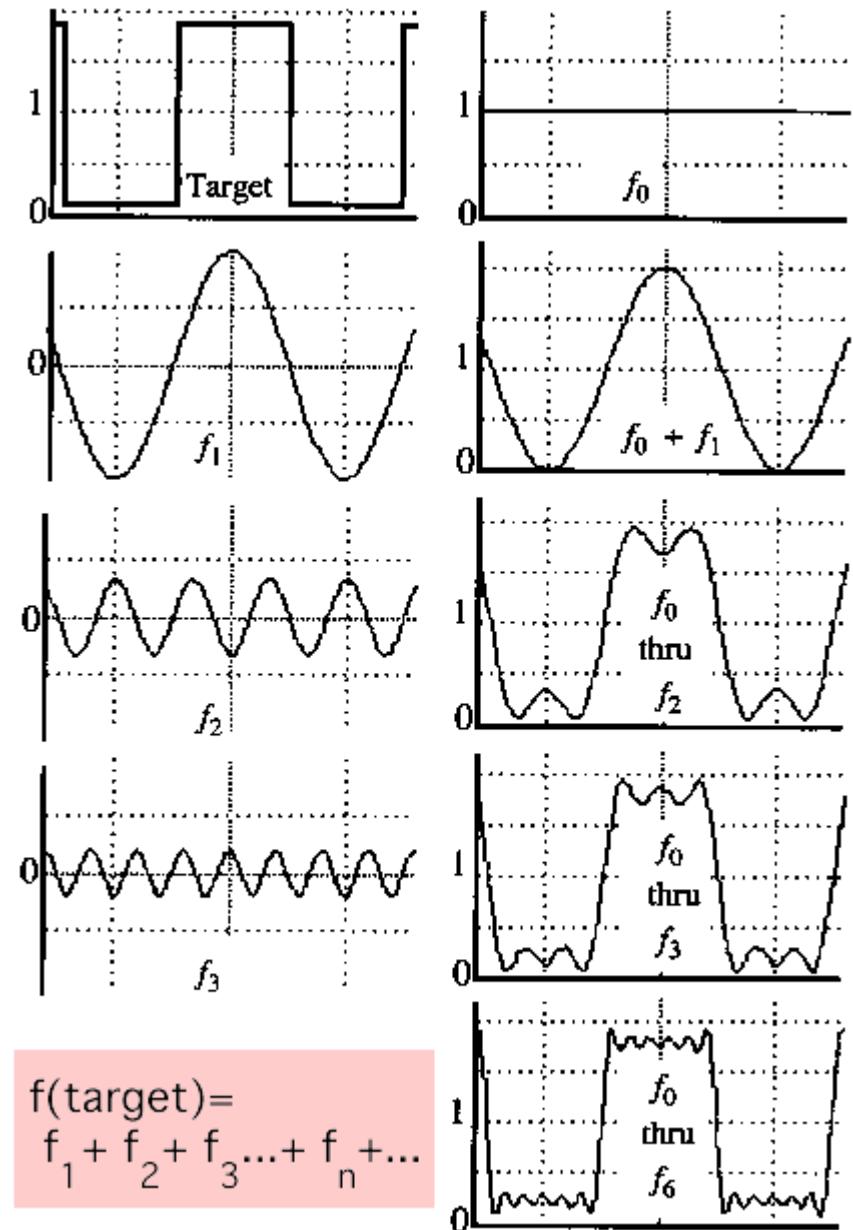
$$A \sin(\omega x + \phi)$$

Add enough of them to get any signal  $g(x)$  you want!

The Fourier transform  $F(\omega)$  stores the magnitude and phase at each frequency

Magnitude  $A$  encodes how much signal there is at a particular frequency  $\omega$

Phase  $\phi$  encodes spatial information (indirectly)



Credit: J. Hays

## Fourier Transform

- We want to understand the frequency  $\omega$  of our signal. So, let's reparametrize the signal by  $\omega$  instead of  $x$ :



For every  $\omega$  from 0 to inf,  $F(\omega)$  holds the amplitude  $A$  and phase  $\phi$  of the corresponding sine  $A \sin(\omega x + \phi)$

- How can  $F$  hold both? Complex number trick!

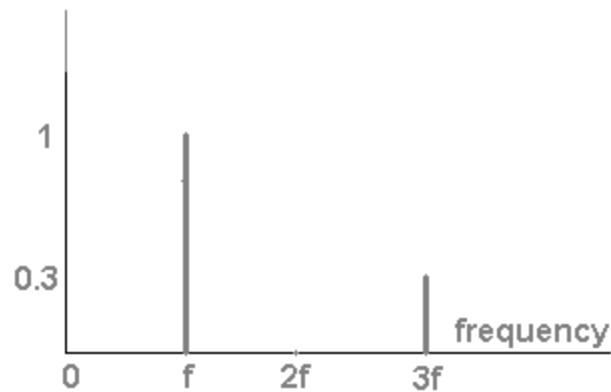
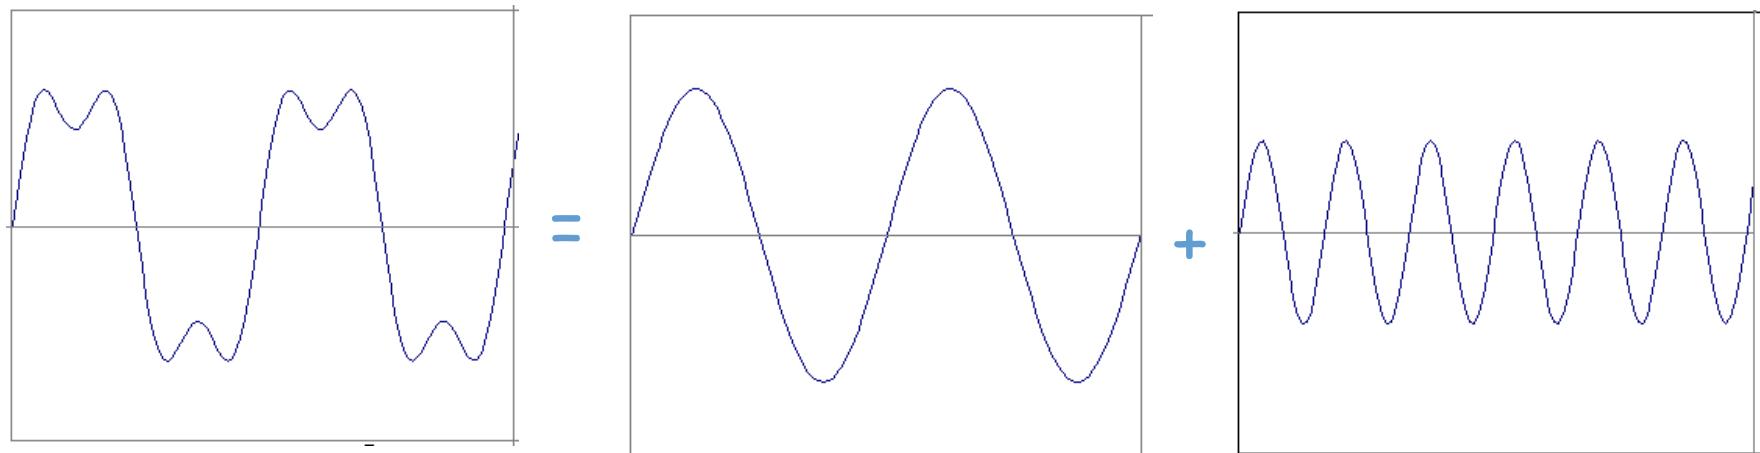
$$F(\omega) = R(\omega) + iI(\omega)$$
$$A = \pm \sqrt{R(\omega)^2 + I(\omega)^2} \quad \phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

We can always go back:

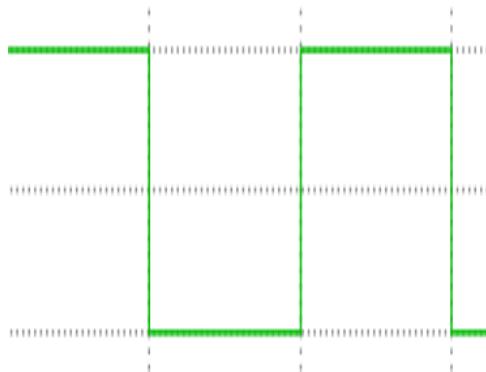


# Frequency Spectra

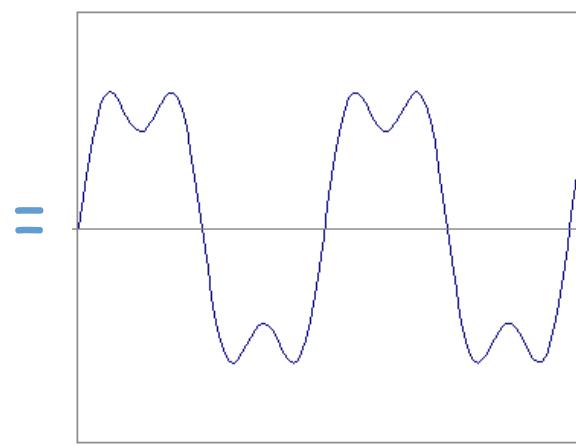
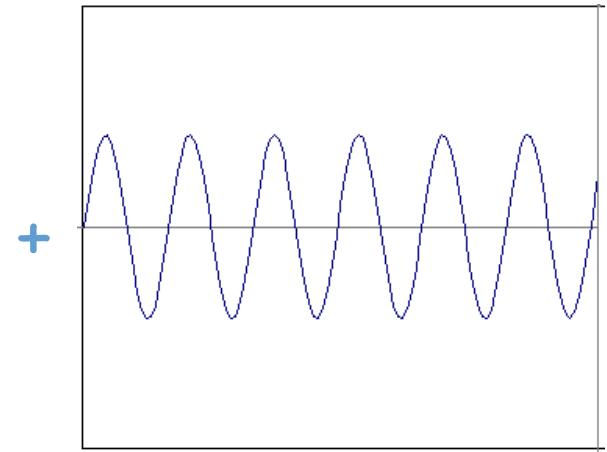
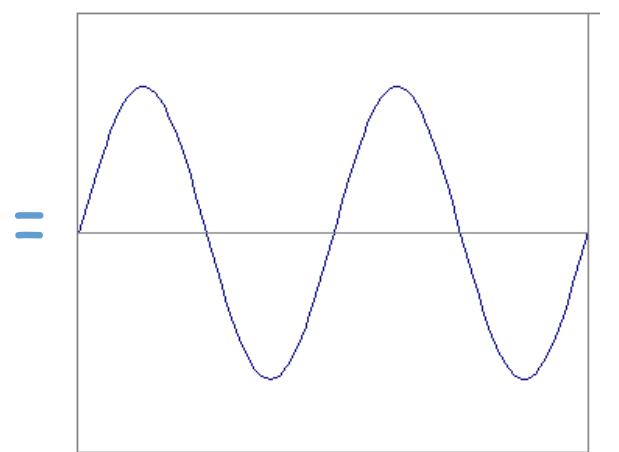
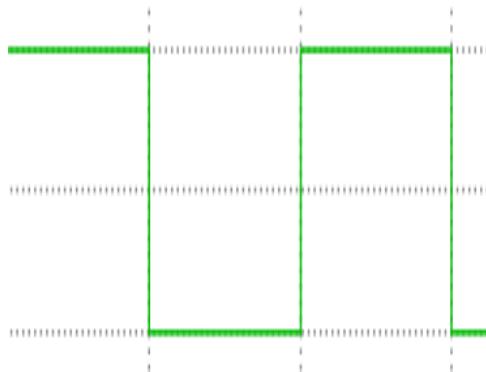
- example :  $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f) t)$



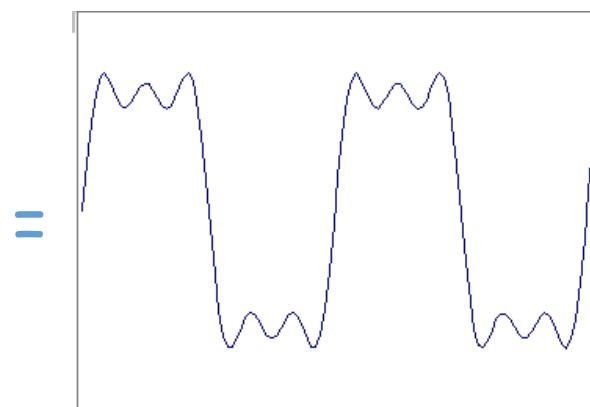
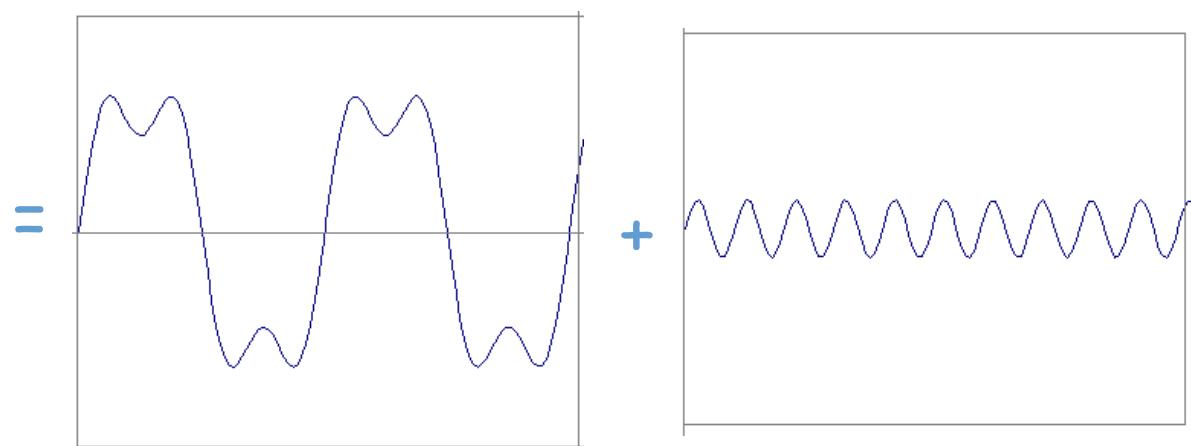
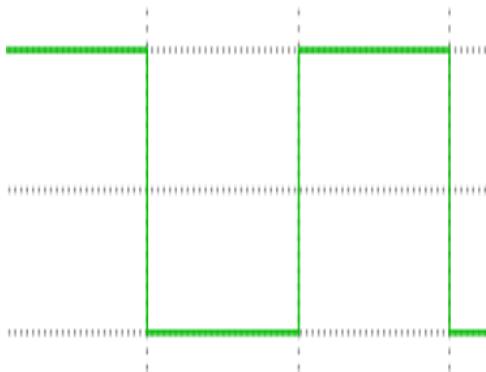
# Frequency Spectra



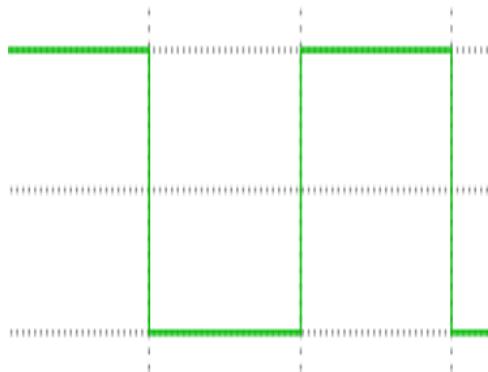
# Frequency Spectra



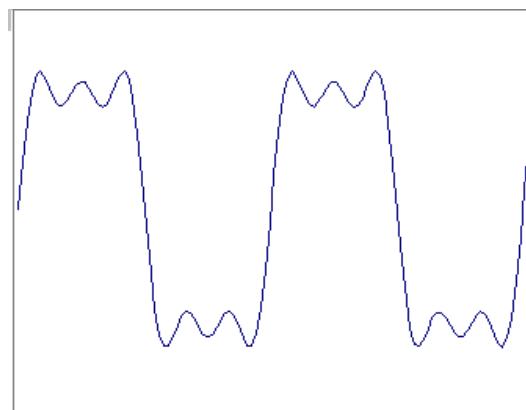
# Frequency Spectra



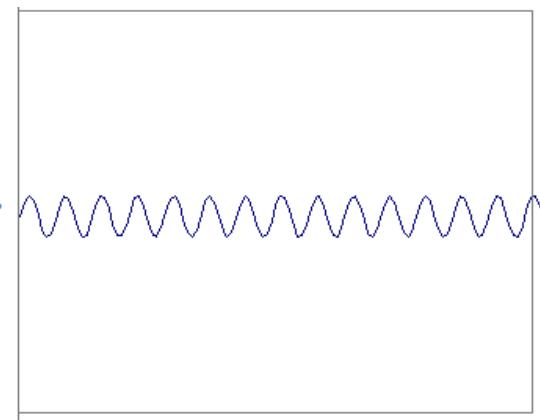
# Frequency Spectra



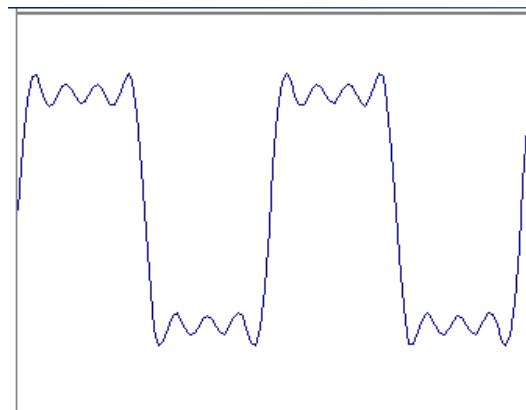
=



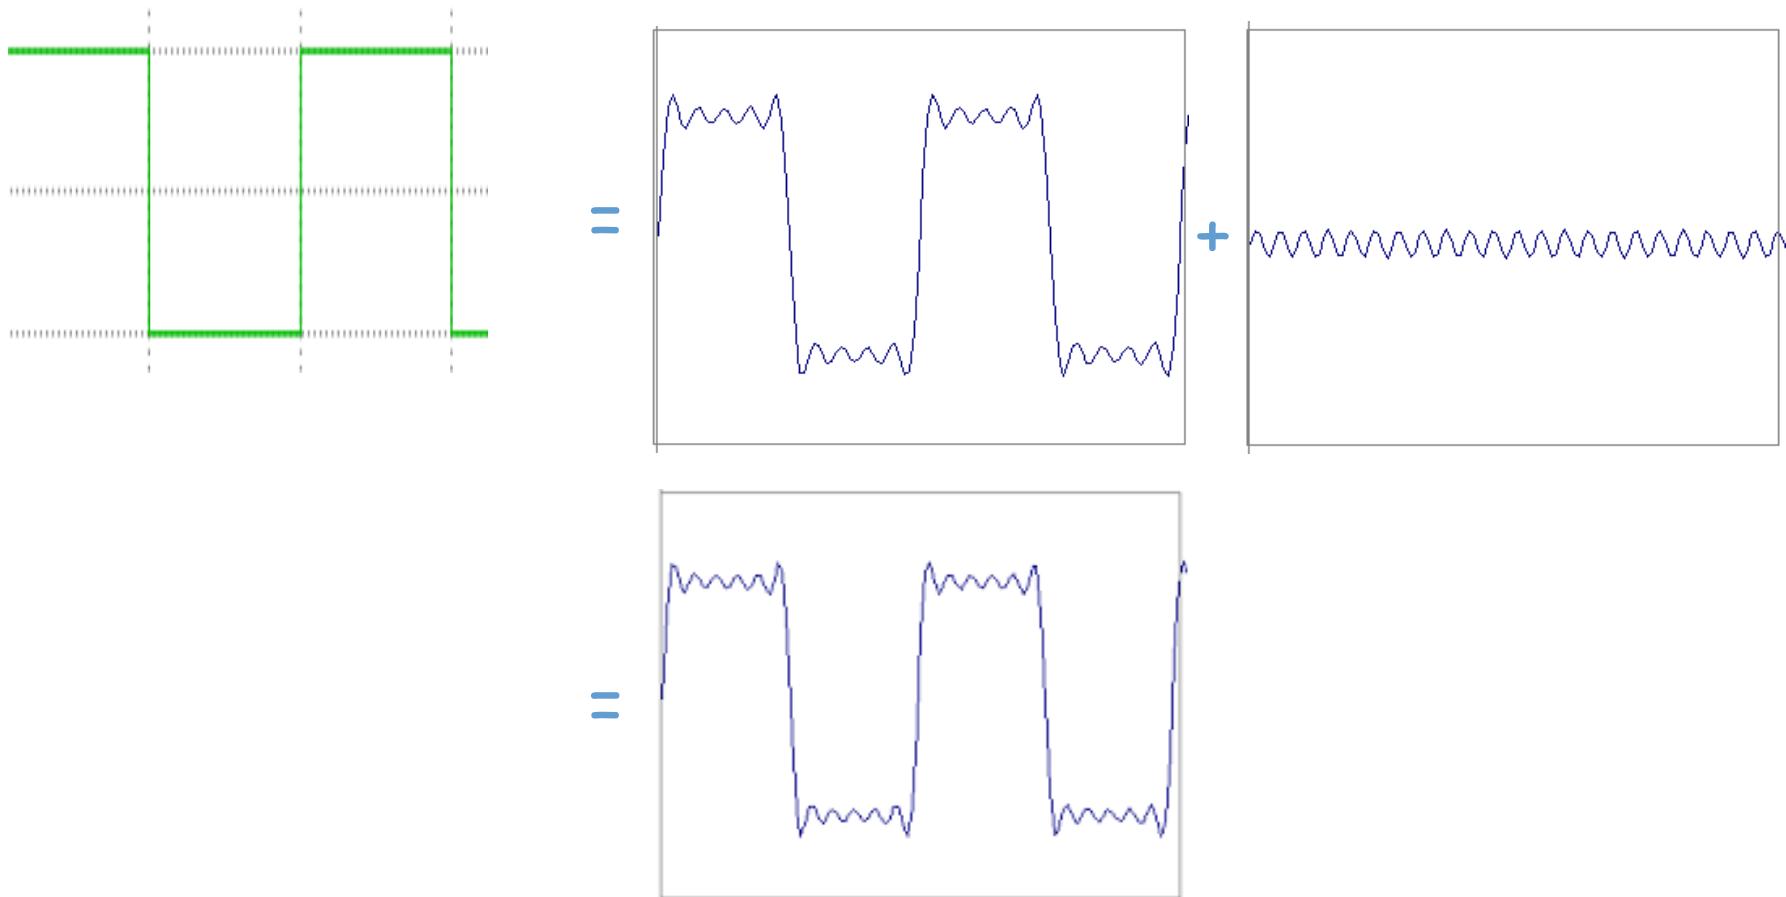
+



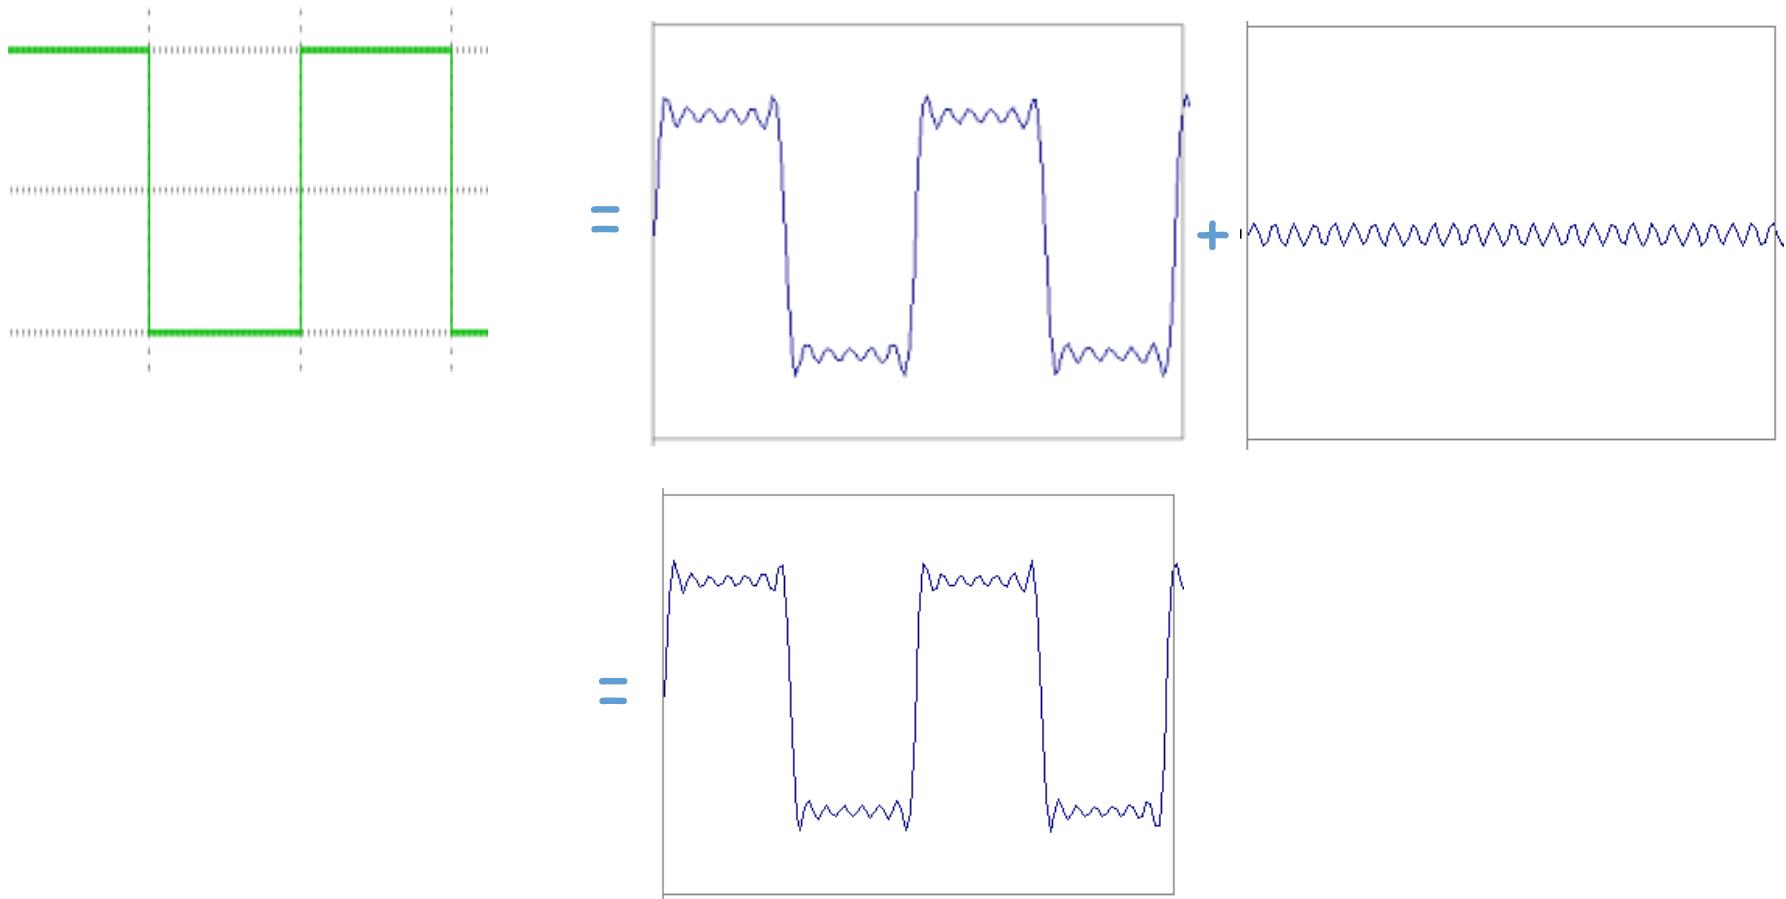
=



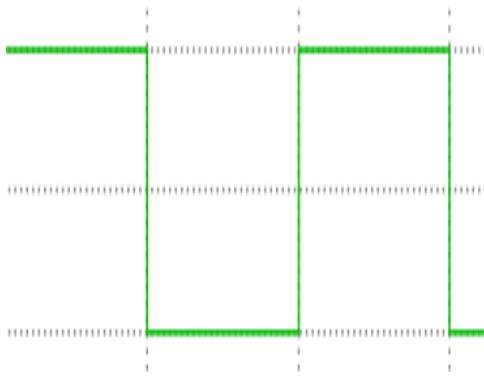
# Frequency Spectra



# Frequency Spectra

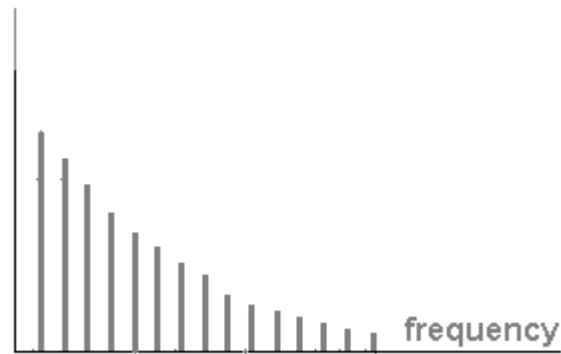


# Frequency Spectra



=

$$A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$



# Basics: The Fourier transform

## Eigenfunctions

- An **eigenfunction** of a system is one that is simply multiplied by another factor in the output.



- We think of this as analogous to the case of eigenvectors from linear algebra.

# Basics: The Fourier transform

## Eigenfunctions

- An **eigenfunction** of a system is one that is simply multiplied by another factor in the output.



- We think of this as analogous to the case of eigenvectors from linear algebra.

## Remark

- Notation

$$e^{iwt} = \exp(iwt)$$

with the imaginary number

$$i = \sqrt{-1}$$

# Basics: The Fourier transform

## Eigenfunctions

- An **eigenfunction** of a system is one that is simply multiplied by another factor in the output.



- We think of this as analogous to the case of eigenvectors from linear algebra.
- For the case of 1D Linear Shift Invariant (LSI) systems we find that  $\exp(iwt)$  is an eigenfunction of convolution.



- Here  $A(w)$  is the (possibly complex) factor by which the input signal is multiplied.
- So, from the input exponential we obtain another exponential; but, scaled and shifted in phase.

# Basics: The Fourier transform

## Eigenfunctions

- An **eigenfunction** of a system is one that is simply multiplied by another factor in the output.



- We think of this as analogous to the case of eigenvectors from linear algebra.
  - For the case of 1D LSI systems we find that  $\exp(iwt)$  is an eigenfunction of convolution.
- 
- A block diagram showing a signal flow. An input signal  $\exp(iwt)$  enters a blue rectangular block from the left. An output signal  $A(w) \exp(iwt)$  exits the block to the right. The block represents a system that multiplies the input by a factor  $A(w)$ .
- Here  $A(w)$  is the (possibly complex) factor by which the input signal is multiplied.
  - So, from the input exponential we obtain another exponential; but, scaled and shifted in phase.

## Frequency

- We call  $w$  the **frequency** (or **wave number**) of the eigenfunction.
- In practice, we use real waveforms, like  $\cos(wt)$  and  $\sin(wt)$ , with the relationship

$$\exp(iwt) = \cos(wt) + i \sin(wt)$$

which is known as **Euler's relation**.

- The complex exponential is used in derivations simply because it provides a compact notation.

# Basics: The Fourier transform

## 1D frequency

- We consider functions of the form

$$f(x) = A \cos(ux + d)$$

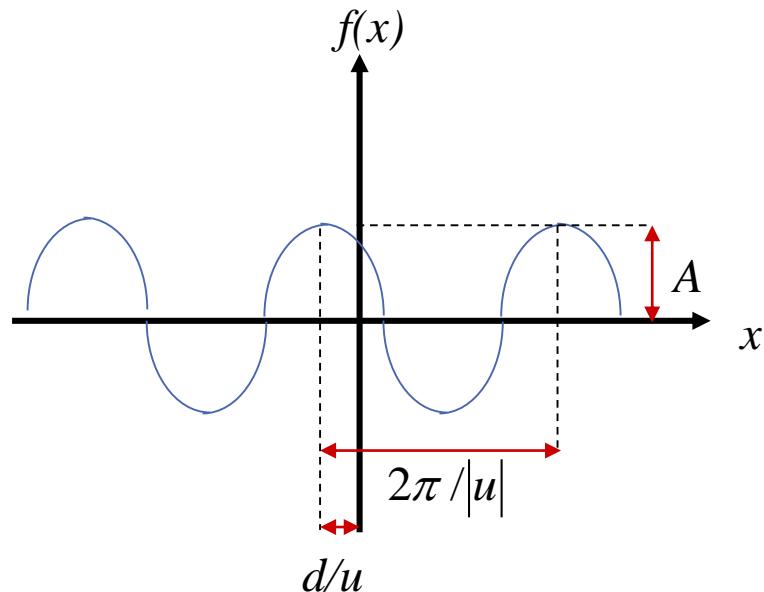
where

$A$  is the amplitude

$u$  is the (angular) frequency

$d$  is the phase constant.

- Notice that the function repeats its value when  $ux + d$  increases by  $2\pi$
- For example, when  $d = 0$ , the maxima and minima occur when  $ux = k\pi$ , for  $k$  an integer.



# Fourier transform: 1D case

## Consider

- A graphical interpretation

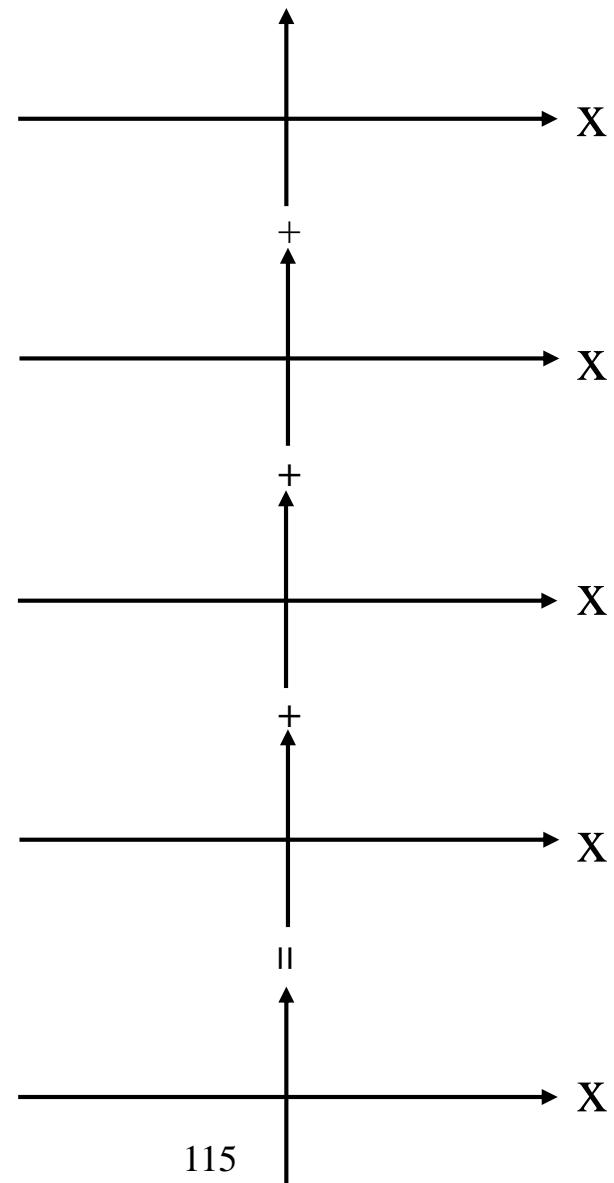
$$1$$

$$+ \frac{1}{2} \cos 2\pi x$$

$$+ \cos 4\pi x$$

$$+ \frac{2}{3} \cos 6\pi x$$

$$= f(x)$$



# Fourier transform: 1D case

Consider

- A graphical interpretation

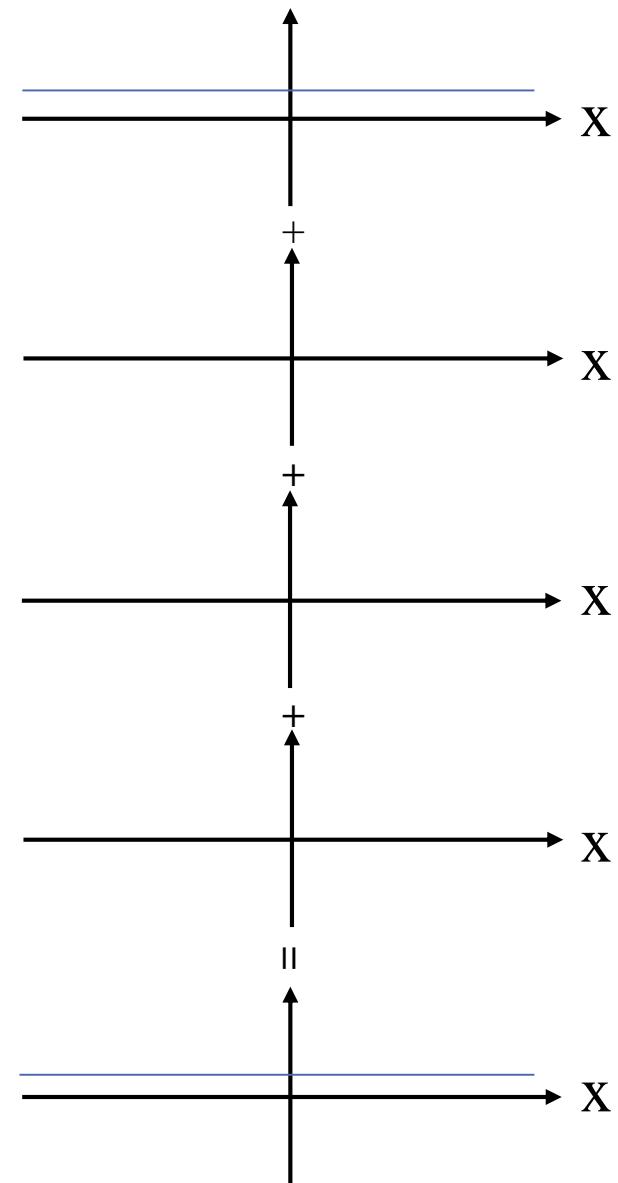
1

$$+ \frac{1}{2} \cos 2\pi x$$

$$+ \cos 4\pi x$$

$$+ \frac{2}{3} \cos 6\pi x$$

$$= f(x)$$



# Fourier transform: 1D case

Consider

- A graphical interpretation

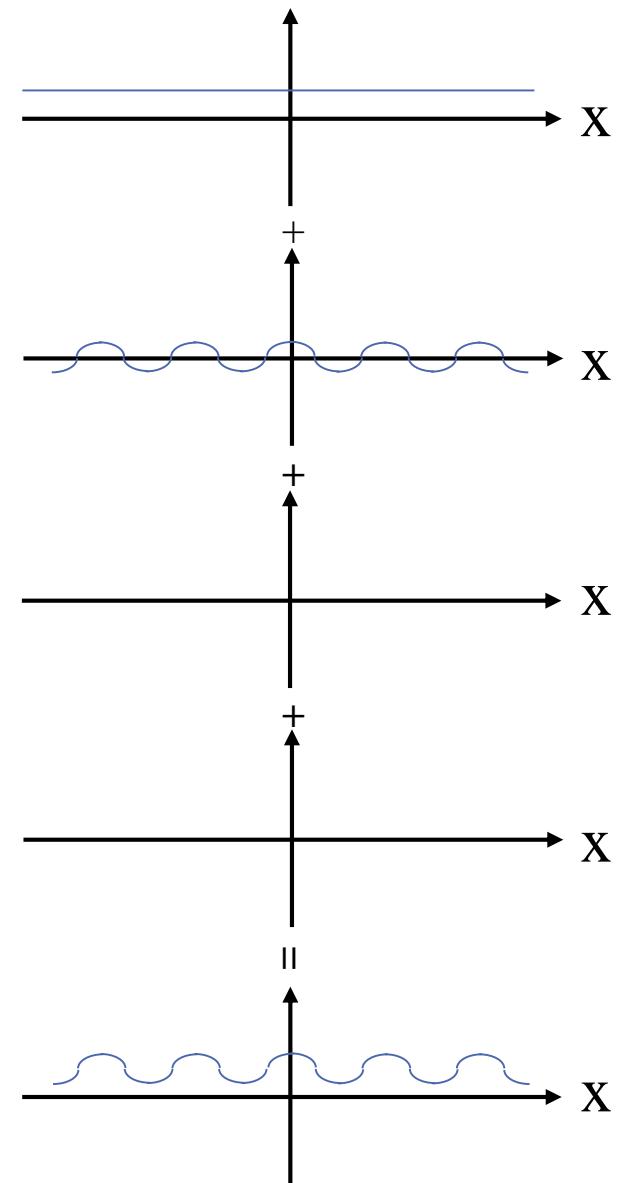
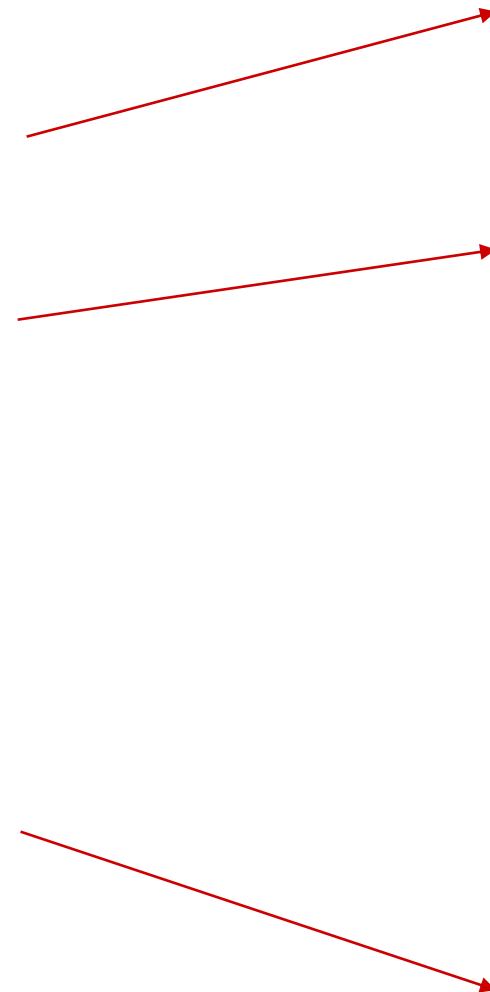
$$1$$

$$+ \frac{1}{2} \cos 2\pi x$$

$$+ \cos 4\pi x$$

$$+ \frac{2}{3} \cos 6\pi x$$

$$= f(x)$$



# Fourier transform: 1D case

Consider

- A graphical interpretation

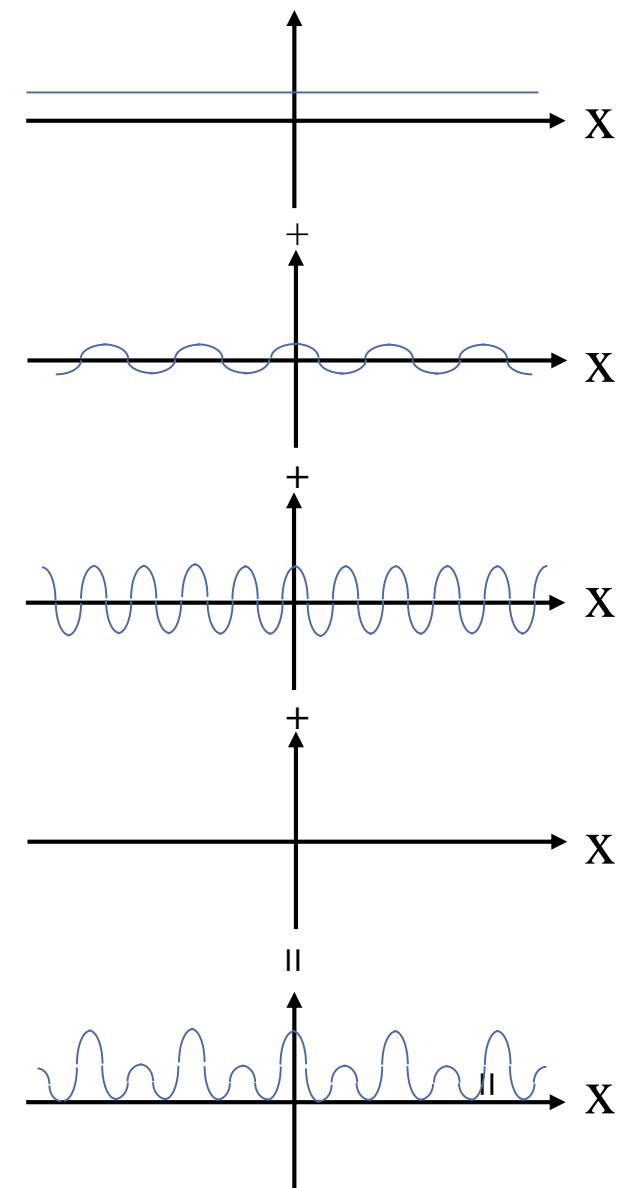
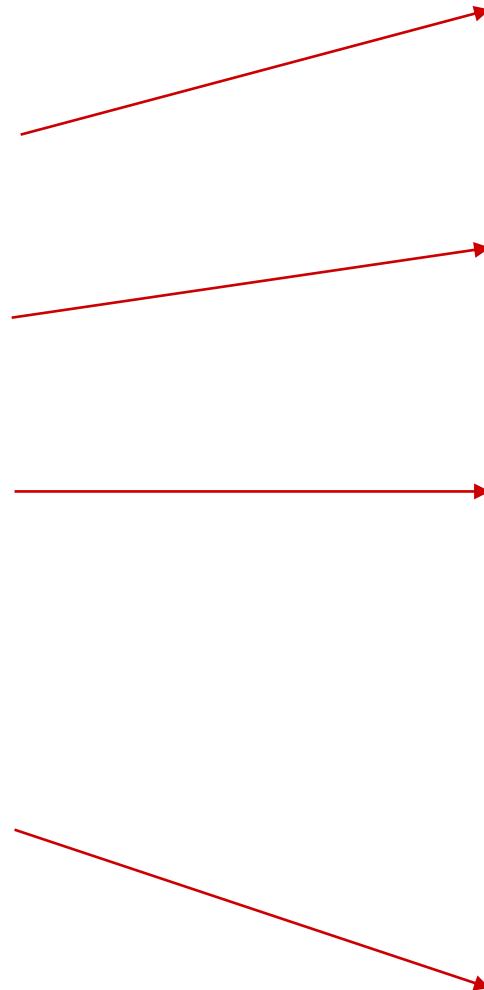
$$1$$

$$+ \frac{1}{2} \cos 2\pi x$$

$$+ \cos 4\pi x$$

$$+ \frac{2}{3} \cos 6\pi x$$

$$= f(x)$$



# Fourier transform: 1D case

## Consider

- A graphical interpretation

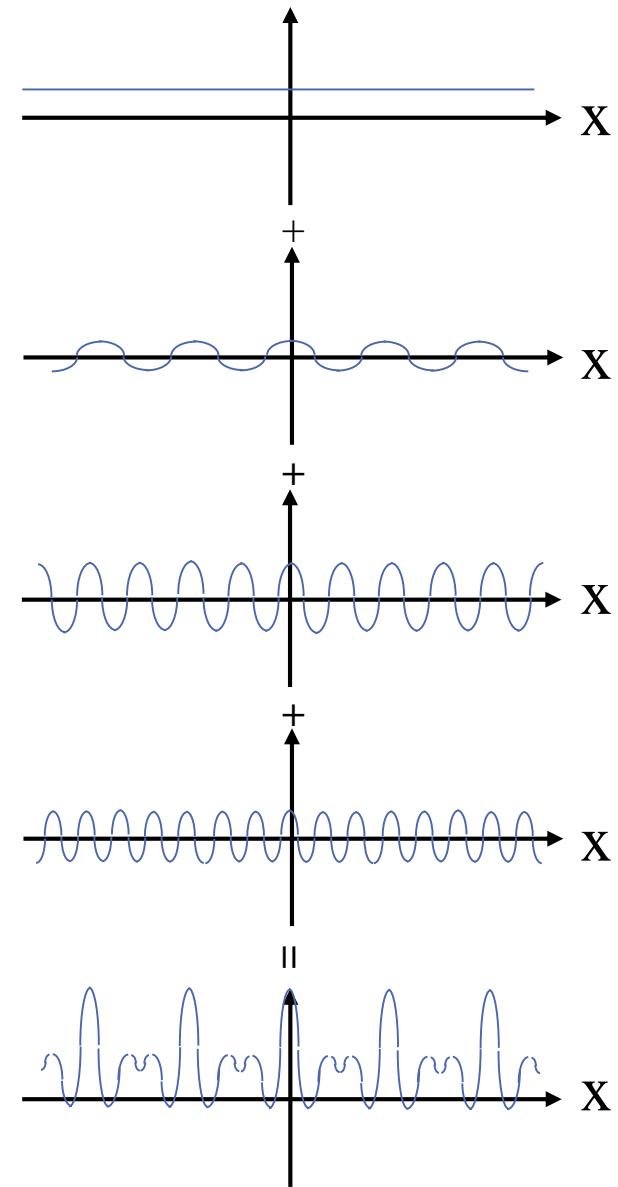
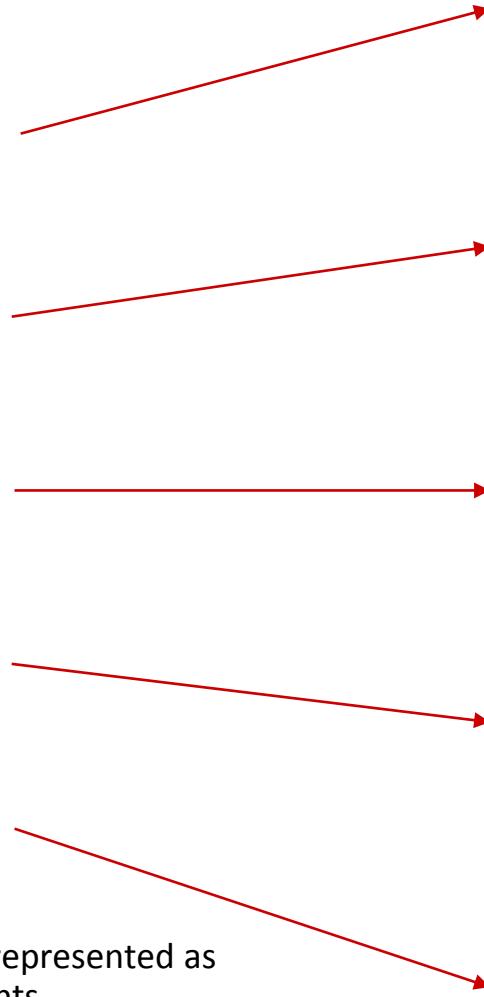
$$1$$

$$+ \frac{1}{2} \cos 2\pi x$$

$$+ \cos 4\pi x$$

$$+ \frac{2}{3} \cos 6\pi x$$

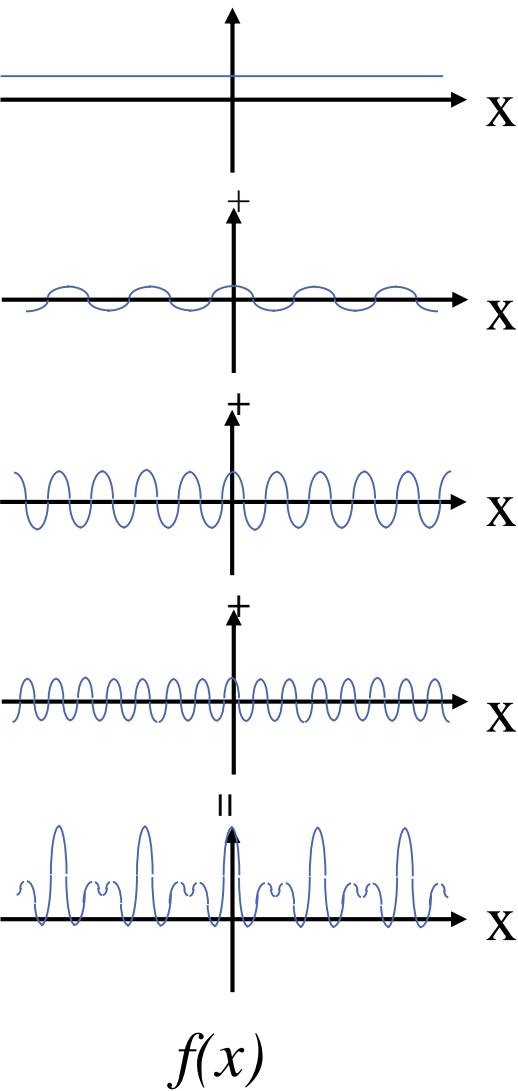
$$= f(x)$$



## Observation

- Complicated signals can be represented as the sum of simple components.

## Fourier transform: 1D case

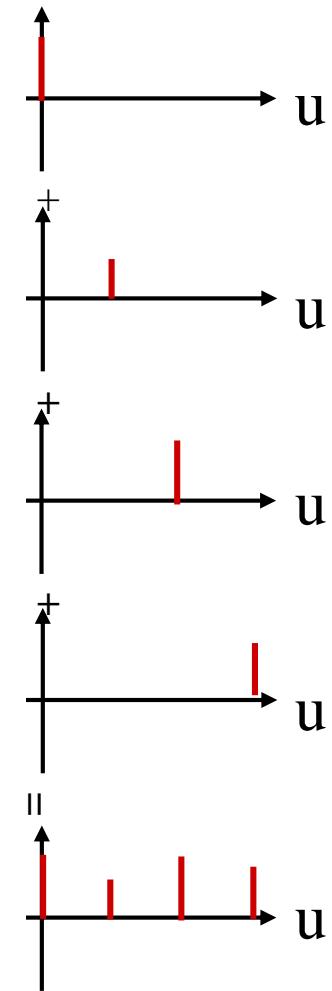


$$1 \cos 0\pi x \rightarrow 1 \cos ux; u = 0$$

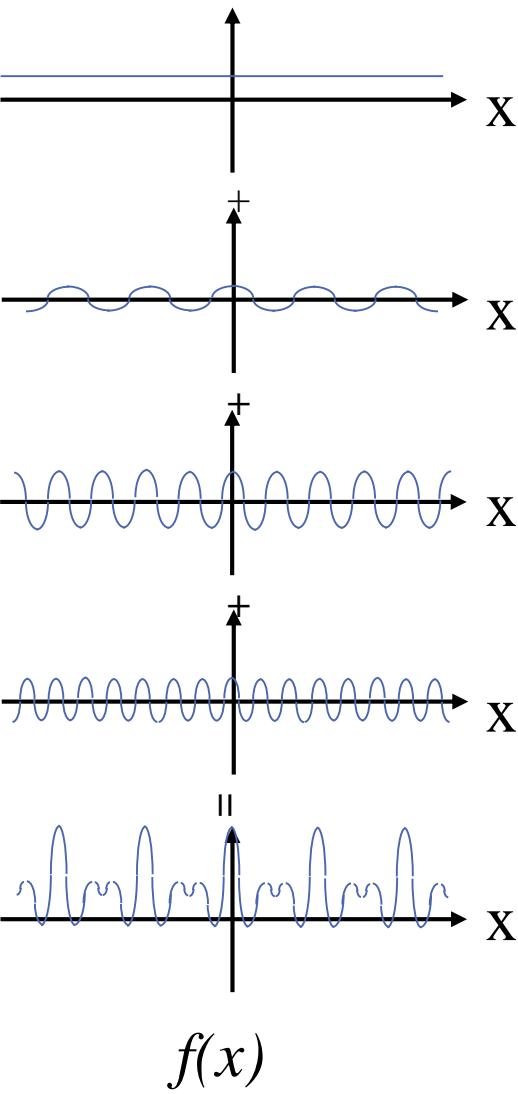
$$+ \frac{1}{2} \cos 2\pi x \rightarrow + \frac{1}{2} \cos ux; u = 2\pi$$

$$+ \cos 4\pi x \rightarrow + 1 \cos ux; u = 4\pi$$

$$+ \frac{2}{3} \cos 6\pi x \rightarrow + \frac{2}{3} \cos ux; u = 6\pi$$



## Fourier transform: 1D case



$$1 \cos 0\pi x \rightarrow 1 \cos ux; u = 0$$

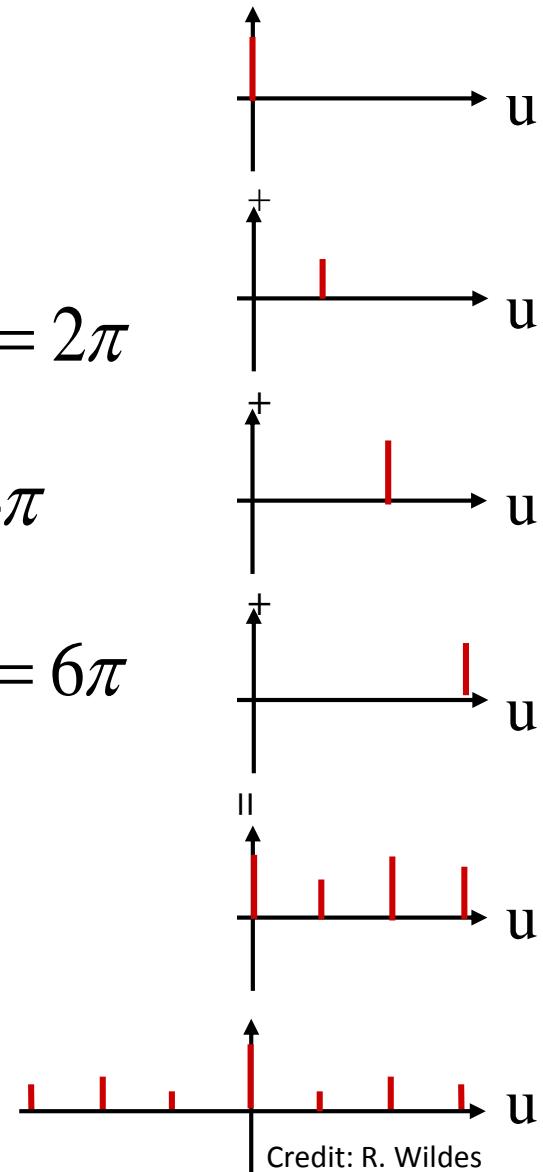
$$+ \frac{1}{2} \cos 2\pi x \rightarrow + \frac{1}{2} \cos ux; u = 2\pi$$

$$+ \cos 4\pi x \rightarrow + 1 \cos ux; u = 4\pi$$

$$+ \frac{2}{3} \cos 6\pi x \rightarrow + \frac{2}{3} \cos ux; u = 6\pi$$

### Note

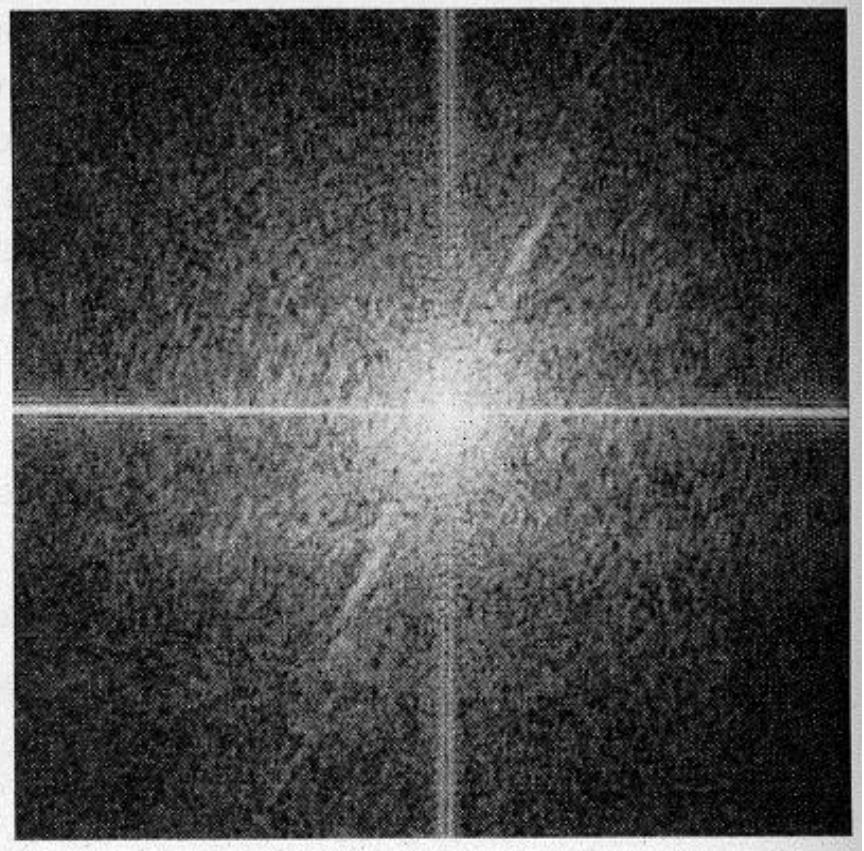
- By symmetry, we may choose to represent this as



# The Fourier transform of an image



Source image (J. Fourier)



Fourier power spectrum

Credit: R. Wildes

# Basics: The 2D Fourier transform

## 2D Eigenfunctions

- For the case of 1D LSI systems we found that  $\exp(iwt)$  is an eigenfunction of convolution.



- One can show that  $\exp[i(ux+vy)]$  is an eigenfunction in 2D.



- The 2D Fourier transform  $F(u,v)$  of  $f(x,y)$  is given by

$$F(u,v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) \exp[-i(ux + vy)] dx dy$$

## The 2D discrete Fourier transform

$$F[u, v] = |F[u, v]| e^{i\phi[u, v]}$$

Discrete domain, Image of height M and width N

Forward transform

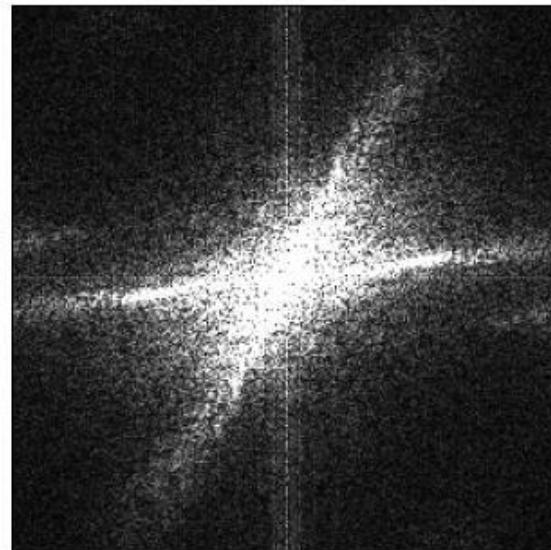
$$F[u, v] = \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f[x, y] e^{-\pi i \left( \frac{ux}{M} + \frac{vy}{N} \right)}$$



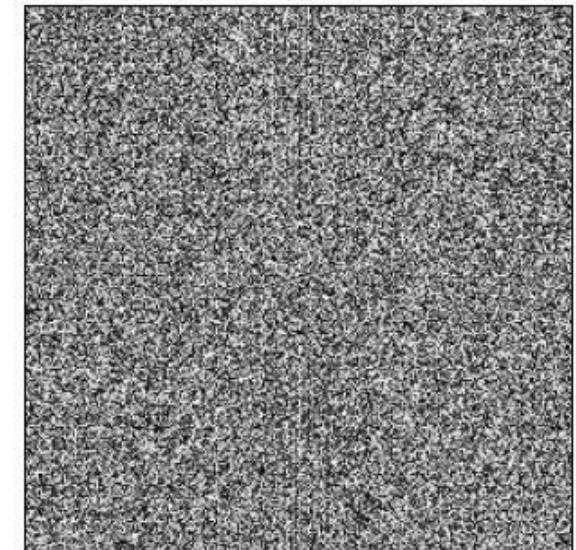
Source image  $f[x, y]$

Inverse transform

$$f[x, y] = \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} F[u, v] e^{+\pi i \left( \frac{xu}{M} + \frac{yv}{N} \right)}$$



Fourier spectrum  $|F[u, v]|$

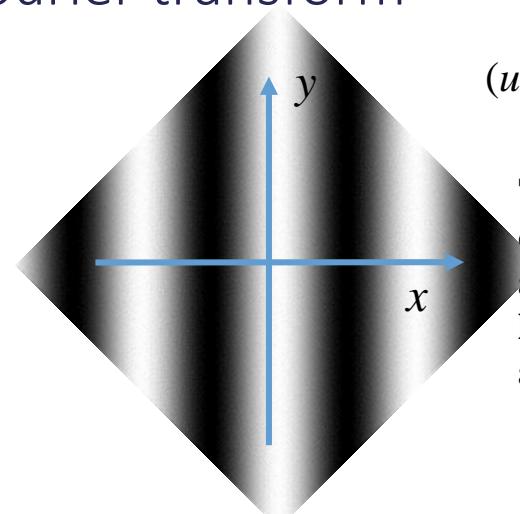


Fourier phase  $\phi[u, v]$

# Basics: The Fourier transform

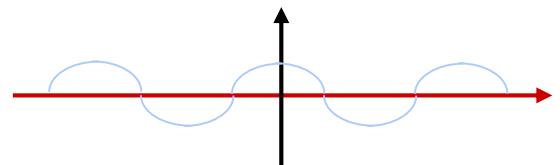
## 2D frequency

- For two spatial dimensions, we see that there are two corresponding frequency components,  $u$  and  $v$ .
- We refer to the  $uv$ -plane as the frequency domain.
- We refer to the  $xy$ -plane as the spatial domain.
- The real waveforms  $\cos(ux+vy)$  and  $\sin(ux+vy)$  correspond to waves in 2D.



$(u, v) = (a, 0)$

**The maxima and minima of the cosinusoids lie along parallel equidistant lines  $ux + vy = k\pi$  for  $k$  an integer.**

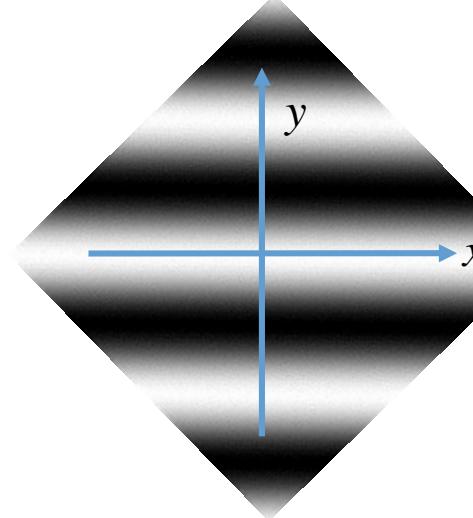


**Cross sections orthogonal to the ridges show a sinusoidal profile**

# Basics: The Fourier transform

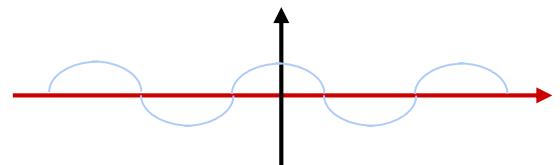
## 2D frequency

- For two spatial dimensions, we see that there are two corresponding frequency components,  $u$  and  $v$ .
- We refer to the  $uv$ -plane as the frequency domain.
- We refer to the  $xy$ -plane as the spatial domain.
- The real waveforms  $\cos(ux+vy)$  and  $\sin(ux+vy)$  correspond to waves in 2D.



$(u,v)=(0,a)$

**The maxima and minima of the cosinusoids lie along parallel equidistant lines  $ux + vy = k\pi$  for  $k$  an integer.**

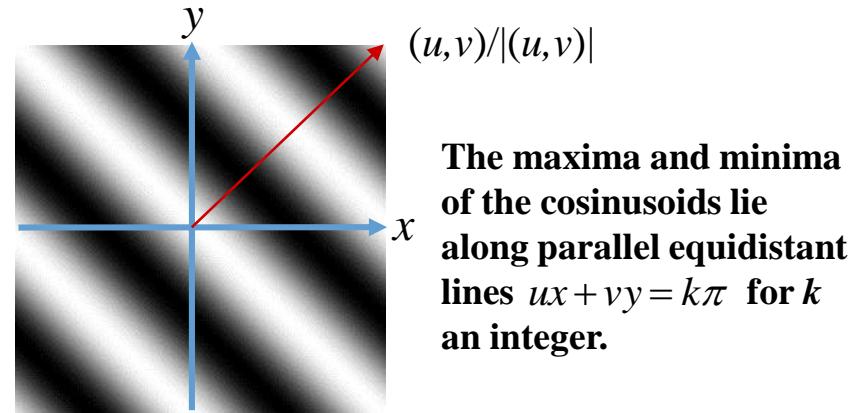


**Cross sections orthogonal to the ridges show a sinusoidal profile**

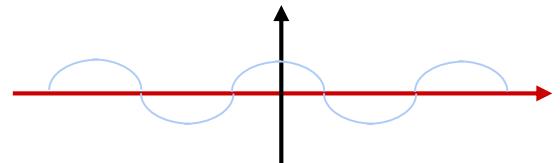
# Basics: The Fourier transform

## 2D frequency

- For two spatial dimensions, we see that there are two corresponding frequency components,  $u$  and  $v$ .
- We refer to the  $uv$ -plane as the frequency domain.
- We refer to the  $xy$ -plane as the spatial domain.
- The real waveforms  $\cos(ux+vy)$  and  $\sin(ux+vy)$  correspond to waves in 2D.



**The maxima and minima of the cosinusoids lie along parallel equidistant lines  $ux + vy = k\pi$  for  $k$  an integer.**

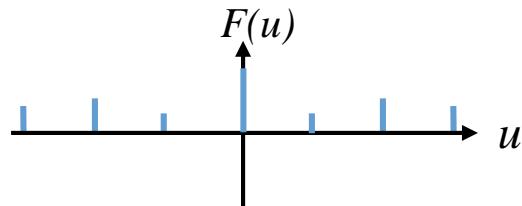
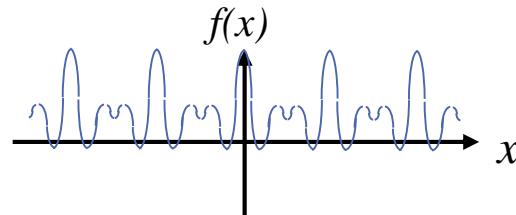


**Cross sections orthogonal to the ridges show a sinusoidal profile**

# The Fourier transform: 2D case

## A 2D example

- Recall the 1D example

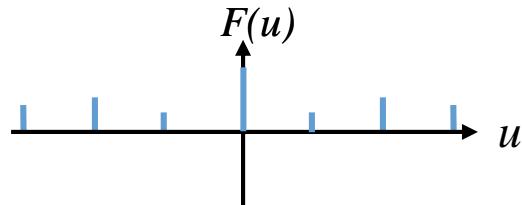
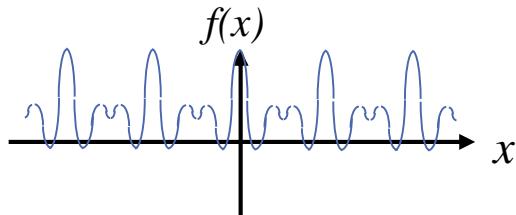


for the cosine component

# The Fourier transform: 2D case

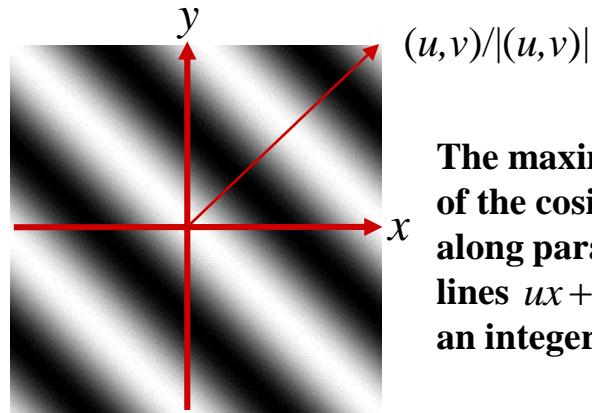
## A 2D example

- Recall the 1D example



for the cosine component

- And the interpretation of 2D spatial frequency



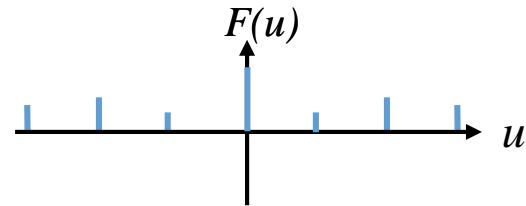
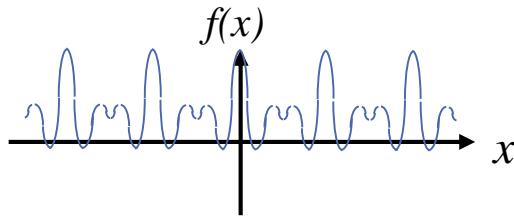
$(u, v)/|(u, v)|$   
**The maxima and minima of the cosinusoids lie along parallel equidistant lines  $ux + vy = k\pi$  for  $k$  an integer.**

$f[x, y]$

# The Fourier transform: 2D case

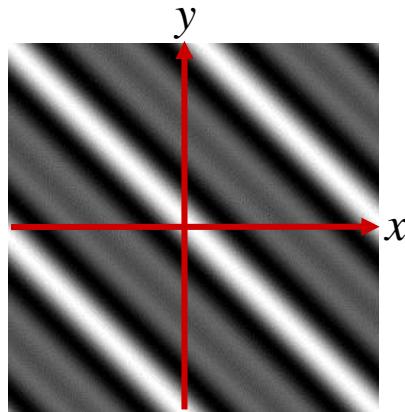
## A 2D example

- Recall the 1D example



for the cosine component

- Then a 2D analogue could be

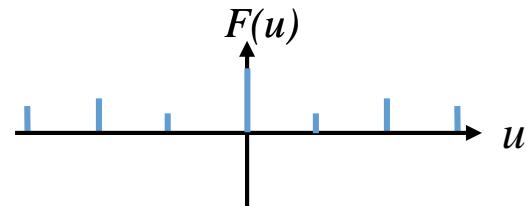
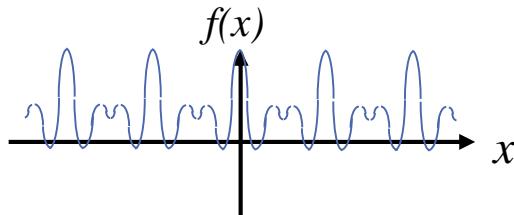


$f[x, y]$

# The Fourier transform: 2D case

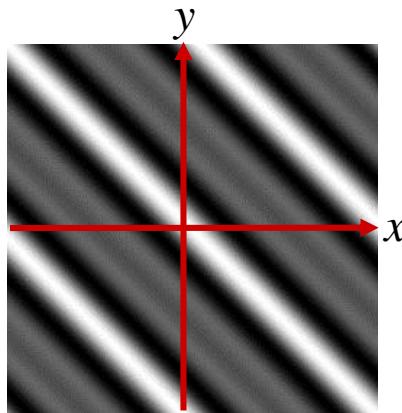
## A 2D example

- Recall the 1D example

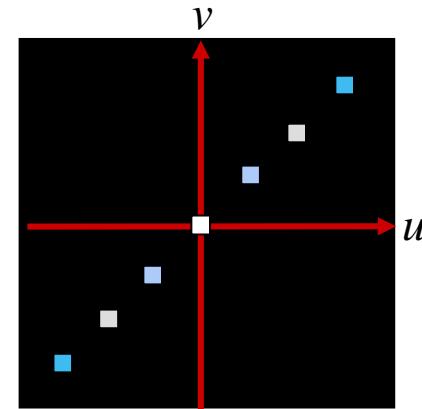


for the cosine component

- Then a 2D analogue could be

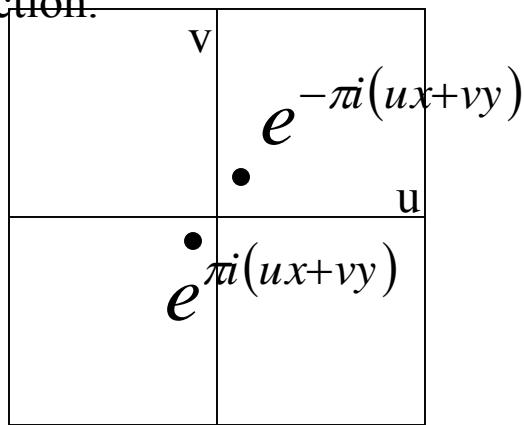


$f[x, y]$



$|F[u, v]|$

To get some sense of what basis elements look like, we plot a basis element --- or rather, its real part --- as a function of  $x, y$  for some fixed  $u, v$ . We get a function that is constant when  $(ux+vy)$  is constant. The magnitude of the vector  $(u, v)$  gives a frequency, and its direction gives an orientation. The function is a sinusoid with this frequency along the direction, and constant perpendicular to the direction.

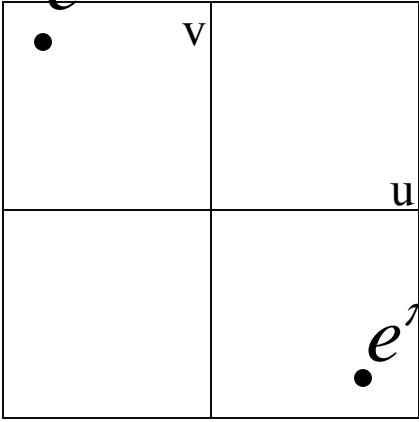


Here  $u$  and  $v$  are larger than in the previous slide.

$$\begin{array}{|c|c|} \hline & v \\ e^{-\pi i(ux+vy)} & \bullet \\ \bullet & u \\ \hline & \\ \bullet & e^{\pi i(ux+vy)} \\ \hline \end{array}$$

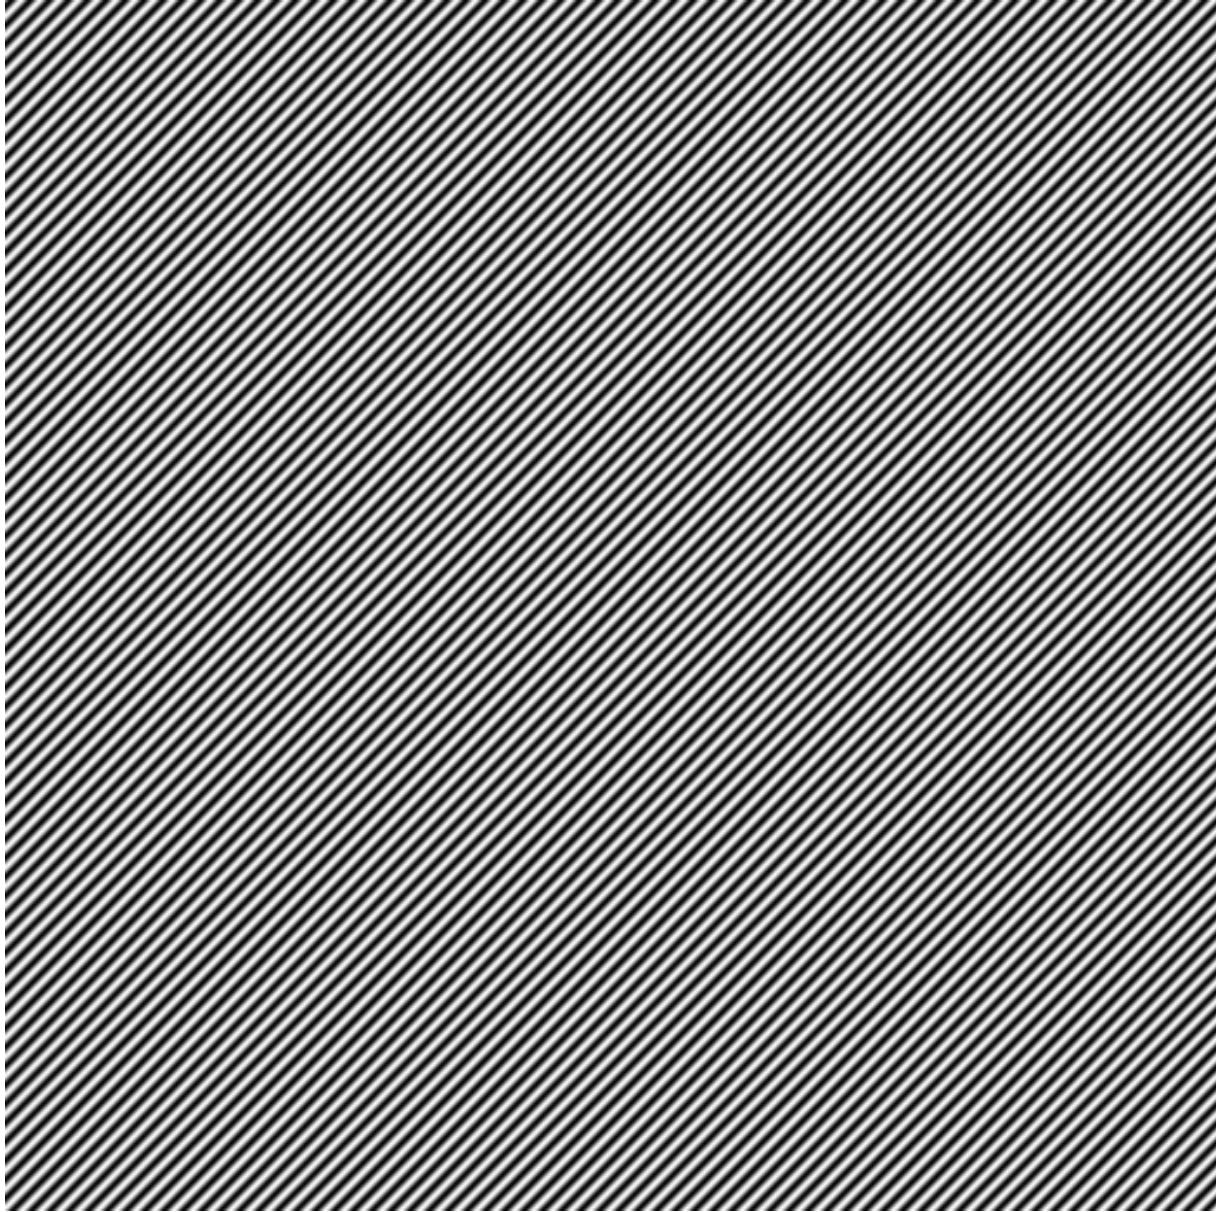


And larger still...

$$e^{-\pi i(ux+vy)}$$


A 2x2 grid with a dot at the origin (0,0). The label 'v' is positioned above the top-right cell, and the label 'u' is positioned to the right of the bottom-right cell.

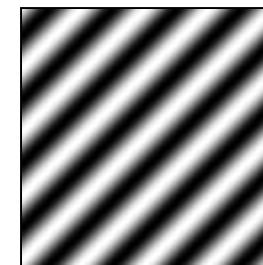
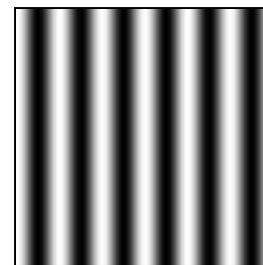
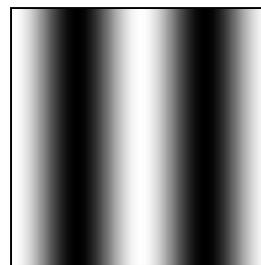
$$e^{\pi i(ux+vy)}$$



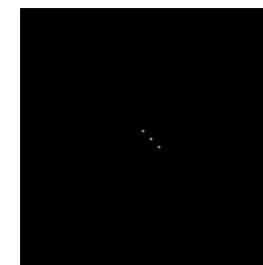
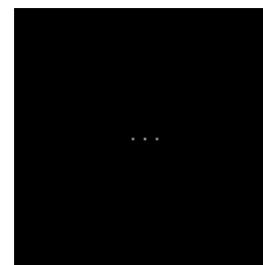
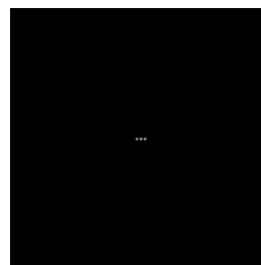
Credit: B. Freeman

# The Fourier transform: Filtering

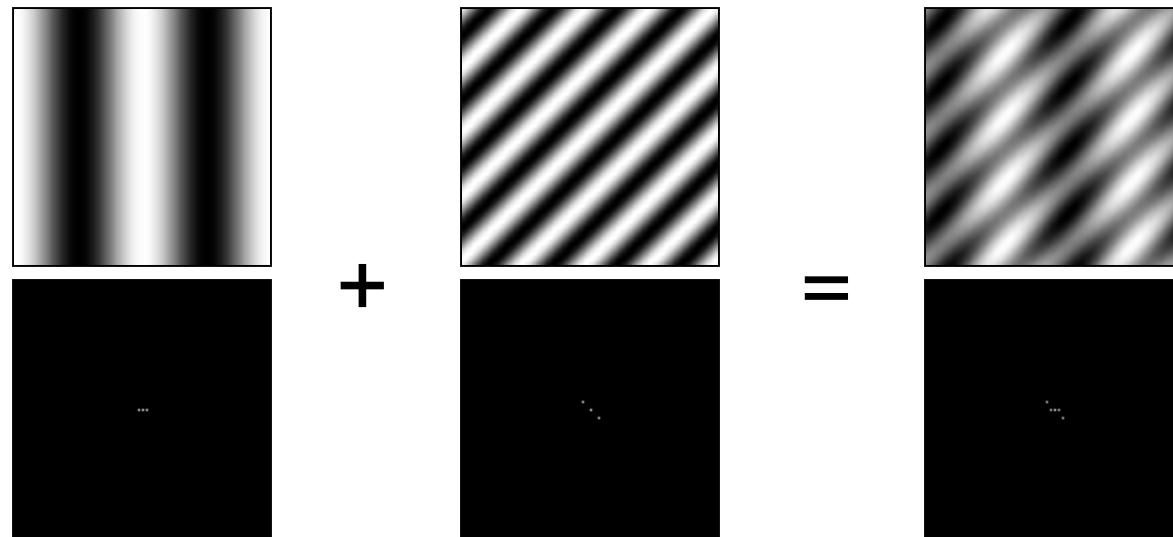
Intensity Image



Fourier Image



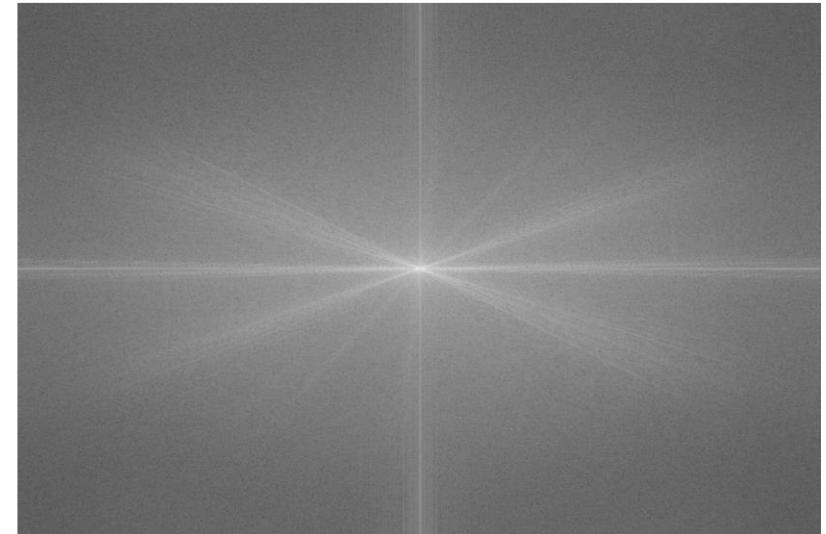
## The Fourier transform: Filtering



# Understanding the Fourier transform of an image



Source image (J. Fourier)



Fourier power spectrum

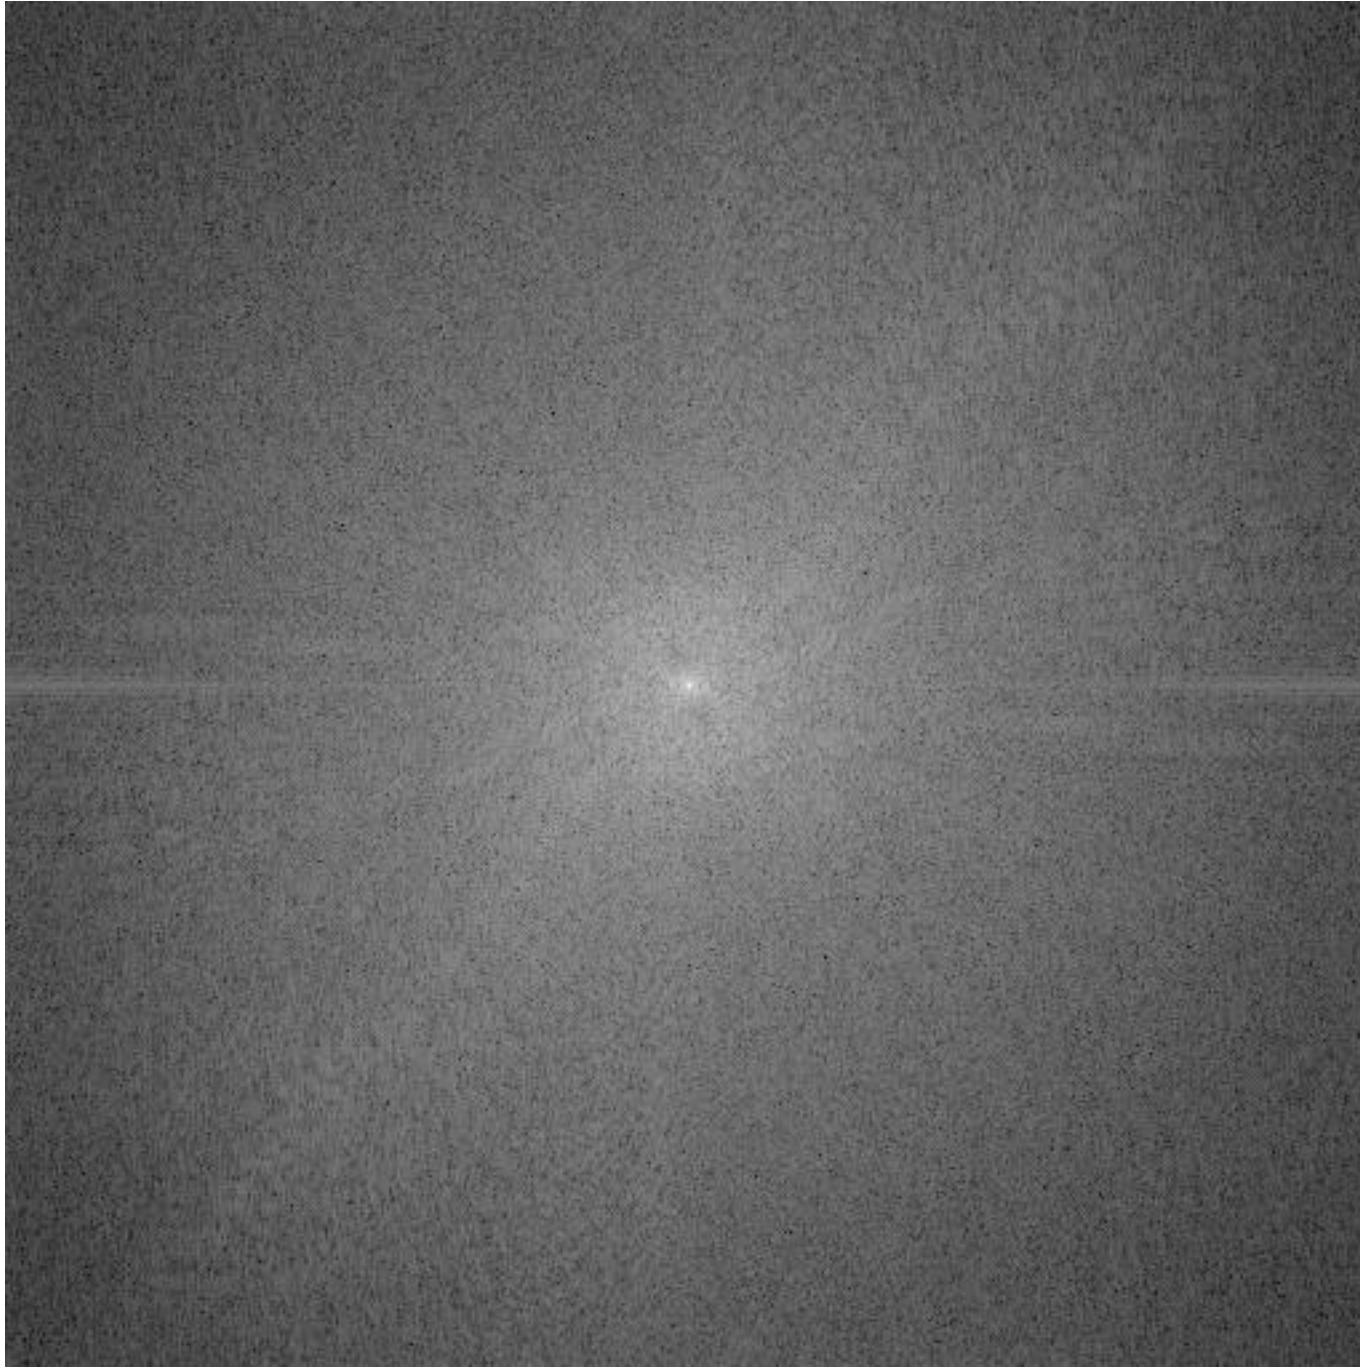
## Phase and Magnitude

- Fourier transform of a real function is complex
  - difficult to plot, visualize
  - instead, we can think of the phase and magnitude of the transform
- Phase is the phase of the complex transform
- Magnitude is the magnitude of the complex transform
- Curious fact
  - all natural images have about the same magnitude transform
  - hence, phase seems to matter, but magnitude largely doesn't
- Demonstration
  - Take two pictures, swap the phase transforms, compute the inverse - what does the result look like?



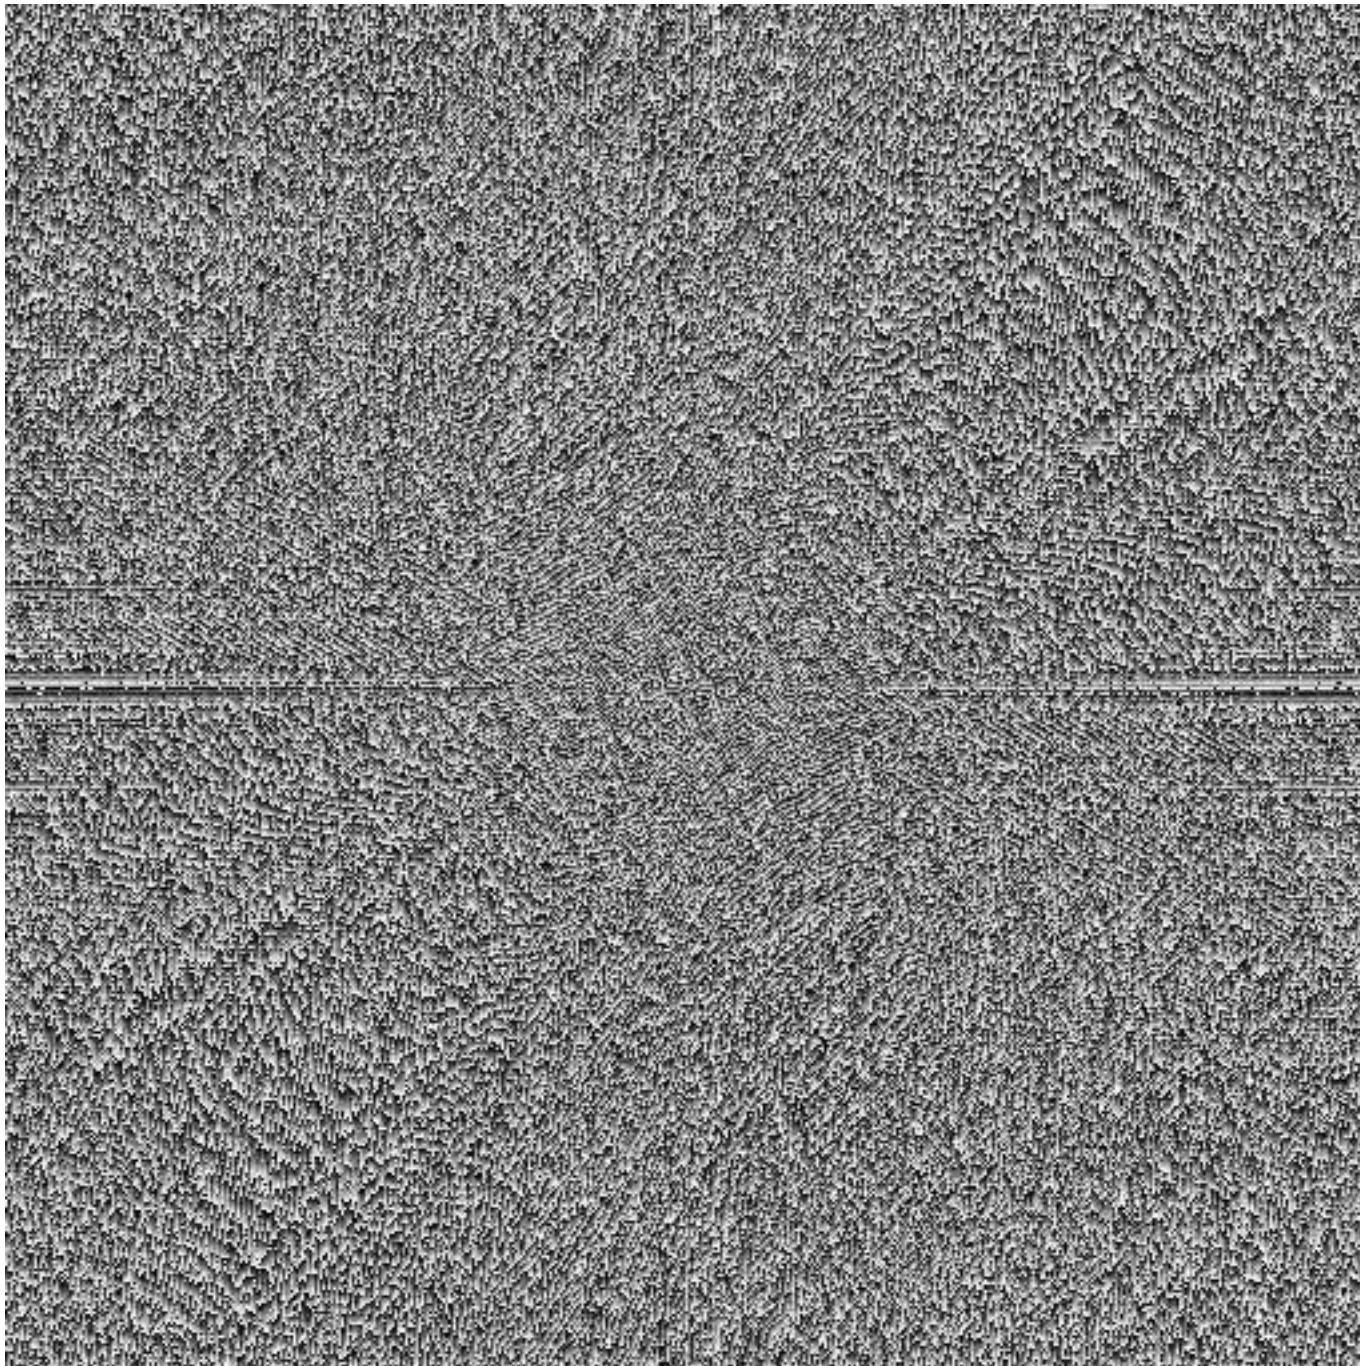
Credit: B. Freeman

This is the  
magnitude  
transform  
of the  
cheetah pic



Credit: B. Freeman

This is the  
phase  
transform  
of the  
cheetah pic

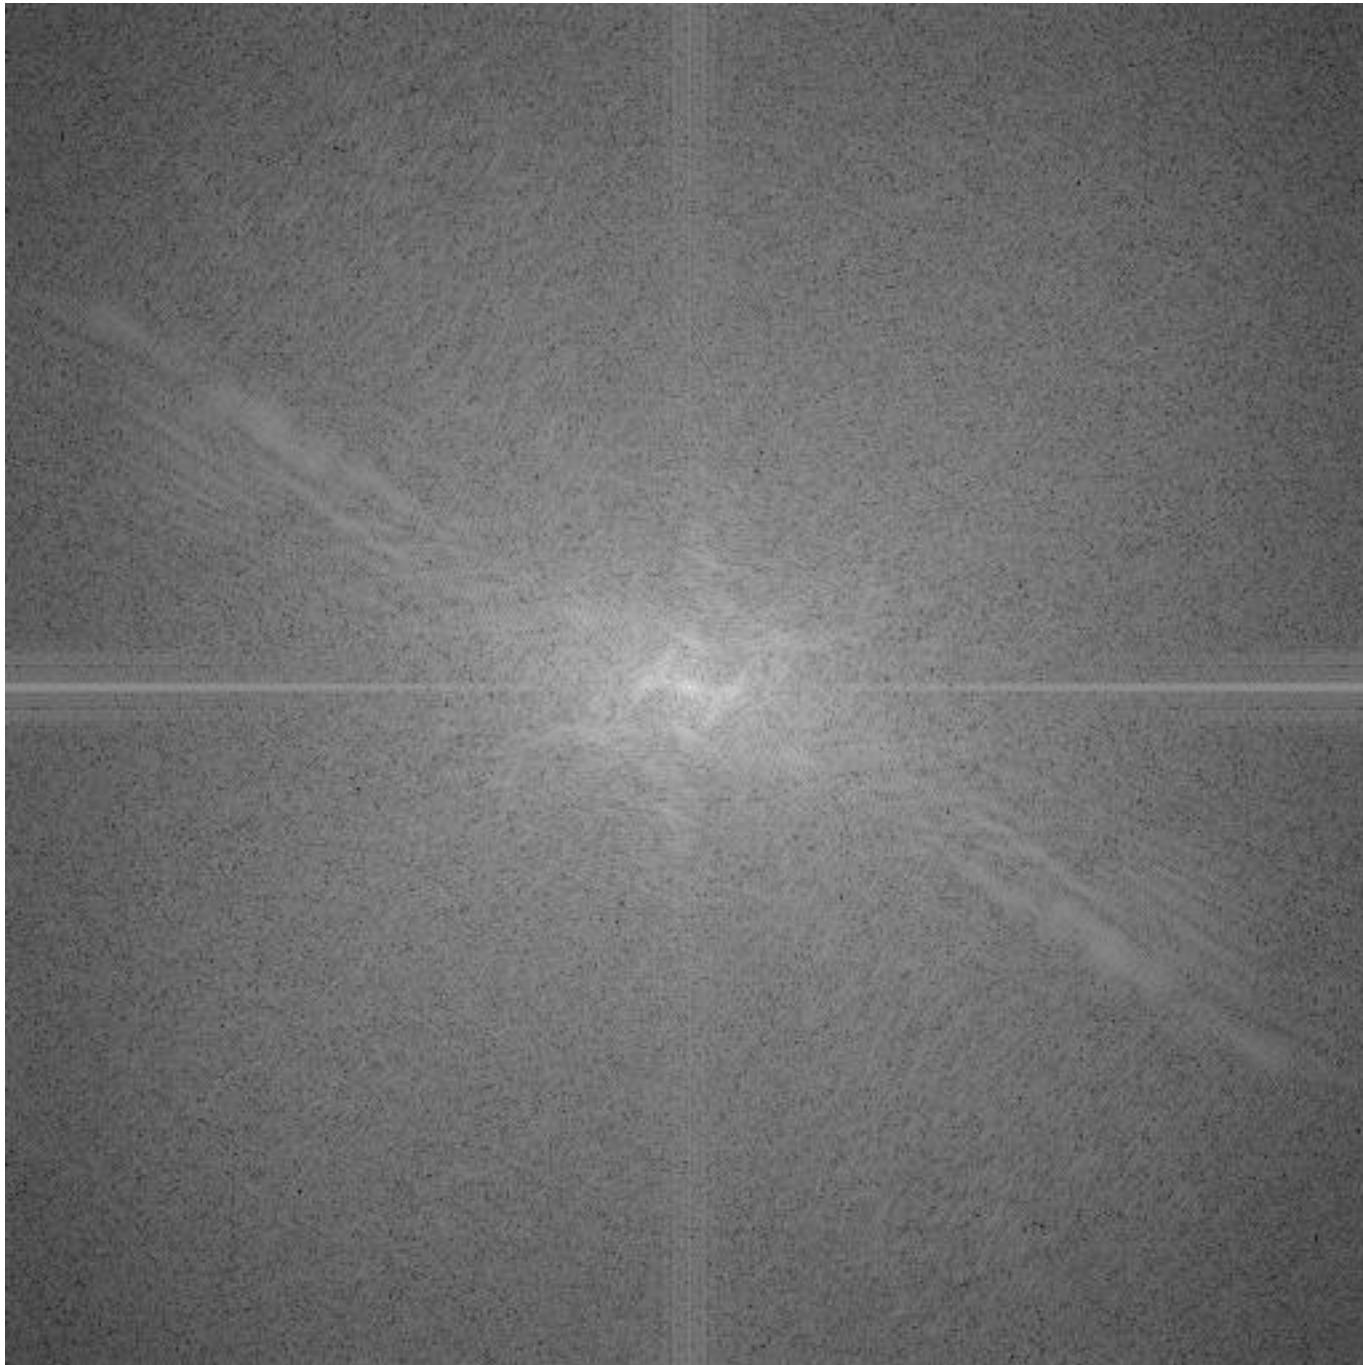


Credit: B. Freeman



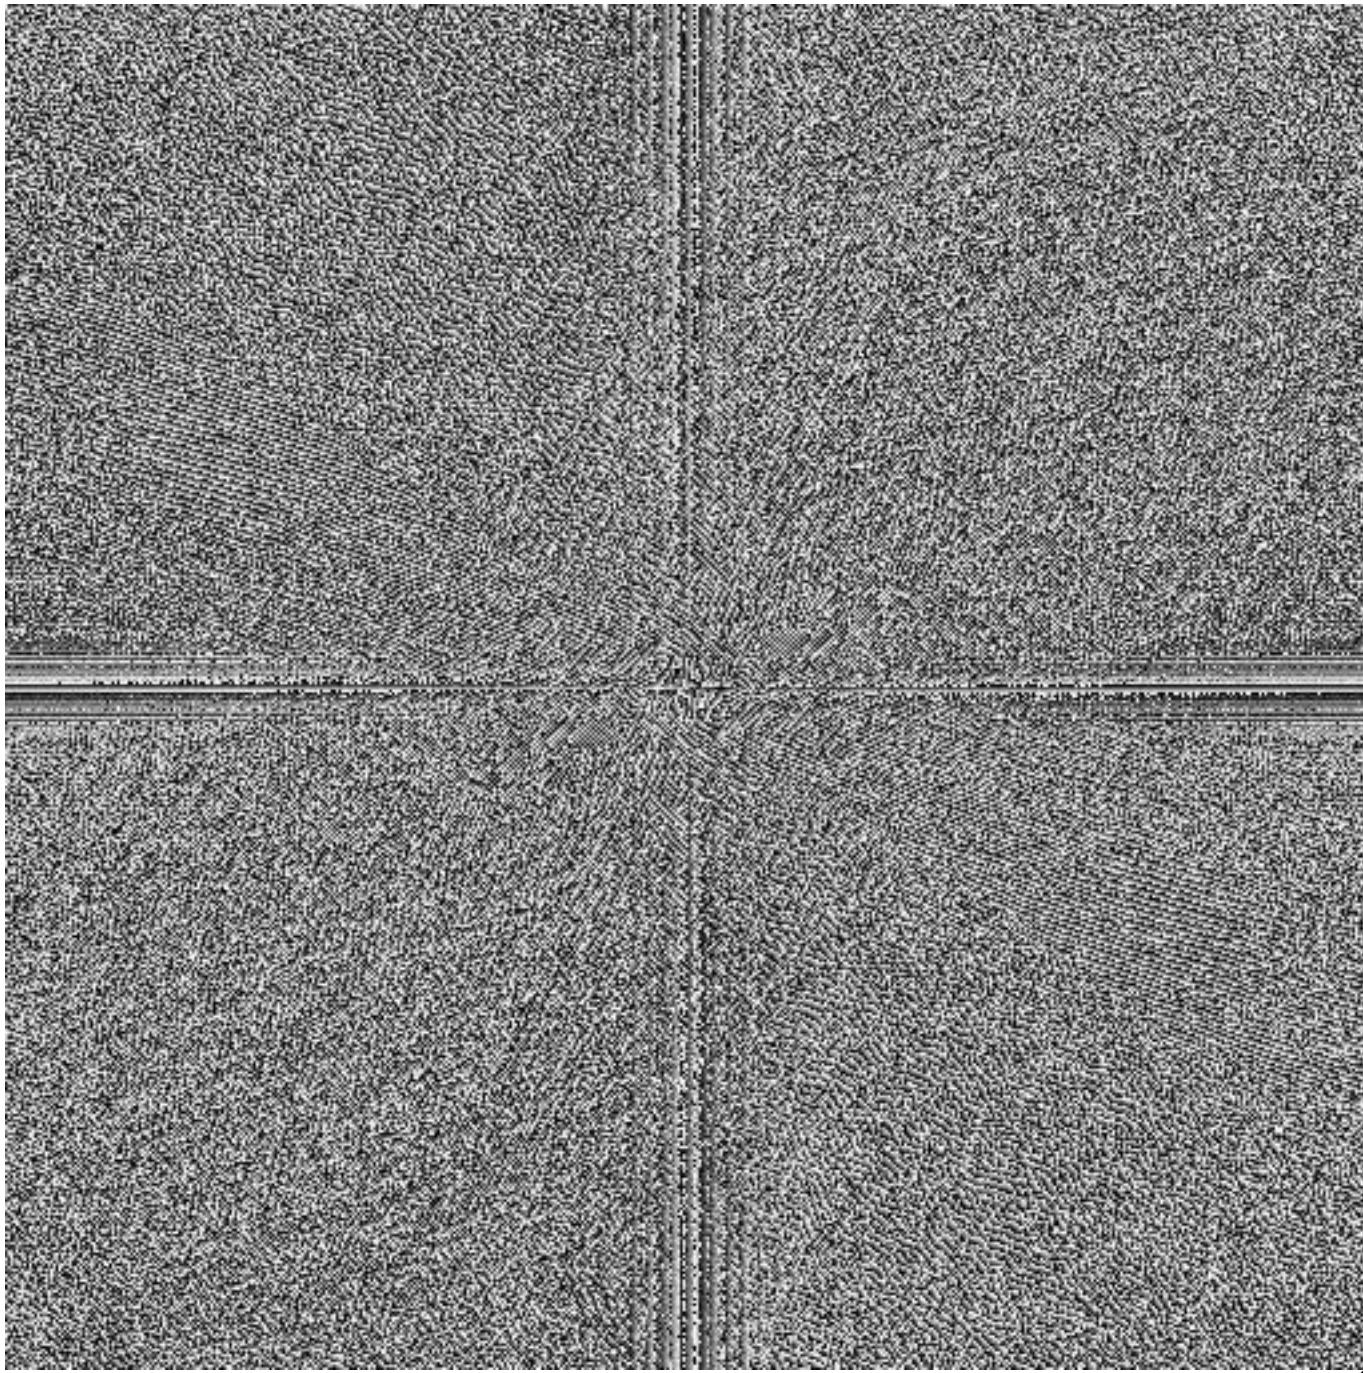
Credit: B. Freeman

This is the  
magnitude  
transform  
of the zebra  
pic



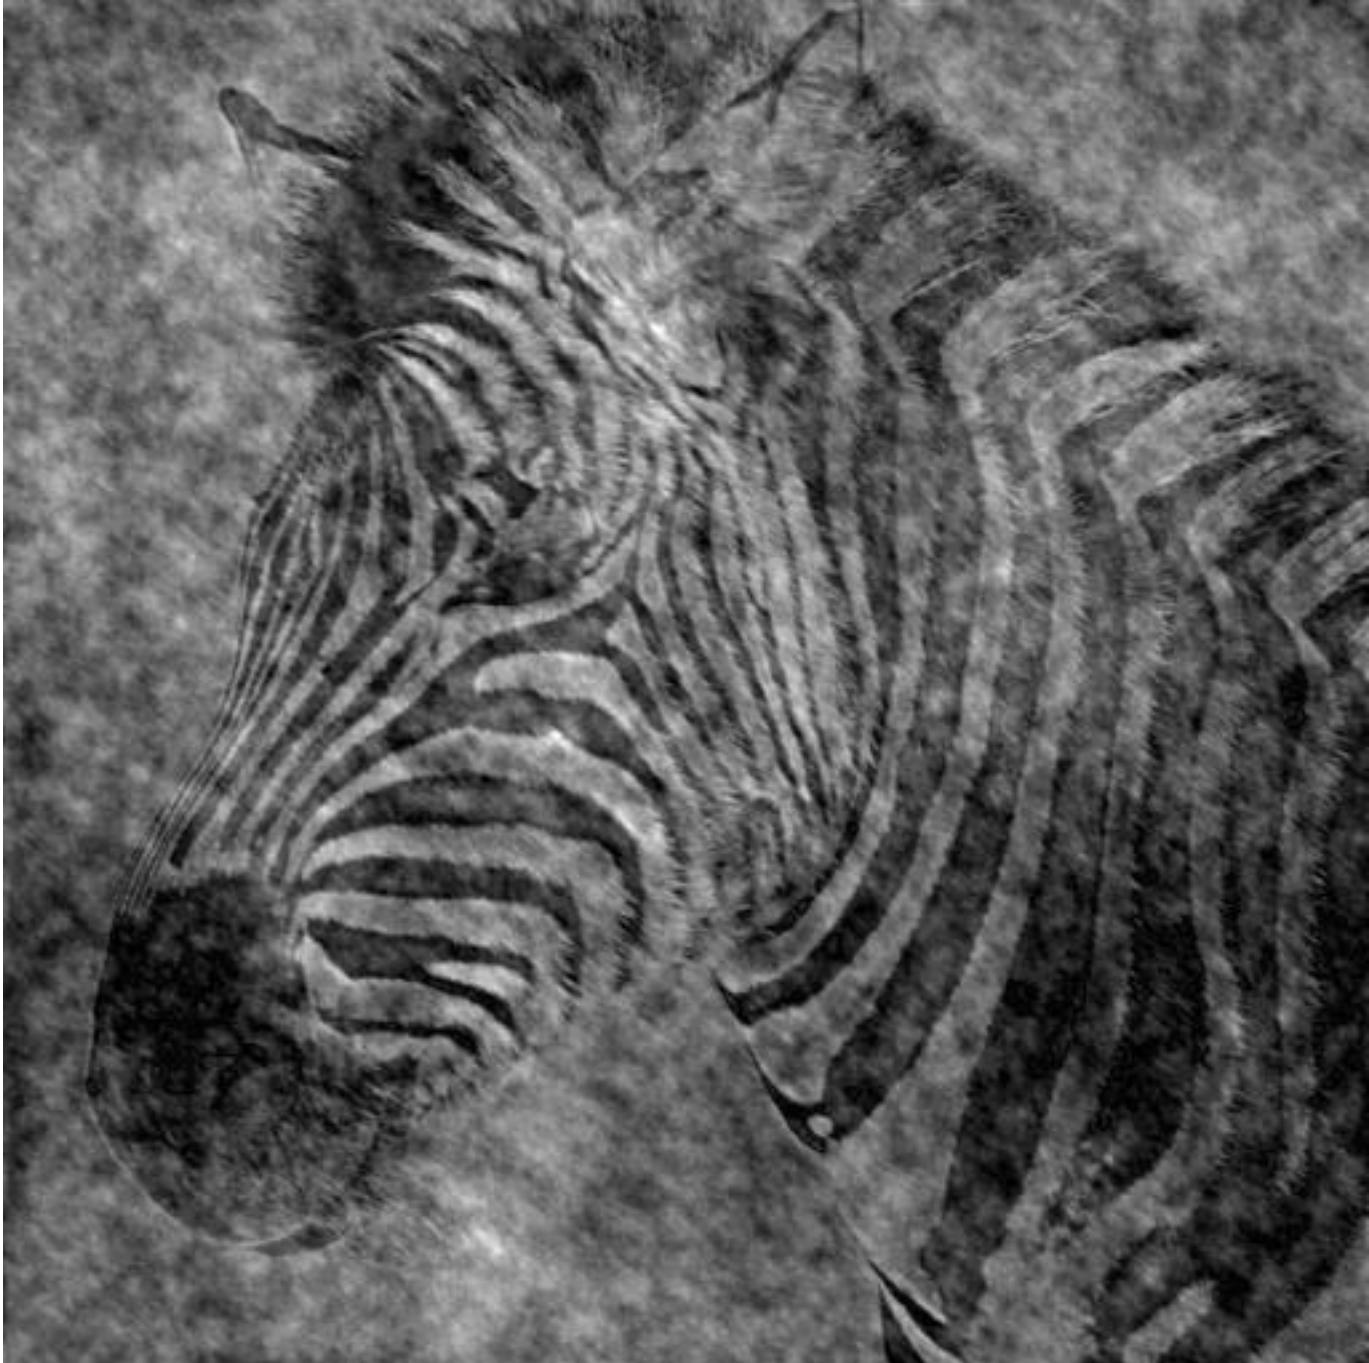
Credit: B. Freeman

This is the  
phase  
transform  
of the zebra  
pic



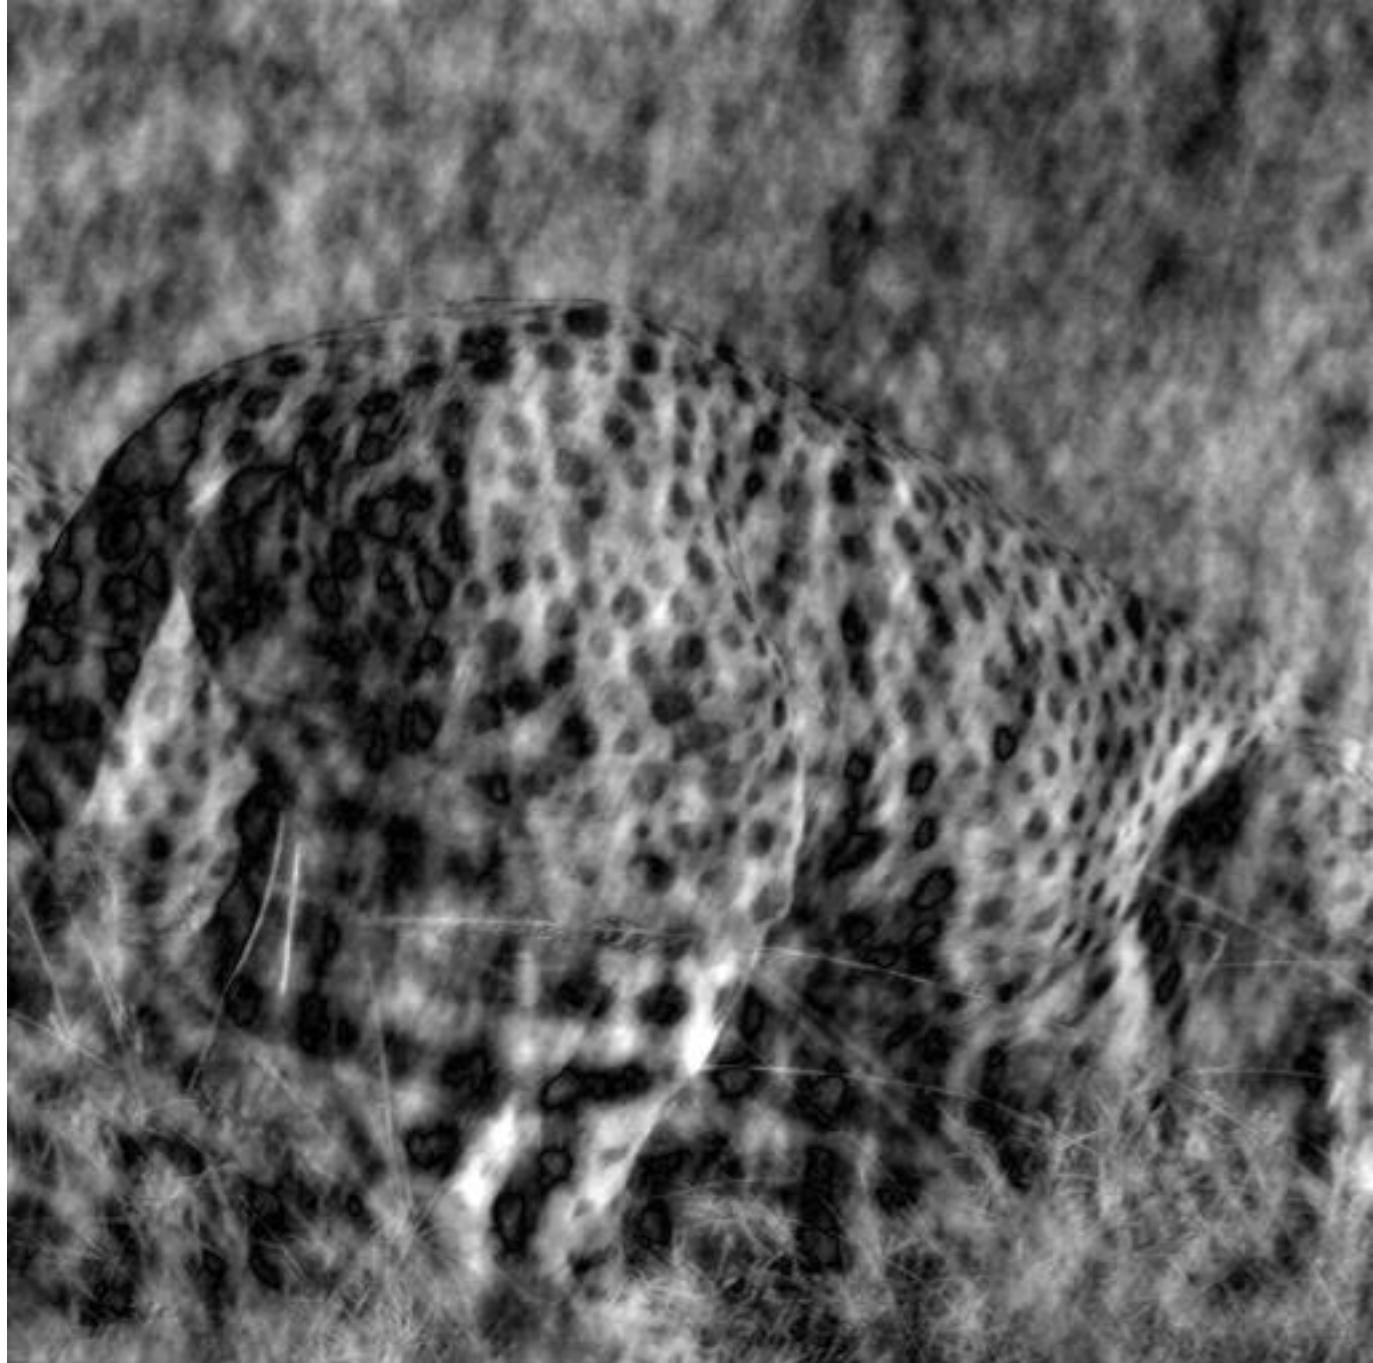
Credit: B. Freeman

Reconstruction  
with zebra  
phase, cheetah  
magnitude



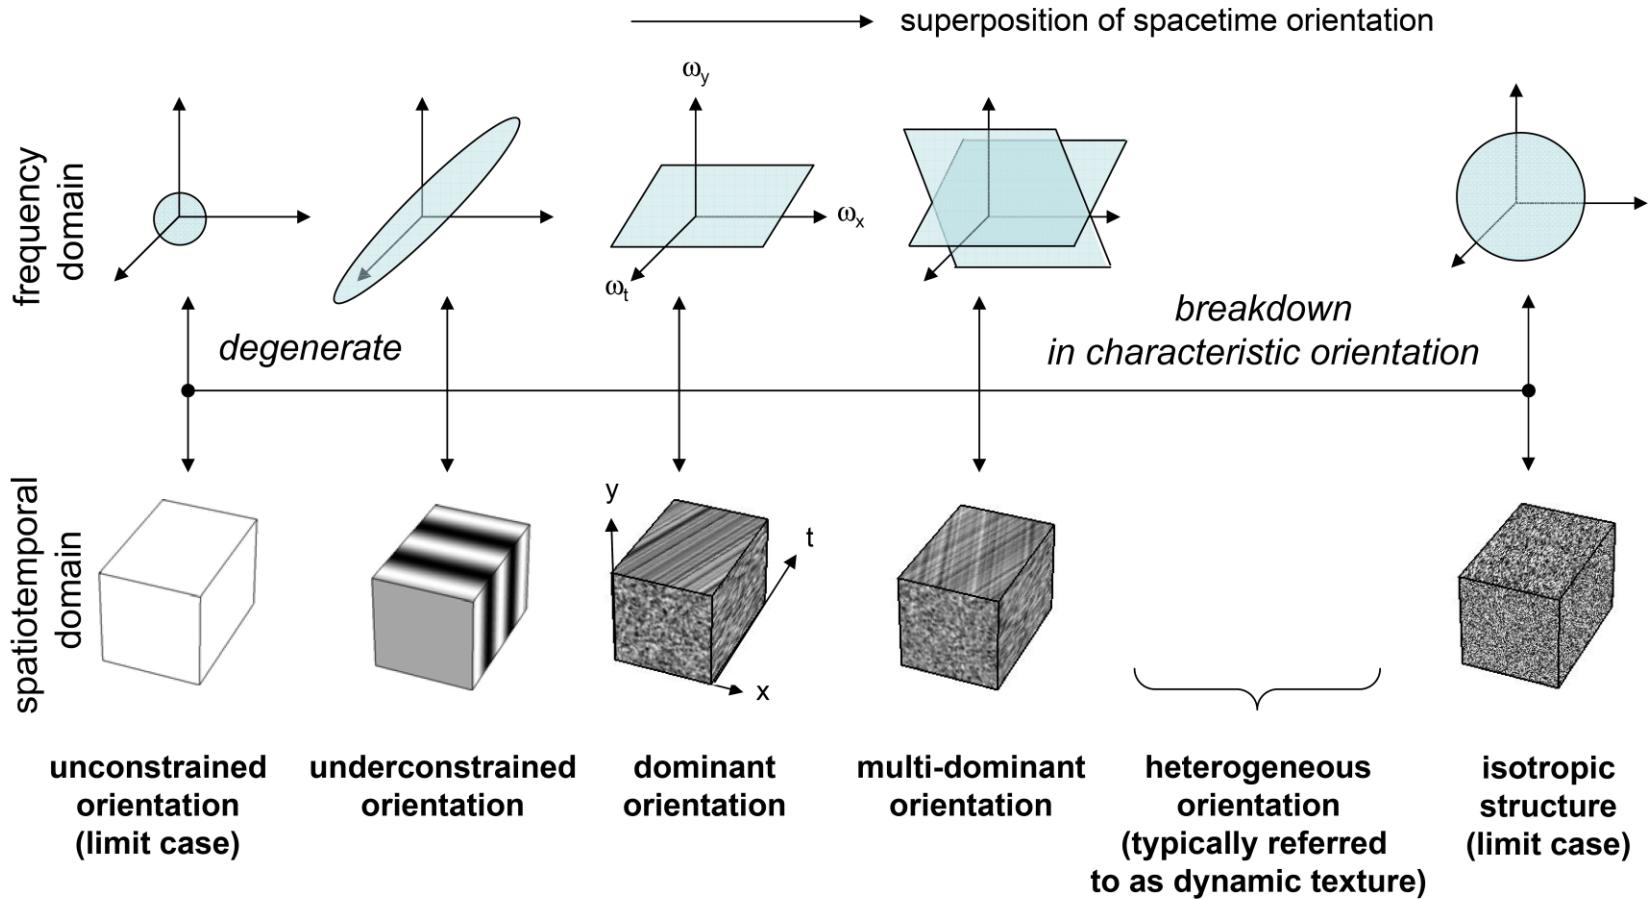
Credit: B. Freeman

Reconstruction  
with cheetah  
phase, zebra  
magnitude



Credit: B. Freeman

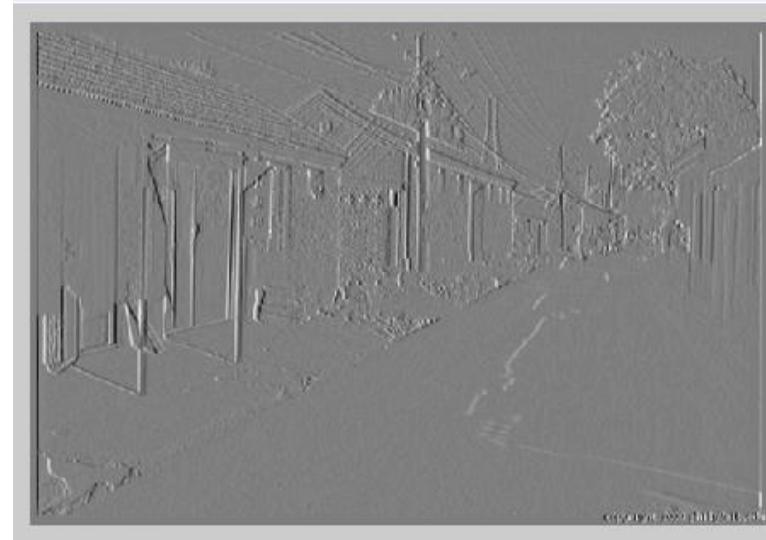
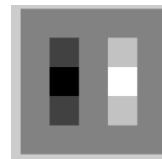
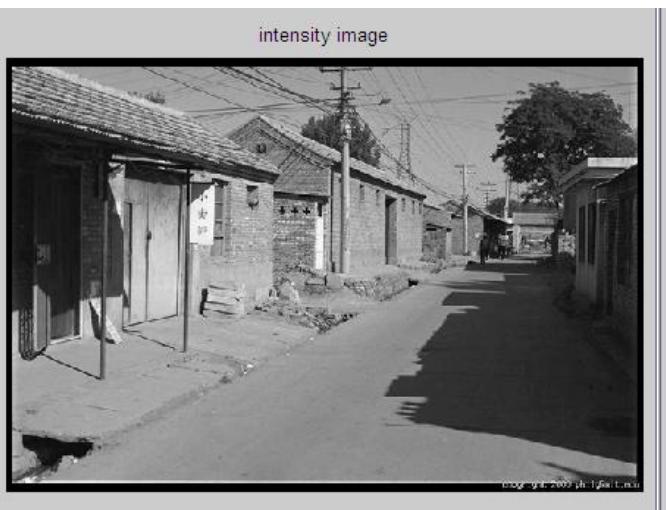
## Extension to 3D



## Filtering in spatial domain

1	0	-1
2	0	-2
1	0	-1

intensity image



## The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

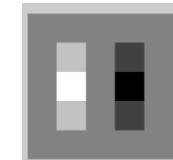
$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

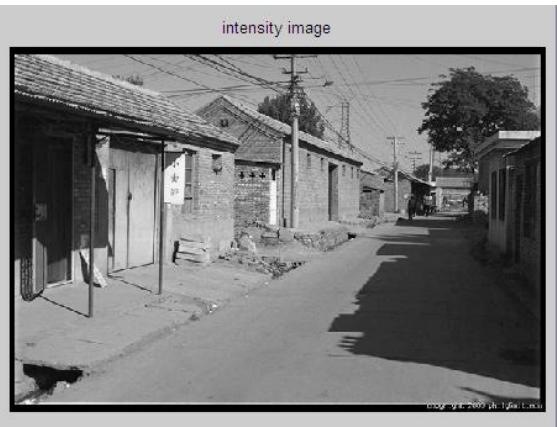
$$g * h = \mathcal{F}^{-1}[\mathcal{F}[g]\mathcal{F}[h]]$$

## Filtering in frequency domain

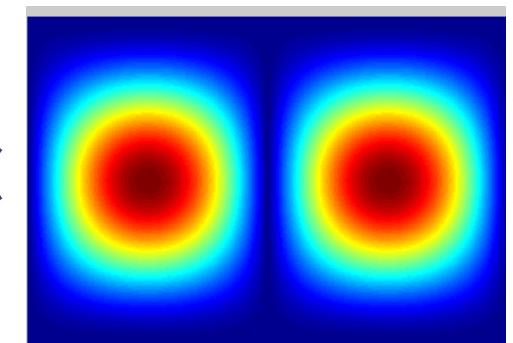
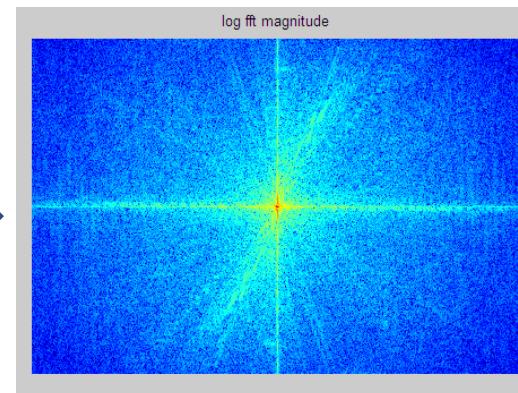
Fast Fourier Transform (FFT) = fast implementation



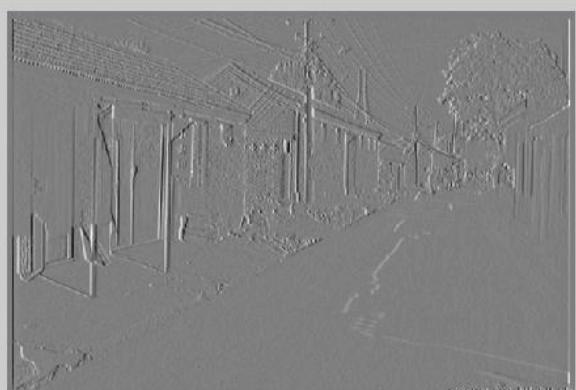
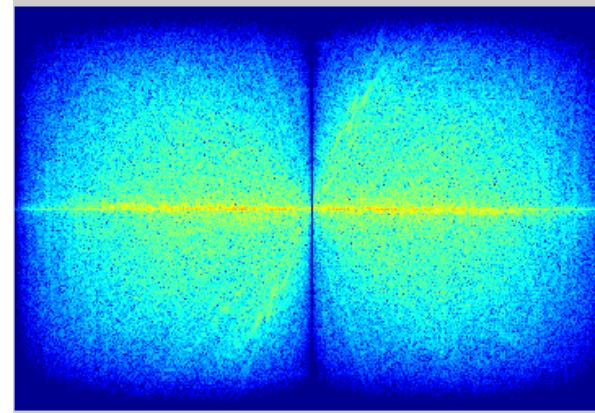
FFT



FFT



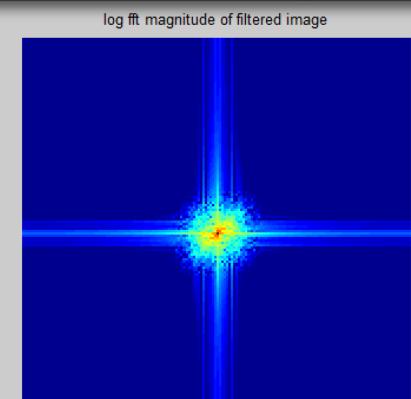
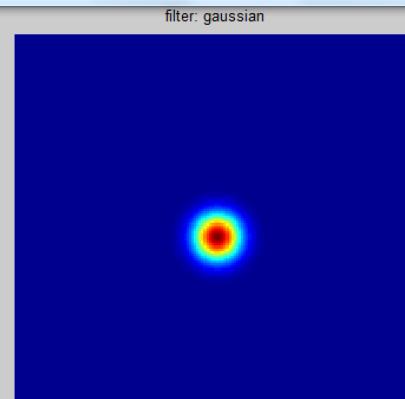
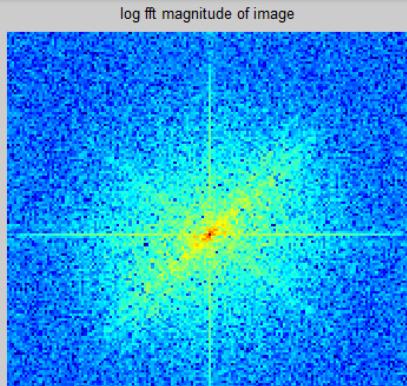
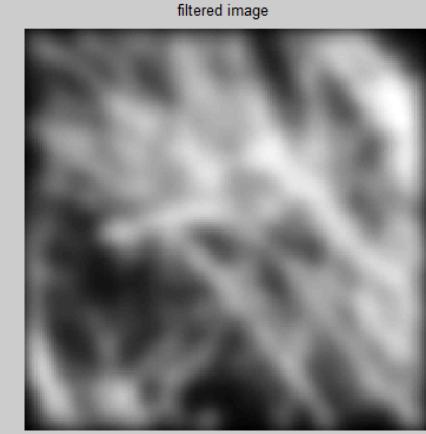
Inverse FFT



Credit: D. Hoiem

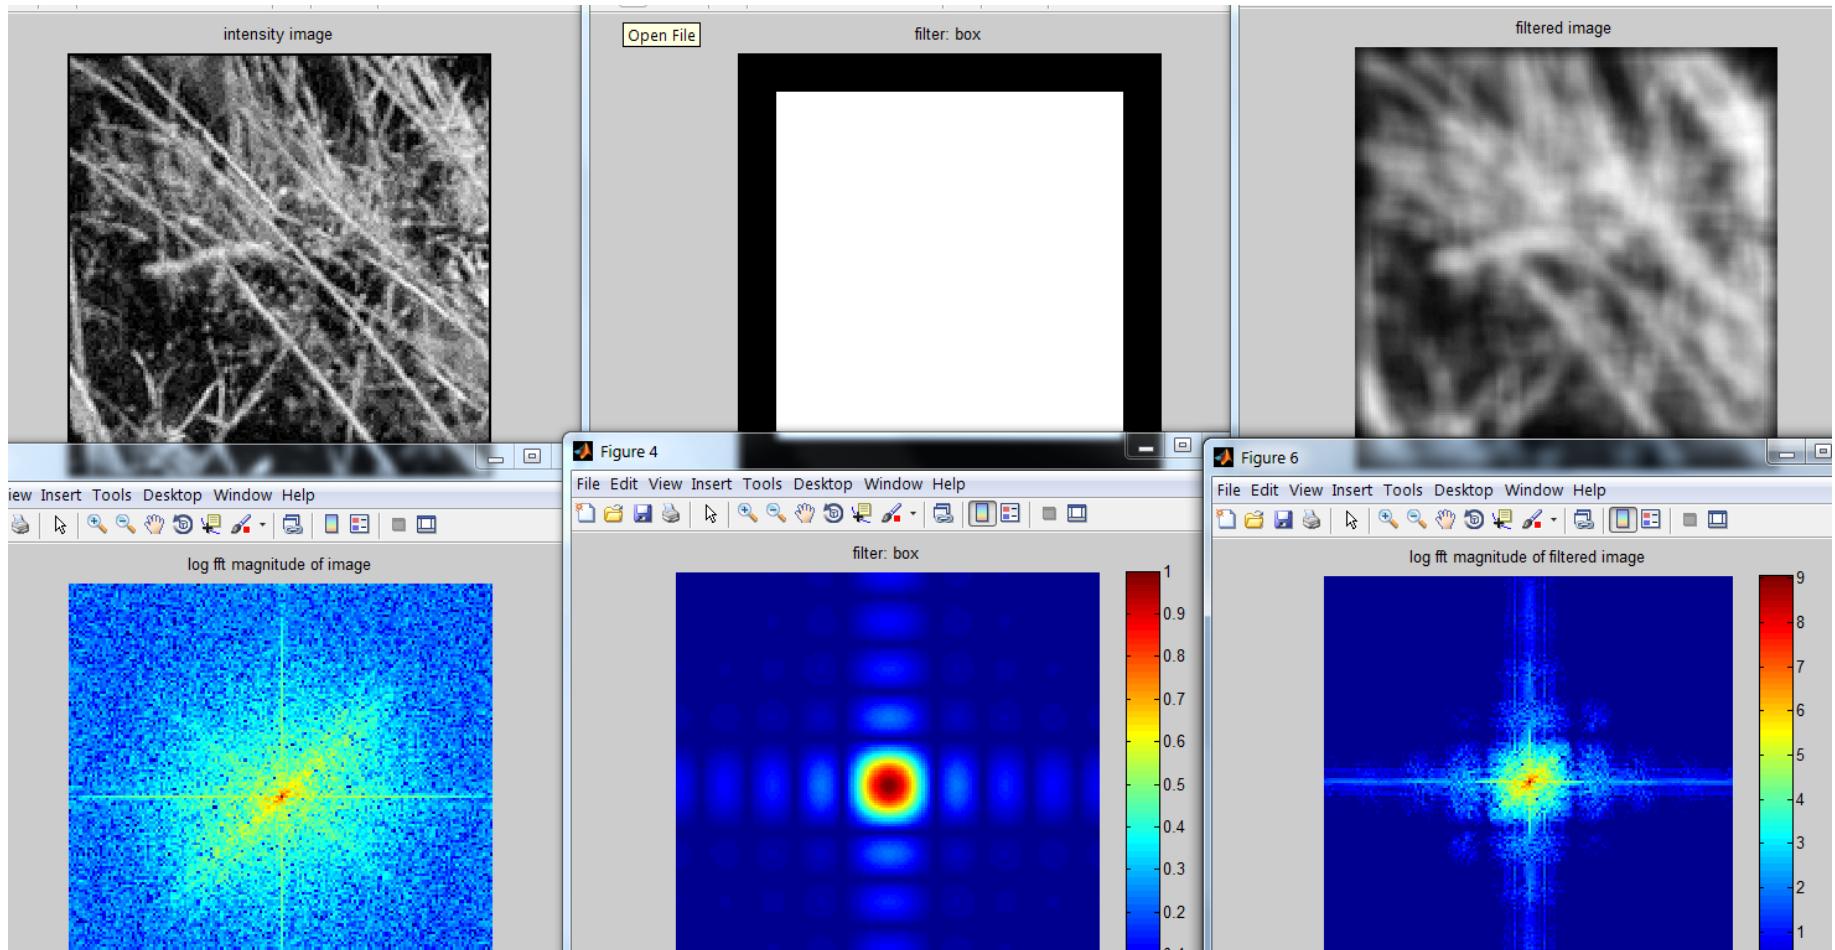
Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

## Gaussian



Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

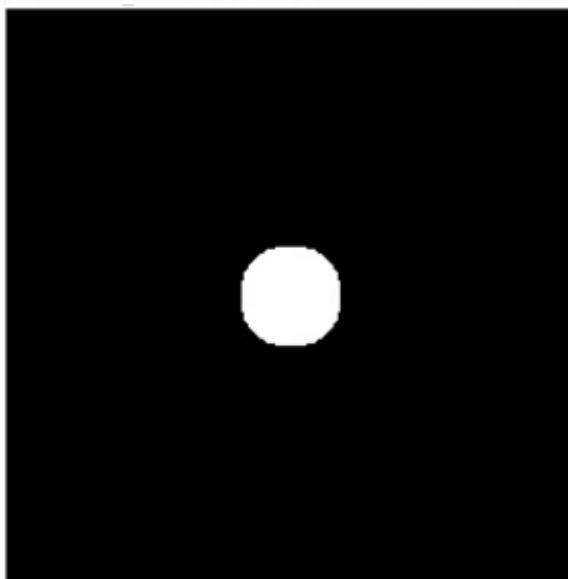
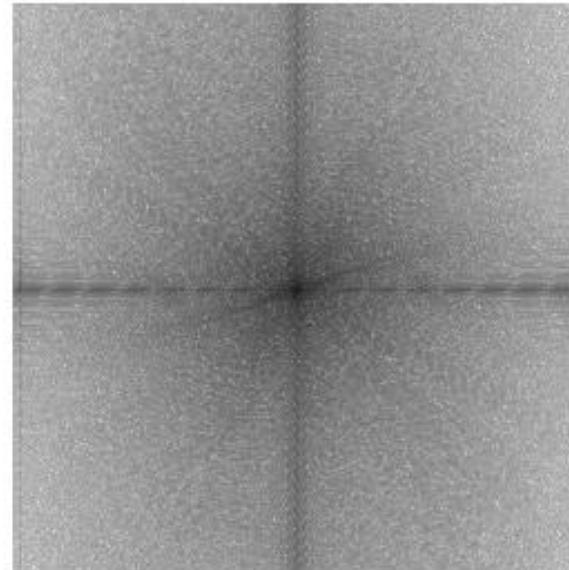
## Box Filter



Credit: D. Hoiem

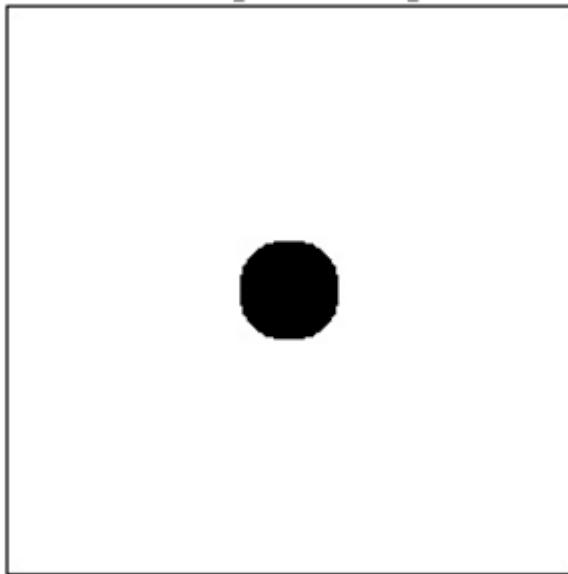
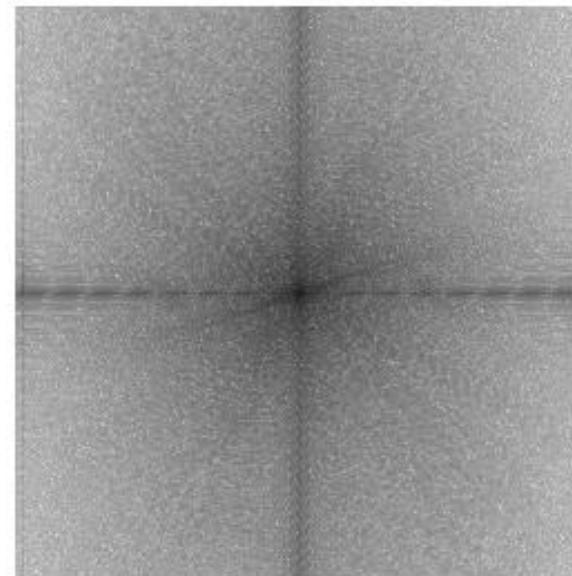
# Low pass filtering

<http://www.reindeergraphics.com>



# High pass filtering

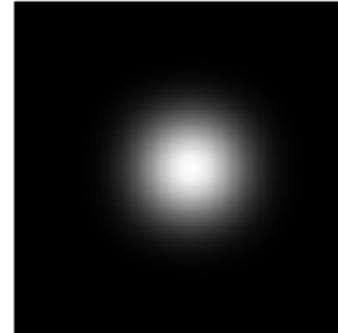
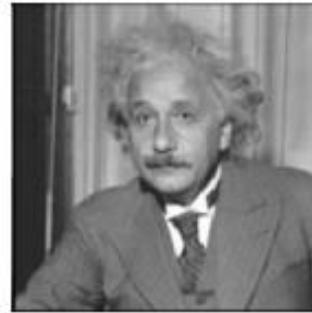
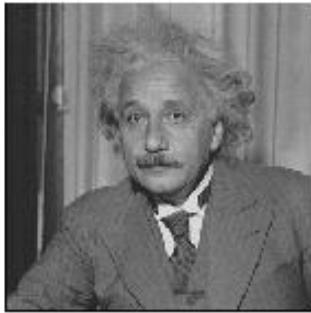
<http://www.reindeergraphics.com>



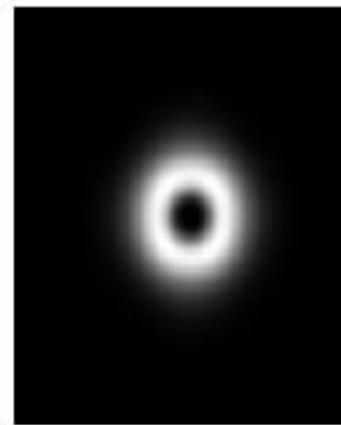
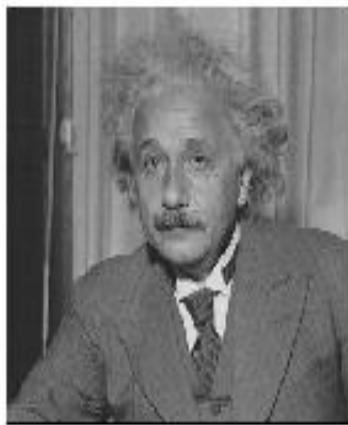
Credit: R. Fergus

Low-pass, Band-pass, High-pass filters

low-pass:

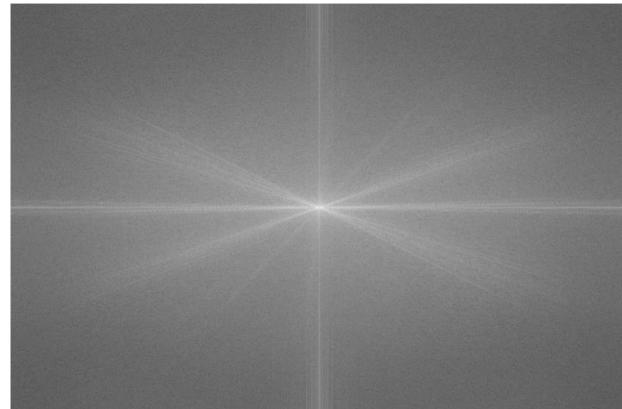


High-pass / band-pass:

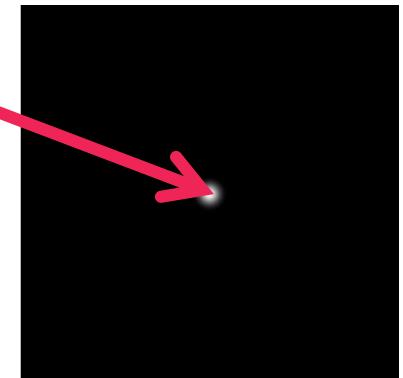
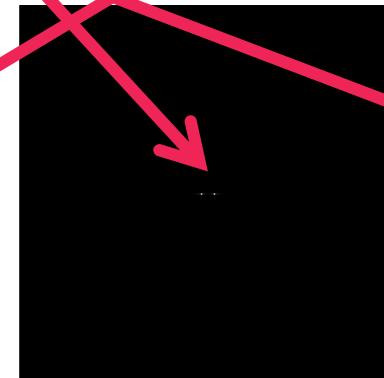
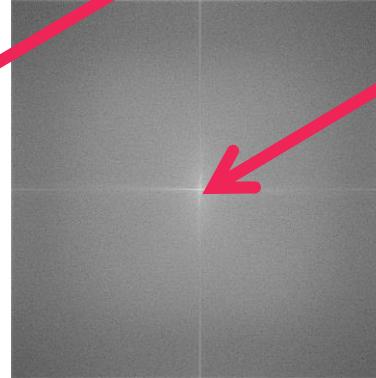
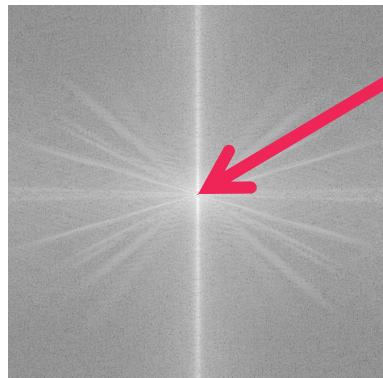
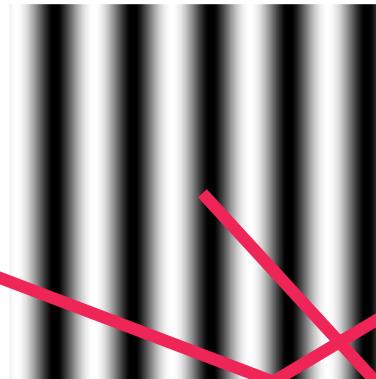
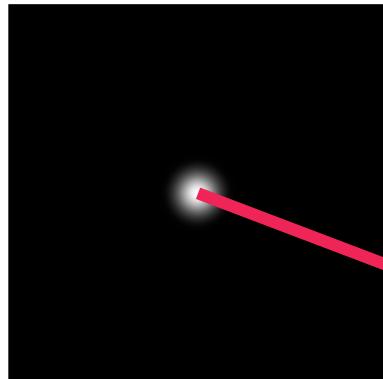


# Why is the Frequency domain useful for us?

- The linear **convolution** operation can be understood from a different angle
- It can be performed very efficient using a clever implementation (FFT)
- The Frequency domain provides an alternative way to understand and manipulate the content of images

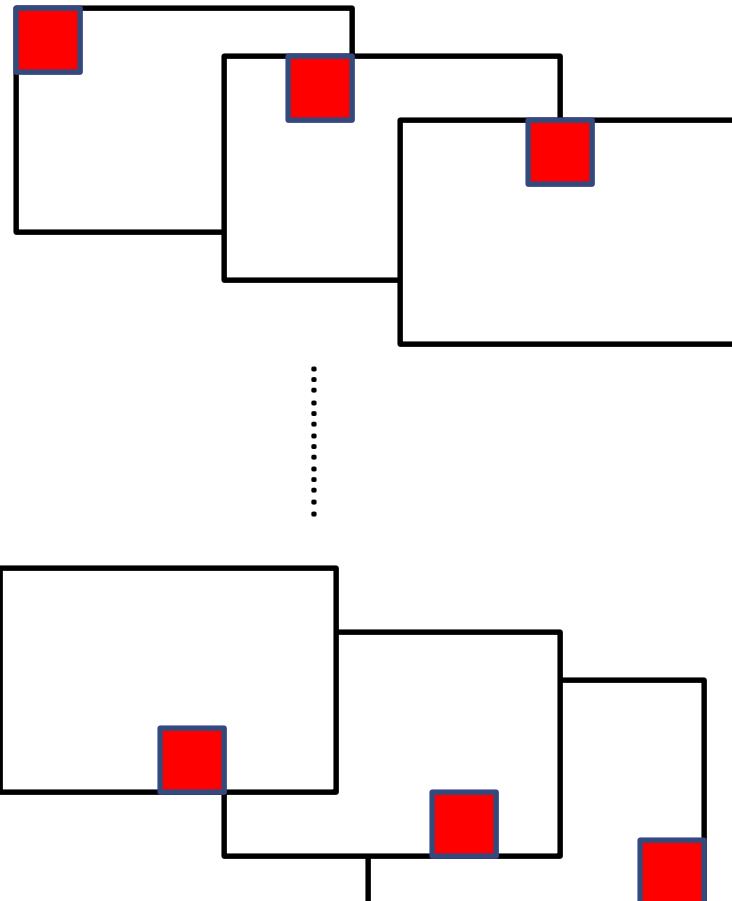


Match the spatial domain image to the Fourier magnitude image



## Spatial Domain

Basis functions:

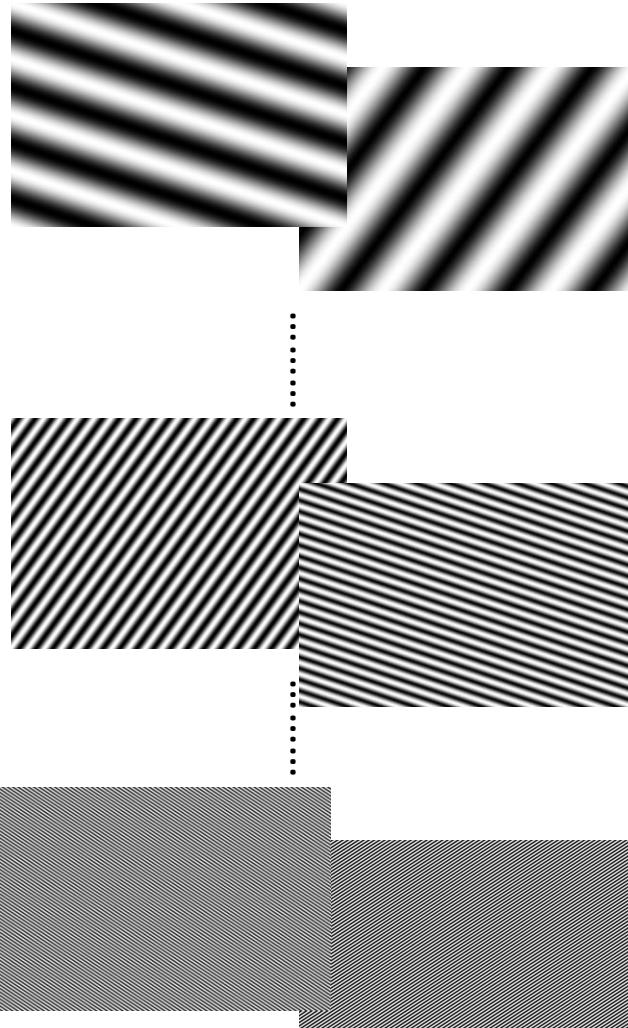


Tells you *where* things are....

... but no concept of *what* it is

## Fourier domain

Basis functions:



Tells you *what* is in the image....

... but not *where* it is

## Modulation property and Gabor filters

Modulation property:

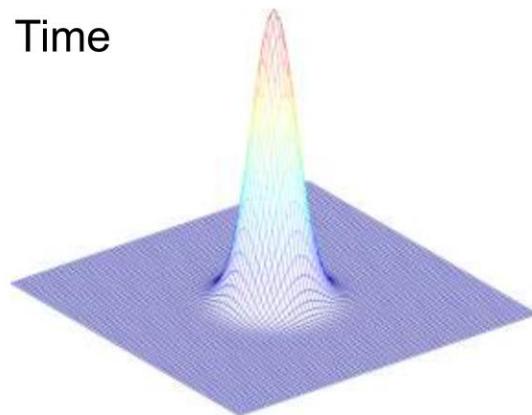
$$f(x) \leftrightarrow F(\omega)$$

$$f(x)e^{j\omega_c x} \leftrightarrow F(\omega - \omega_c)$$

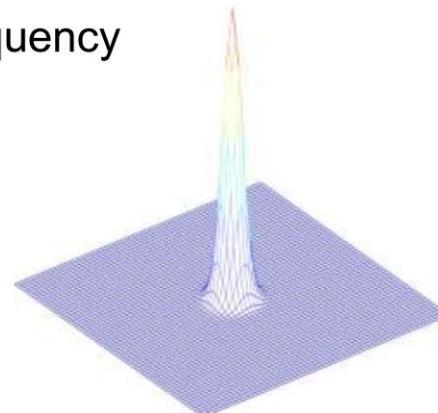
Gaussian:

$$\frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \leftrightarrow e^{-\frac{(\omega_x^2+\omega_y^2)\sigma^2}{2}}$$

Time



Frequency



## Modulation property and Gabor filters

Modulation property:

$$f(x) \leftrightarrow F(\omega)$$

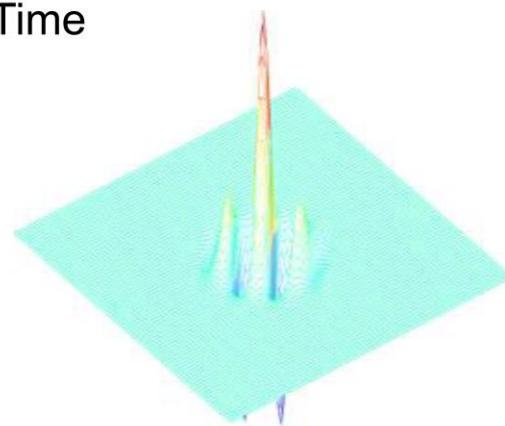
$$f(x)e^{j\omega_c x} \leftrightarrow F(\omega - \omega_c)$$

Gaussian:

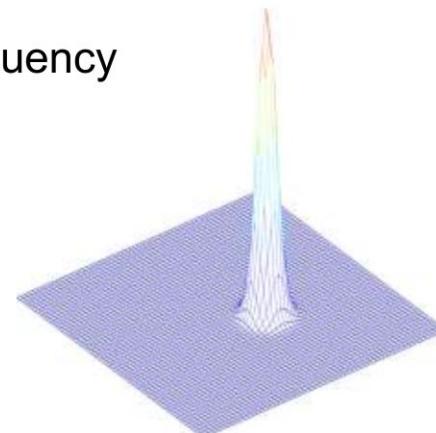
$$\frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \leftrightarrow e^{-\frac{(\omega_x^2+\omega_y^2)\sigma^2}{2}}$$

Gabor:  $\frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} e^{j(\omega_x^c x + \omega_y^c y)} \leftrightarrow e^{-j\frac{\sigma^2((\omega_x - \omega_x^c)^2 + (\omega_y - \omega_y^c)^2)}{2}}$

Time

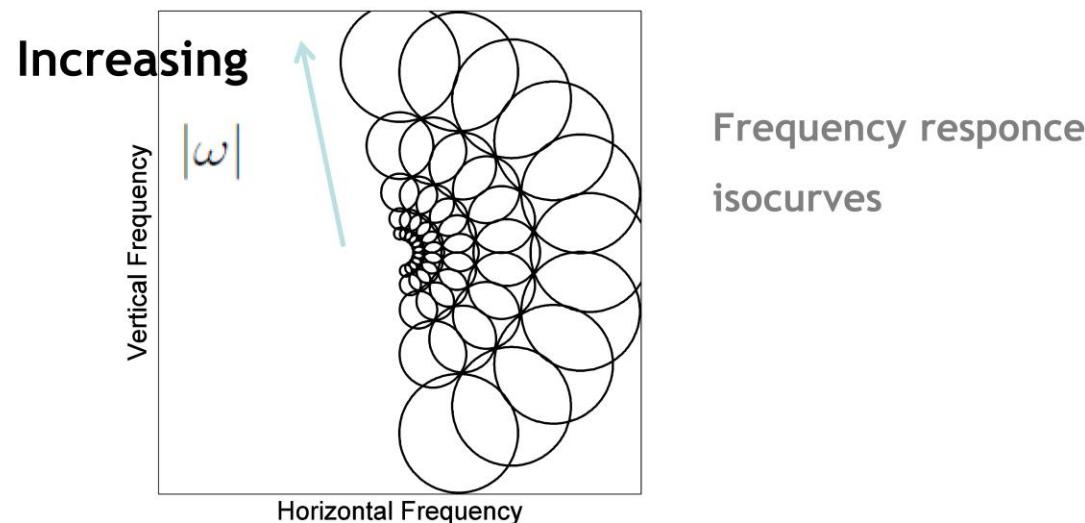
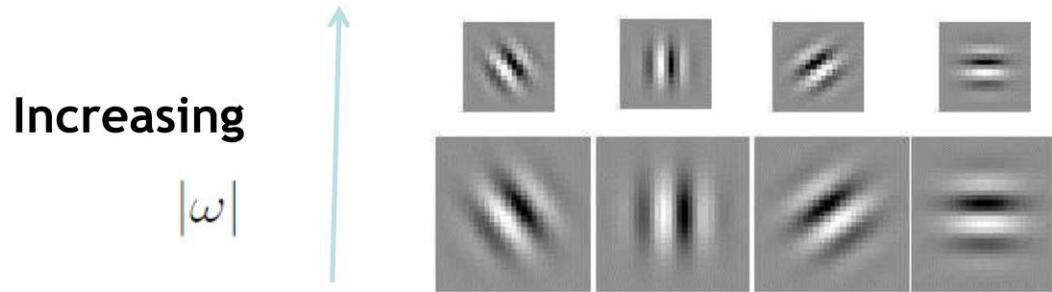


Frequency

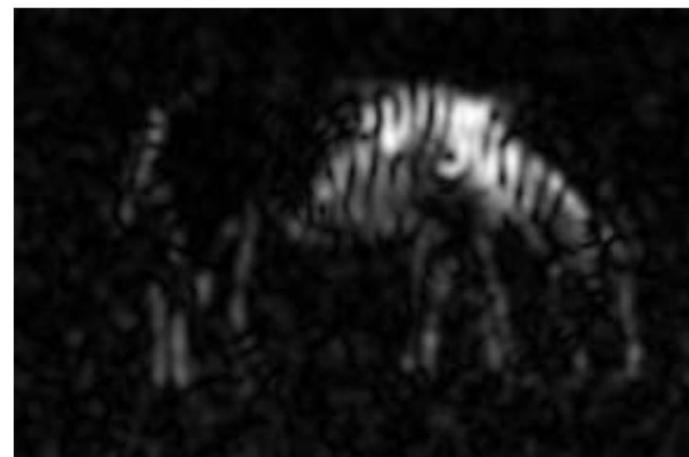
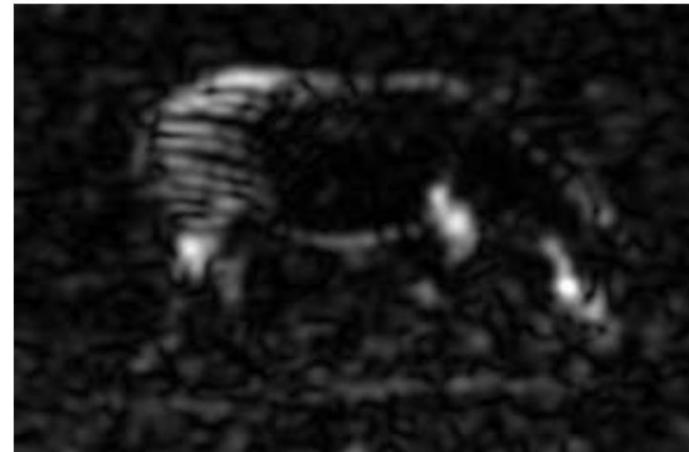
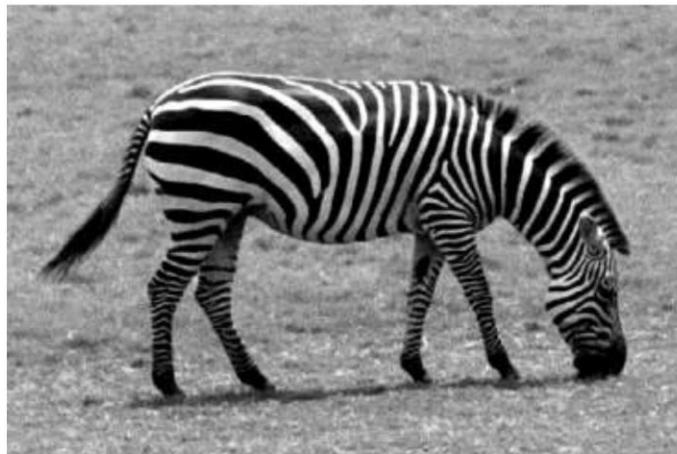


## 2D Gabor filterbank and texture analysis

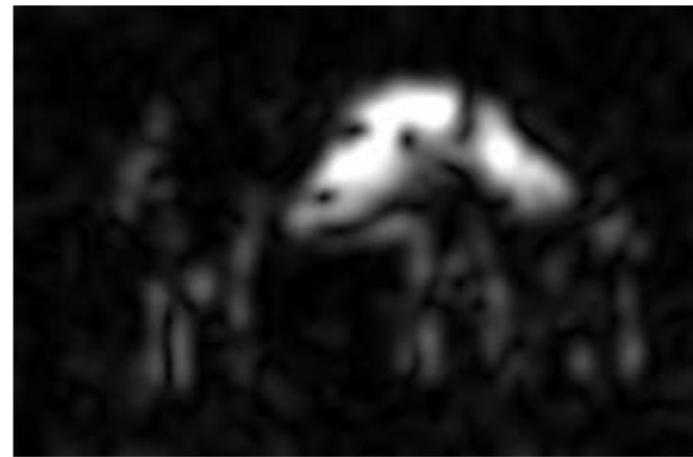
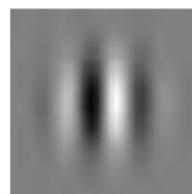
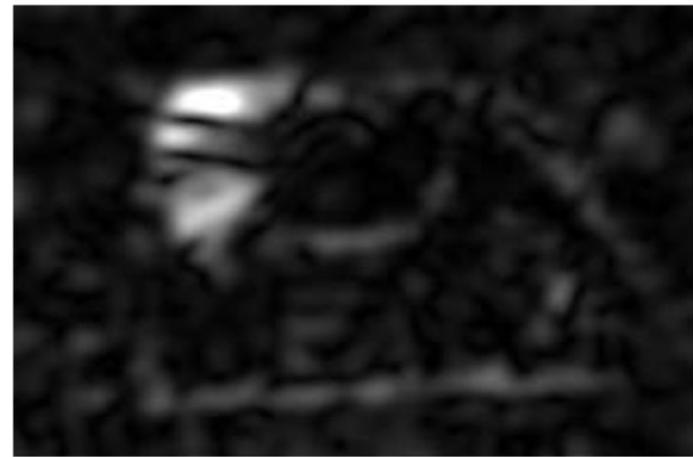
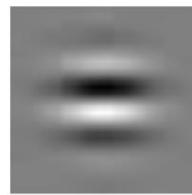
Consider many combinations of  $|\omega|$  and  $\angle\omega$



## 2D Gabor filterbank and texture analysis



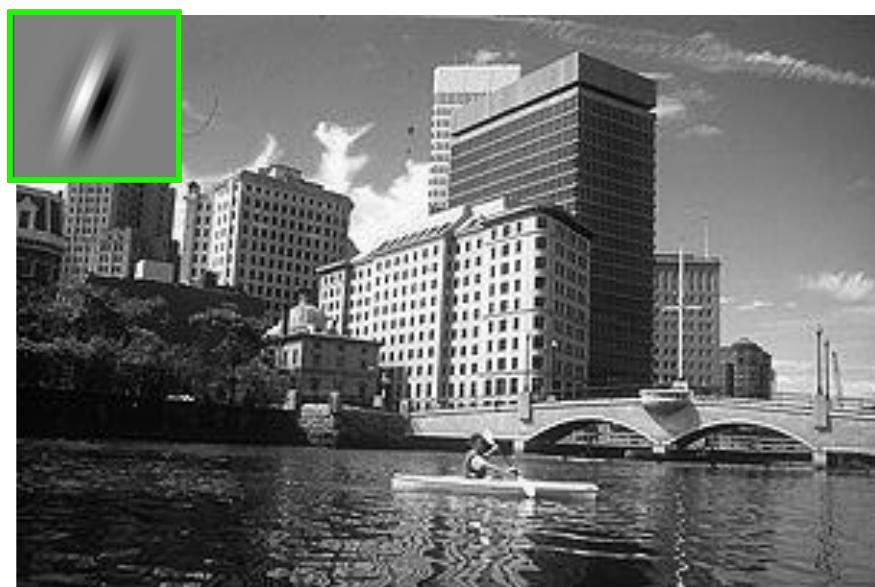
## 2D Gabor filterbank and texture analysis



## Summary: Images as a composition of local parts

### Filtering example

- Why filtering?
  - Statistics of images look similar at different locations
  - Dependencies are very local
  - Filtering is an operation with translation *equivariance*

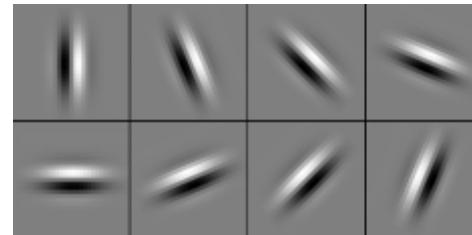


## Summary: Compare: SIFT Descriptor

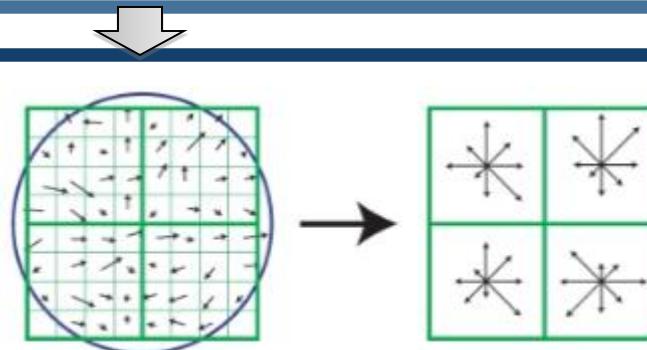
Image  
Pixels



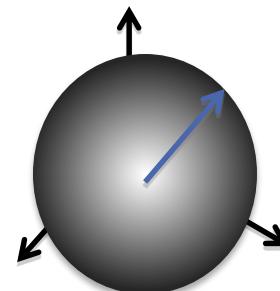
Apply  
Gabor filters



Spatial pool  
(Sum)



Normalize to unit  
length



Feature  
Vector

# Summary

- Linear filtering and the importance of **convolution**
  - Apply a filtermask to the local neighborhood at each pixel in the image
  - The filtermask defines how to combine values from neighbors.
  - Can be used for
    - Extract **intermediate representations** to abstract images by higher-level “**features**”, for further processing (i.e., preserve the useful information only and discard redundancy)
    - Image **modification**, e.g., to reduce noise, resize, increase contrast, etc.
    - Match **template** images (e.g. by correlating two image patches)
- Image filtering in the frequency domain
  - Provides a nice way to illustrate the effect of linear filtering
    - Filtering is a way to modify the frequencies of images
  - Efficient signal filtering is possible in that domain
  - The frequency domain offers an alternative way to understanding and manipulating the image.