

Image and Video Understanding

2VO 710.095 WS

Christoph Feichtenhofer, Axel Pinz

Slide credits:

Many thanks to all the great computer vision researchers on which this presentation relies on.

Most material is taken from tutorials at NIPS, CVPR and BMVC conferences.

Outline

- Convolutional Networks (ConvNets) for Image Classification
 - Operations in each layer
 - Architecture
 - Visualizations
 - Results
- Krizhevsky, A., Sutskever, I. and Hinton, G. E.,
ImageNet Classification with Deep Convolutional Neural Networks, NIPS 2012
- M. Zeiler & R. Fergus, *Visualizing and Understanding Convolutional Networks*, ECCV, 2014
- Representations for Video Classification
 - Hand-designed features
 - Spatiotemporal ConvNets
 - Two-stream ConvNets
- Wang et al., *Action Recognition by Dense Trajectories*, CVPR 2011.
- Karpathy et al., *Large-scale Video Classification with Convolutional Neural Networks*, CVPR 2014
- K. Simonyan & A. Zisserman, *Two-Stream Convolutional Networks for Action Recognition in Videos*, NIPS 2014

Goal: Scene Understanding

Car

Person

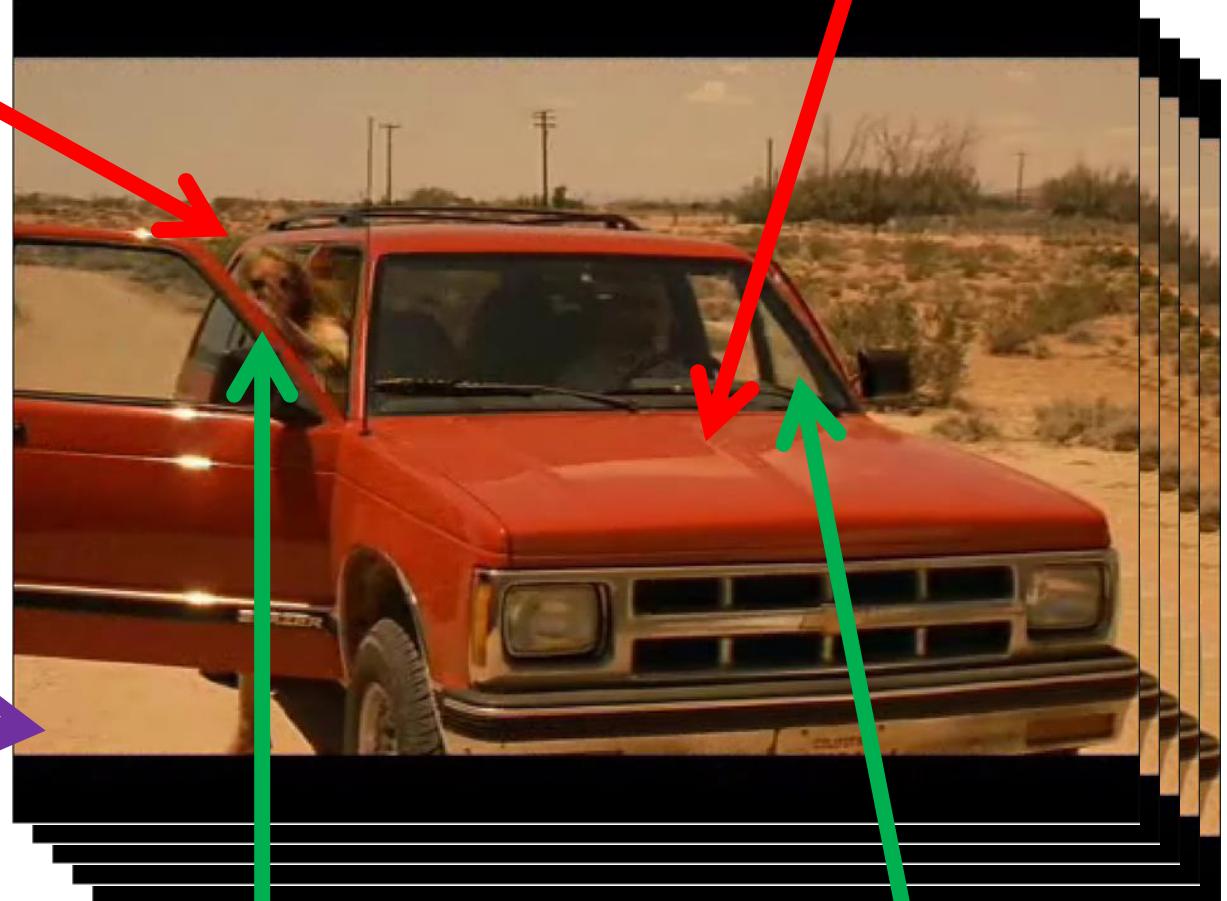
Scenes Objects

Actions

Activities

Desert

Leave a car in the desert

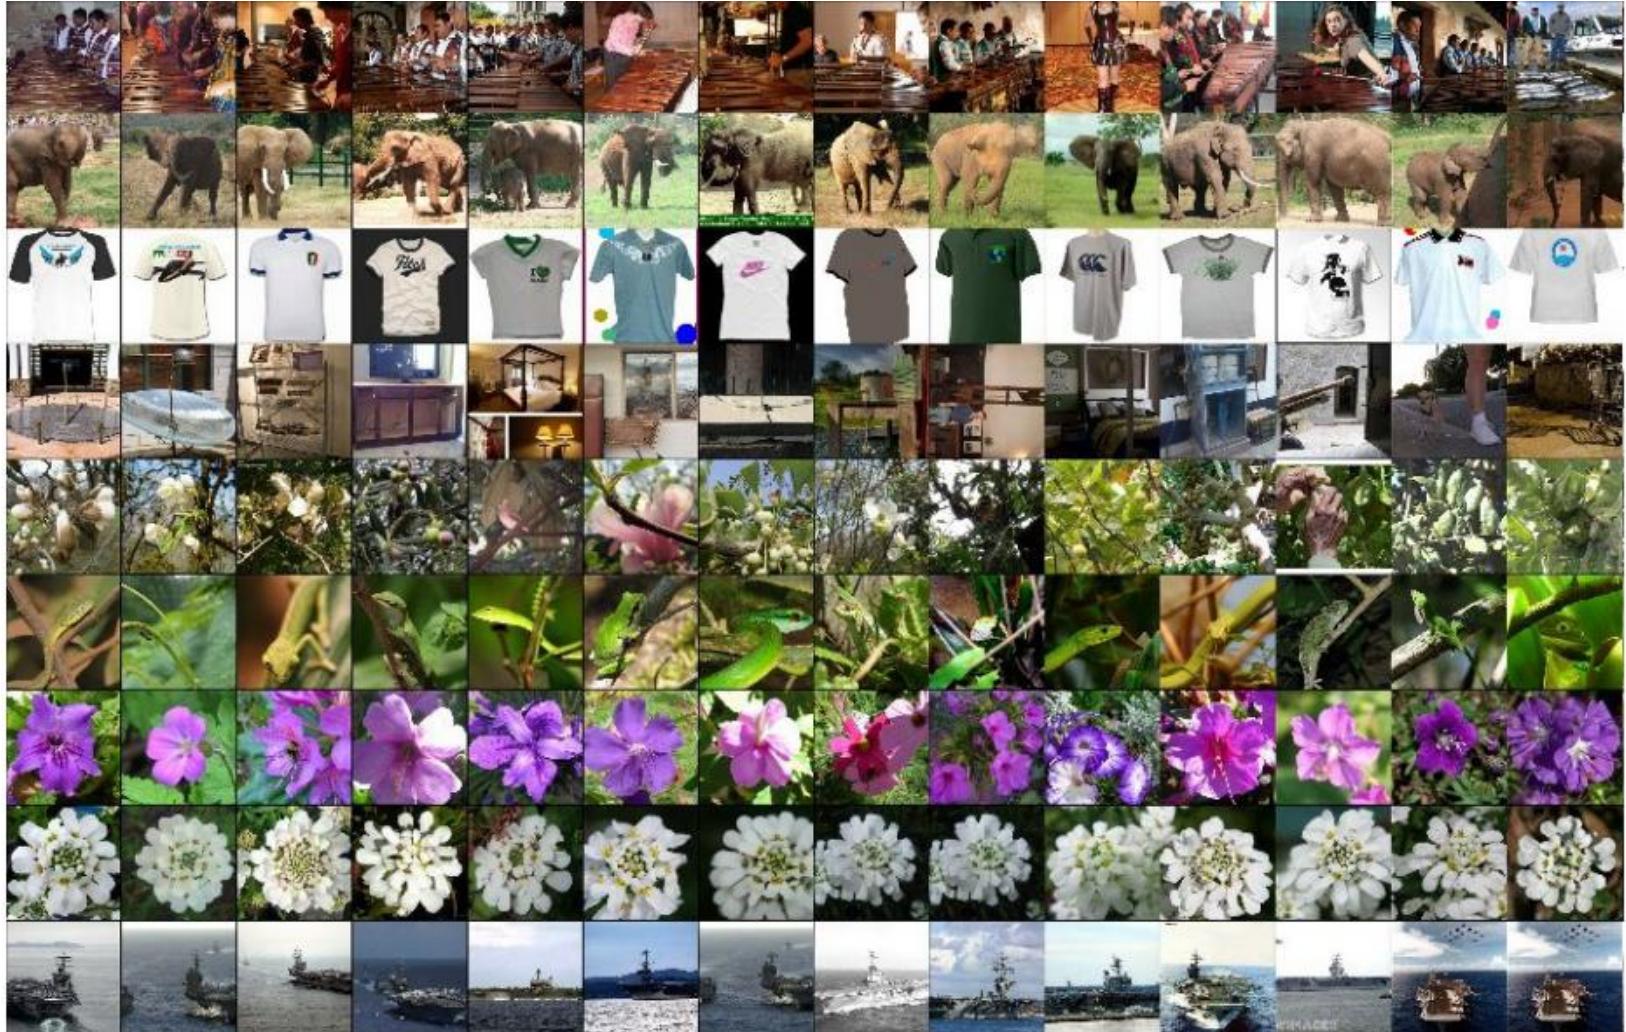


Temporal input sequence

Get out of car

Open door

One application: Image retrieval



Deep Learning - breakthrough in visual and speech recognition



Credit: B. Ginzburg

A lot of buzz about Deep Learning



[Microsoft On Deep Learning for Speech](#) goto 3:00-5:10

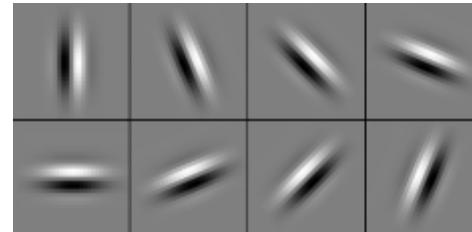
Credit: B. Ginzburg

Summary: Compare: SIFT Descriptor

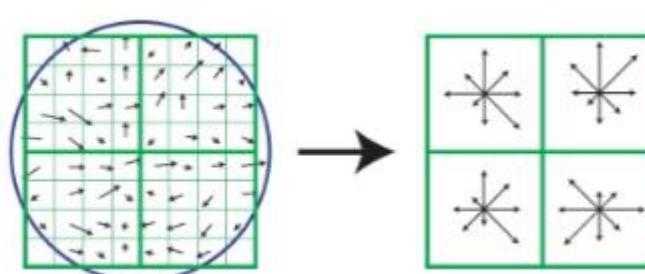
Image
Pixels



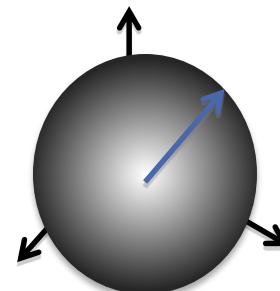
Apply
Gabor filters



Spatial pool
(Sum)



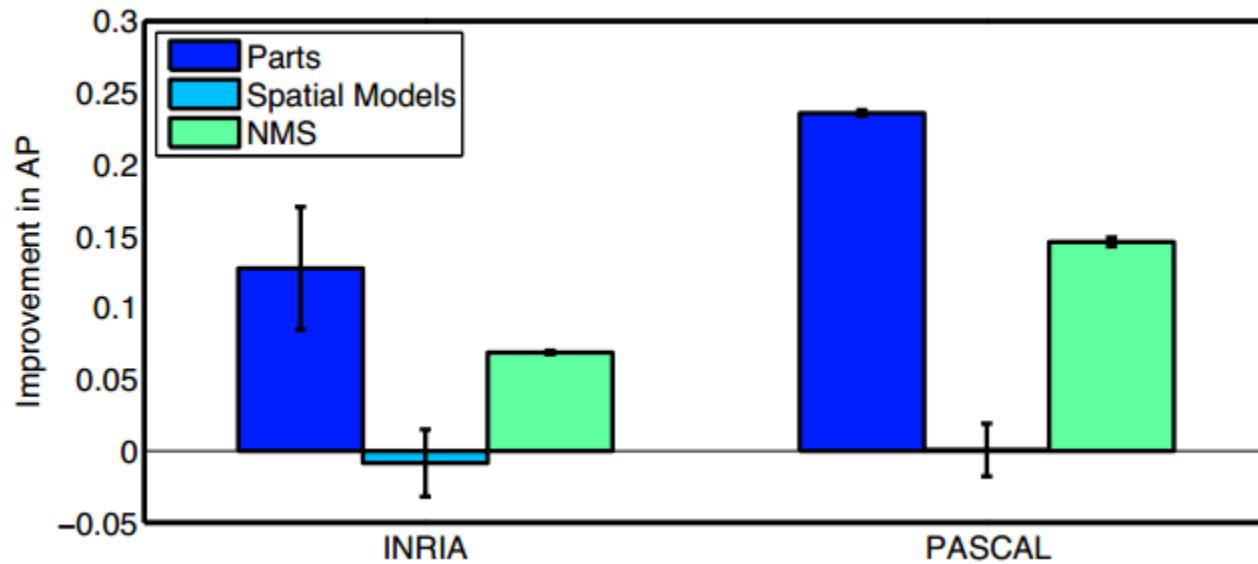
Normalize to unit
length



Feature
Vector

What are the weakest links limiting performance?

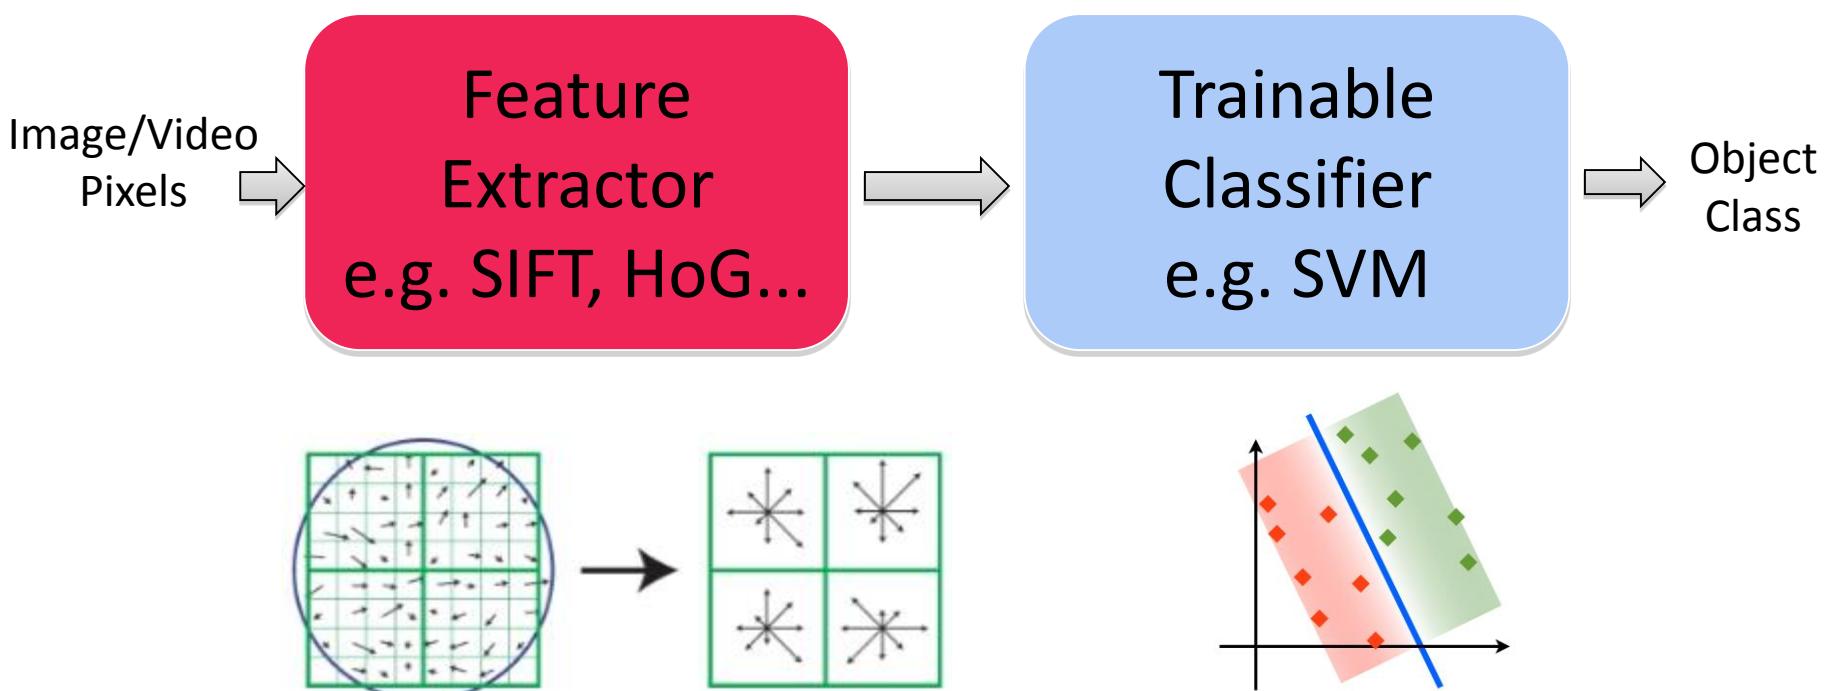
- Replace each component of the deformable part model detector with humans
- Good Features (part detection) and accurate Localization (NMS) are most important



[Parikh & Zitnick CVPR'10]

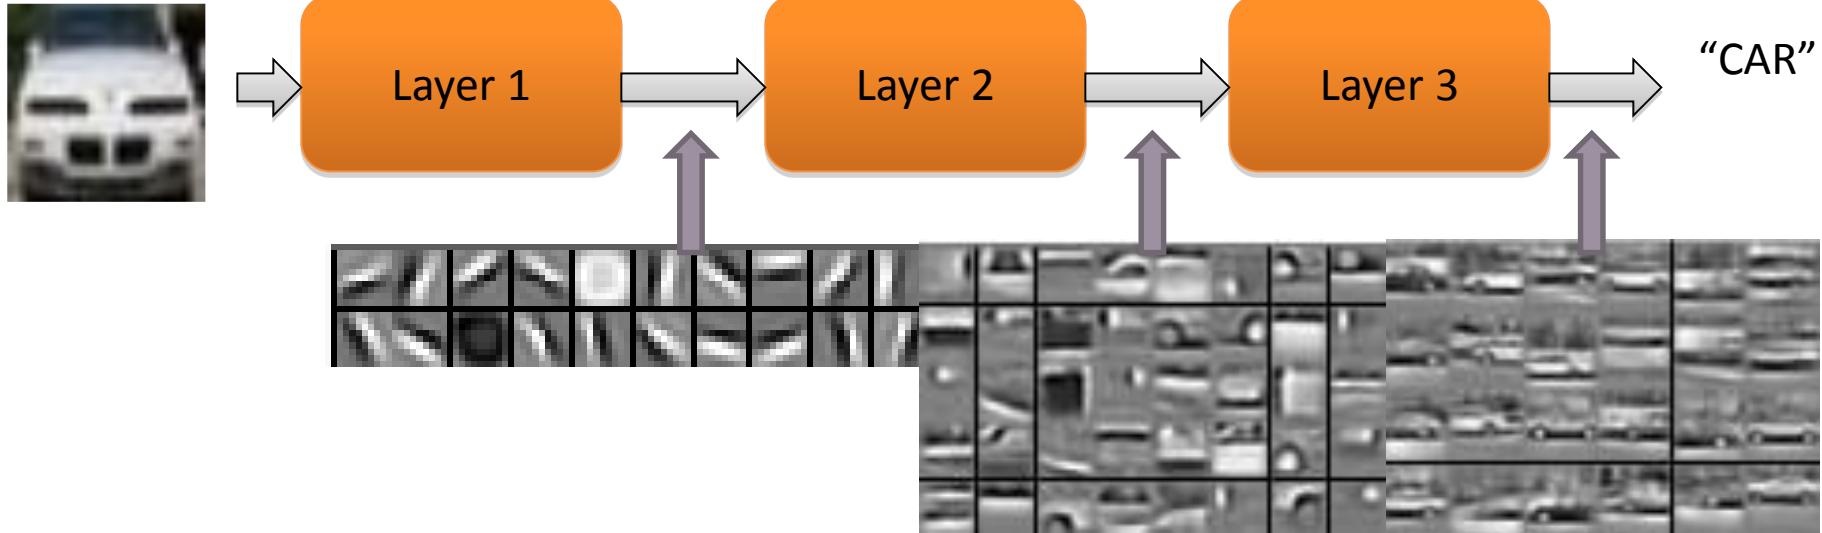
Typical visual recognition pipeline

- Select / develop features
- Add on top of this Machine Learning for multi-class recognition and train classifier



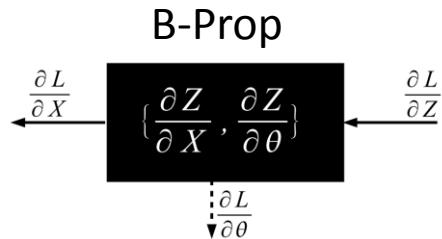
Intuition Behind Deep Neural Nets

- Build features automatically based on training data
- Combine feature extraction and classification
- All the way from pixels → classifier
 - Each box is a feature detector
- One layer extracts features from output of previous layer



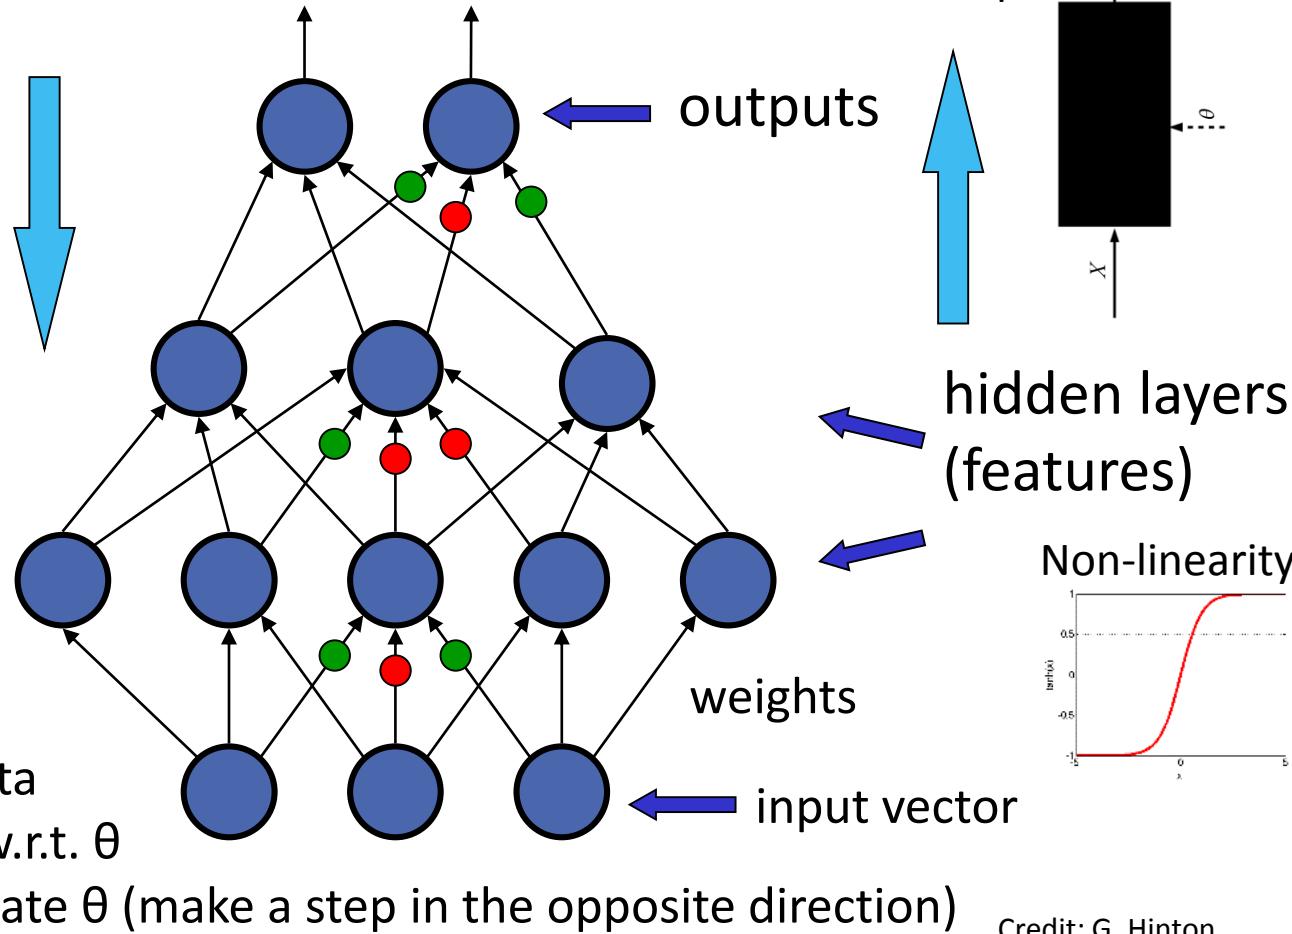
Some Key Ingredients for Convolutional Neural Networks

Neural networks trained via backpropagation



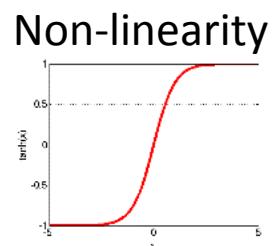
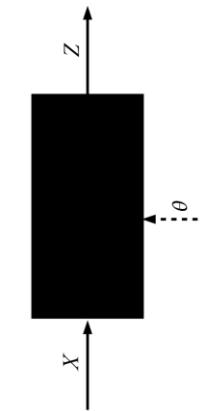
Back-propagate
error signal to get
derivatives for learning
parameters θ

Compare outputs with
correct answer to get
error signal L



Training

- F-Prop / B-Prop
- Learning by stochastic gradient descent (SGD):
 - A) Compute loss L on small mini-batch of data
 - B) Compute gradient w.r.t. θ
 - C) Use gradient to update θ (make a step in the opposite direction)

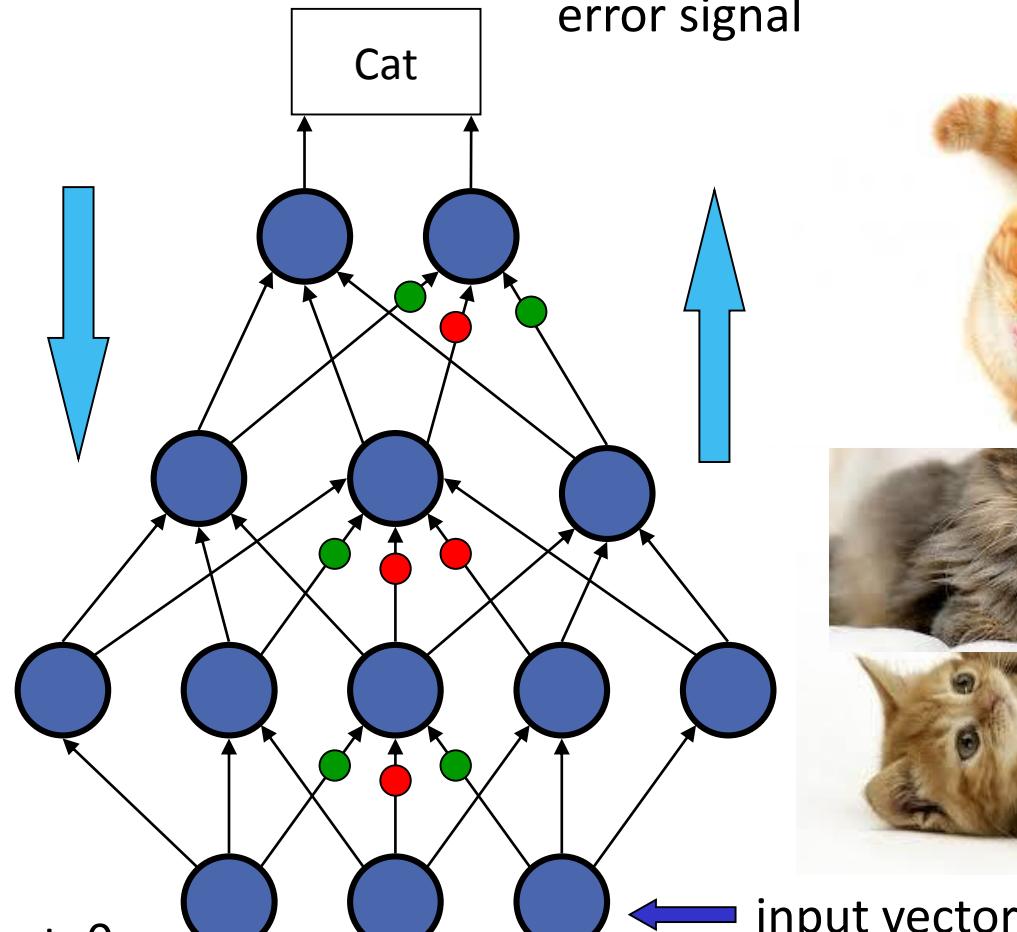


Credit: G. Hinton

Neural networks trained via backpropagation

Back-propagate
error signal to get
derivatives for learning
parameters θ

Compare outputs with
correct answer to get
error signal



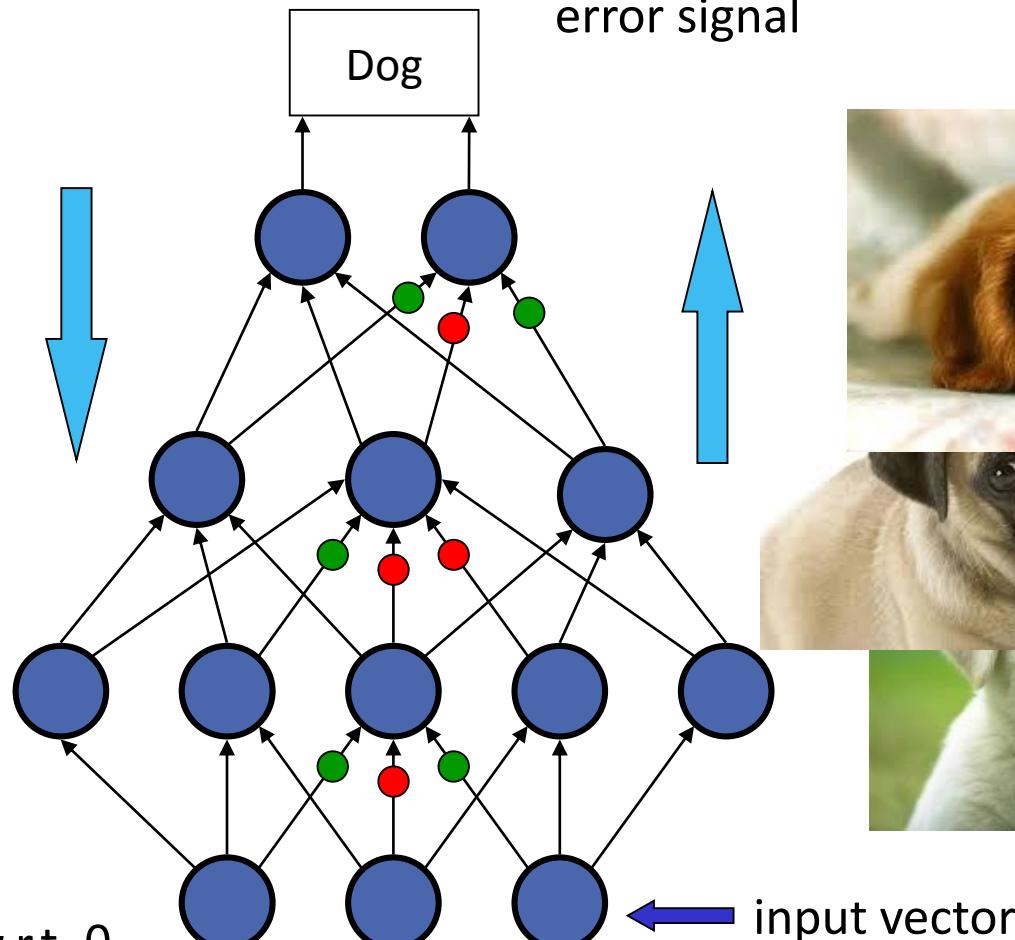
Training

- F-Prop / B-Prop
- Learning by SGD:
- A) Compute loss on small mini-batch
- B) Compute gradient w.r.t. θ
- C) Use gradient to update θ

Neural networks trained via backpropagation

Back-propagate
error signal to get
derivatives for learning
parameters θ

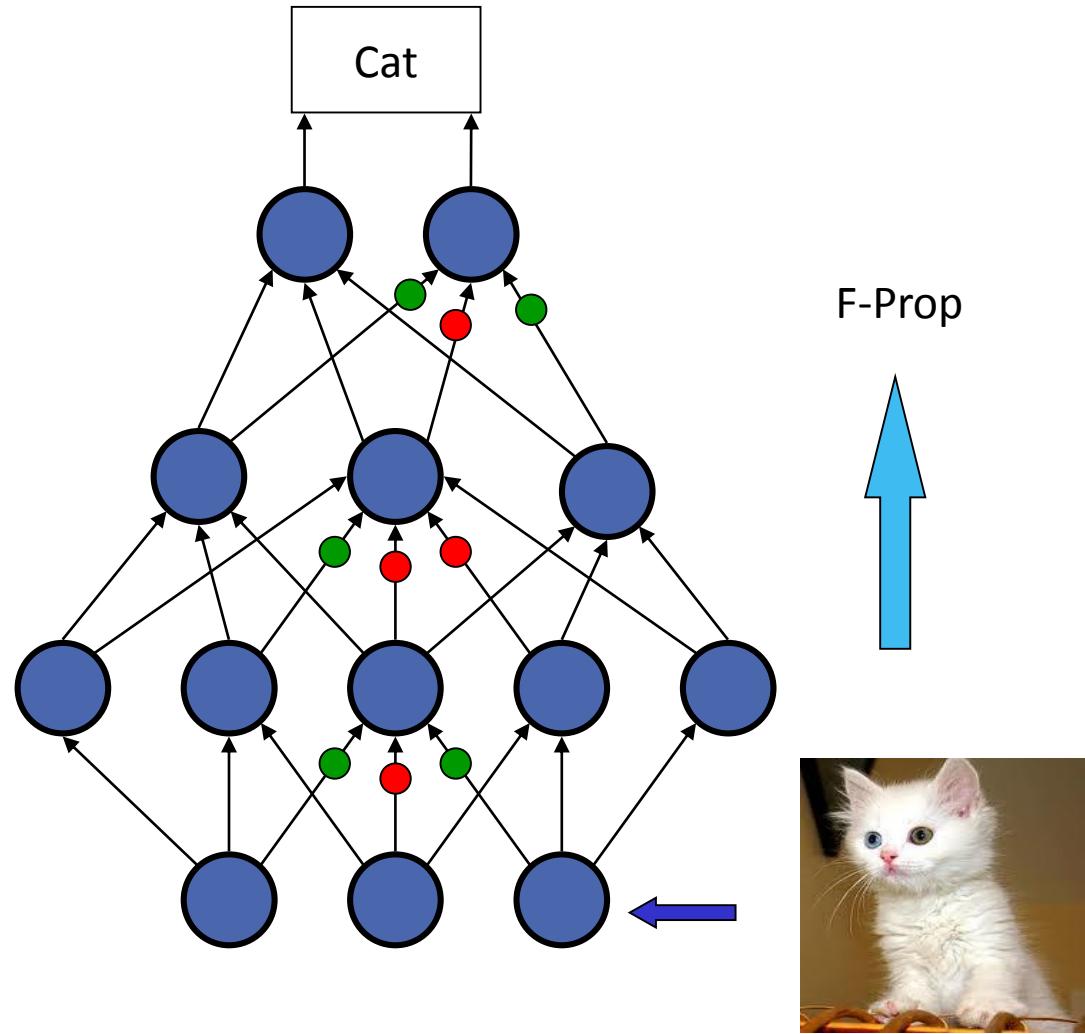
Compare outputs with
correct answer to get
error signal



Training

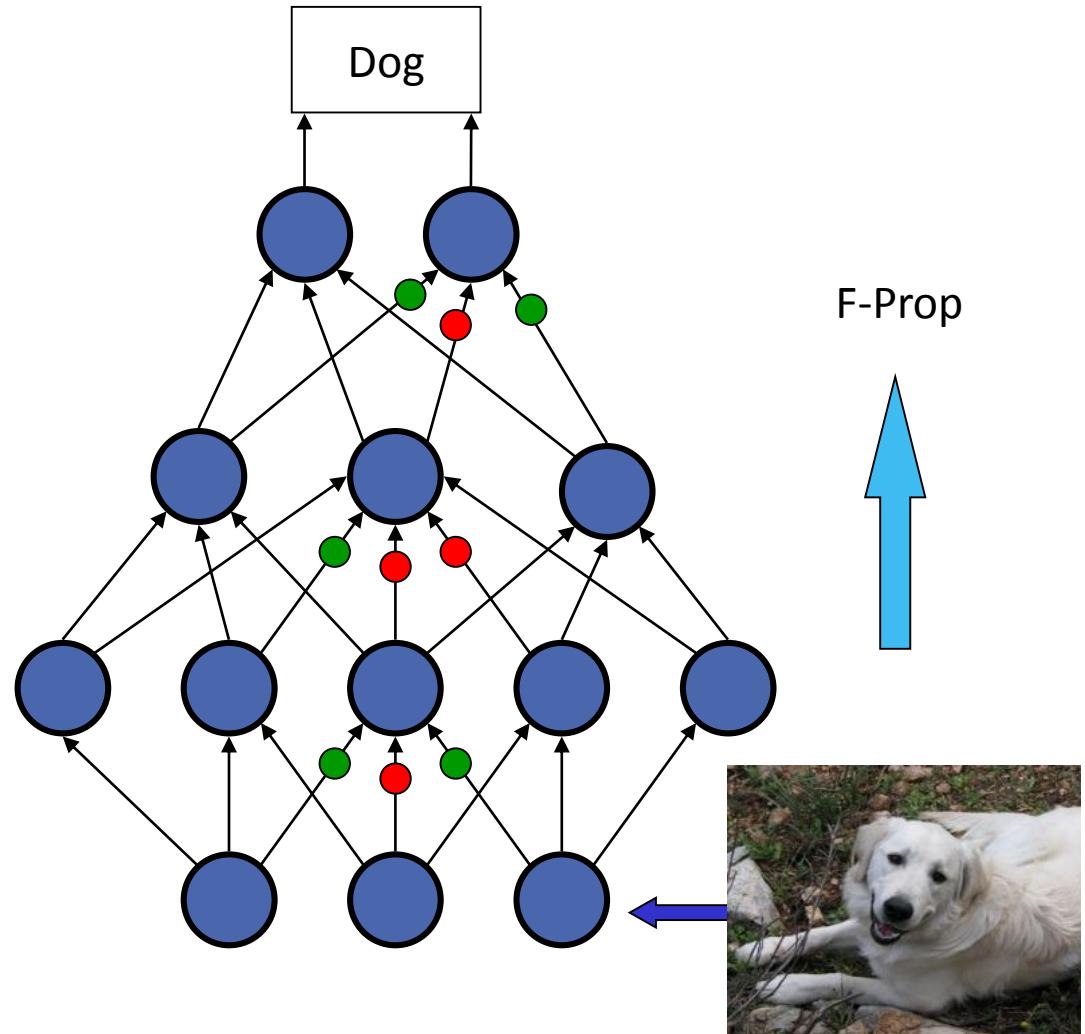
- F-Prop / B-Prop
- Learning by SGD:
- A) Compute loss on small mini-batch
- B) Compute gradient w.r.t. θ
- C) Use gradient to update θ

Neural networks testing



Credit: G. Hinton

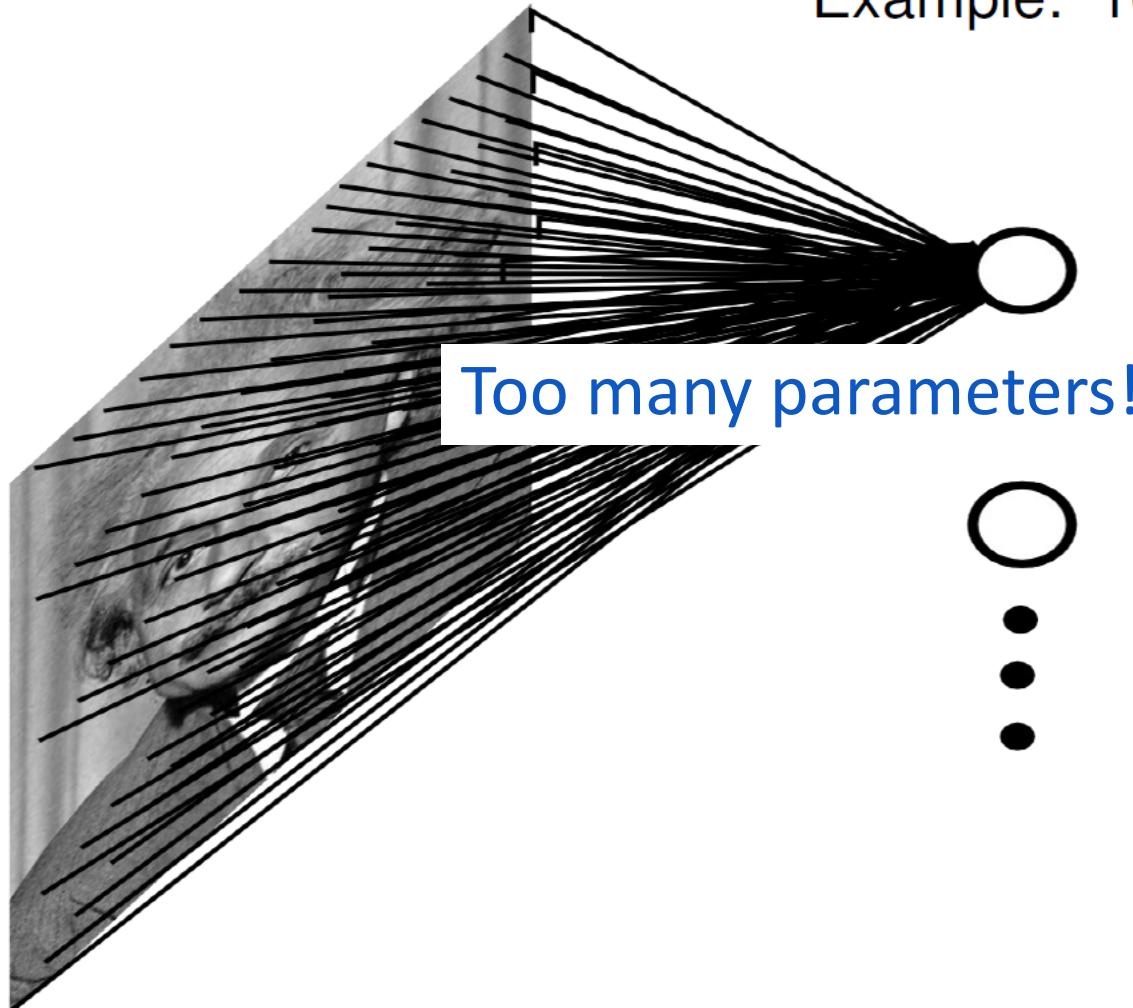
Neural networks testing



Credit: G. Hinton

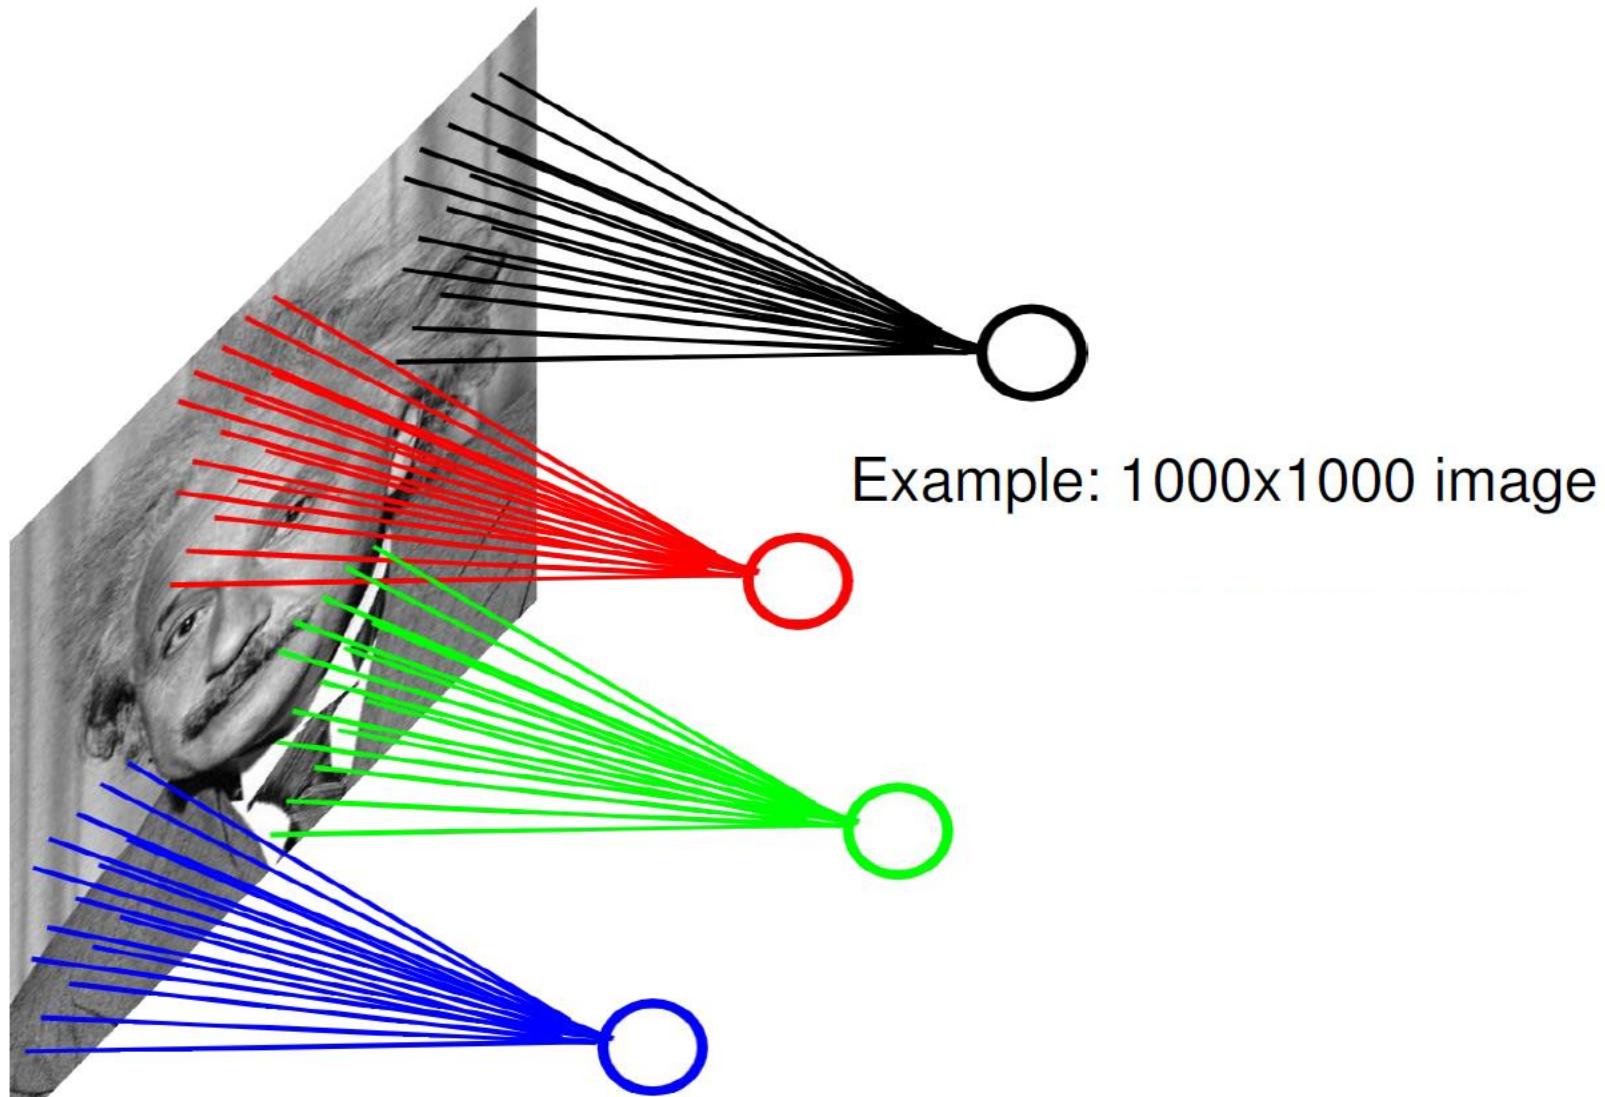
Motivation: Images as a composition of local parts
“Pixel-based” representation

Example: 1000x1000 image



Motivation: Images as a composition of local parts

“Patch-based” representation



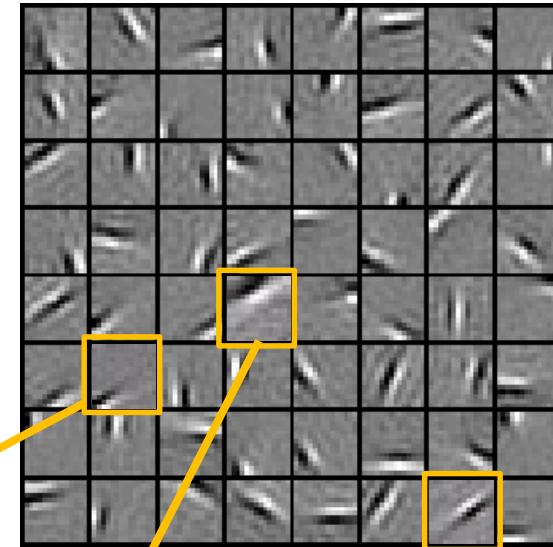
Motivation: Images as a composition of local parts

Sparse coding example

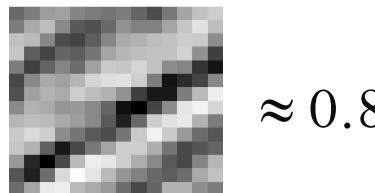
Natural Images



Learned bases (ϕ_1, \dots, ϕ_{64}): "Edges"



Test example

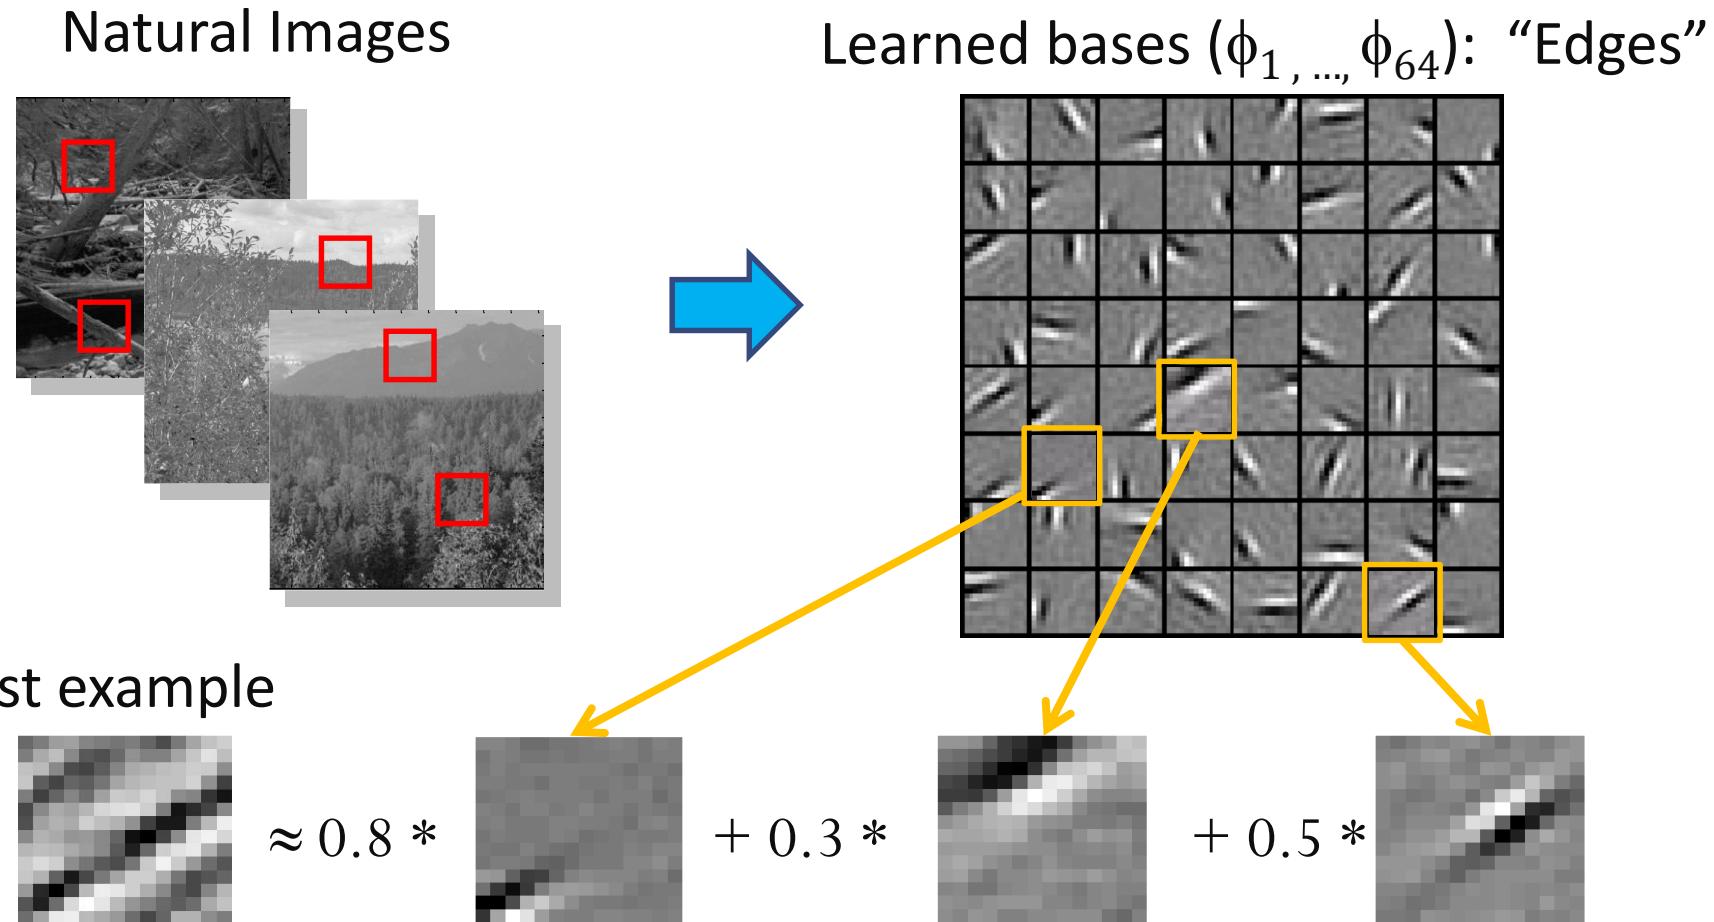


$$x \approx 0.8 * \phi_{36} + 0.3 * \phi_{42} + 0.5 * \phi_{63}$$

$[a_1, \dots, a_{64}] = [0, 0, \dots, 0, \mathbf{0.8}, 0, \dots, 0, \mathbf{0.3}, 0, \dots, 0, \mathbf{0.5}, 0]$
(feature representation)

Motivation: Images as a composition of local parts

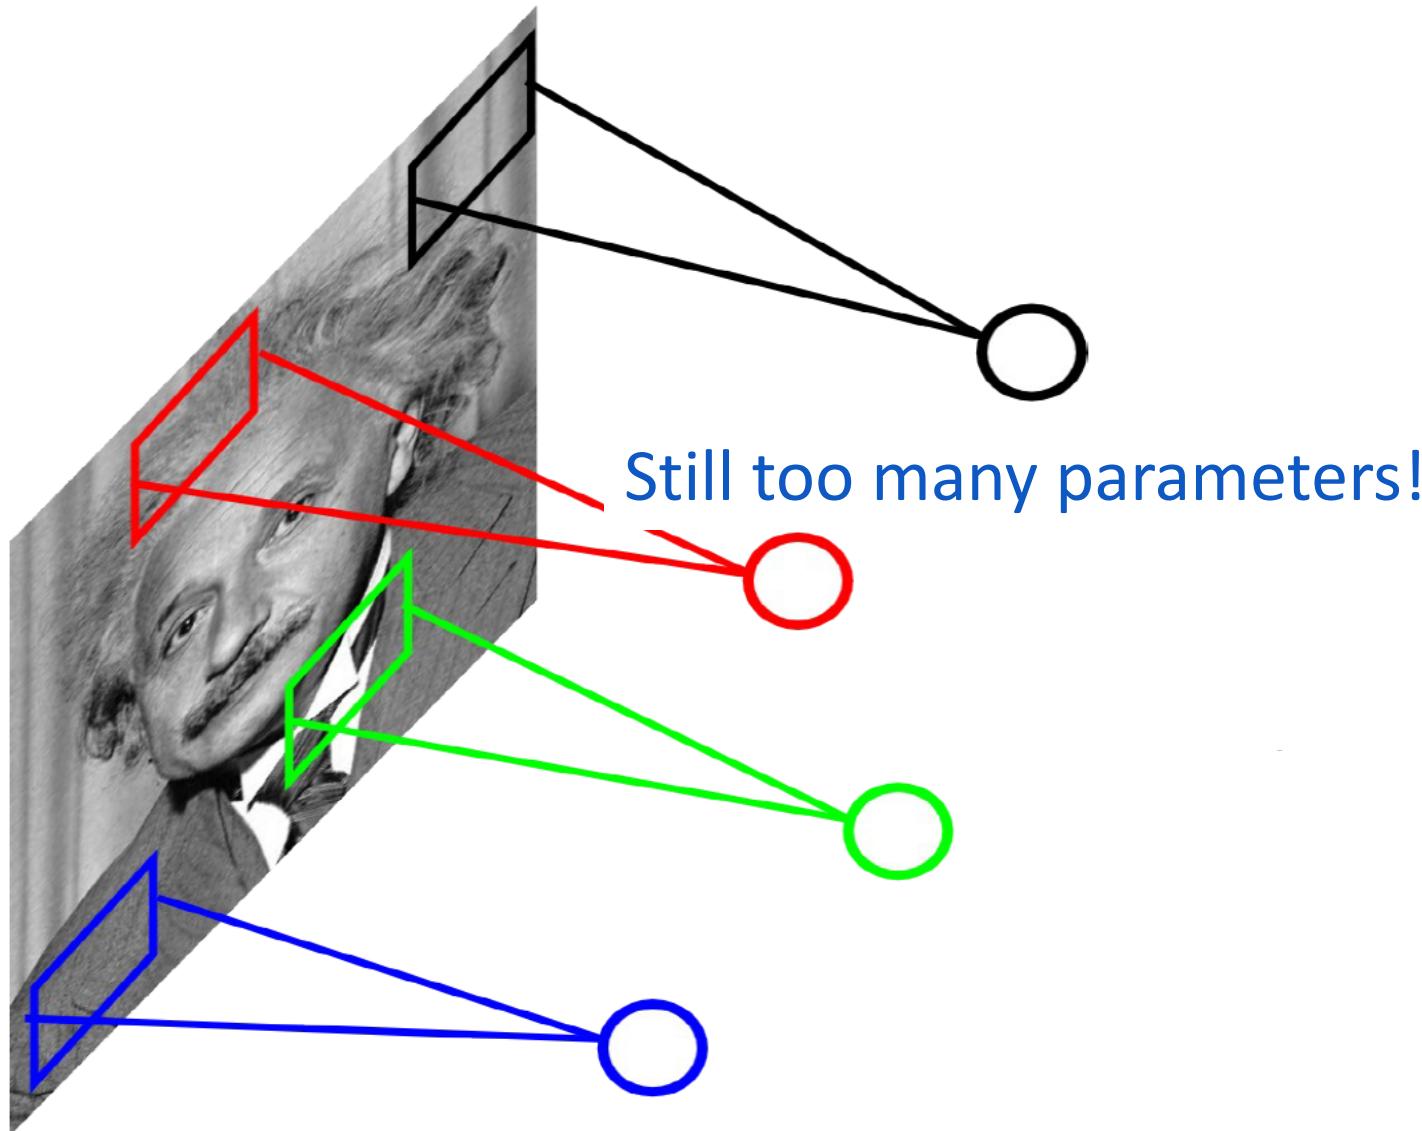
Sparse coding example



- Method “invents” edge detection
- Automatically learns to represent an image in terms of the edges that appear in it
- Gives a more succinct, higher-level representation than the raw pixels
- Quantitatively similar to primary visual cortex (area V1) in brain

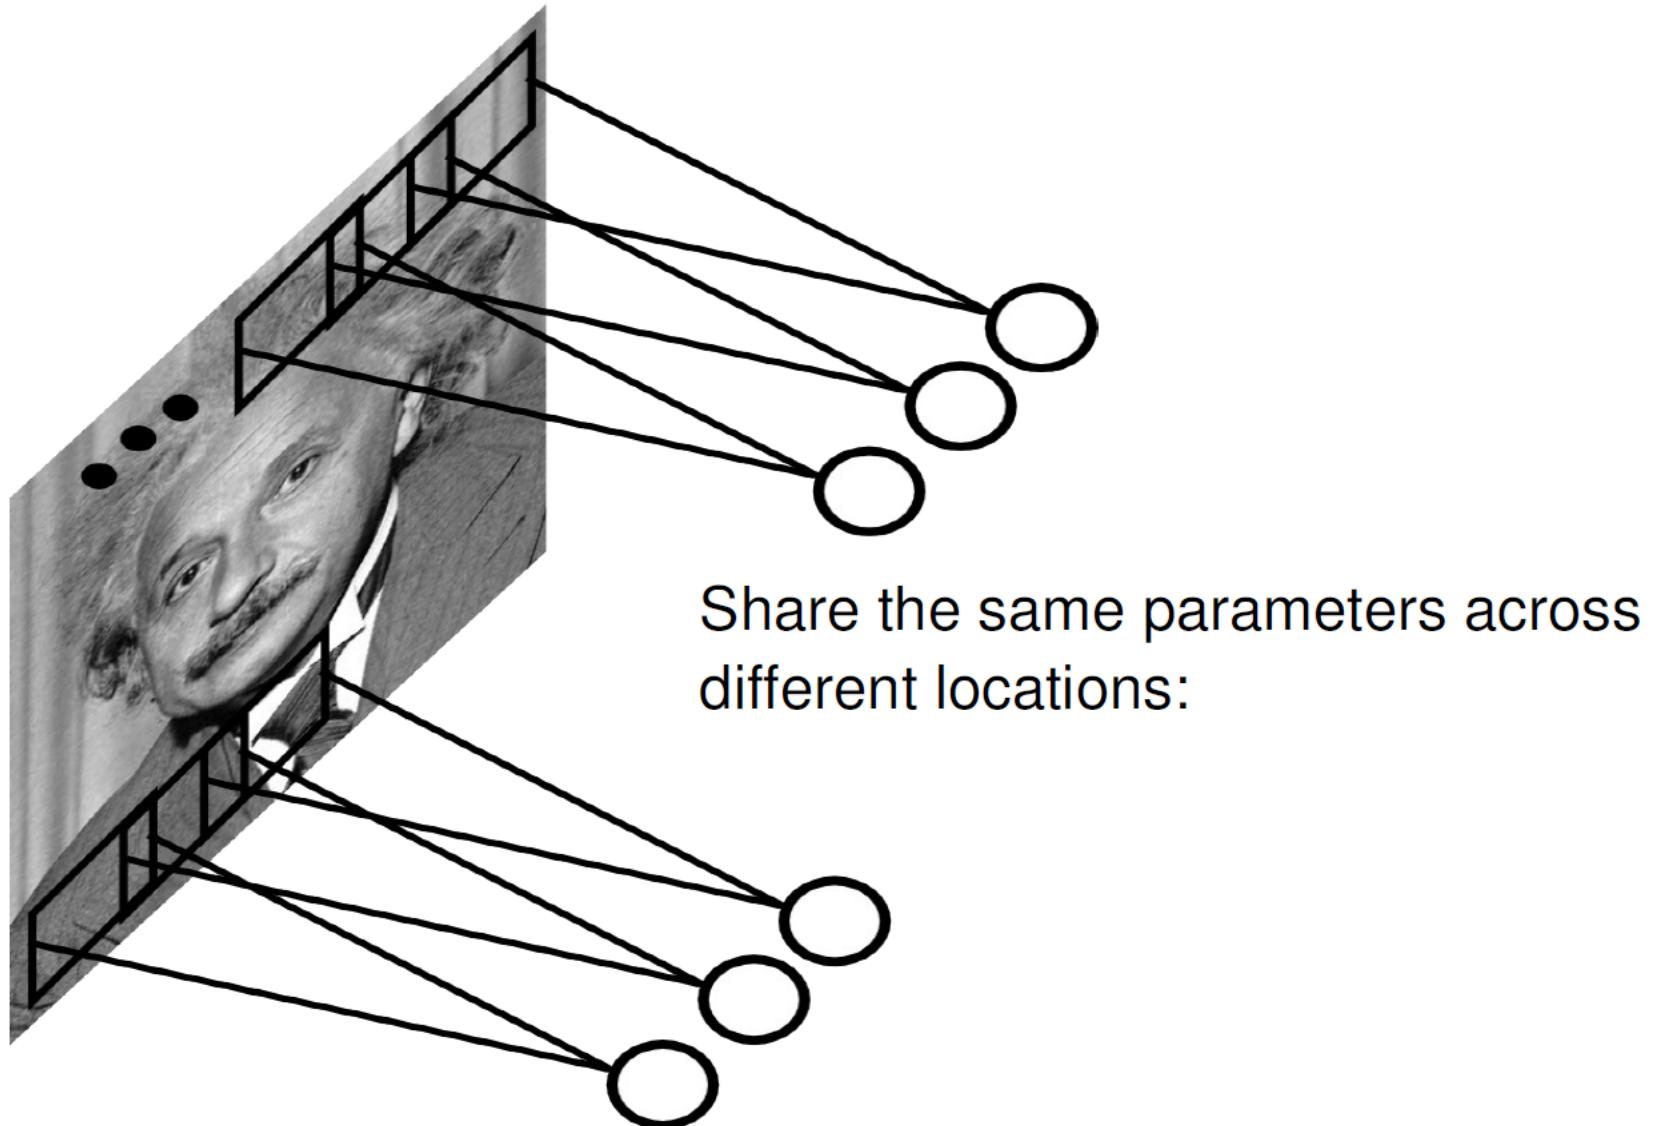
Motivation: Images as a composition of local parts

“Patch-based” representation



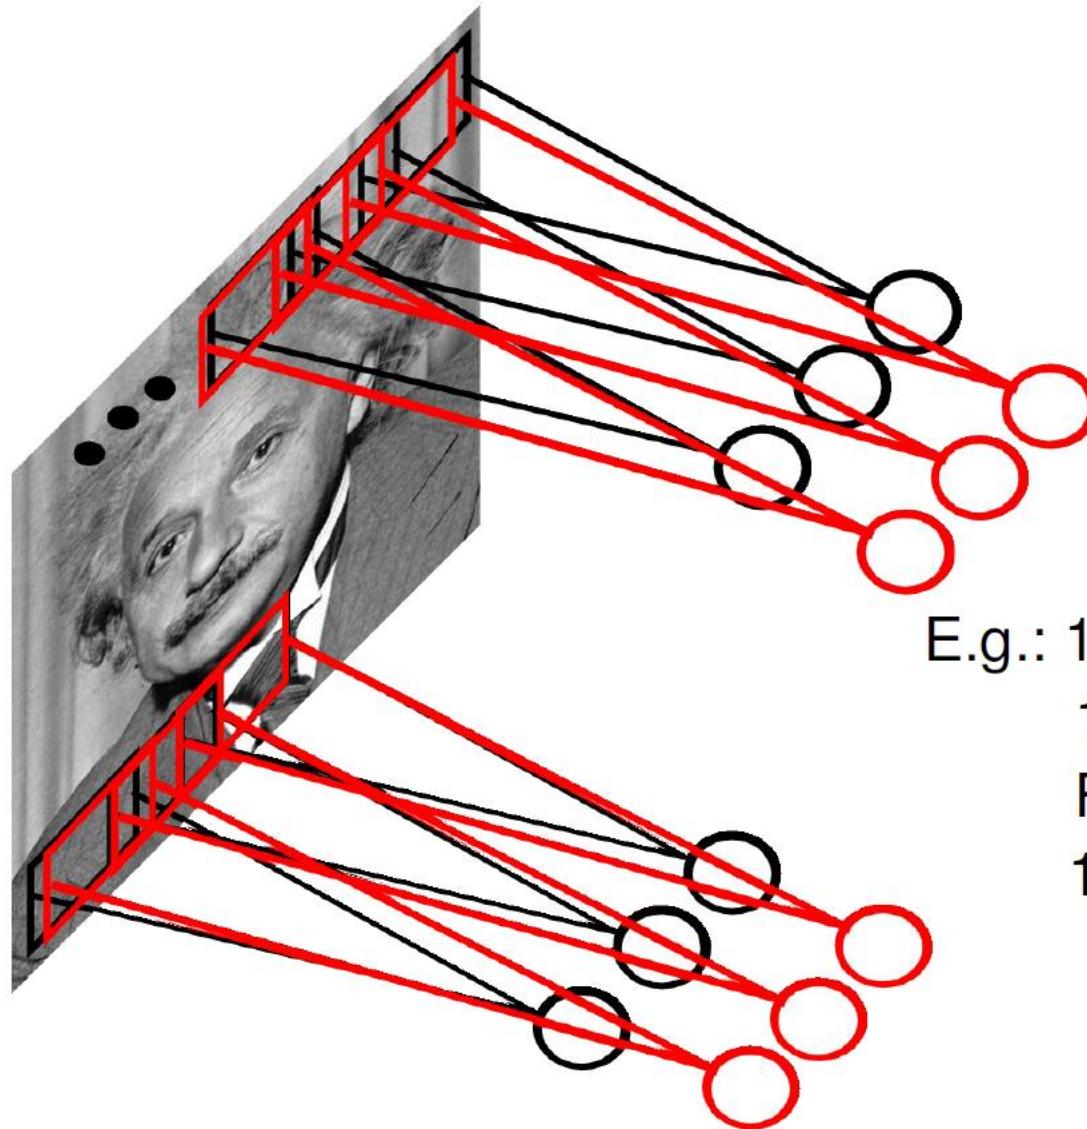
Motivation: Images as a composition of local parts

Convolution example



Motivation: Images as a composition of local parts

Convolution example

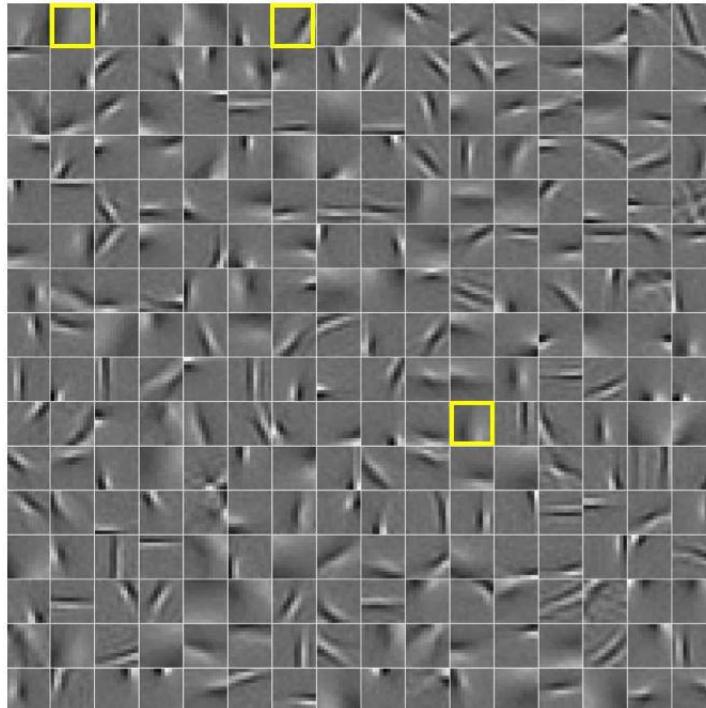


E.g.: 1000x1000 image
100 Filters
Filter size: 10x10
10K parameters

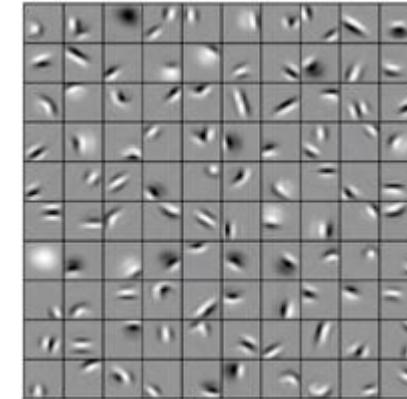
Motivation: Images as a composition of local parts

Filtering example

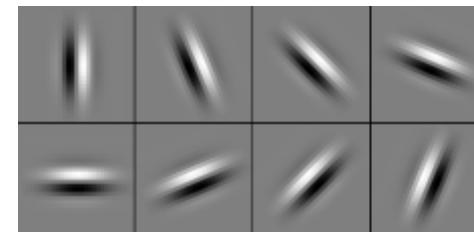
- Why translation *equivariance*?
 - Input translation leads to a translation of features
 - Fewer filters needed: no translated replications
 - But still need to cover orientation/frequency



Patch-based



Patch-based

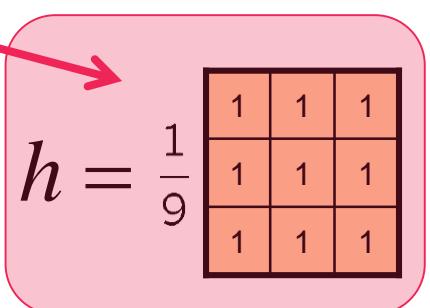


Convolutional

3x3 box average filter to blur an image (e.g., to remove noise)

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9}$$


1	1	1
1	1	1
1	1	1

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0

$g =$

Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0						

Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0	10					

Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0	10	20				

Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0	10	20	30	30	30	30

Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$g =$

			0	10	20	30	30	30	

Convolution: The 2D case

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$

$$h = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

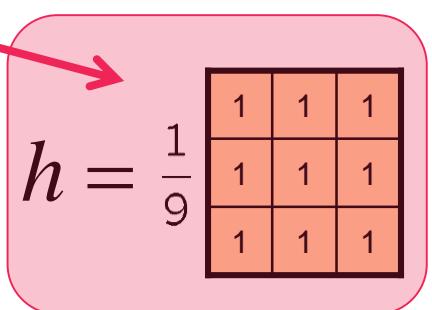
$g =$

			0	10	20	30	30	30	

Think of the filter as a feature detector now
(e.g., how smooth is a region?)

Numerical calculation: Smoothing via local averaging

$$g(x, y) = \sum_{k,l} f(x-k, y-l)h(k, l)$$



$h = \frac{1}{9}$

1	1	1
1	1	1
1	1	1

$f =$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0

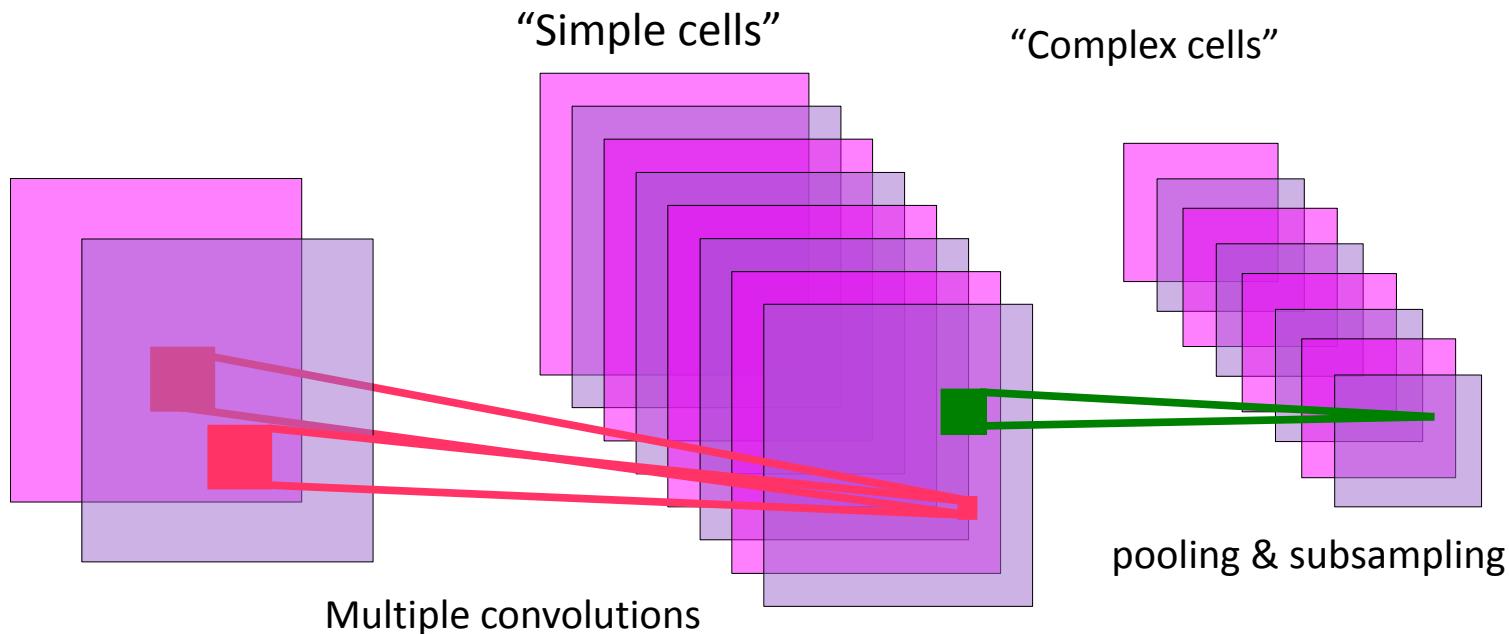
$g =$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Convolutional Neural Networks

Multistage HubelWiesel Architecture: An Old Idea for Local Shift Invariance

- [Hubel & Wiesel 1962]
 - Simple cells detect local features
 - Complex cells “pool” the outputs of simple cells within a retinotopic neighborhood.



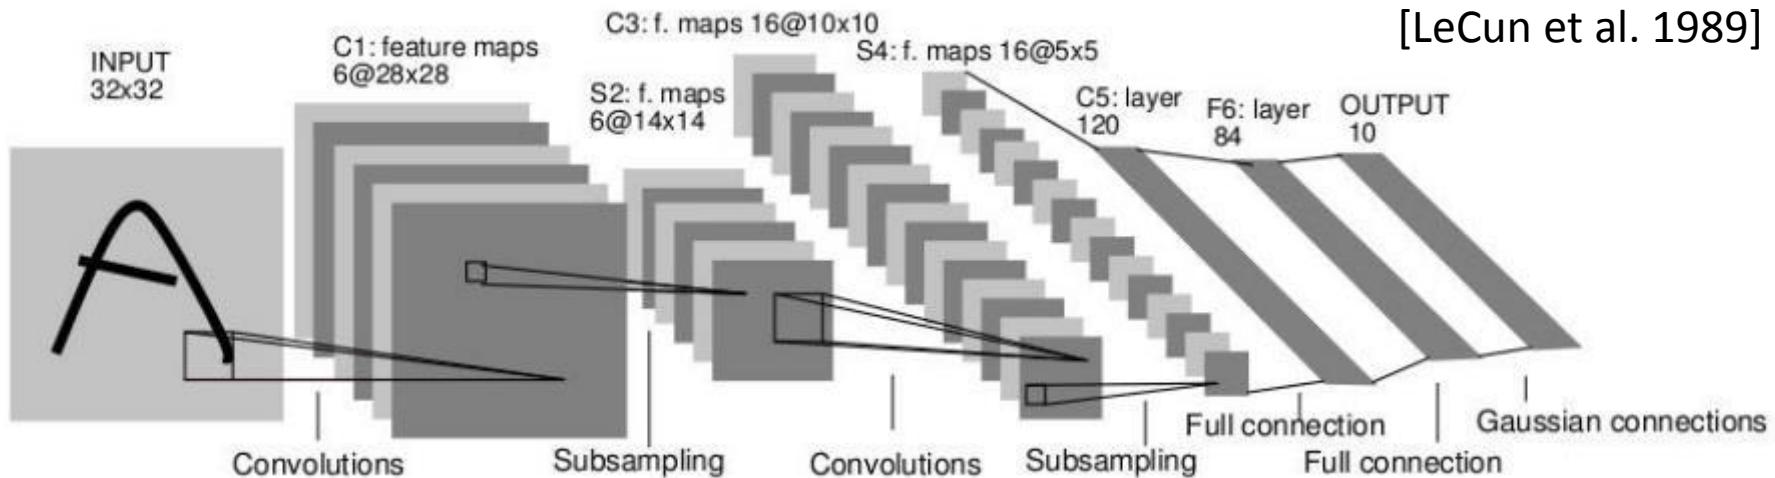
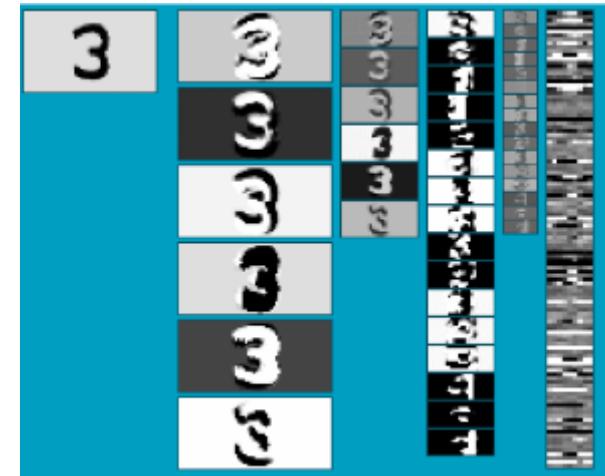
Convolutional Networks
[LeCun 1988-present]

Retinotopic Feature Maps

Credit: Y. LeCun

Convolutional Neural Networks

- Neural network with specialized connectivity structure
- After a few convolution and subsampling stages, spatial resolution is very small
- Use fully connected layers up to classification at output layer



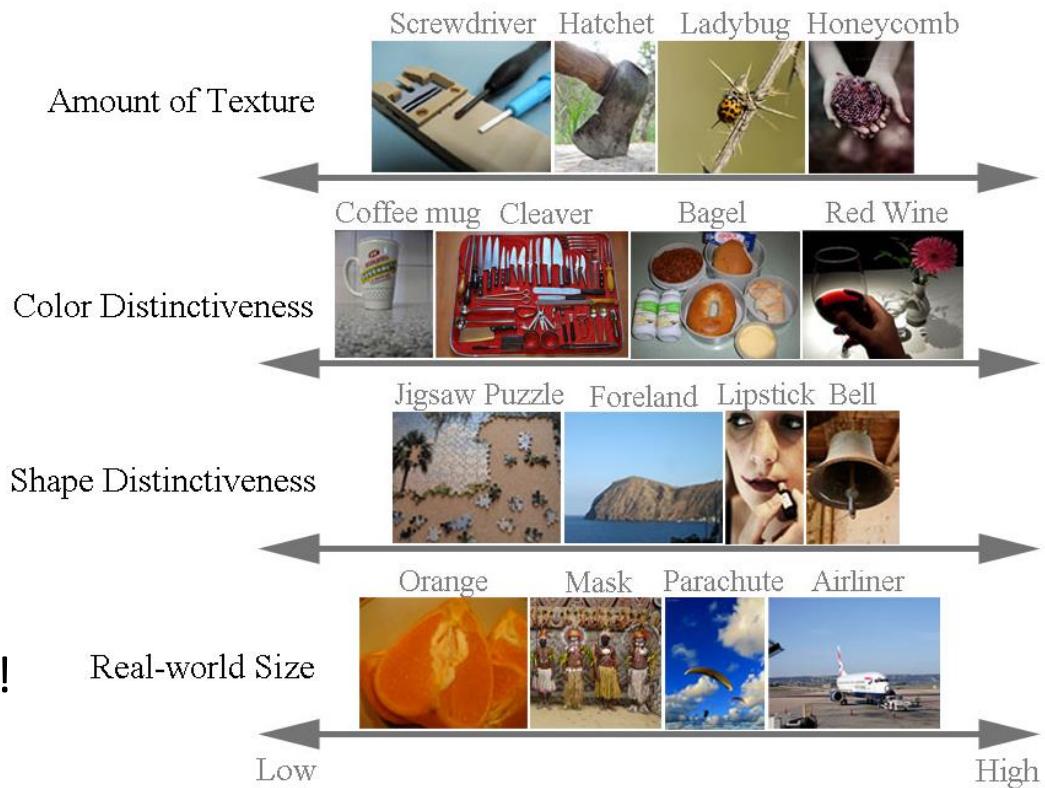
[LeCun et al. 1989]

Training large ConvNets on



Key ingredients for CNNs

- **Large annotated dataset**
- Strong regularization (dropout)
- GPU(s) for fast processing
 - ~ 150 images/sec
 - days—weeks of training



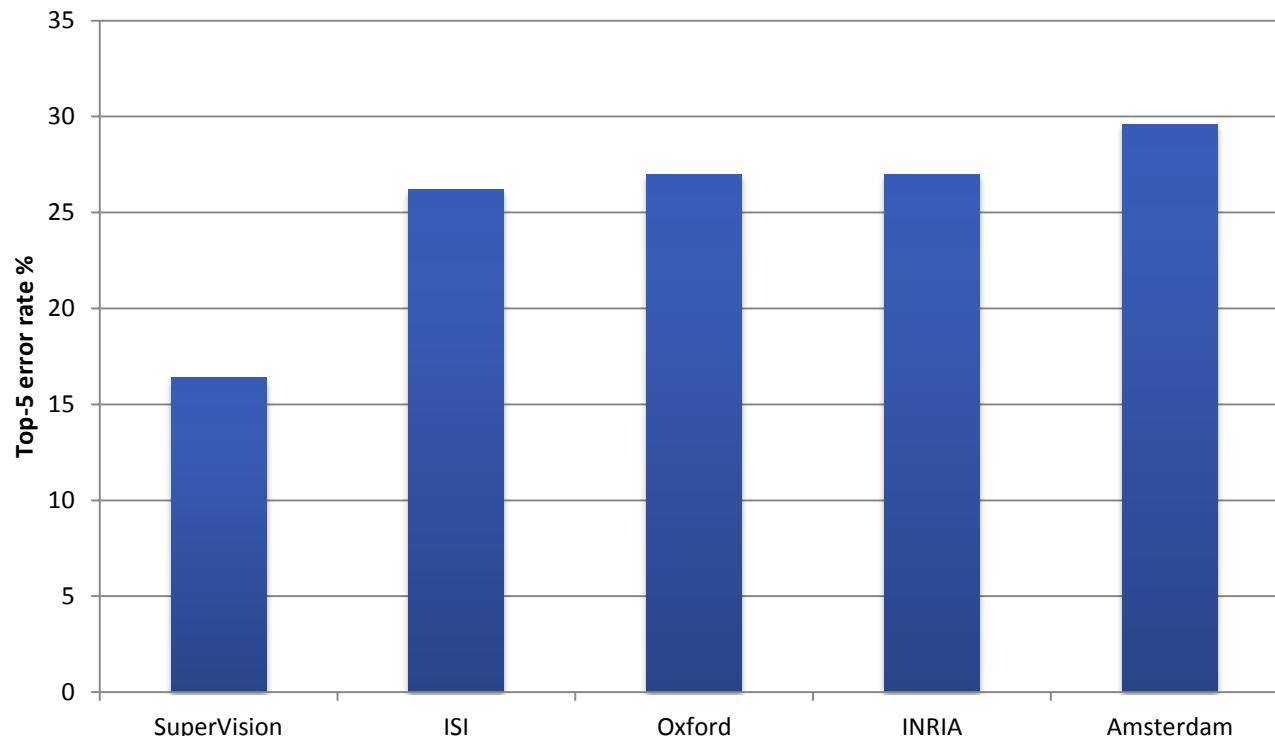
Imagenet database:

- 1K classes
- ~ 1K training images per class!
- ~ 1M training images

[Russakovsky et al. ICCV'13]

ImageNet Classification 2012

- Krizhevsky et al. -- 16.4% error (top-5)
- Next best (non-convnet) – 26.2% error



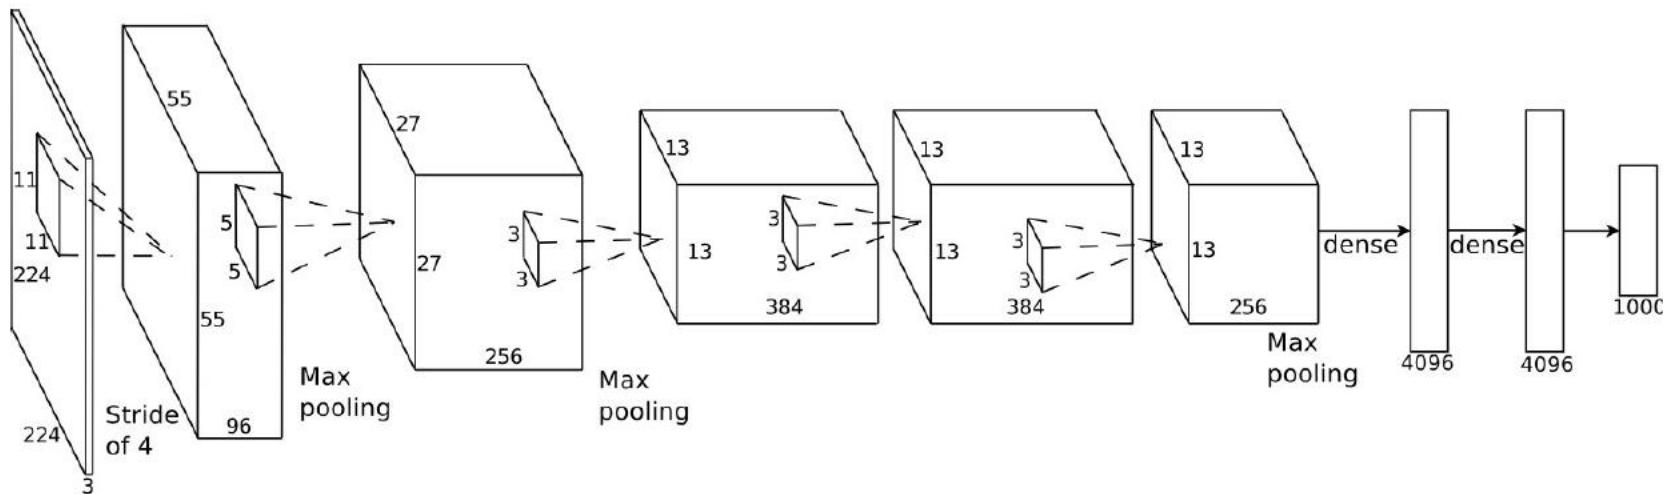
Architecture of [Krizhevsky et al. NIPS'12]

ImageNet Classification 2012

- 16.4% error (top-5)
- next best (variant of SIFT + Fisher Vectors) – 26.2% error

Same idea as in [LeCun'98] but on a larger scale:

- more training images (10^6 vs 10^3)
- more hidden layers

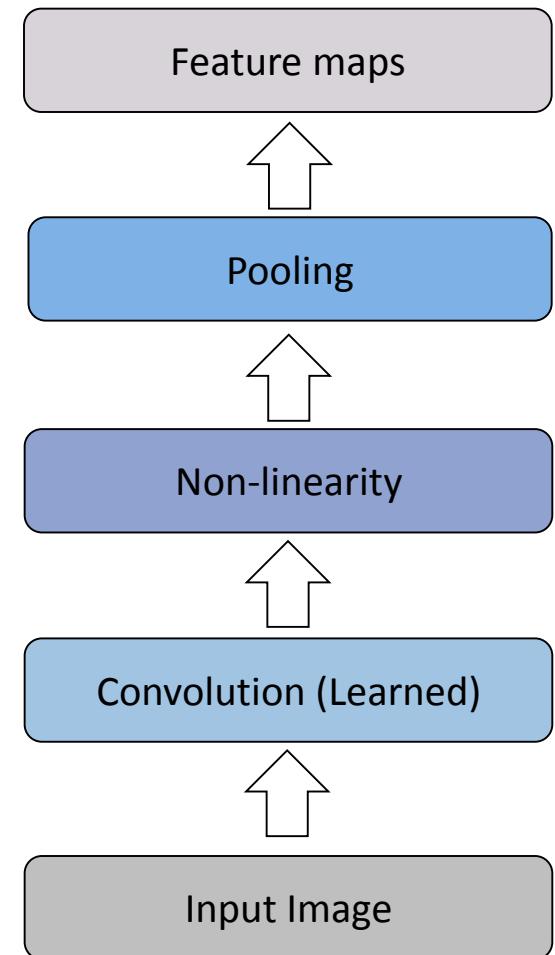
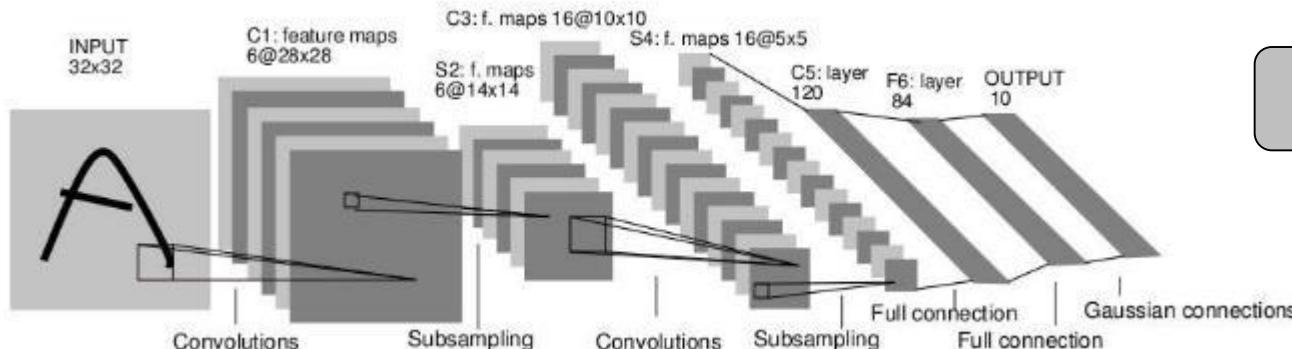


Overall: 60,000,000 parameters which are trained on 2 GPUs for a week with several tricks to reduce overfitting

- Data augmentation
- DropOut (new regularization)

ConvNet Architecture

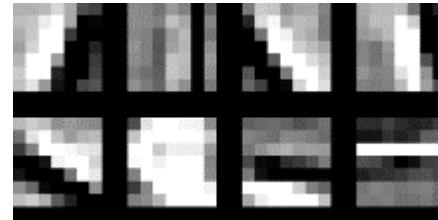
- Feed-forward:
 - Convolve input
 - Non-linearity (rectified linear)
 - Pooling (local max)
- Supervised
- Train convolutional filters by back-propagating classification error



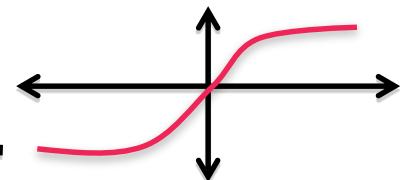
Components of Each Layer

Pixels /
Features

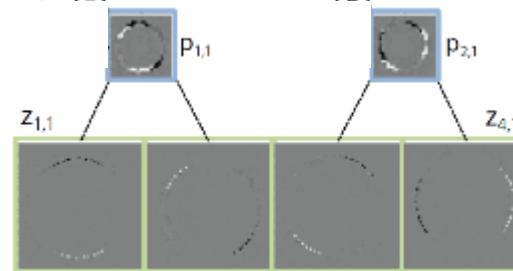
Filter with
Dictionary
(convolutional
or tiled)



+ Non-linearity



Spatial/Feature
(Sum or Max)



Normalization
between
feature responses



Output Features

[Optional]

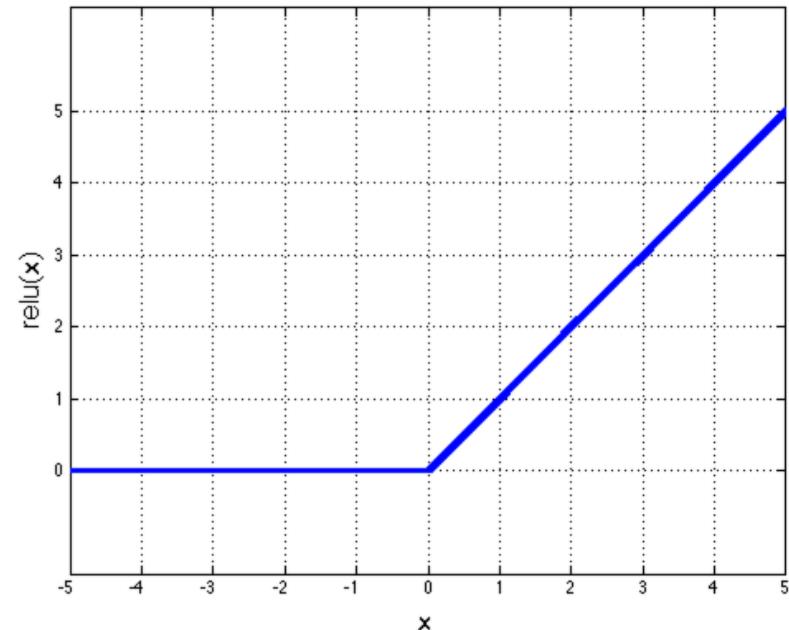
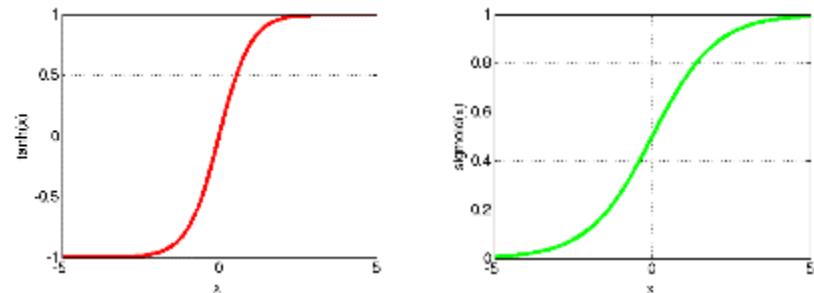
Convolutional filtering

- Why convolution?
 - Statistics of images look similar at different locations
 - Dependencies are very local
 - Filtering is an operation with translation *equivariance*



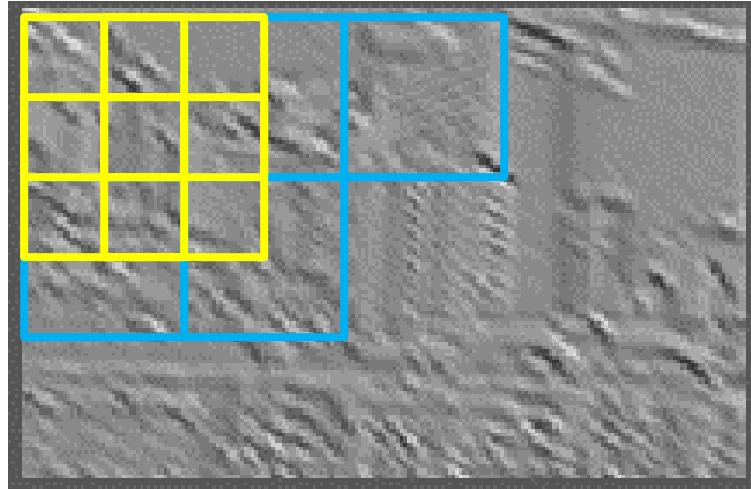
Non-Linearity

- Non-linearity applied to each response
 - Per-feature independent
 - Tanh
 - Sigmoid: $1/(1+\exp(-x))$
 - Rectified linear : $\max(0,x)$
 - Simplifies backprop
 - Makes learning faster
 - Avoids saturation issues
(much higher dynamic range)
- Preferred option

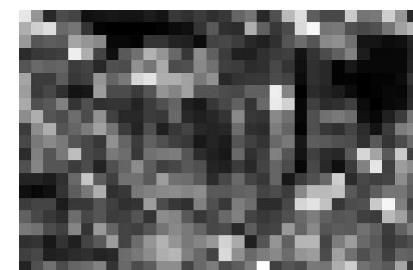


Pooling

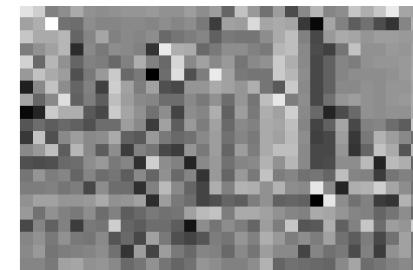
- Spatial Pooling
 - Non-overlapping / overlapping regions (-0.3% in error)
 - Sum or max
 - Provides invariance to local transformations



Max

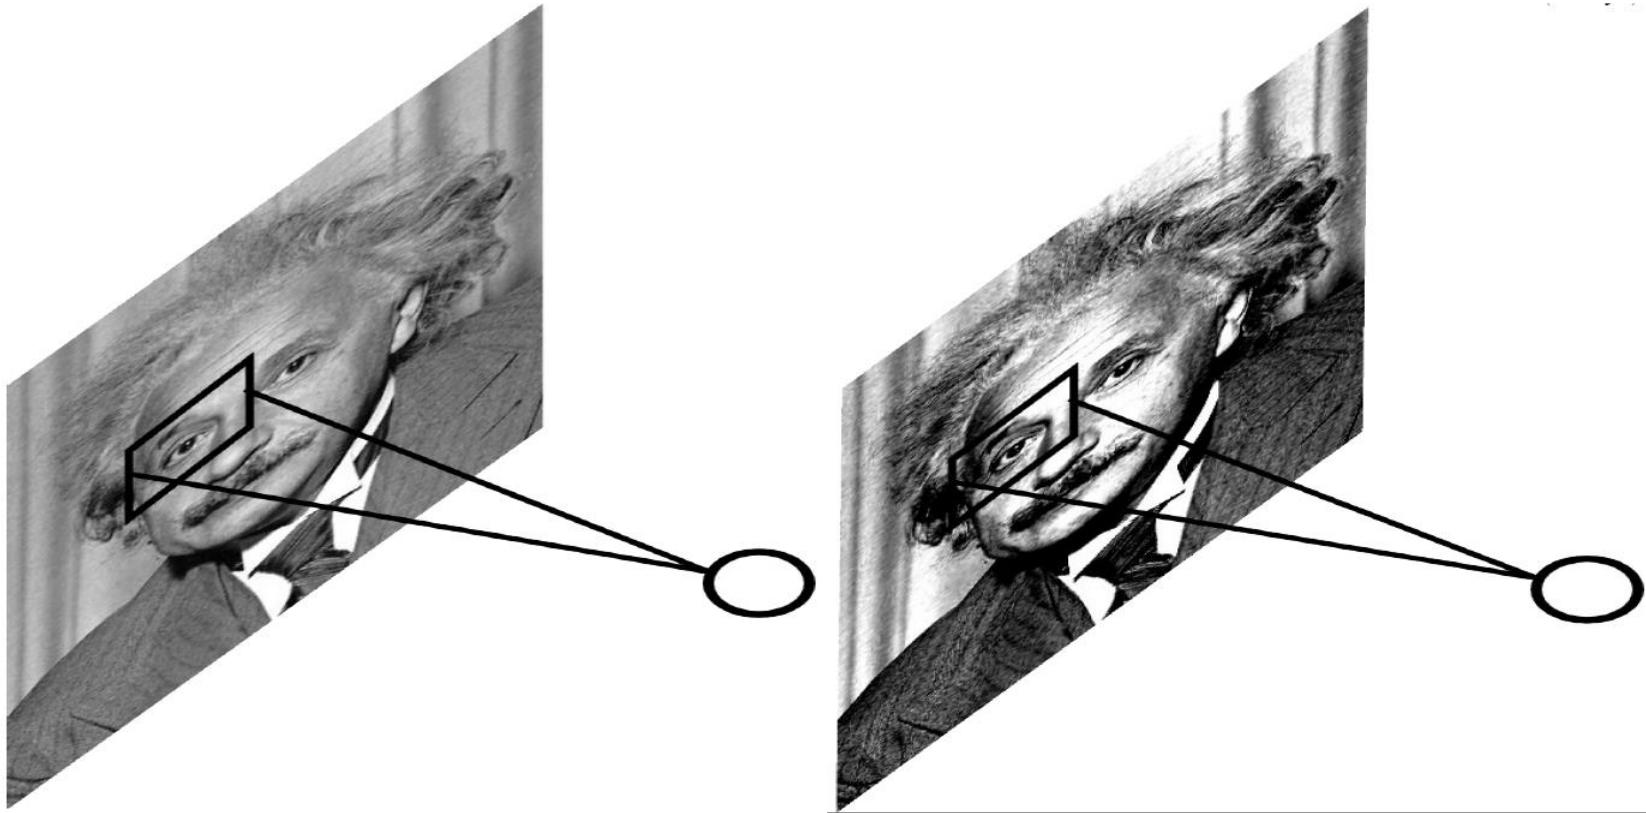


Sum



Normalization

- Contrast normalization

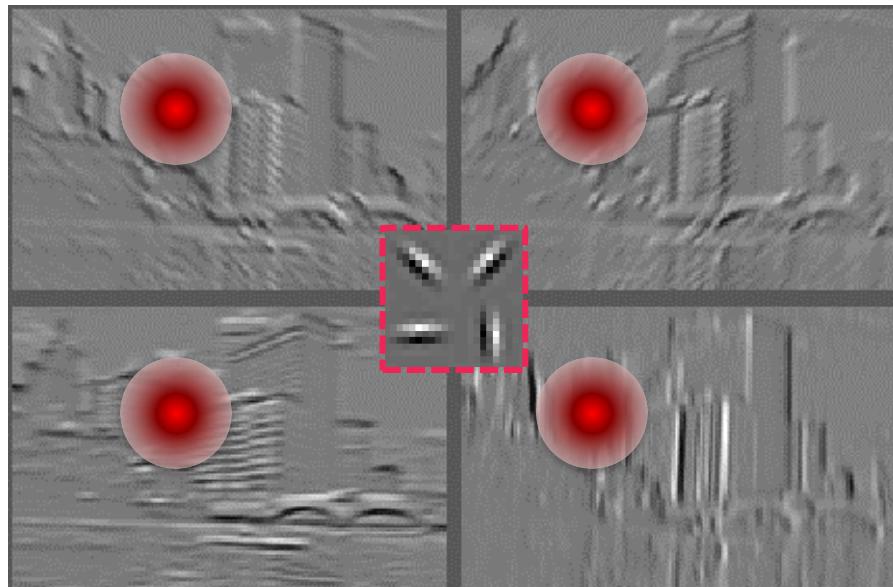


The same response for different contrasts is desired

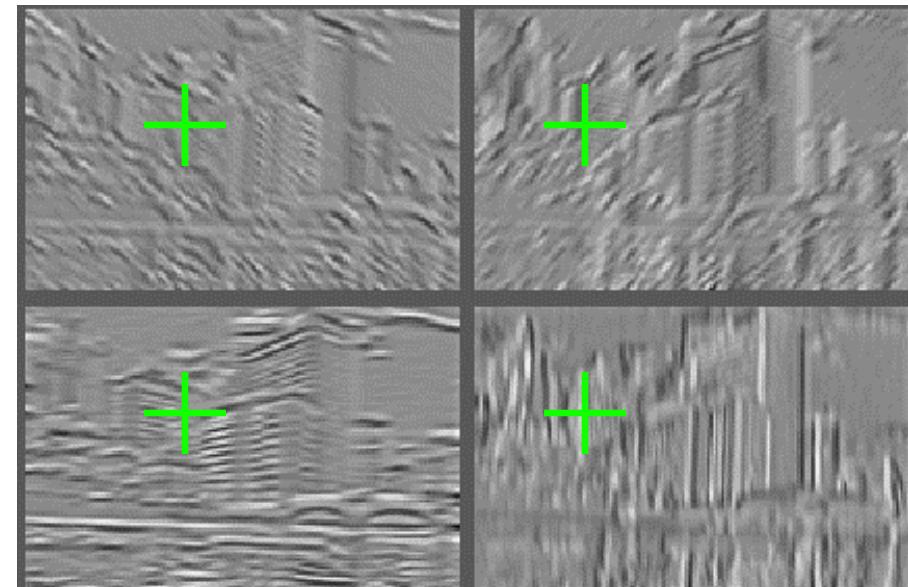
- Contrast normalization
 - Local mean
 - Equalizes the distribution of pixel values



(feature maps)
Gaussian



Feature Maps

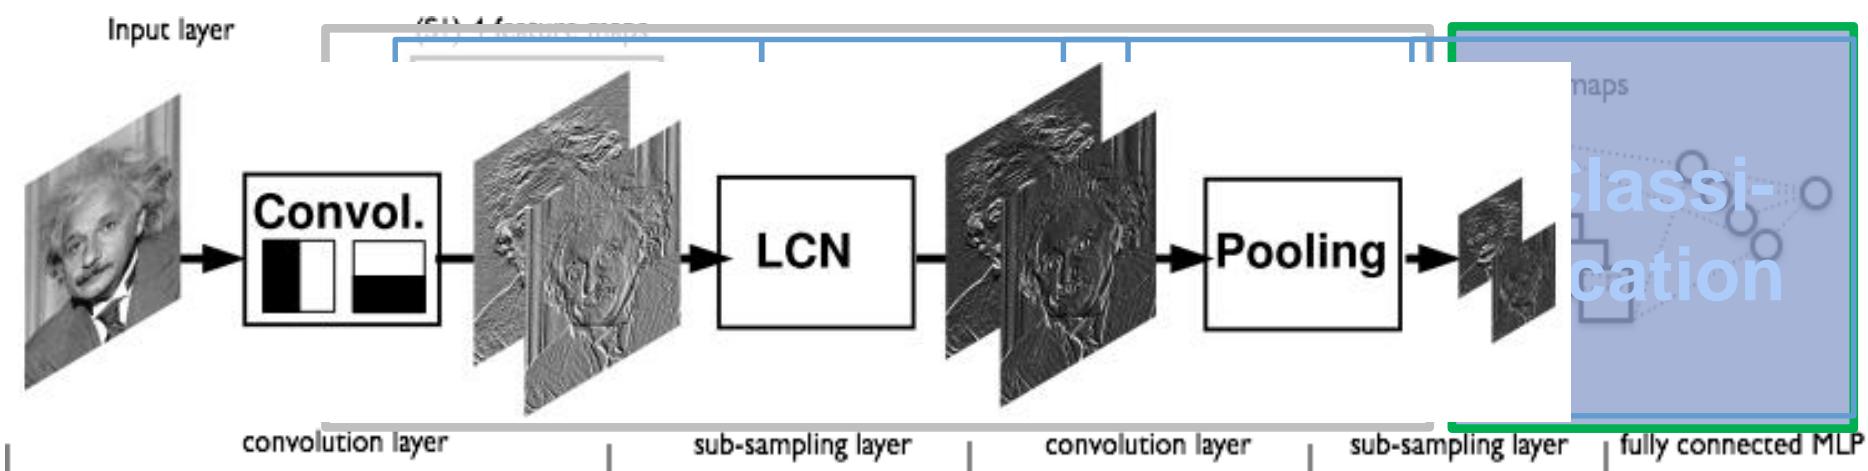


Feature Maps
After Contrast Normalization

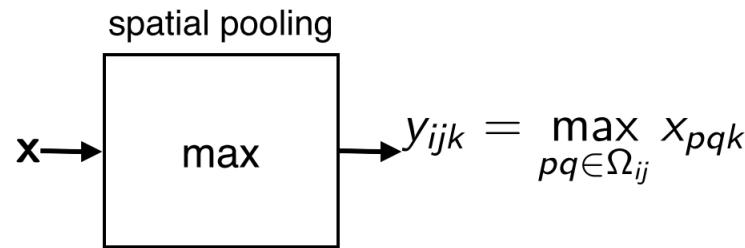
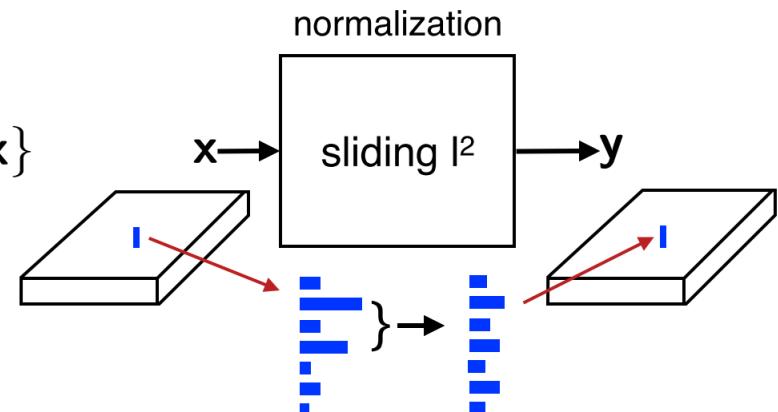
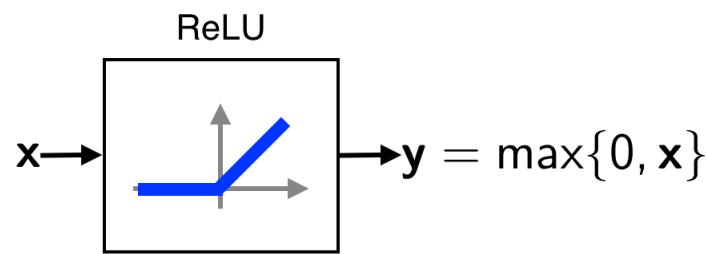
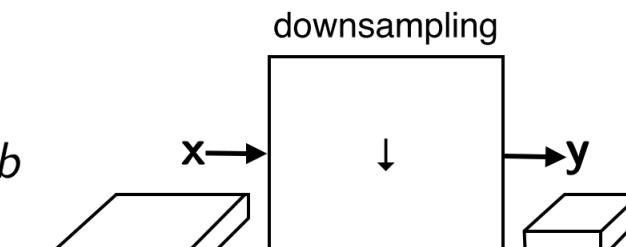
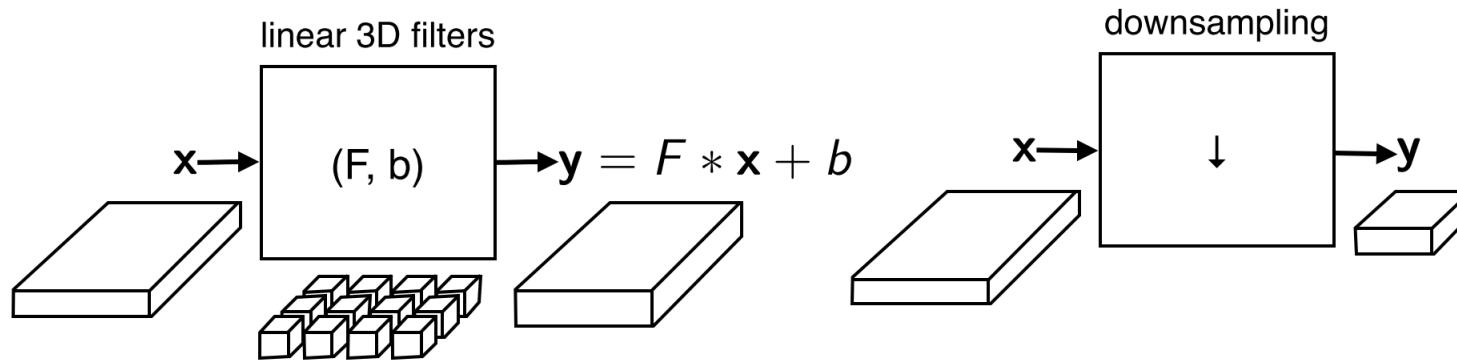
Recap: Components of a CNN

CNN - multi-layer NN architecture

- Convolutional + Non-Linear Layer
- Sub-sampling Layer
- Convolutional +Non-Linear Layer
- Fully connected layers
- Supervised



Summary: Components of Each Layer



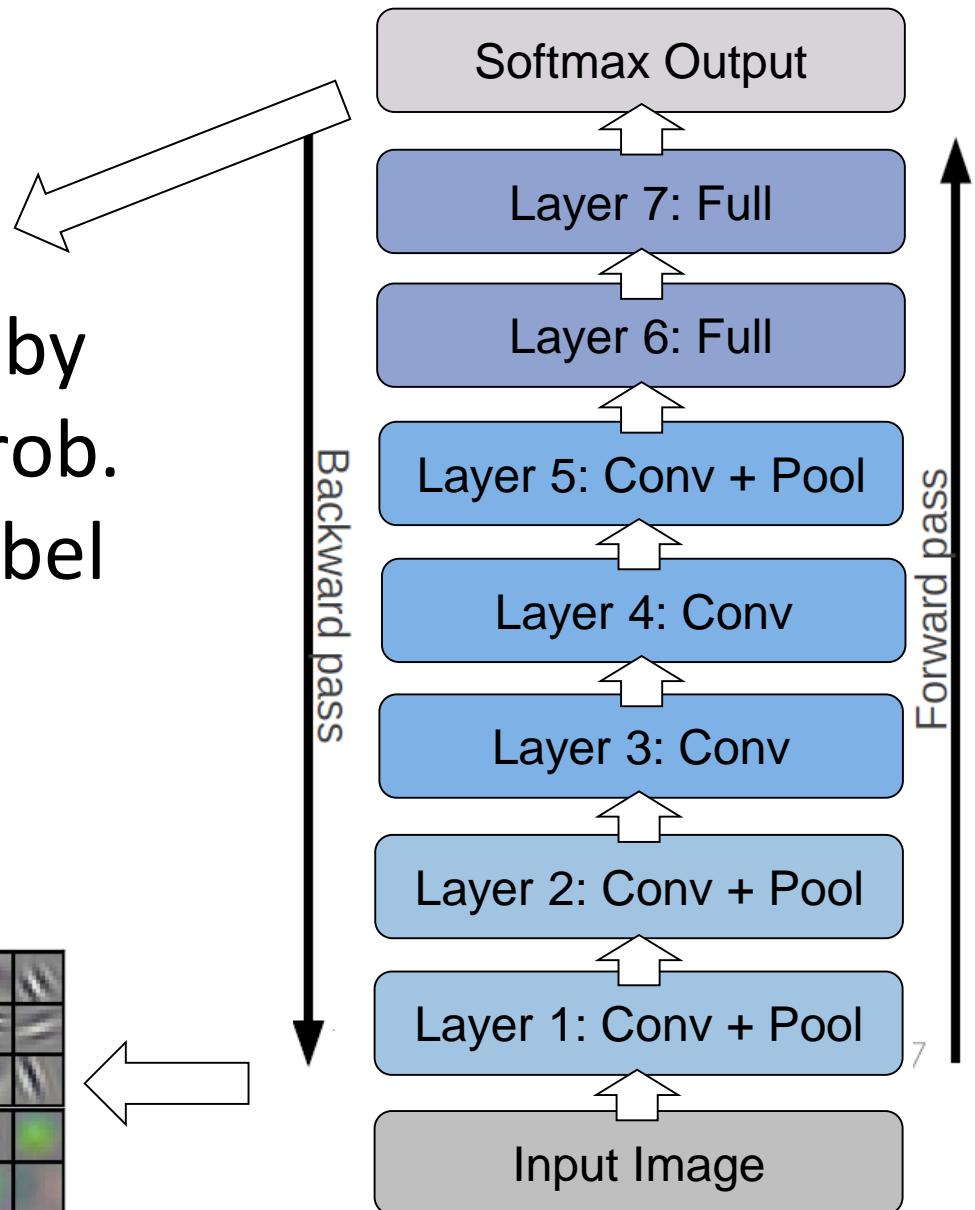
Architecture of Krizhevsky et al.

One output unit per class

x_i = total input to output unit i

$$f(x_i) = \frac{\exp(x_i)}{\sum_{j=1}^{1000} \exp(x_j)}$$

- Trained via backprop by maximizing the log-prob. of the correct class-label
- 8 layers total
- Trained on ImageNet dataset

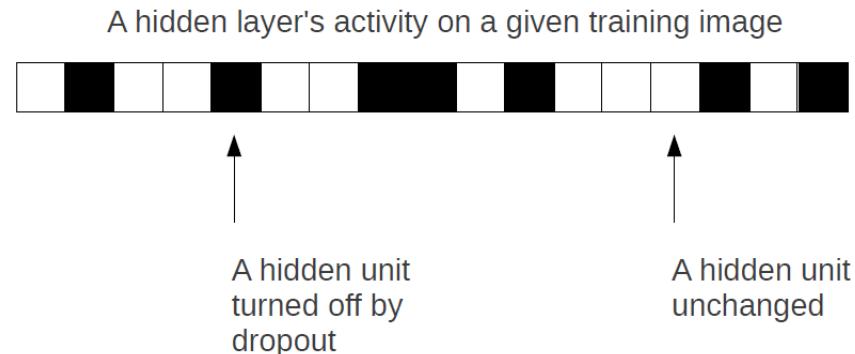


Improving Generalization: DropOut

[Hinton et al. NIPS'12]

Motivation:

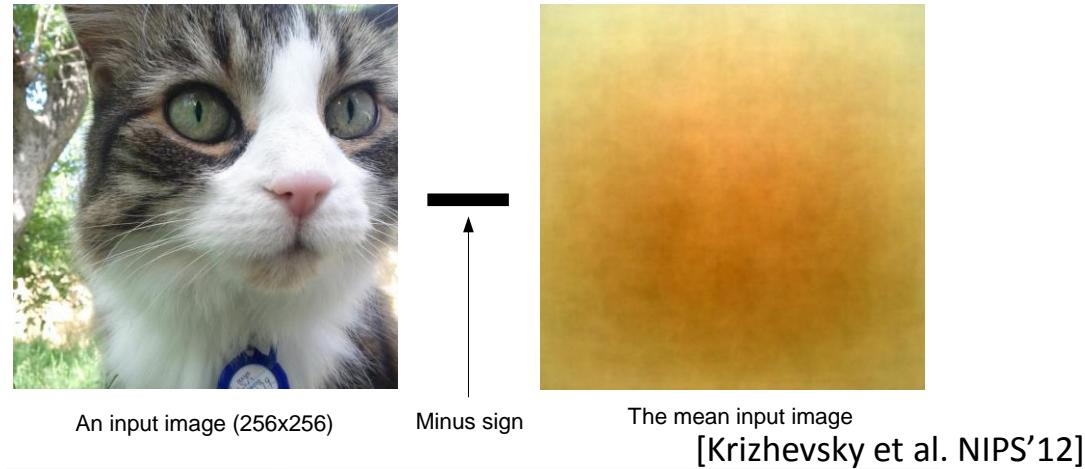
- Random Forests generalize well due to averaging of many models
- Decision Trees are **fast** - ConvNets are **slow** – many models are not feasible
- Similar to random forest bagging [Breiman'94], but differs in that parameters are shared
- For fully connected layers only:
 - In training: Independently set each hidden unit activity to zero with 0.5 probability
 - In testing multiply neuron output by 0.5
- Corresponds to averaging over exponentially many samples from different model architectures



Further pre-processing tricks

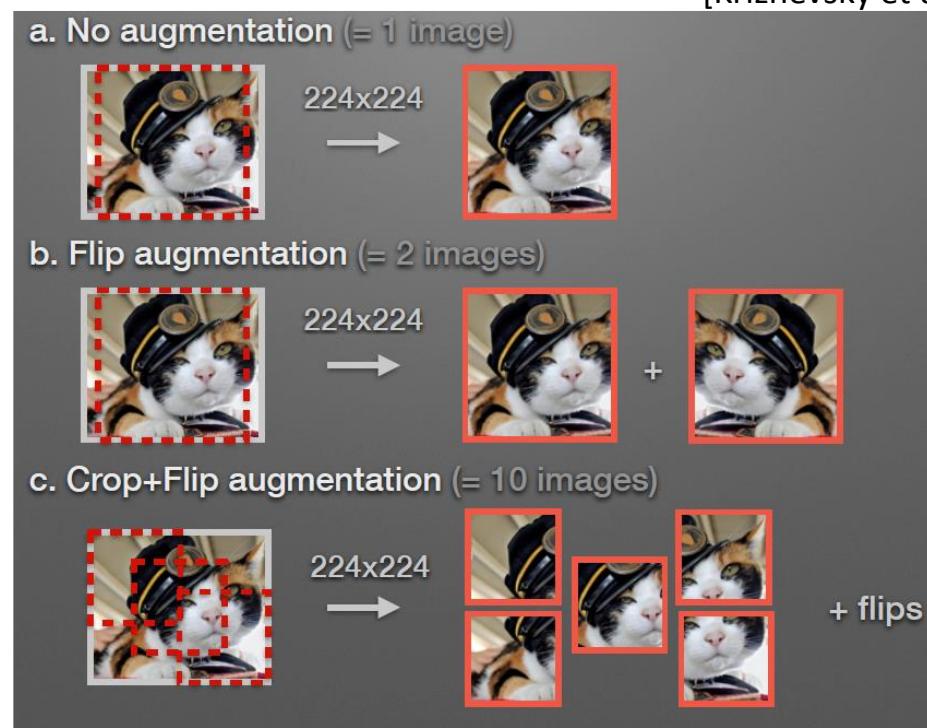
- Mean removal

Centered (0-mean) RGB values.



- Data augmentation

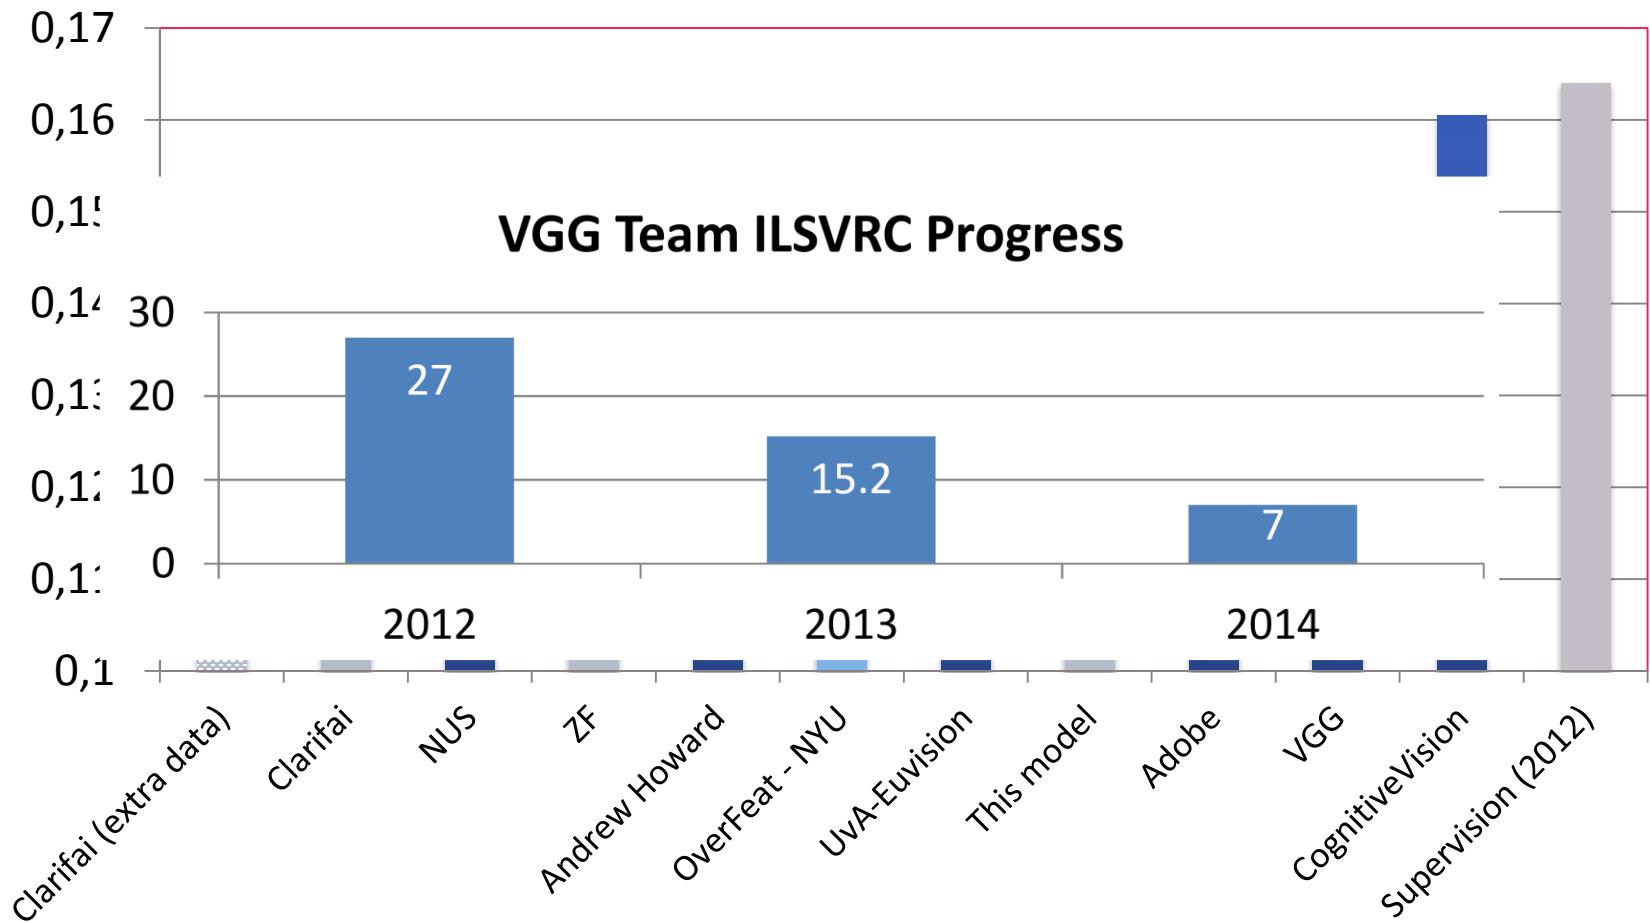
Train on 224x224 patches extracted randomly from images, and also their horizontal reflections



[Chatfield et al. BMVC'14]

ImageNet Classification 2013/2014 Results

- <http://www.image-net.org/challenges/LSVRC/2013/results.php>



- Pre-2012: 26.2% error → 2012: 16.5% error → 2013: 11.2% error

ImageNet Sample classifications

[Krizhevsky et al. NIPS'12]



mite	container ship	motor scooter	leopard
mite	container ship	motor scooter	leopard
black widow	lifeboat	go-kart	jaguar
cockroach	amphibian	moped	cheetah
tick	fireboat	bumper car	snow leopard
starfish	drilling platform	golfcart	Egyptian cat
grille	mushroom	cherry	Madagascar cat
convertible	agaric	dalmatian	squirrel monkey
grille	mushroom	grape	spider monkey
pickup	jelly fungus	elderberry	titi
beach wagon	gill fungus	ffordshire bullterrier	indri
fire engine	dead-man's-fingers	currant	howler monkey

ImageNet Sample classifications

[Krizhevsky et al. NIPS'12]



lens cap

reflex camera
Polaroid camera
pencil sharpener
switch
combination lock



abacus

typewriter keyboard
space bar
computer keyboard
accordion



slug

zucchini
ground beetle
common newt
water snake



hen

cock
cocker spaniel
partridge
English setter



tiger

tiger
tiger cat
tabby
boxer
Saint Bernard



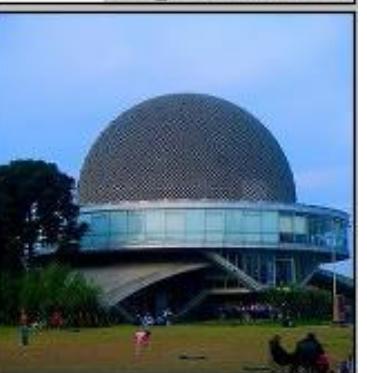
chambered nautilus

lampshade
throne
goblet
table lamp
hamper



tape player

cellular telephone
slot
reflex camera
dial telephone
iPod



planetarium

planetarium
dome
mosque
radio telescope
steel arch bridge

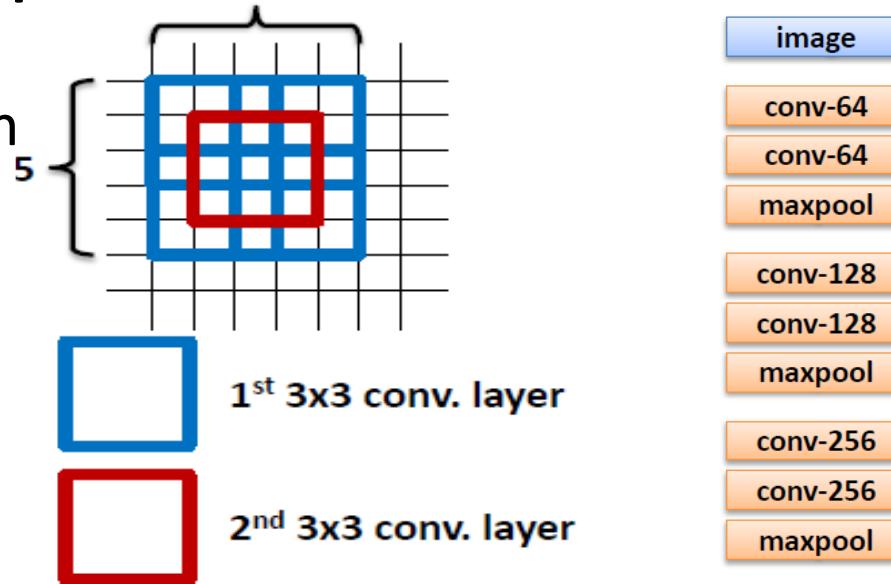
ImageNet Classification Progress from '12-'14

2012 Teams	%error	2013 Teams	%error	2014 Teams	%error
Supervision (Toronto)	15.3	Clarifai (NYU spinoff)	11.7	GoogLeNet	6.6
ISI (Tokyo)	26.1	NUS (singapore)	12.9	VGG (Oxford)	7.3
VGG (Oxford)	26.9	Zeiler-Fergus (NYU)	13.5	MSRA	8.0
XRCE/INRIA	27.0	A. Howard	13.5	A. Howard	8.1
UvA (Amsterdam)	29.6	OverFeat (NYU)	14.1	DeeperVision	9.5
INRIA/LEAR	33.4	UvA (Amsterdam)	14.2	NUS-BST	9.7
		Adobe	15.2	TTIC-ECP	10.2
		VGG (Oxford)	15.2	XYZ	11.2
		VGG (Oxford)	23.0	UvA	12.1

Better (\approx deeper) architectures exist now

ILSVRC14 Winners: ~7.3% Top-5 error₅

- VGG: 16 layers of stacked 3x3 convolution with stride 1



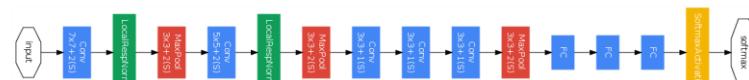
Other details:

- Rectification (ReLU) non-linearity
- 5 max-pool layers (x2 reduction)
- no normalisation
- 3 fully-connected (FC) layers

Better (\approx deeper) architectures exist now

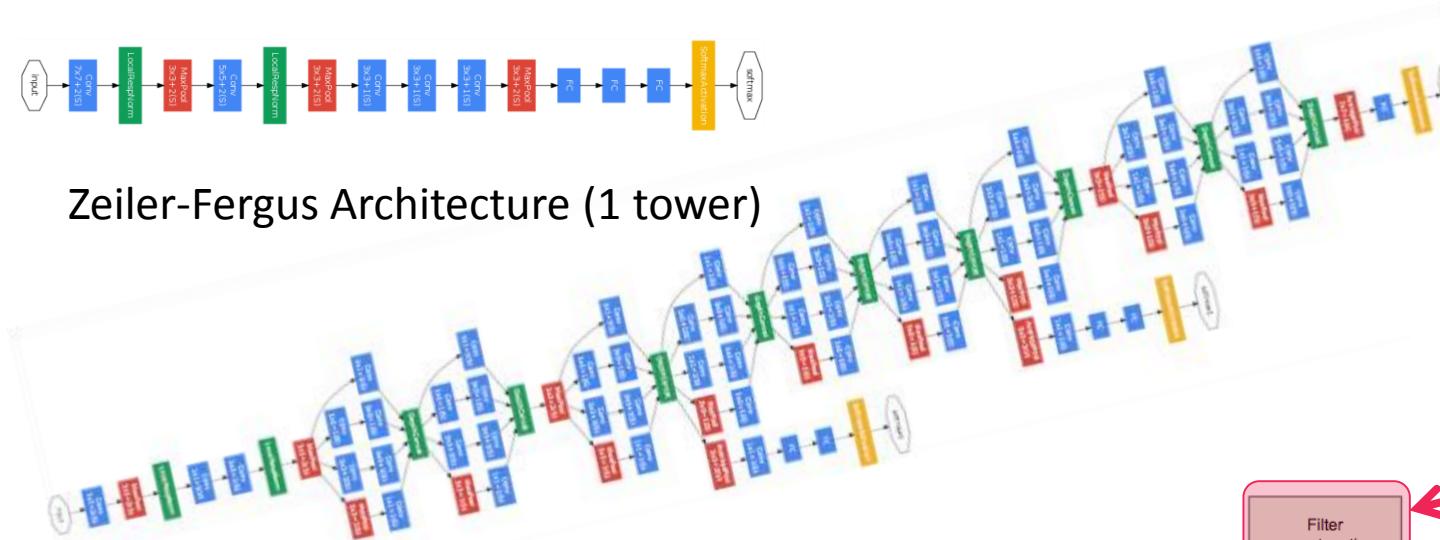
ILSVRC14 Winners: ~6.6% Top-5 error

- GoogLeNet: composition of multi-scale dimension-reduced modules:



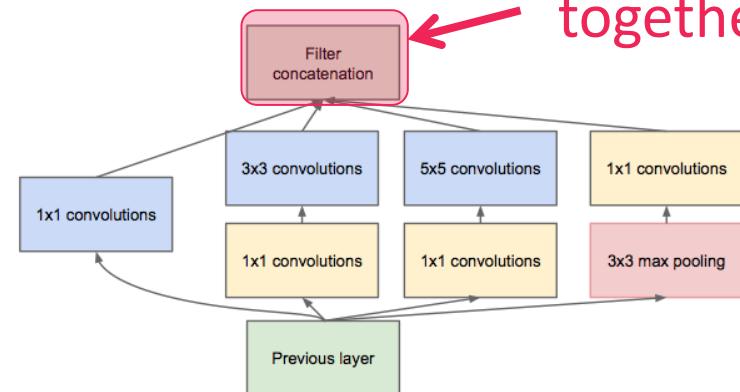
Convolution
Pooling
Softmax
Other

Zeiler-Fergus Architecture (1 tower)



“wire together
what fires
together”

- 1x1 convolutions serve as dimensionality reduction



Credit: C. Szegedy

Classification failure cases



Groundtruth: **coffee mug**

GoogLeNet:

- **table lamp**
- **lamp shade**
- **printer**
- **projector**
- **desktop computer**

Classification failure cases



Groundtruth: **hay**

GoogLeNet:

- sorrel (horse)
- hartebeest
- Arabian camel
- warthog
- gaselle

Classification failure cases



Groundtruth: **Police car**

GoogLeNet:

- laptop
- hair drier
- binocular
- ATM machine
- seat belt

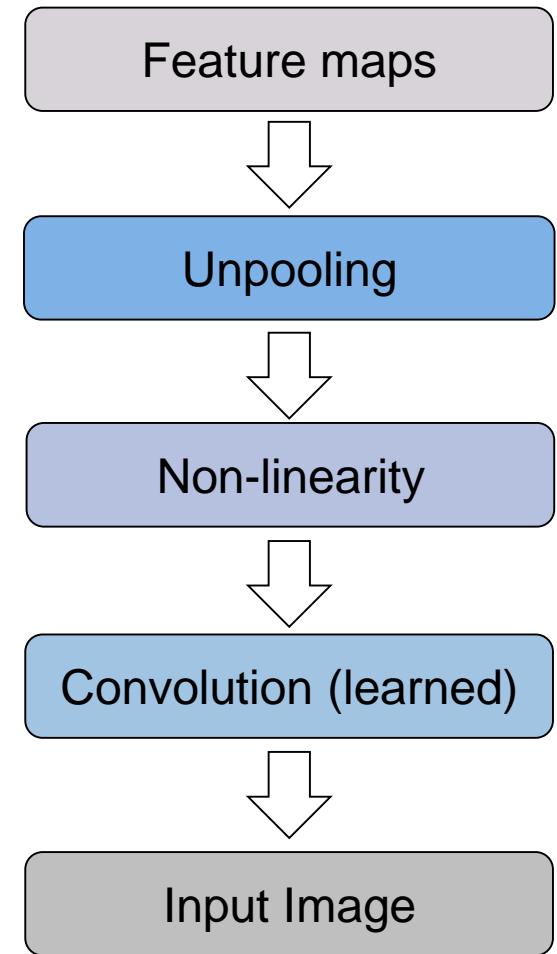
What is learned? Visualizing CNNs

M. Zeiler & R. Fergus, *Visualizing and Understanding Convolutional Networks*, ECCV, 2014

Visualization using Deconvolutional Networks

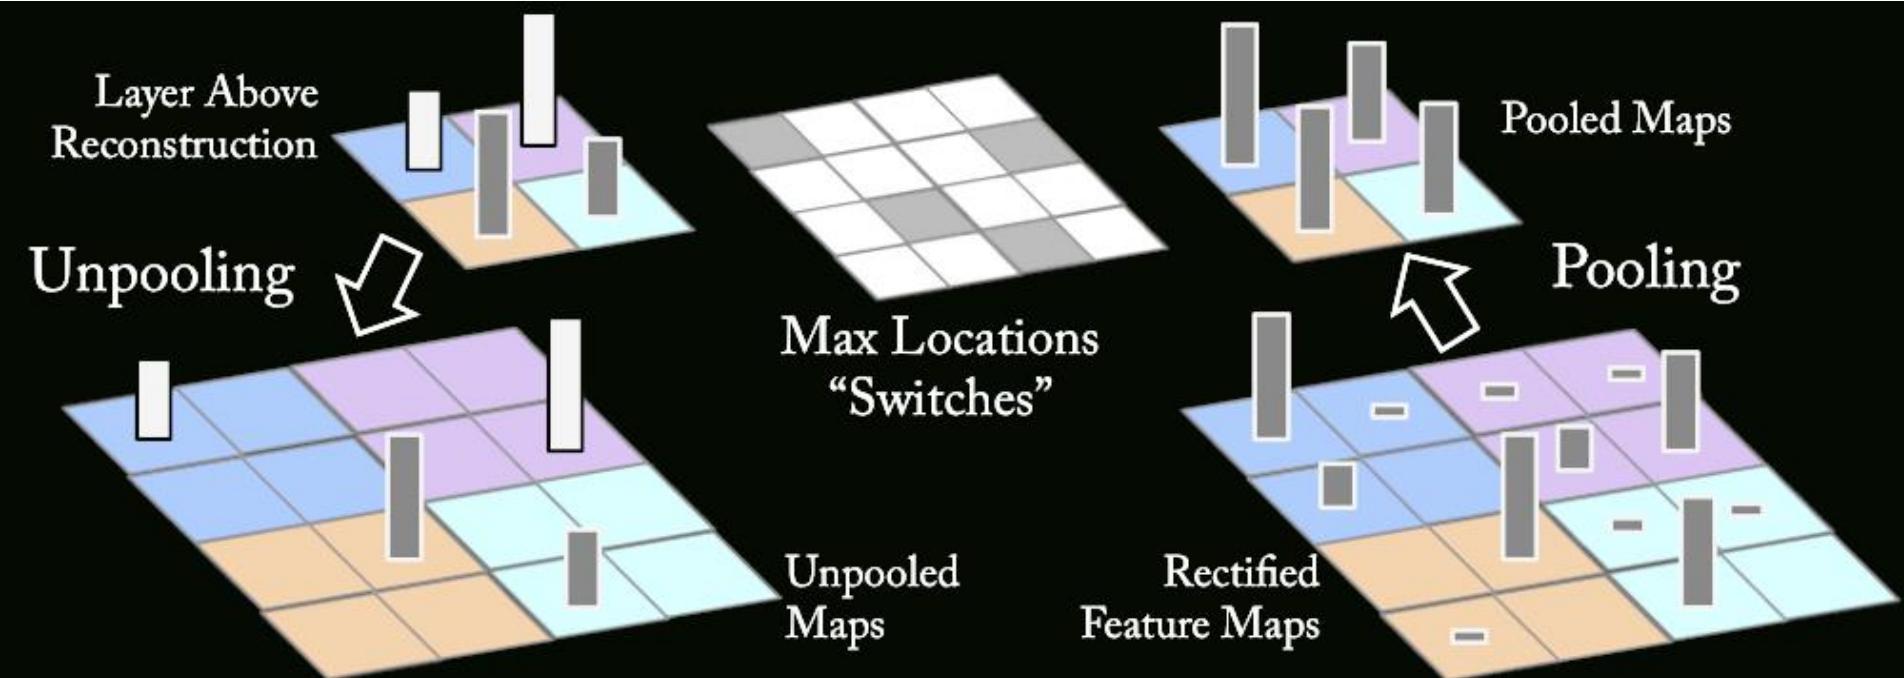
[Zeiler et al. CVPR'10, ICCV'11, ECCV'14]

- Provides way to map activations at high layers back to the input
- Same operations as Convnet, but in reverse:
 - Unpool feature maps
 - Convolve unpooled maps
 - Filters copied from Convnet
- Used here purely as a probe
 - Originally proposed as unsupervised learning method
 - No inference, no learning



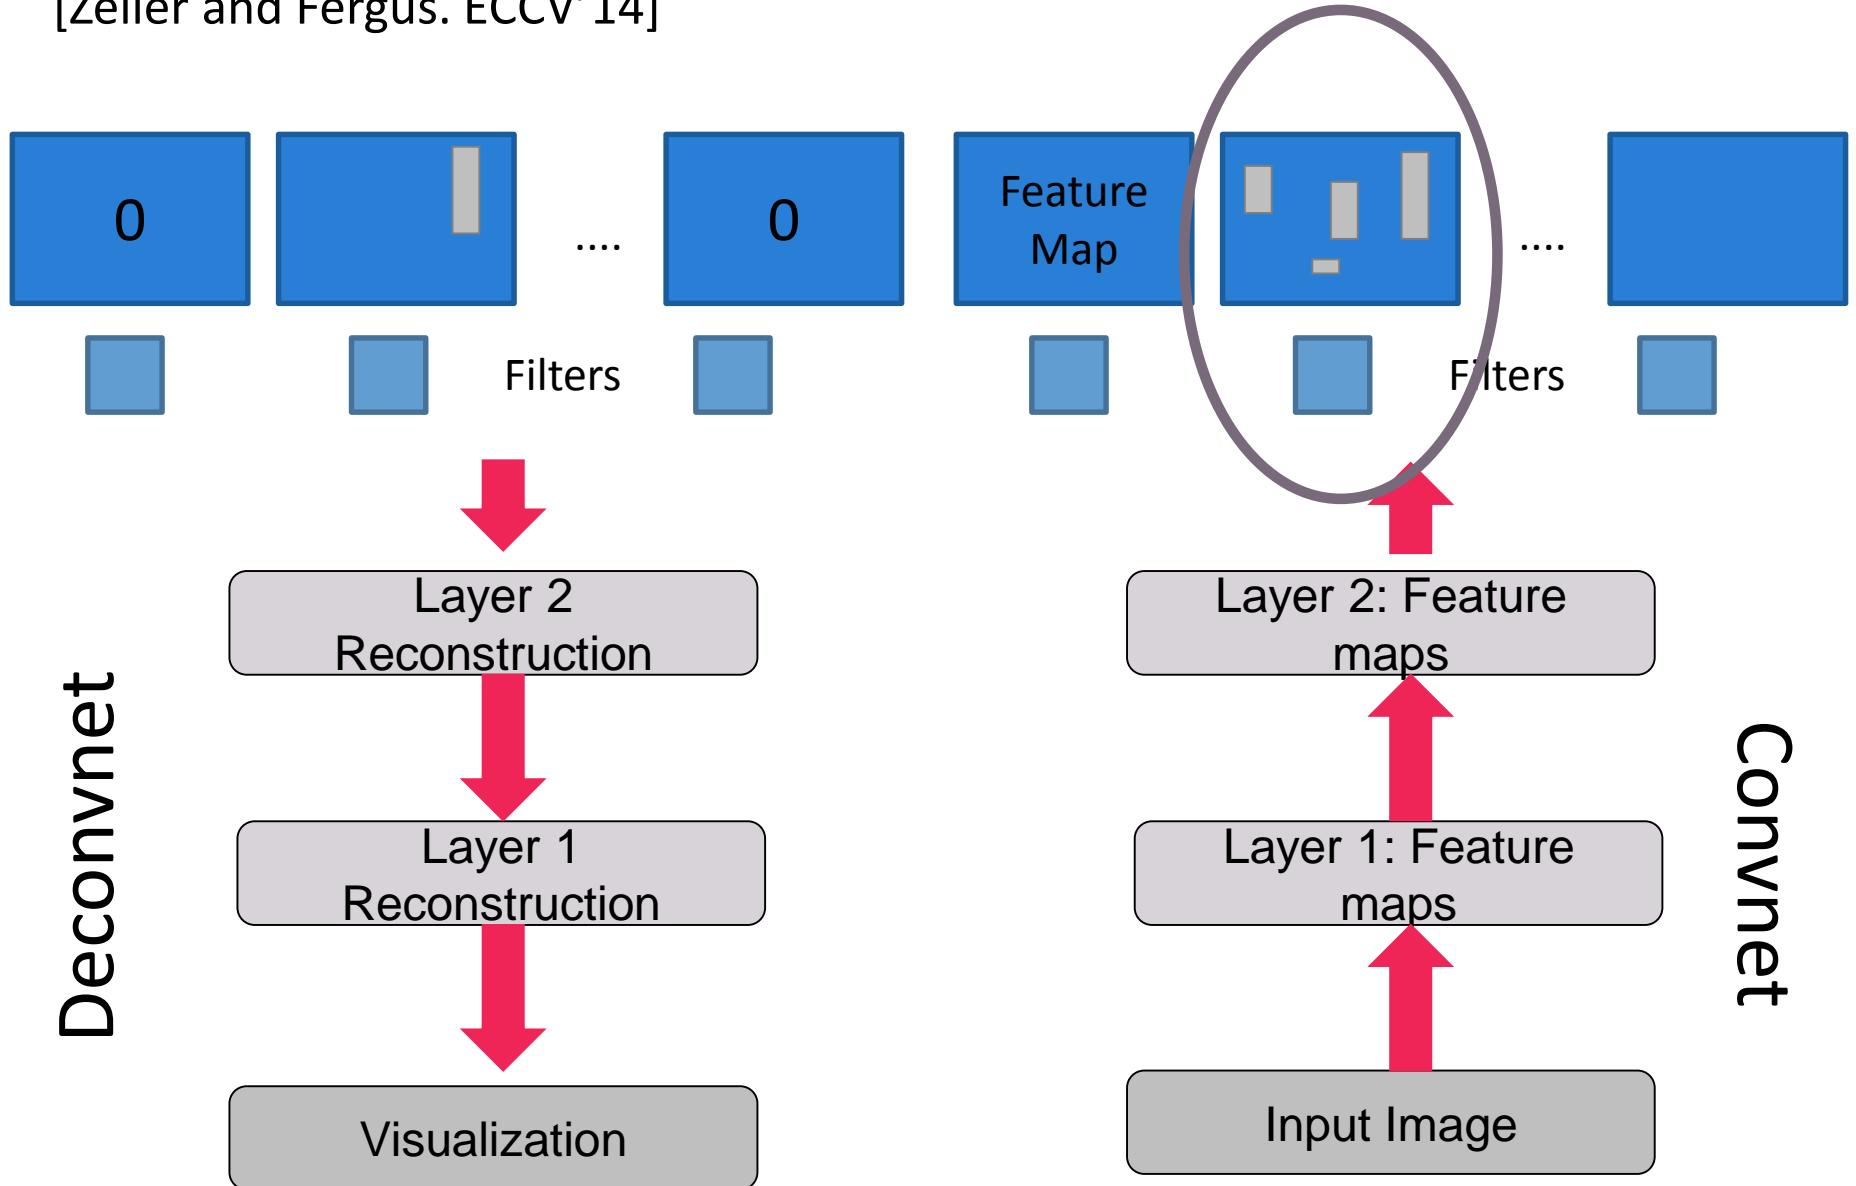
Unpooling Operation

- Switches record where the pooled activations came from to “unpool” the reconstructed layer above



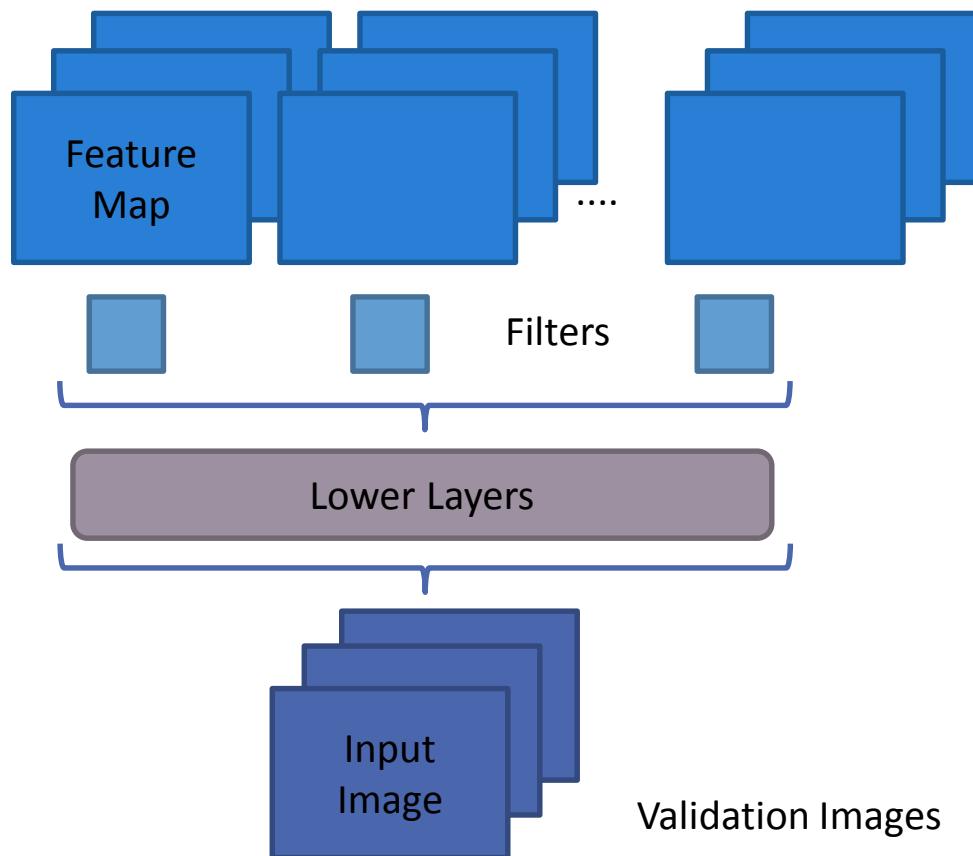
Deconvnet Projection from Higher Layers

[Zeiler and Fergus. ECCV'14]



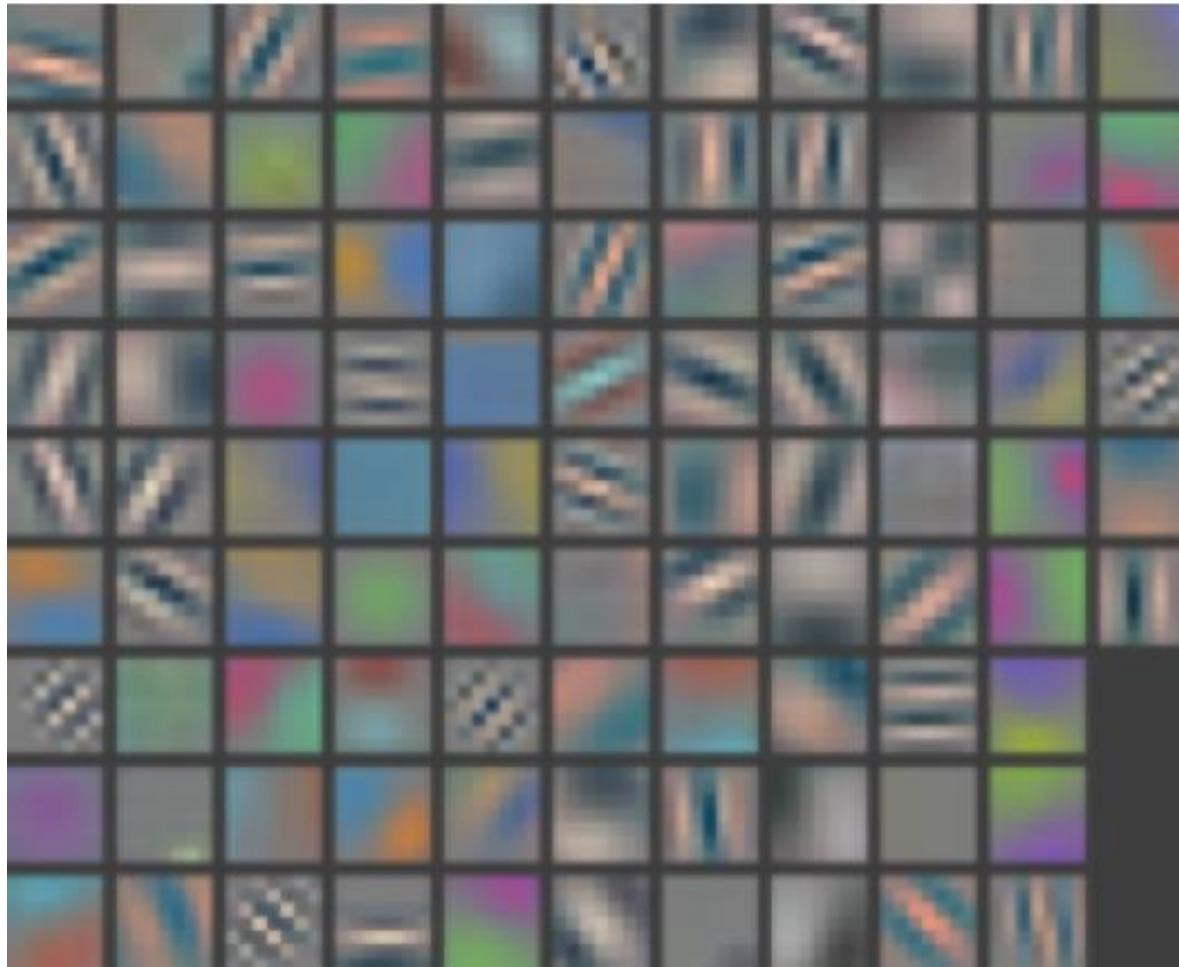
Visualizations of Higher Layers

- Use ImageNet 2012 validation set (stack of images)
- Push each image through network and look for image with the strongest activation for each feature map



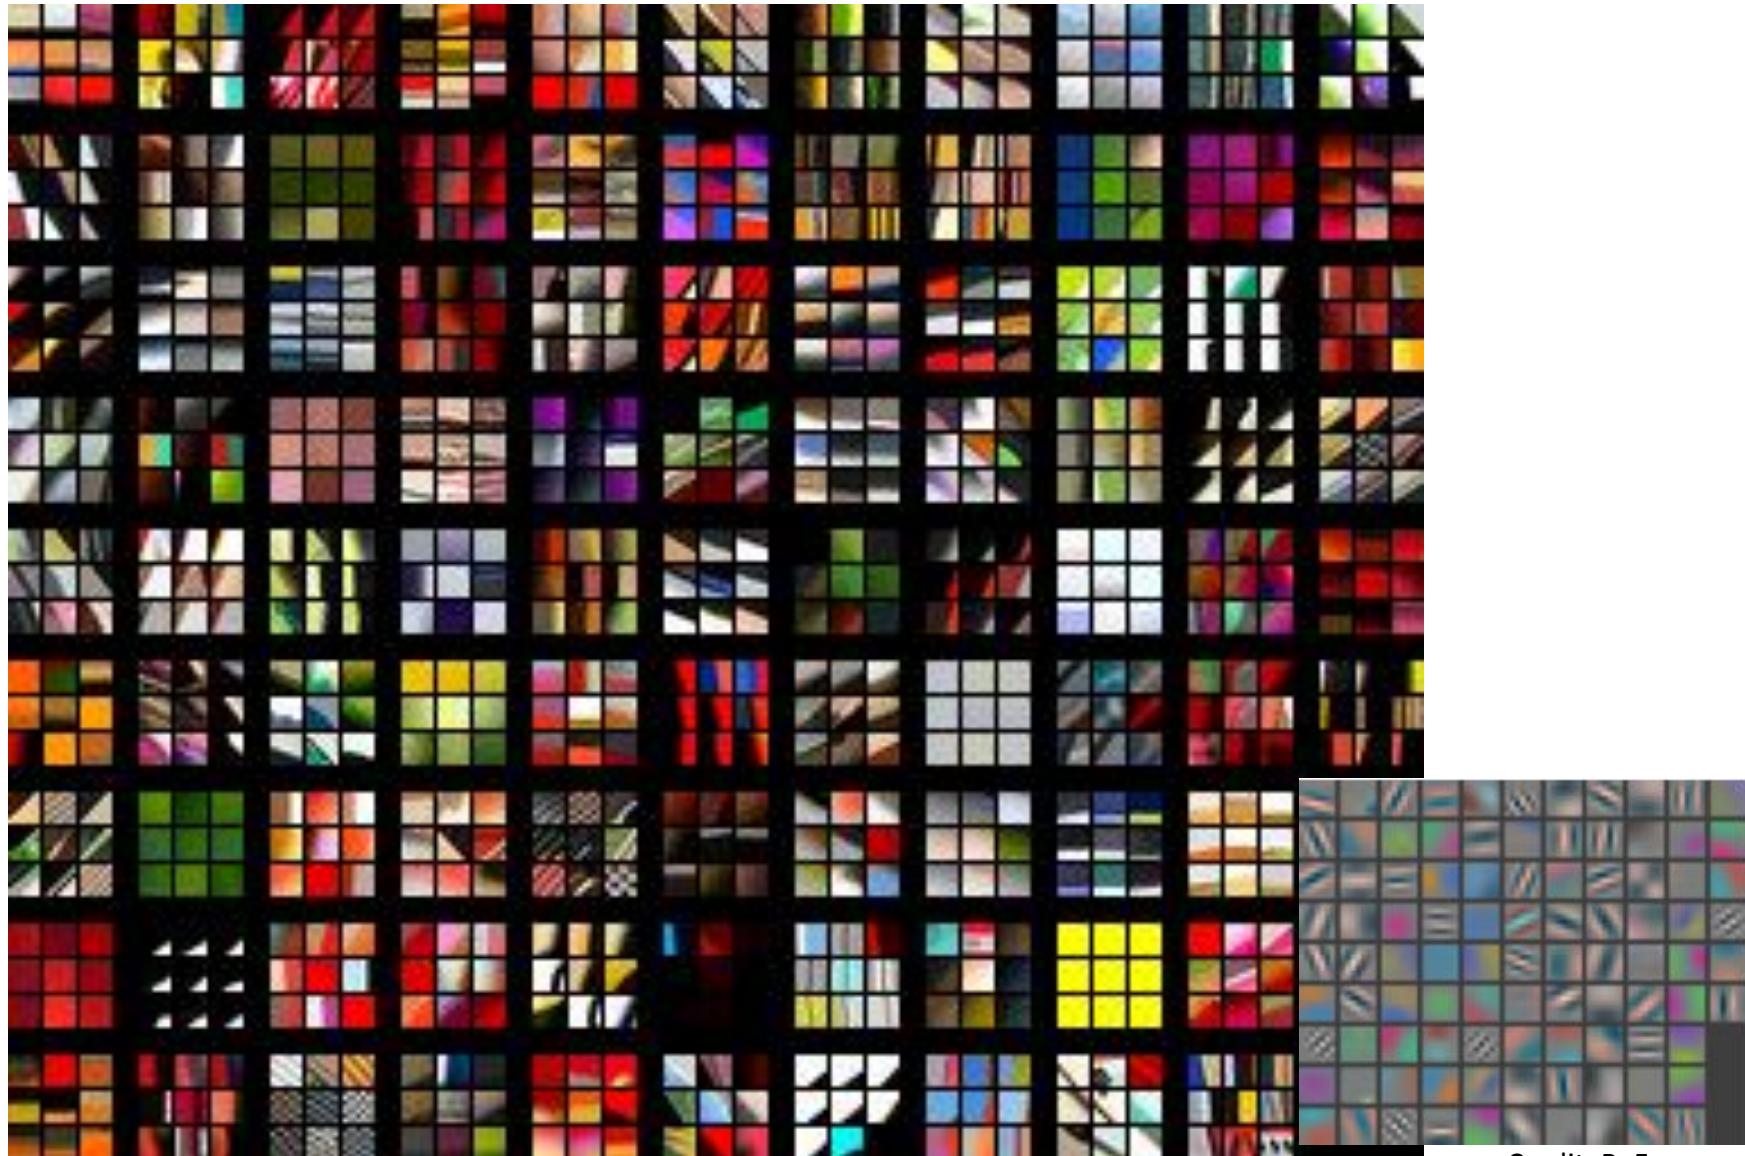
- Take max activation from feature map associated with each filter
- Use Deconvnet to project back to pixel space
- Use pooling “switches” distinctive to that activation

Layer 1 Filters



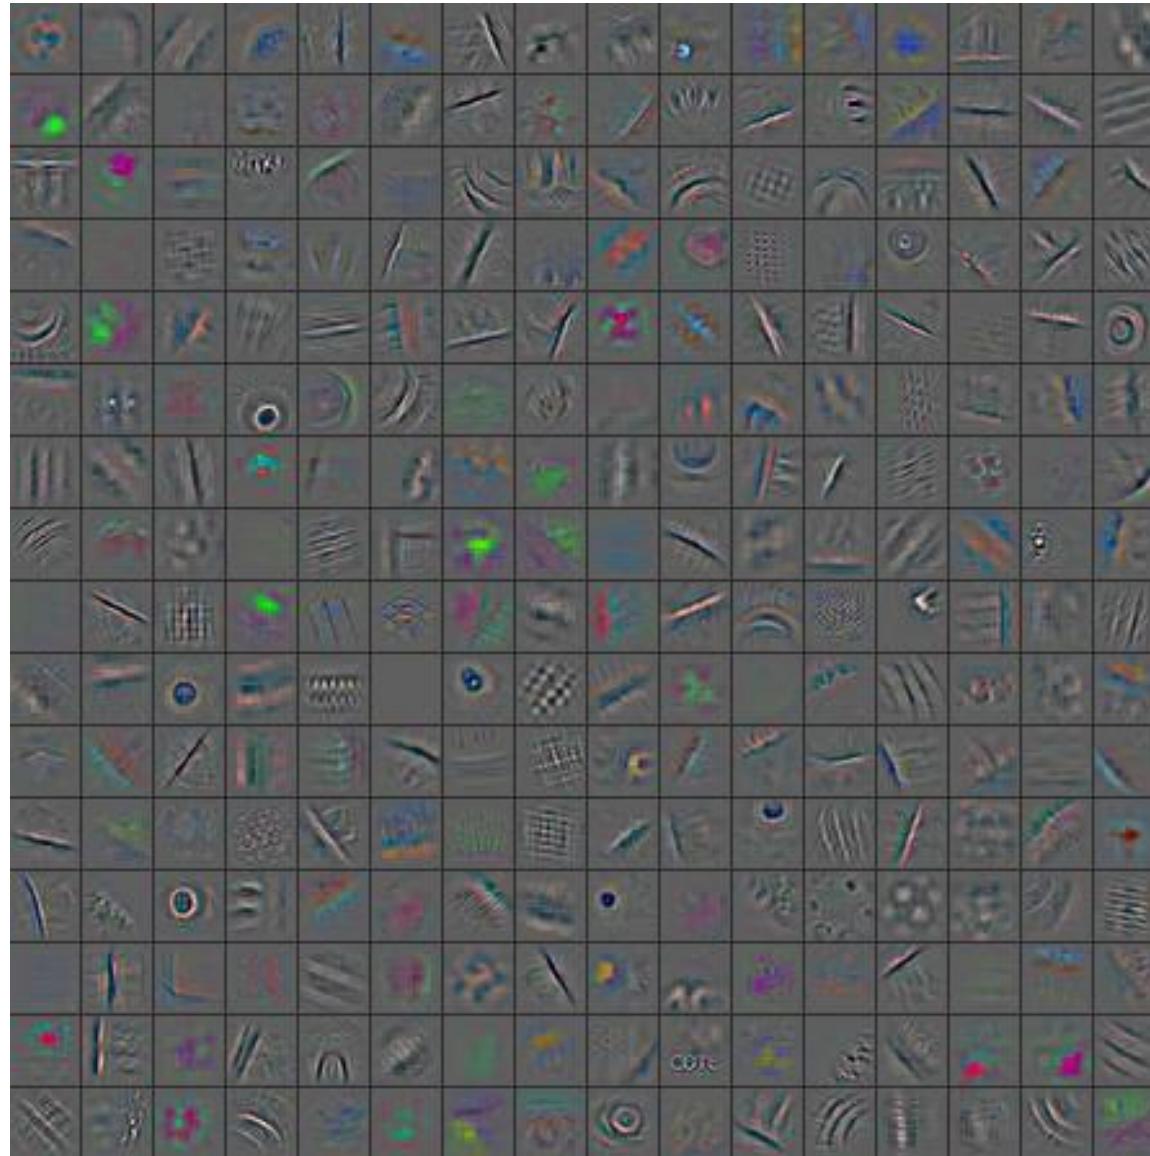
Credit: R. Fergus

Layer 1: Top-9 Patches



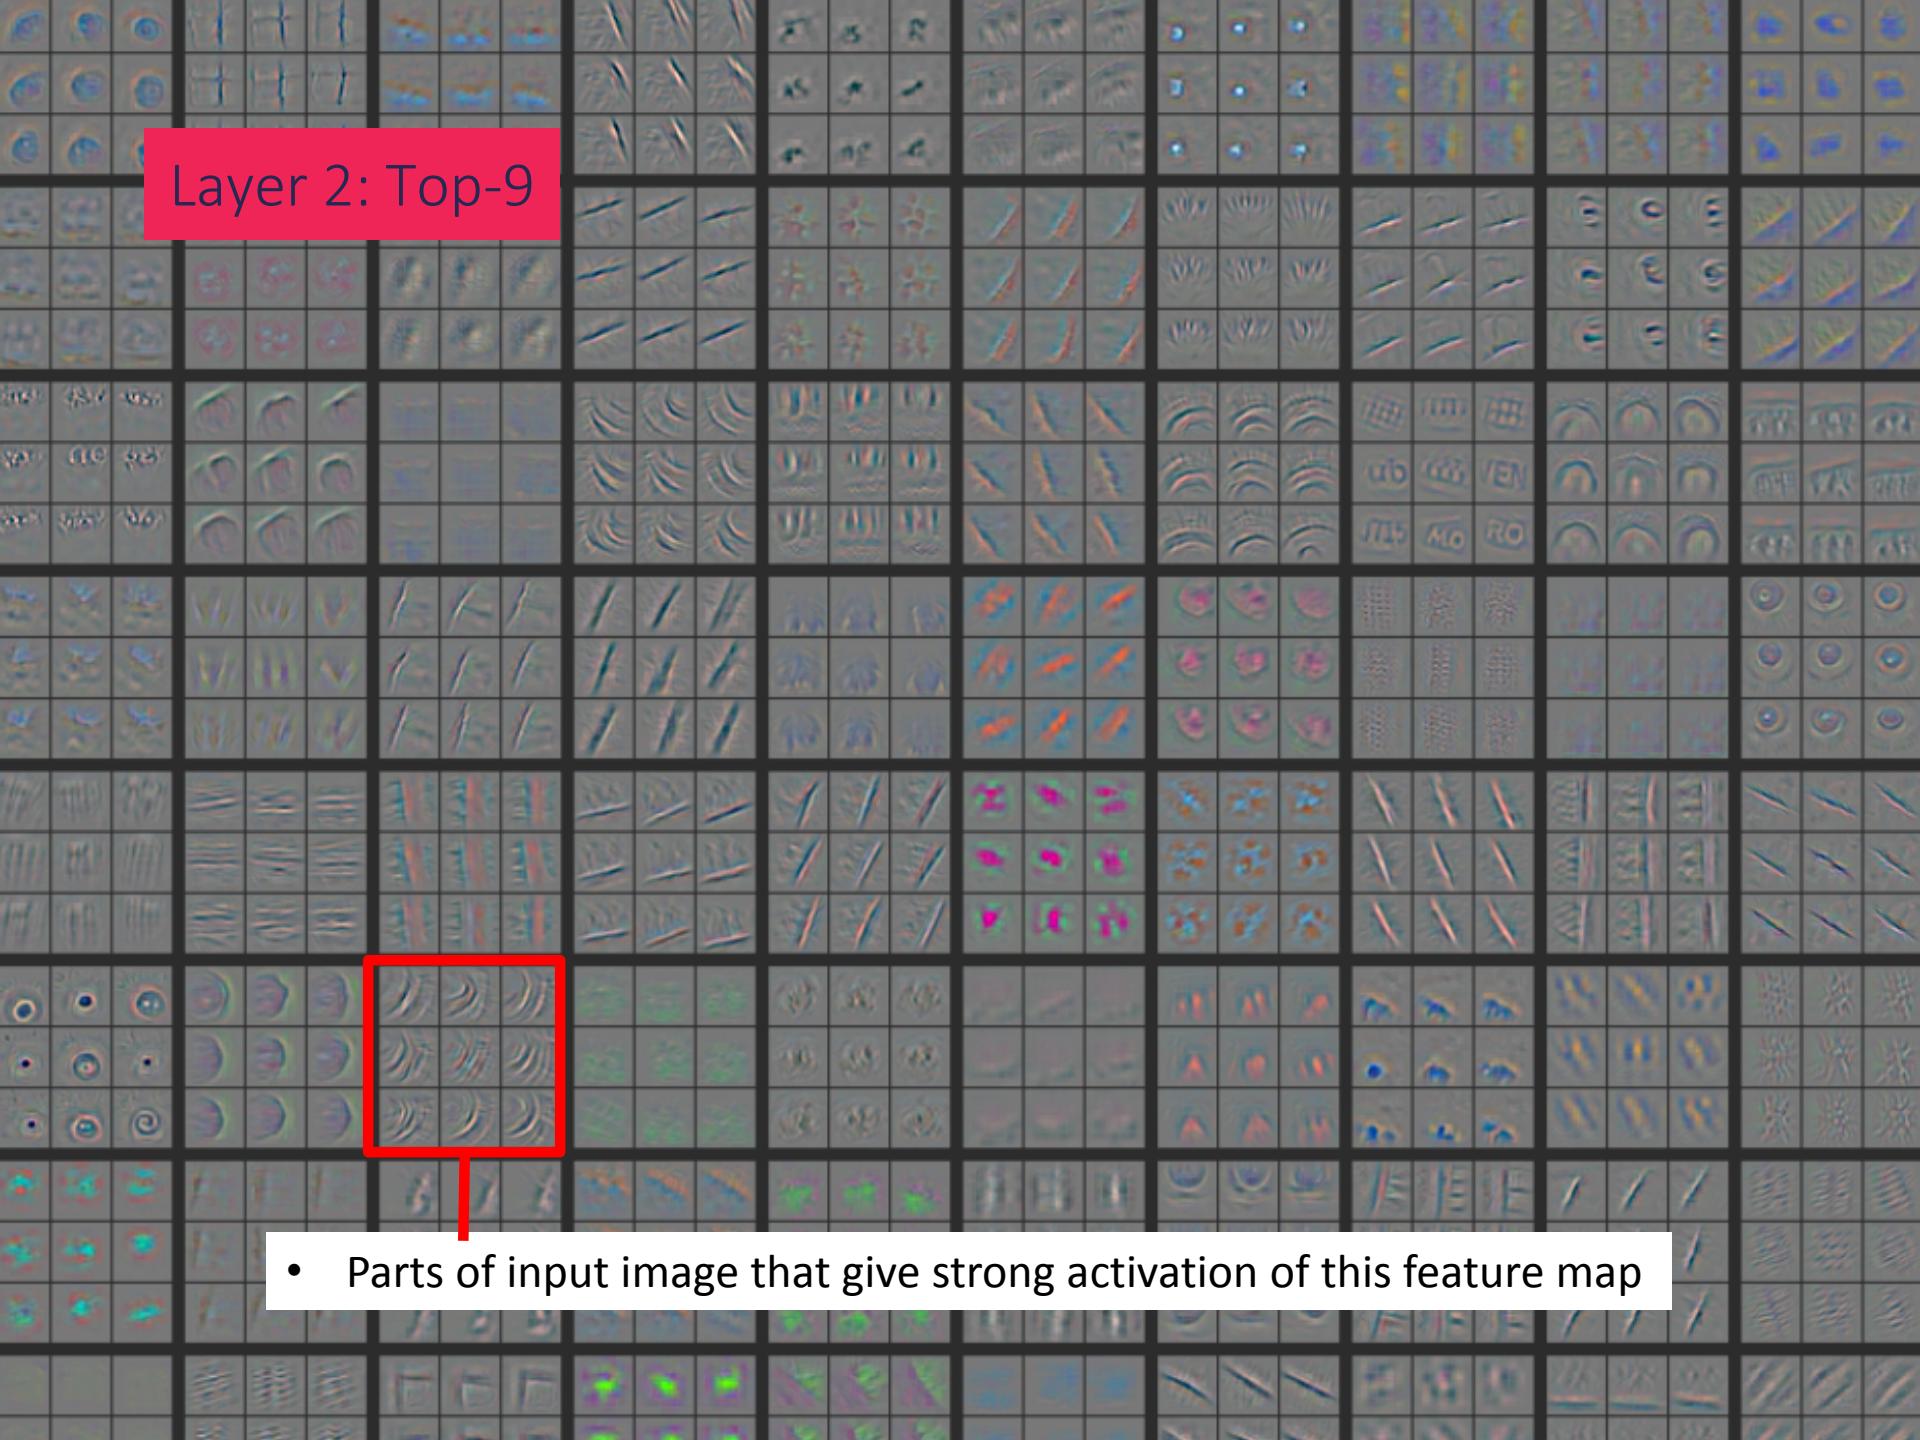
Credit: R. Fergus

Layer 2: Top-1



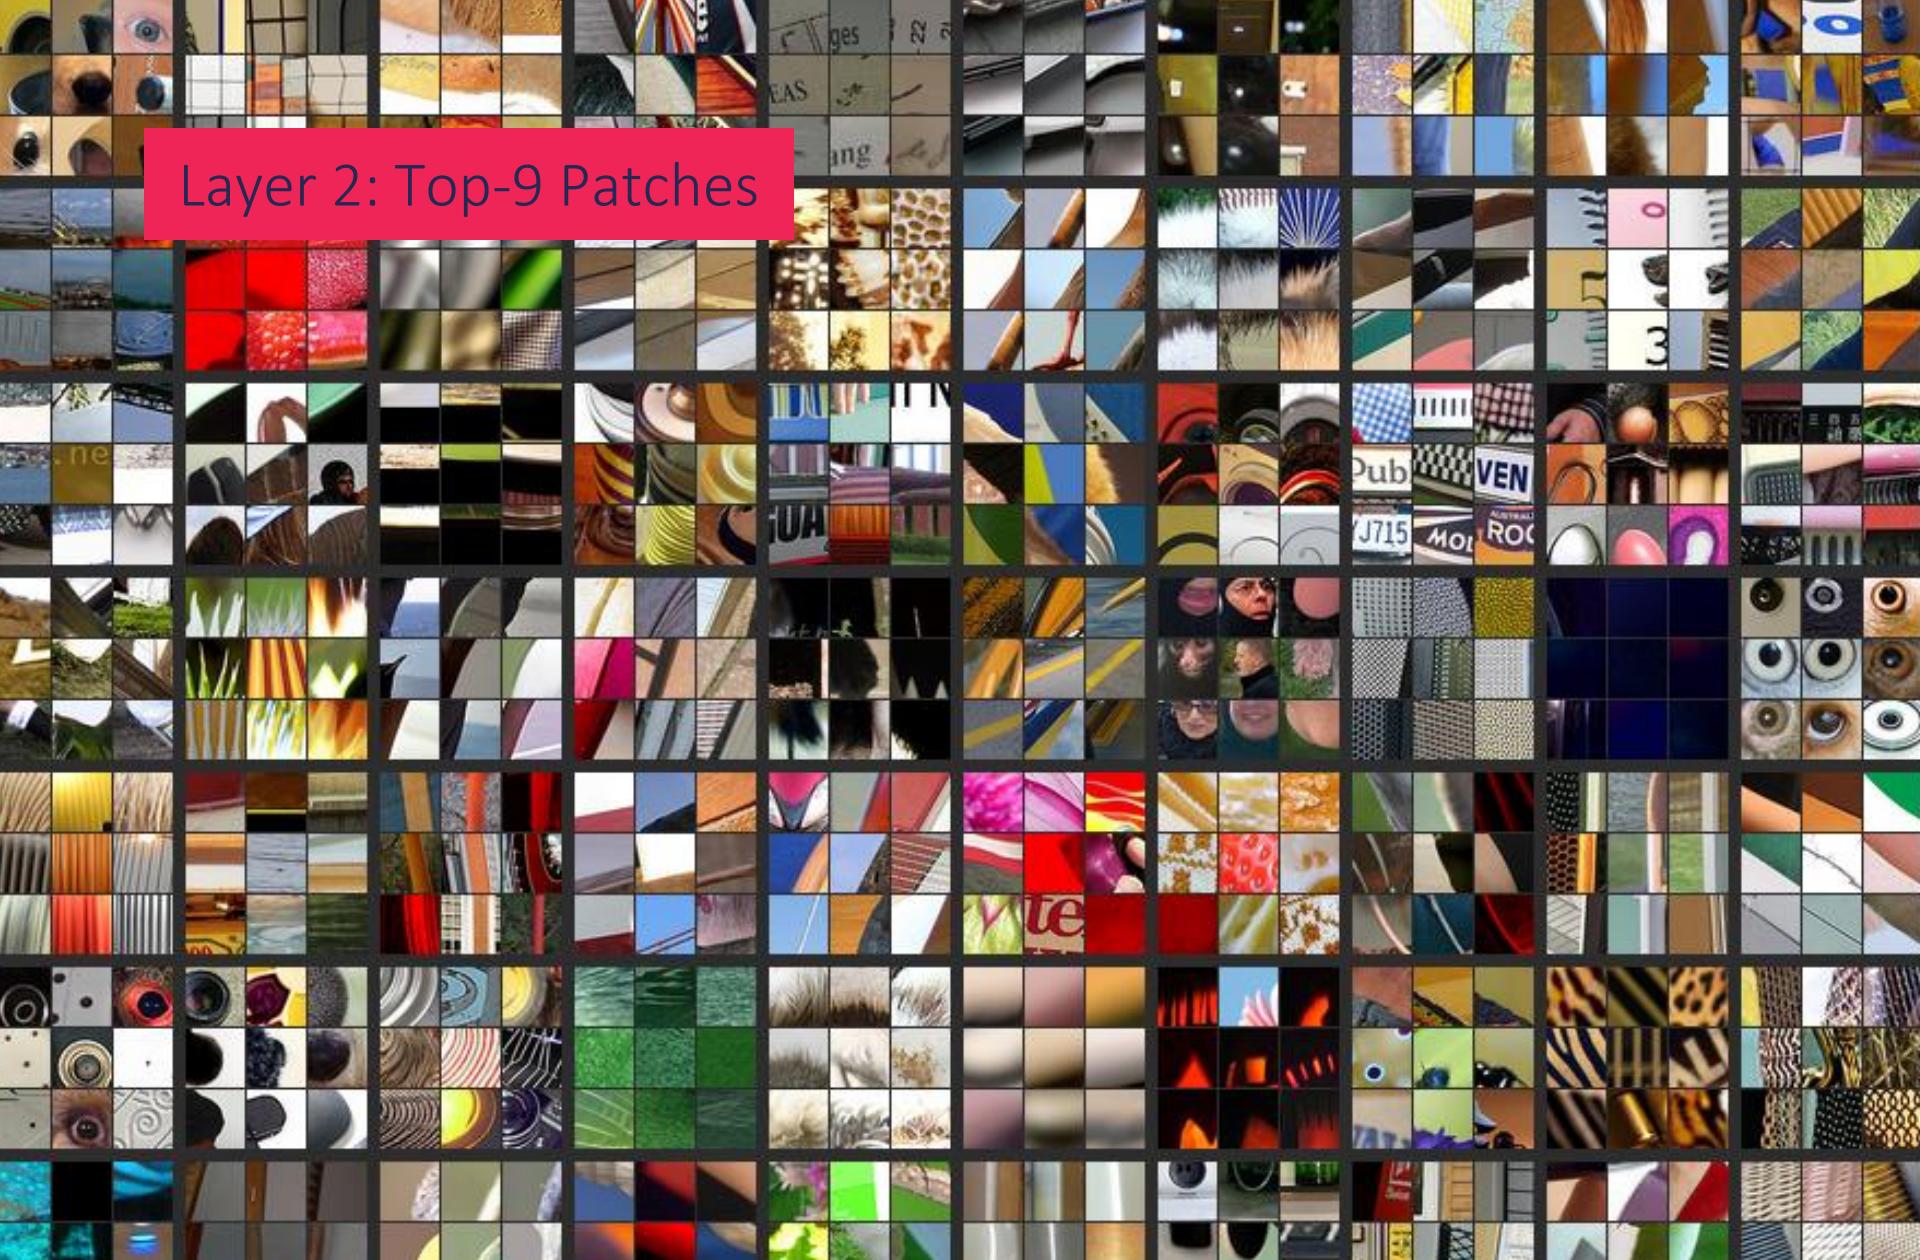
Credit: R. Fergus

Layer 2: Top-9



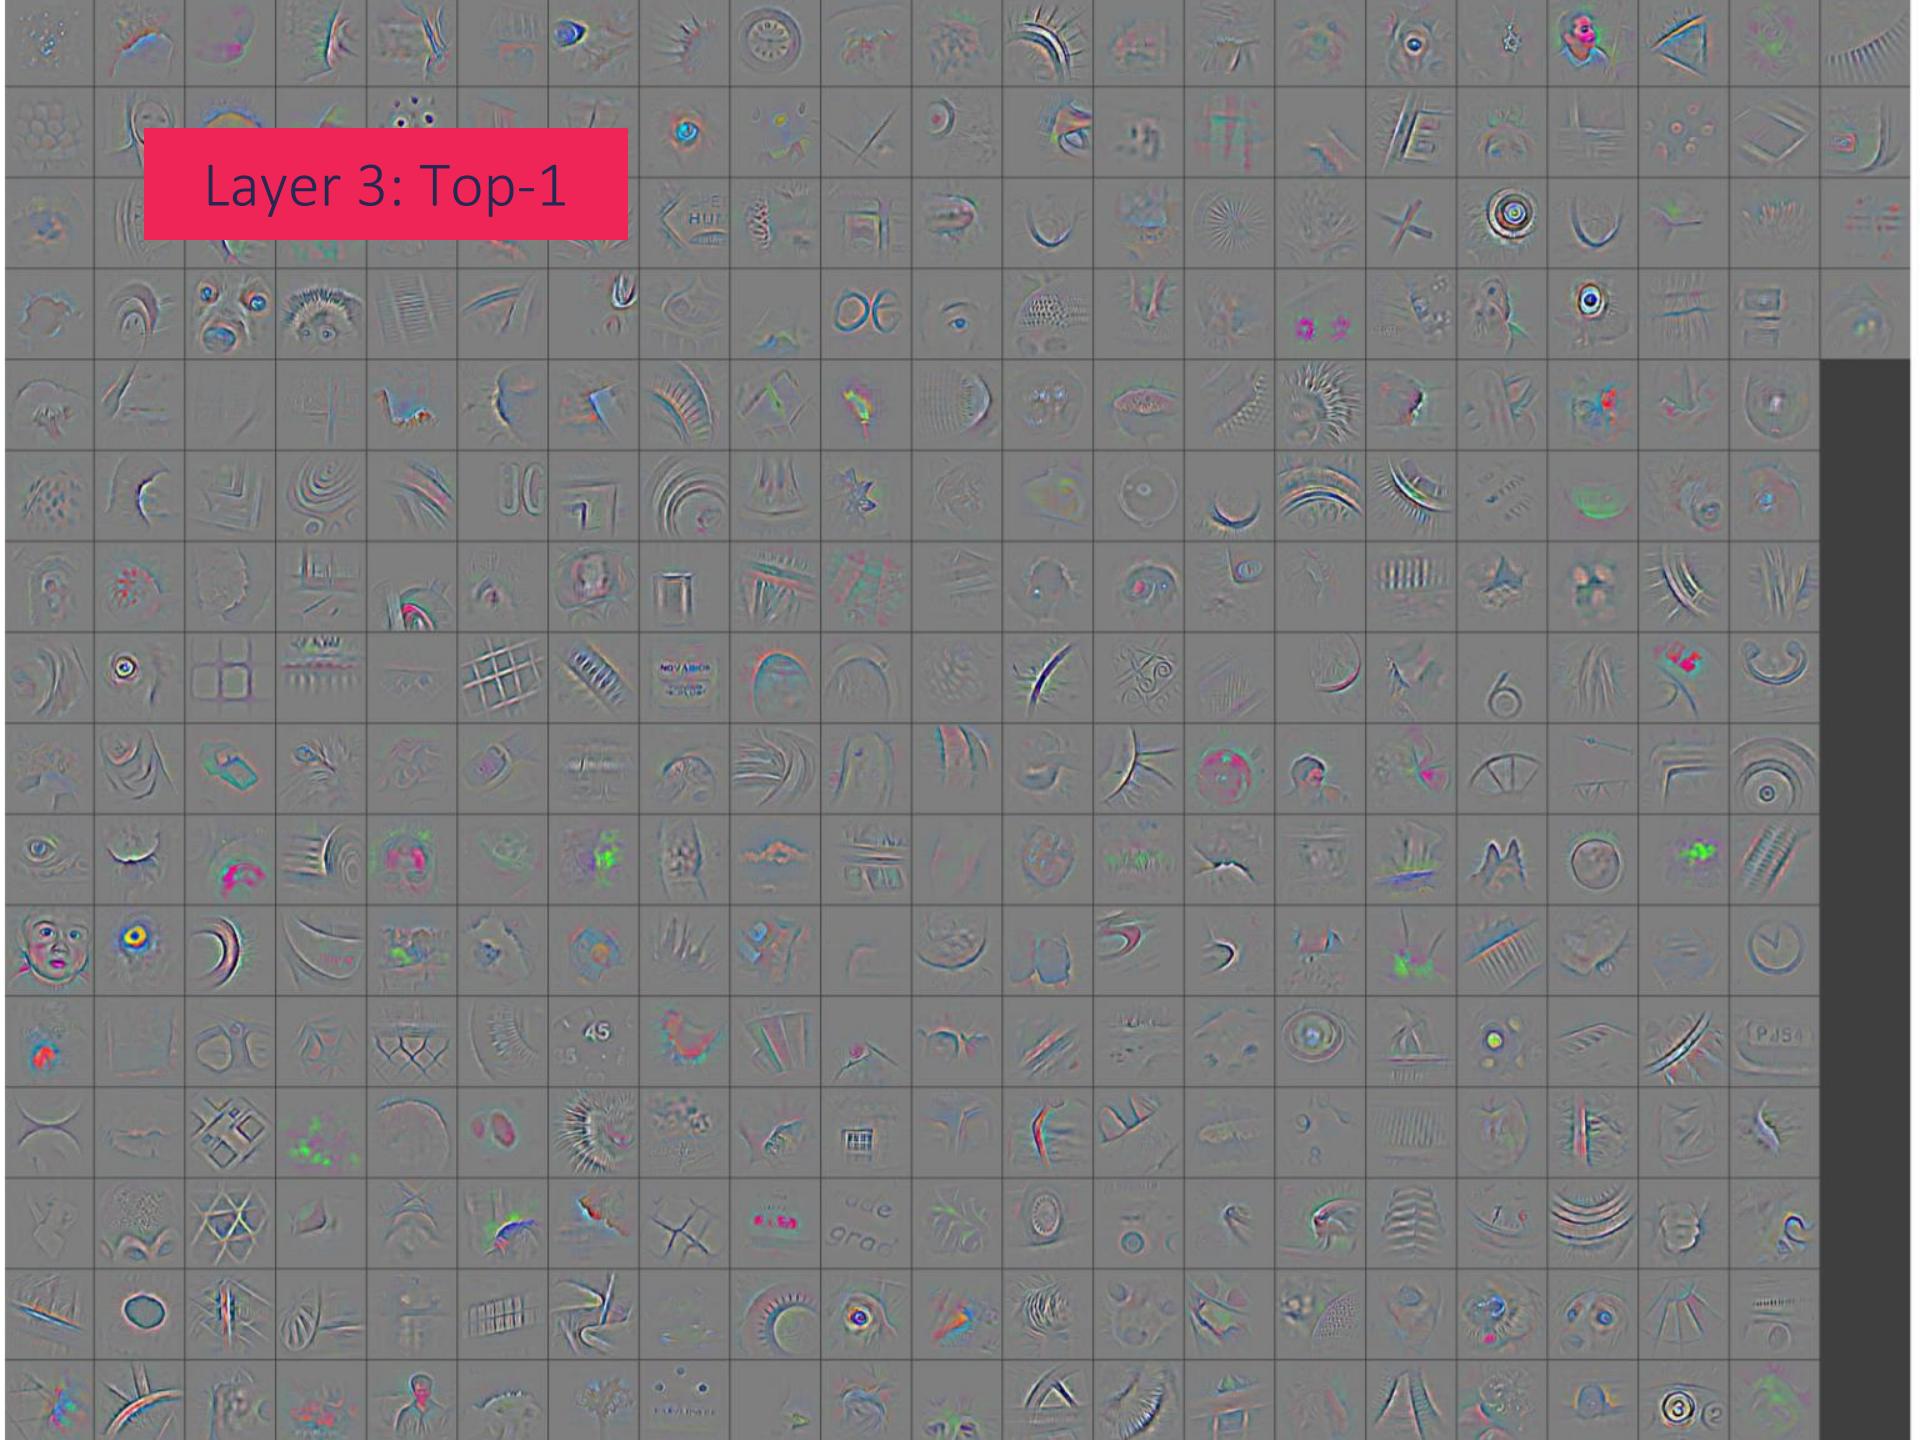
- Parts of input image that give strong activation of this feature map

Layer 2: Top-9 Patches



- Patches from validation images that give maximal activation of a given feature map

Layer 3: Top-1



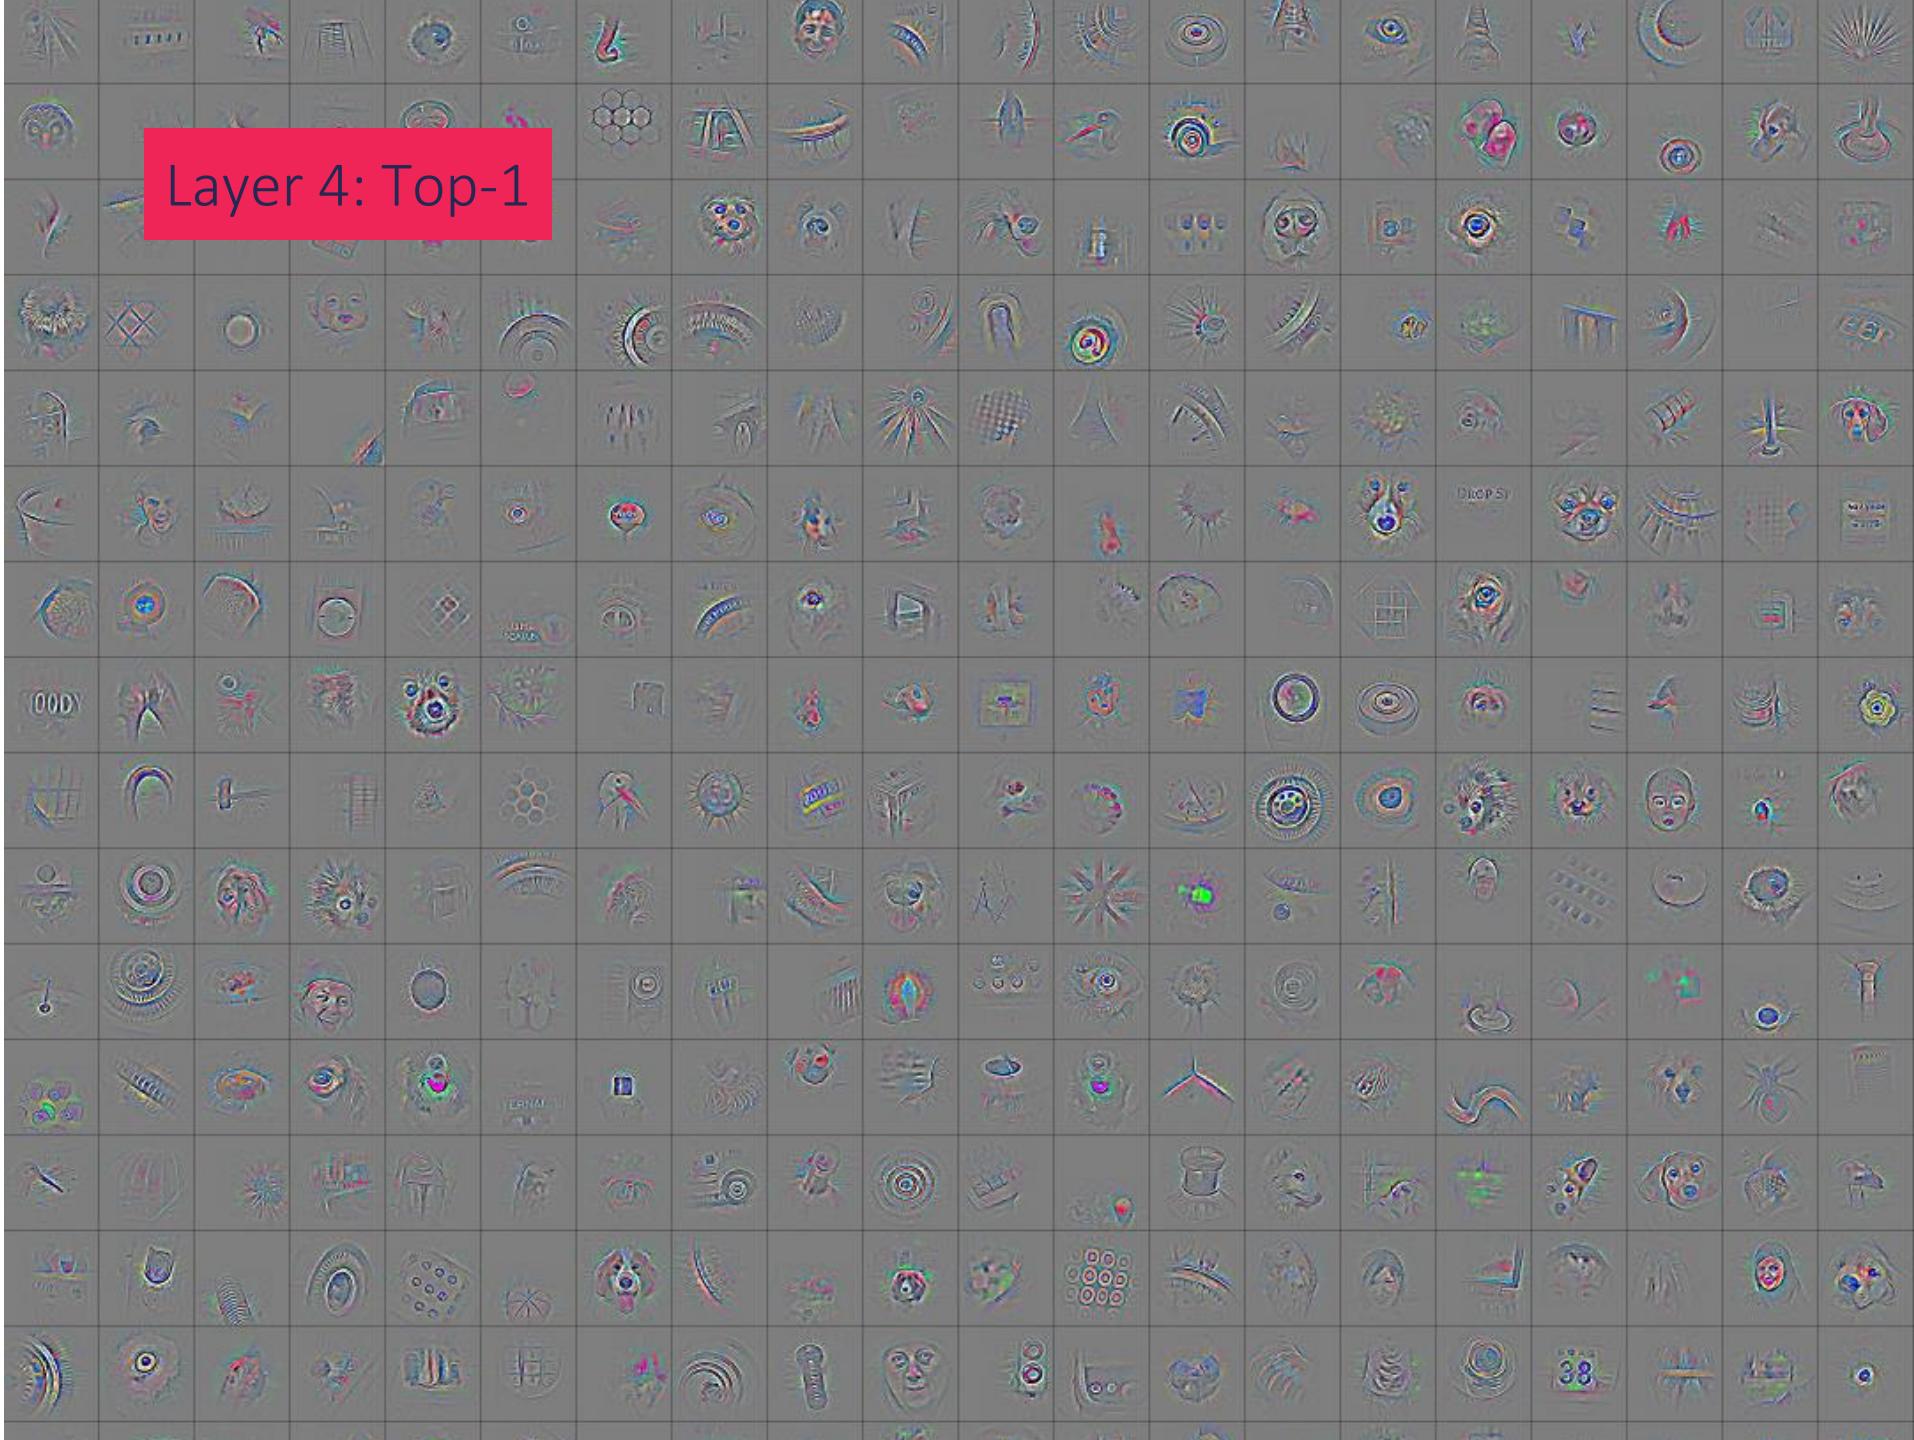
Layer 3: Top-9



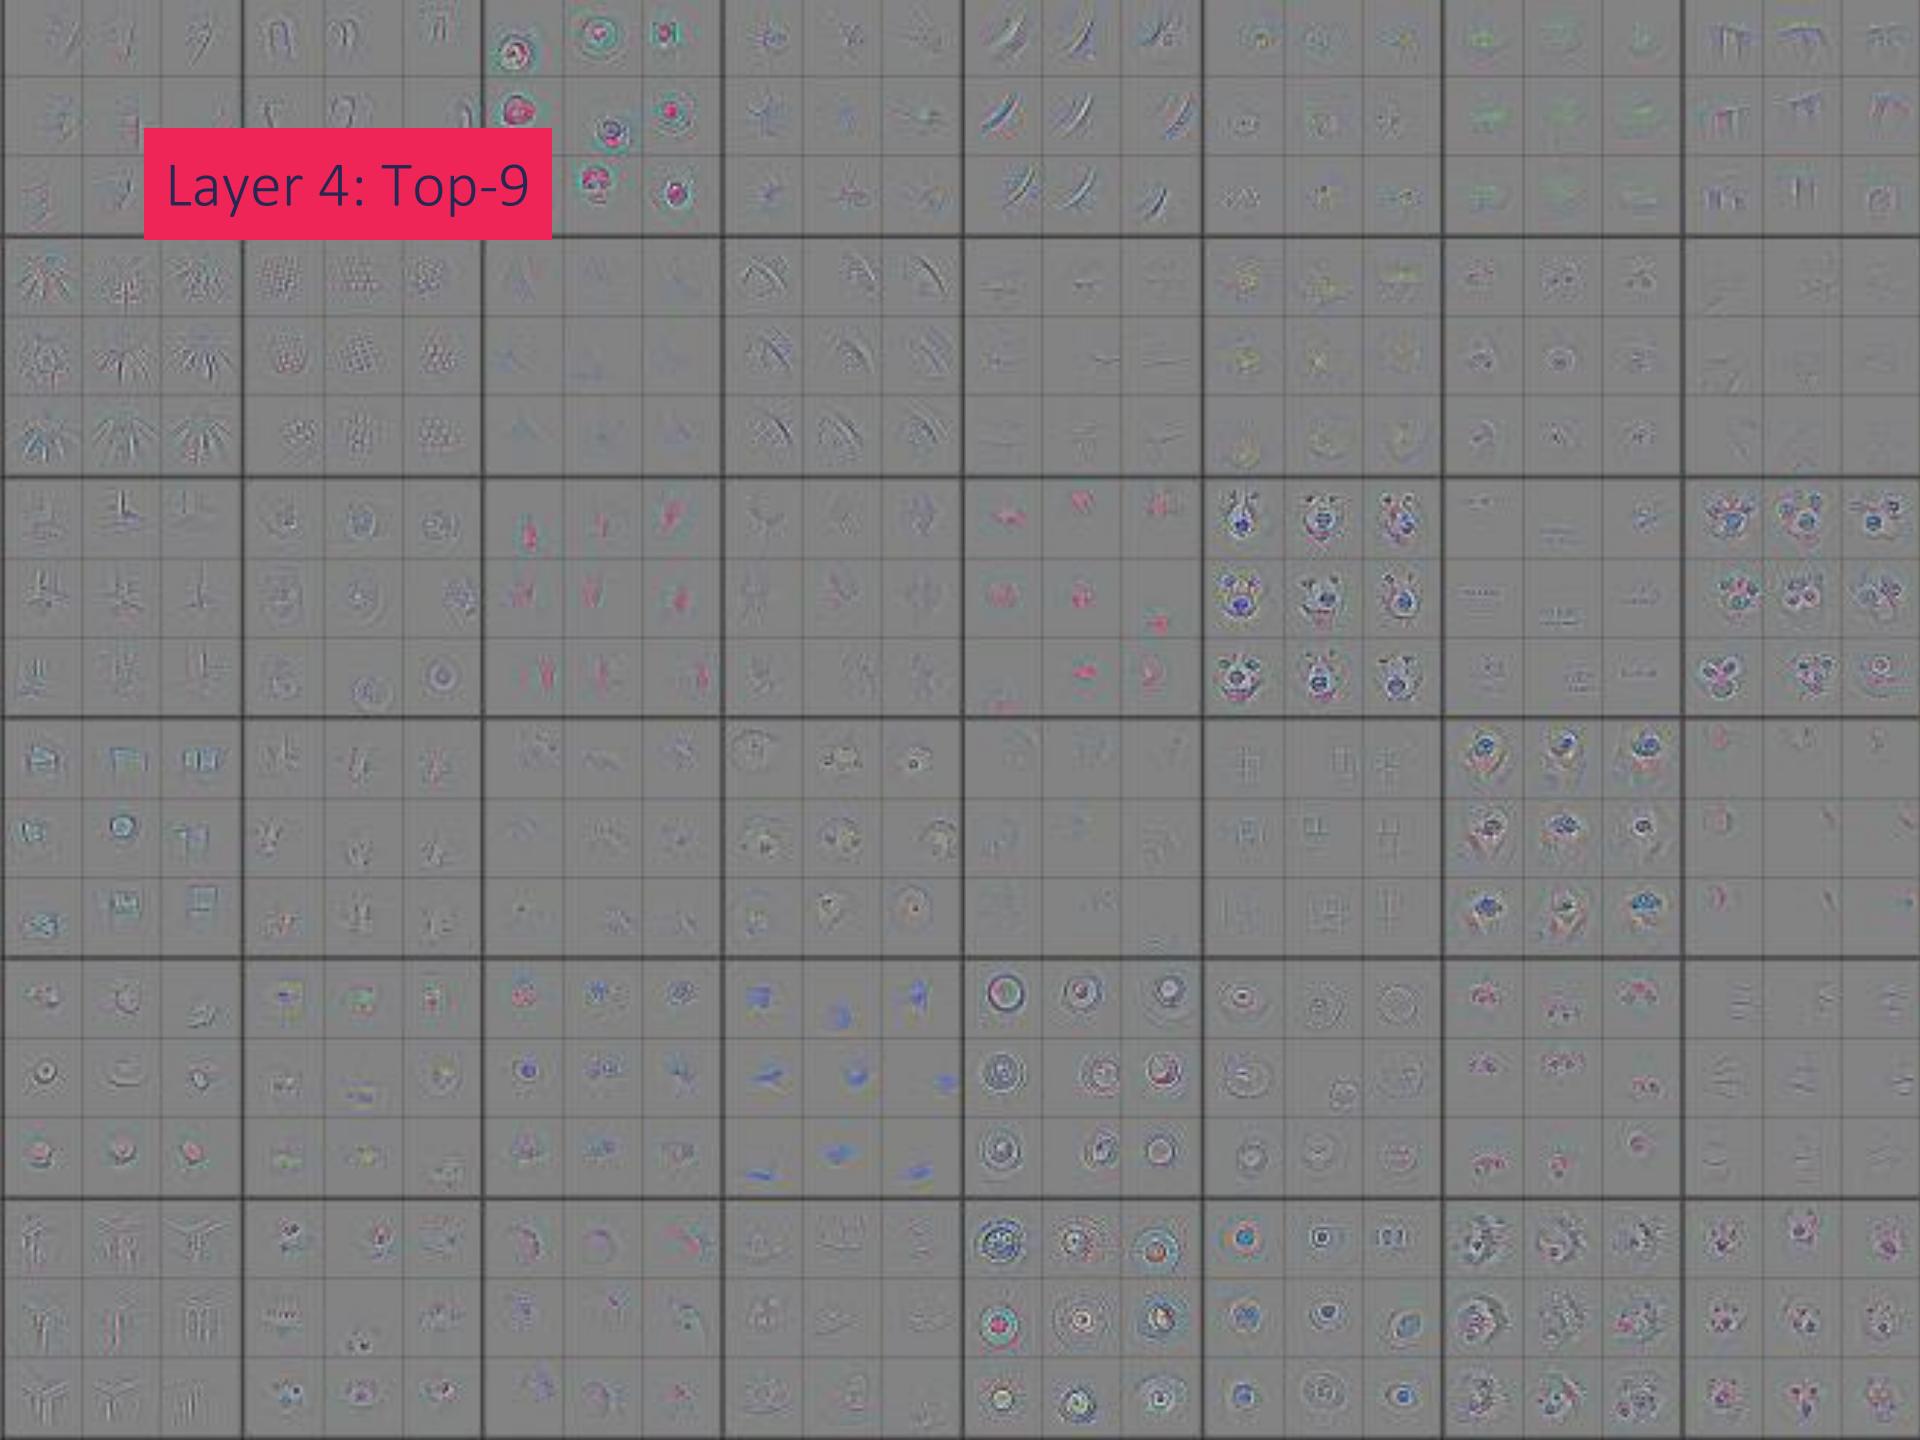
Layer 3: Top-9



Layer 4: Top-1



Layer 4: Top-9



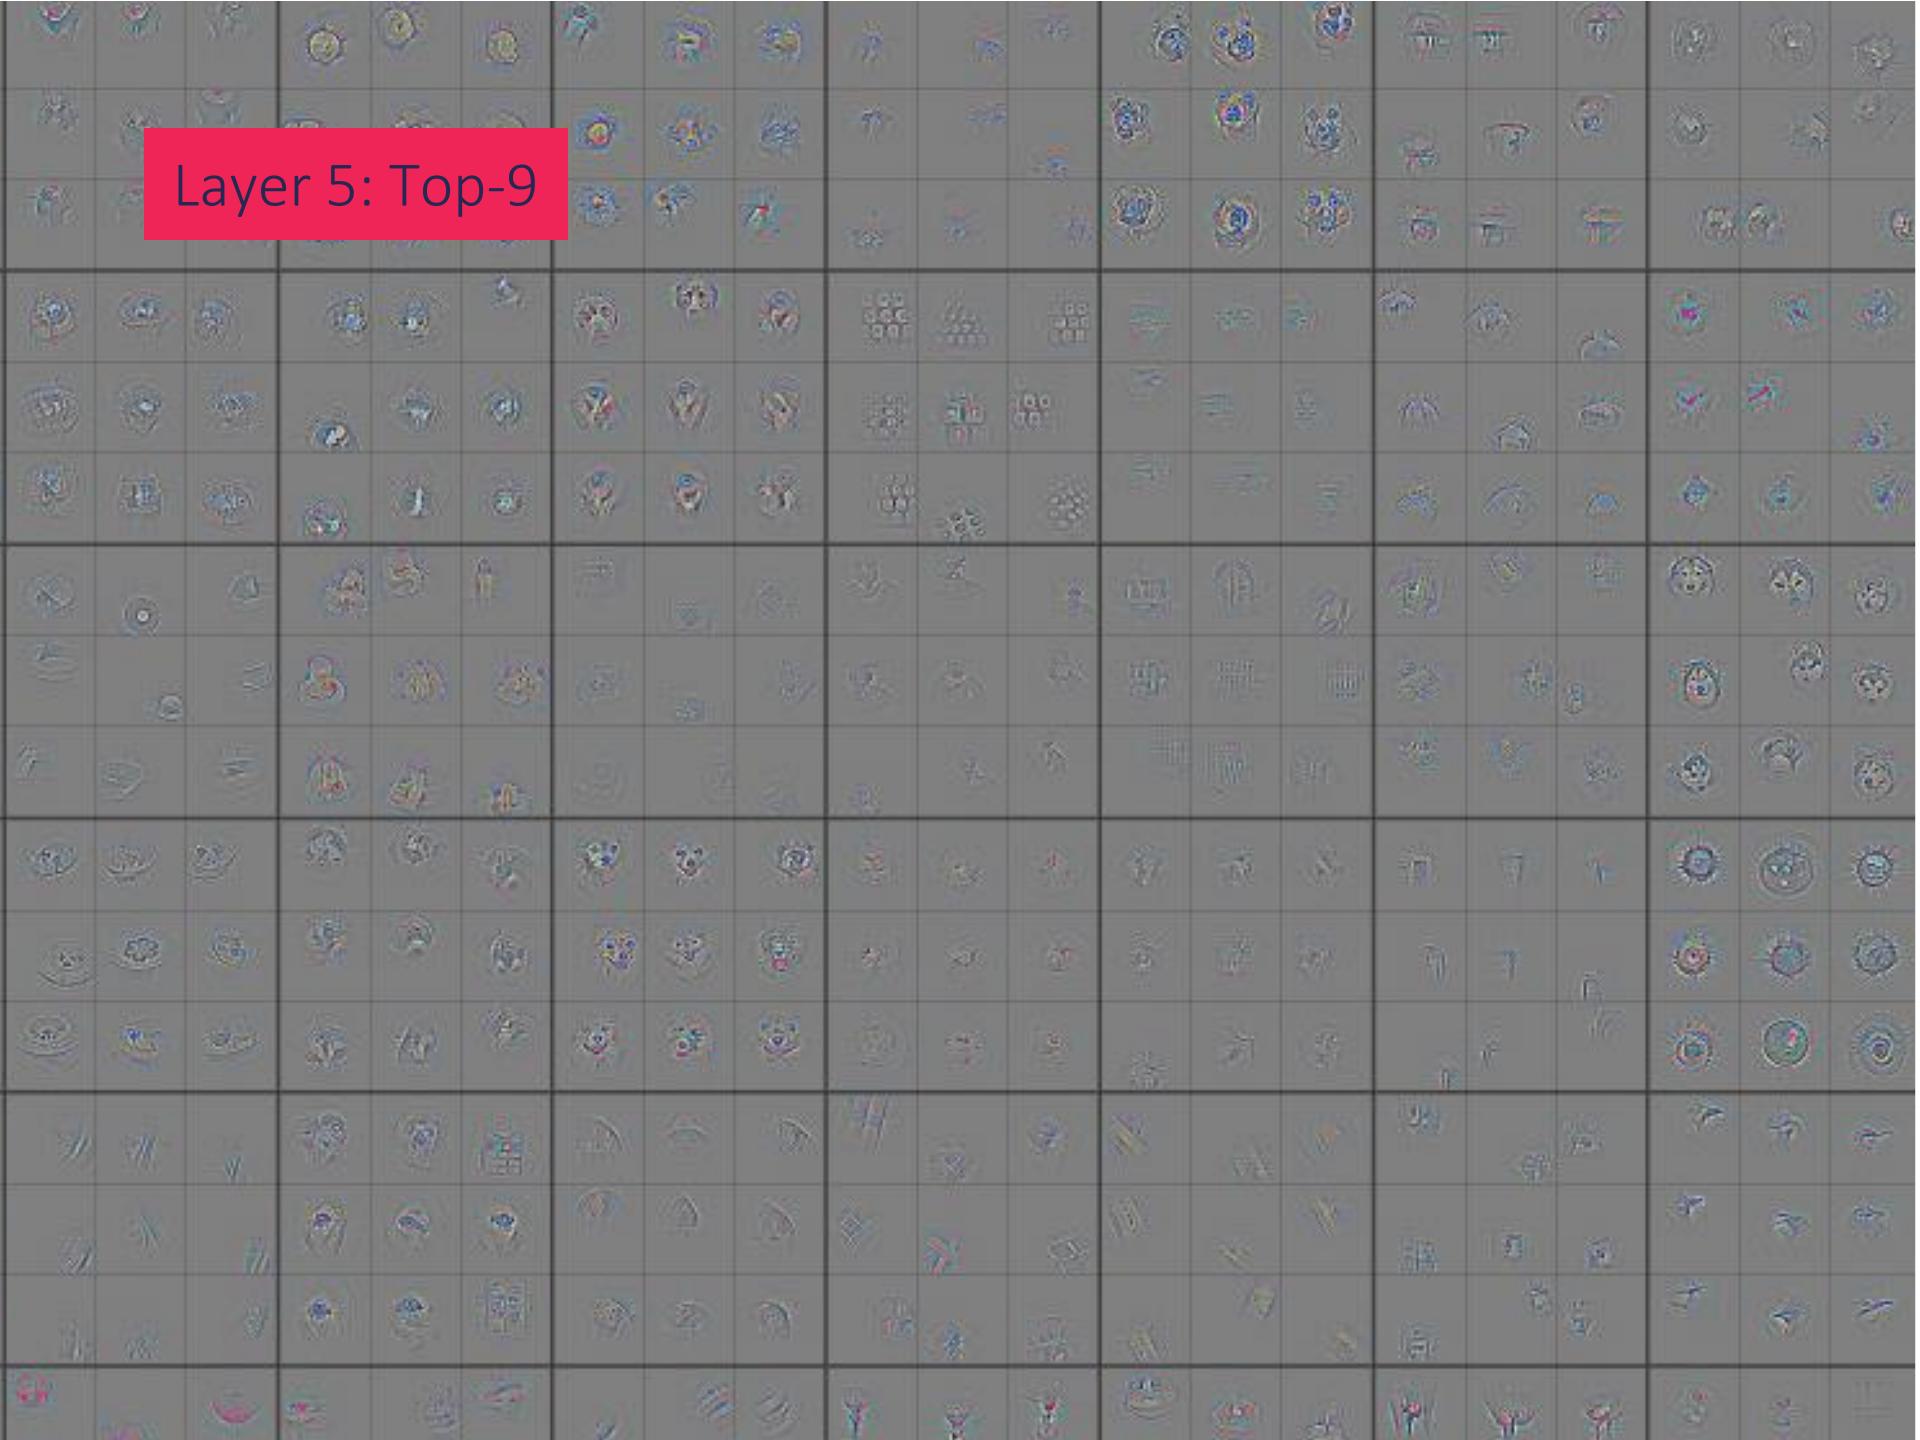
Layer 4: Top-9 patches



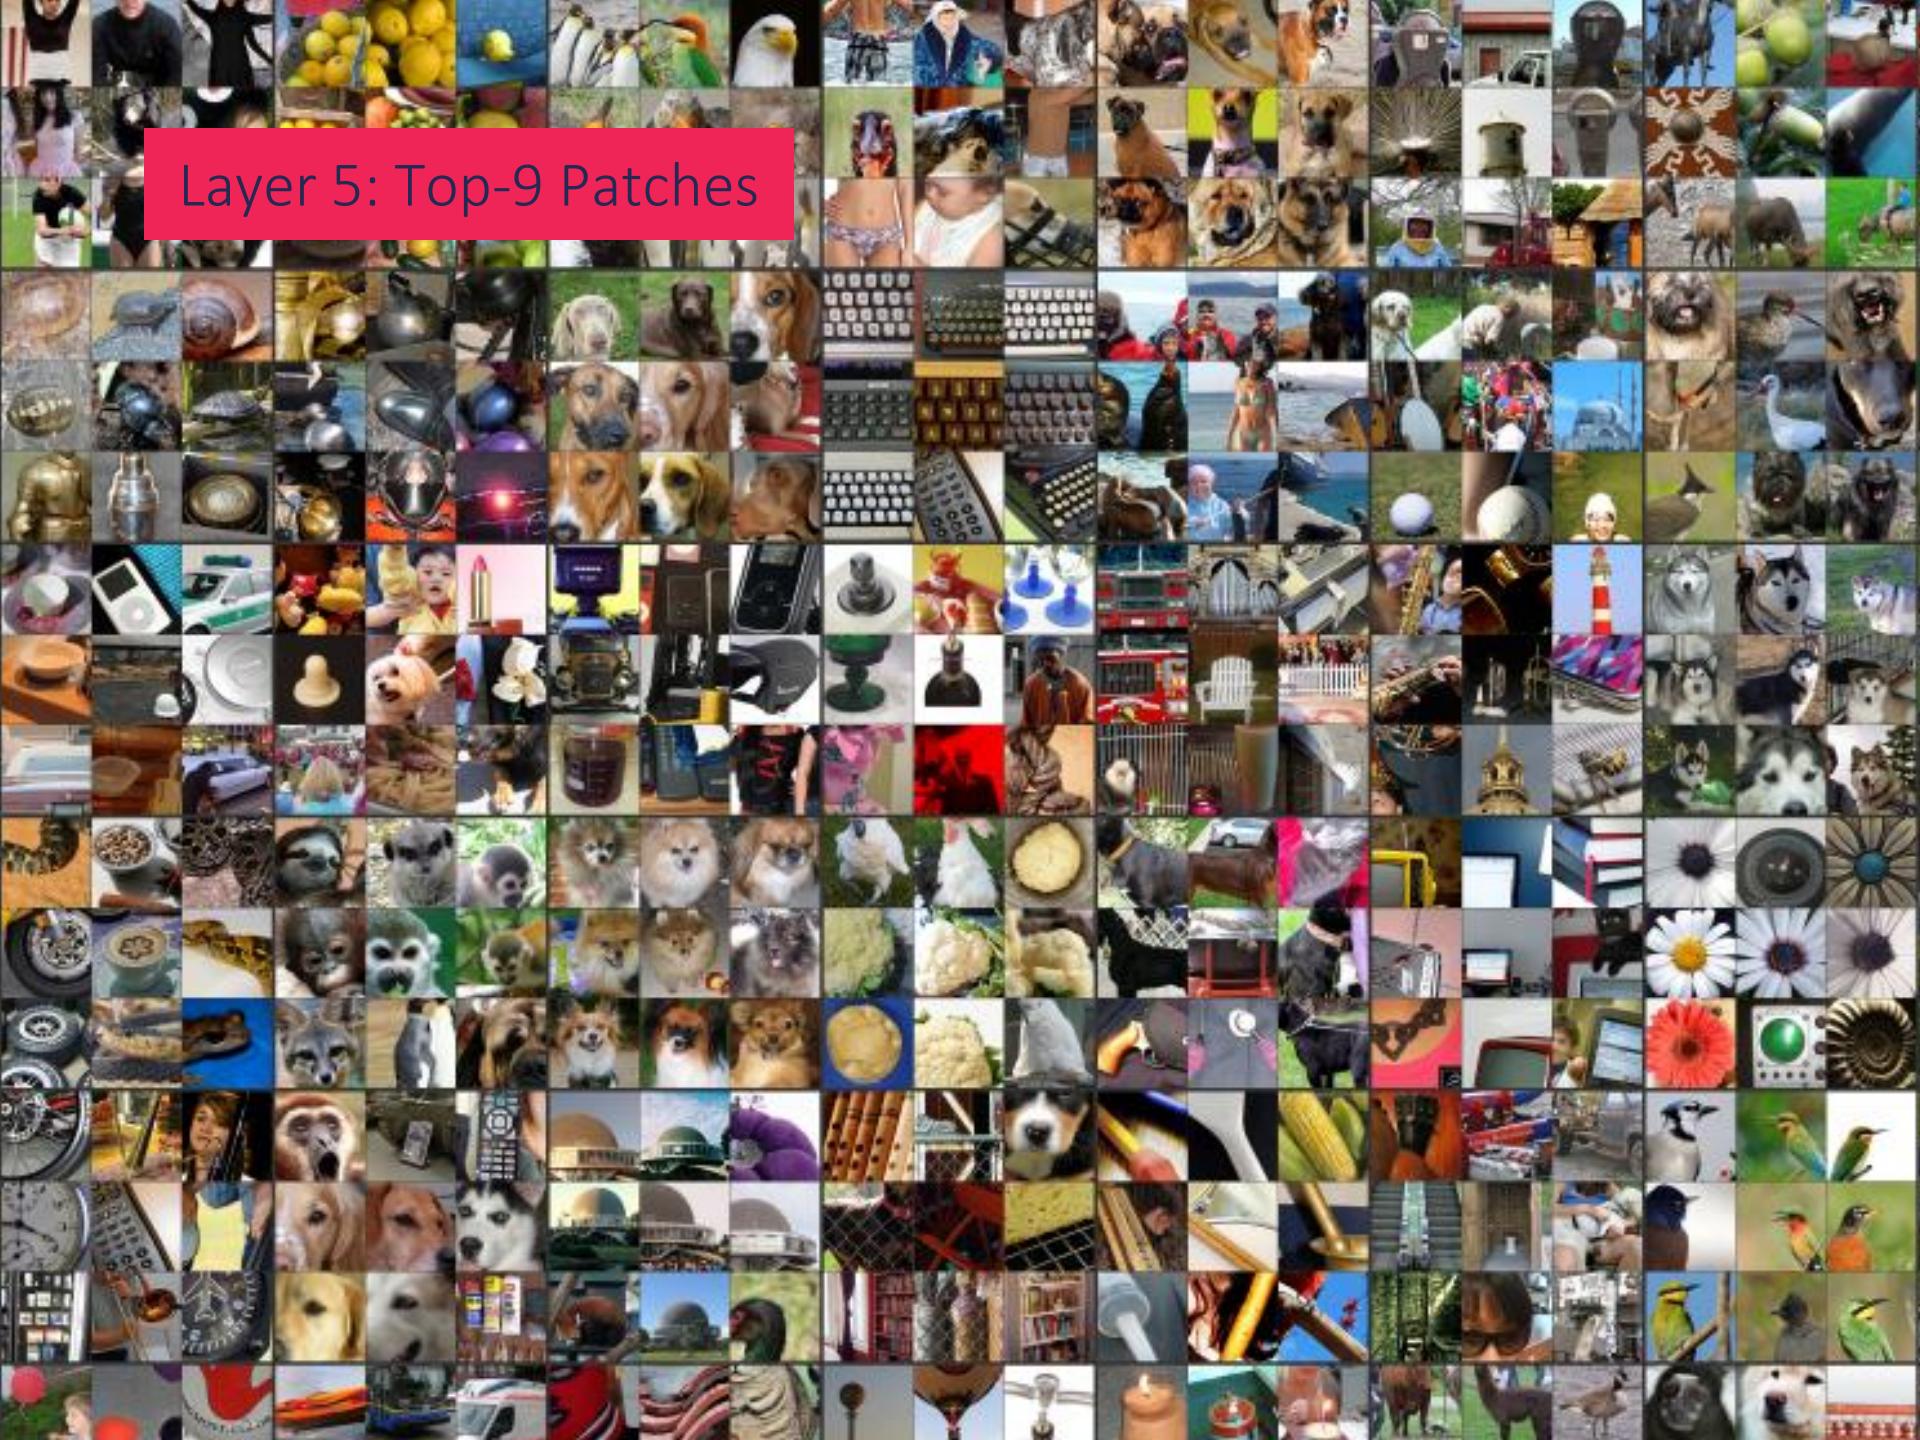
Layer 5: Top-1



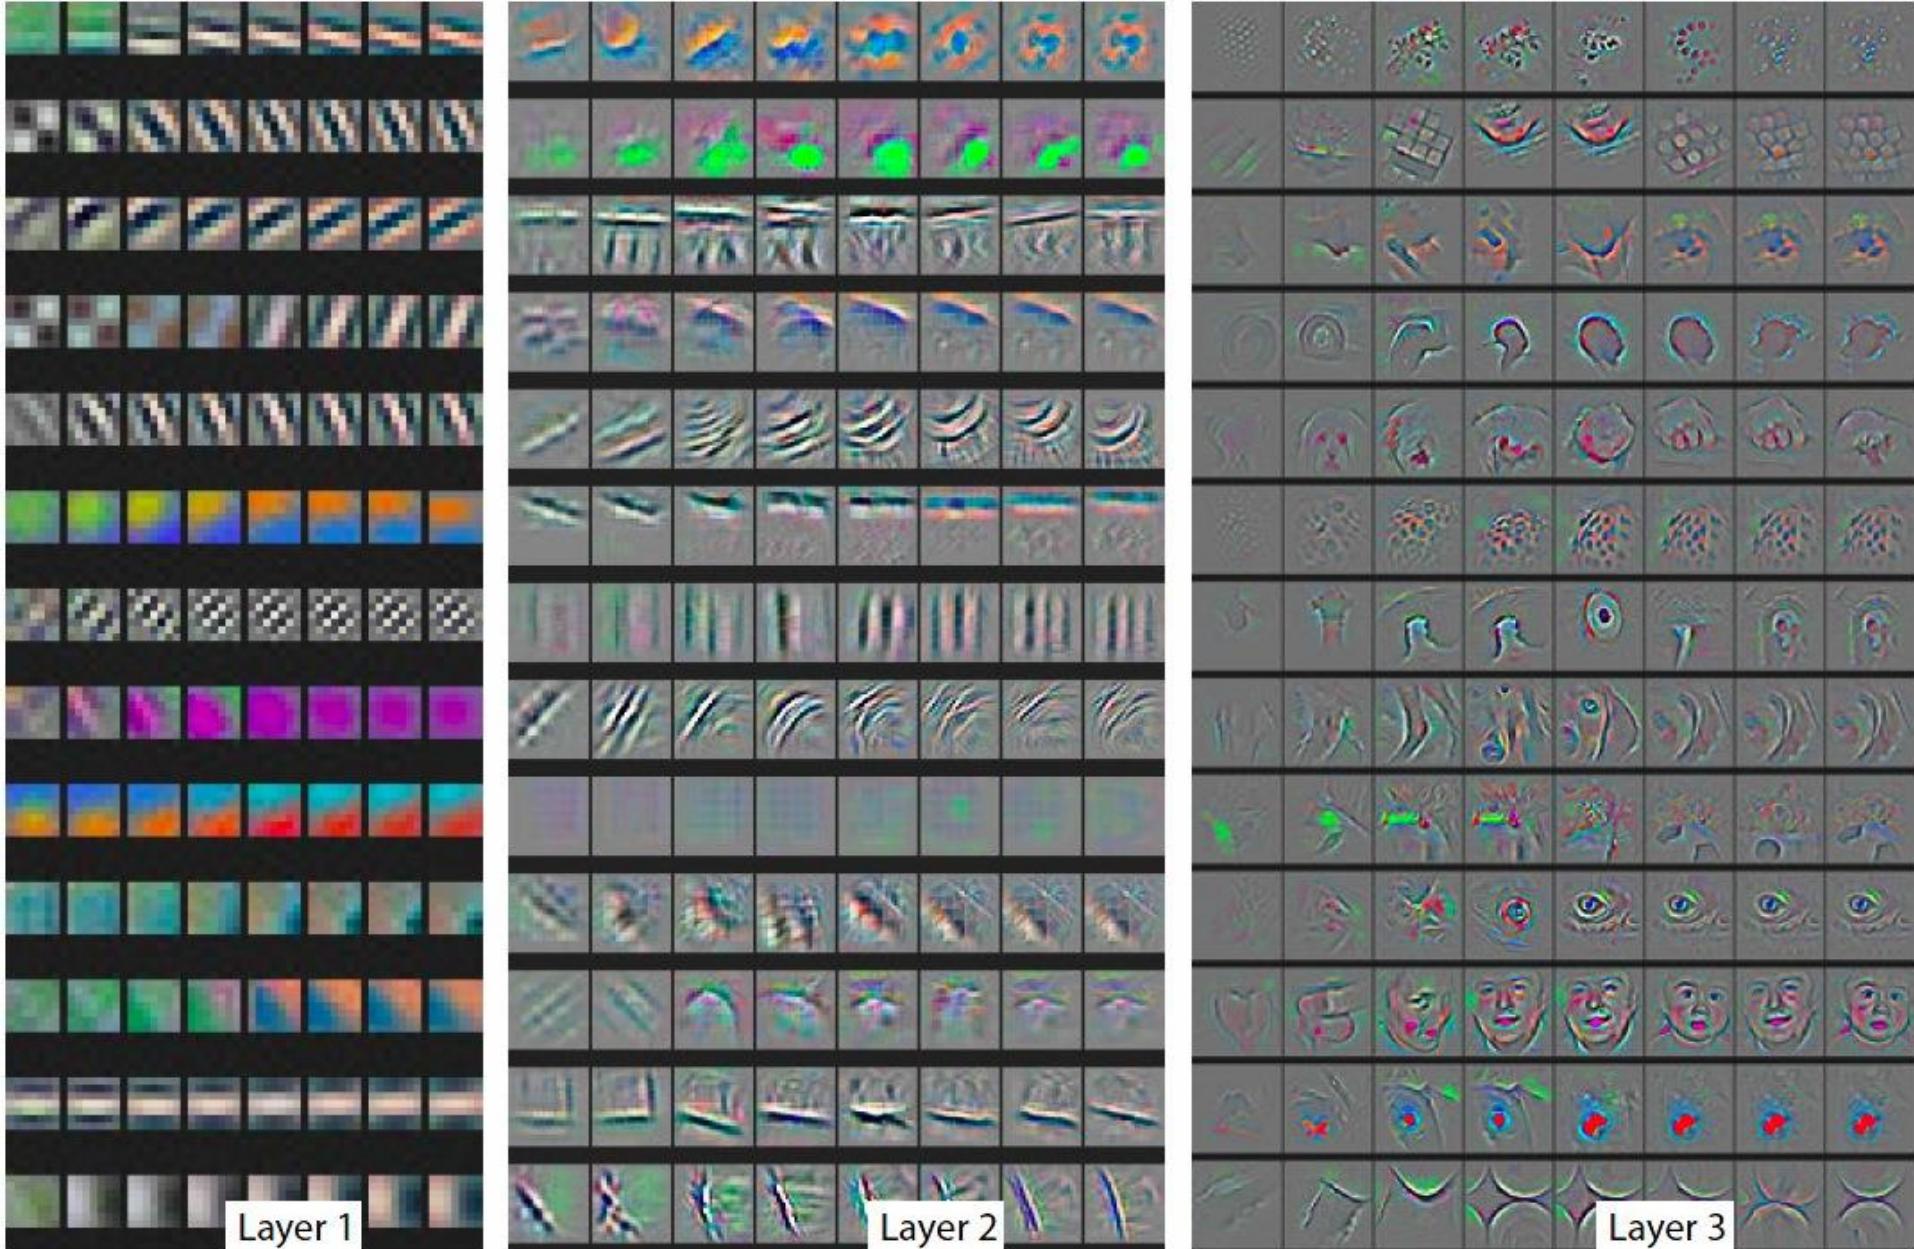
Layer 5: Top-9



Layer 5: Top-9 Patches

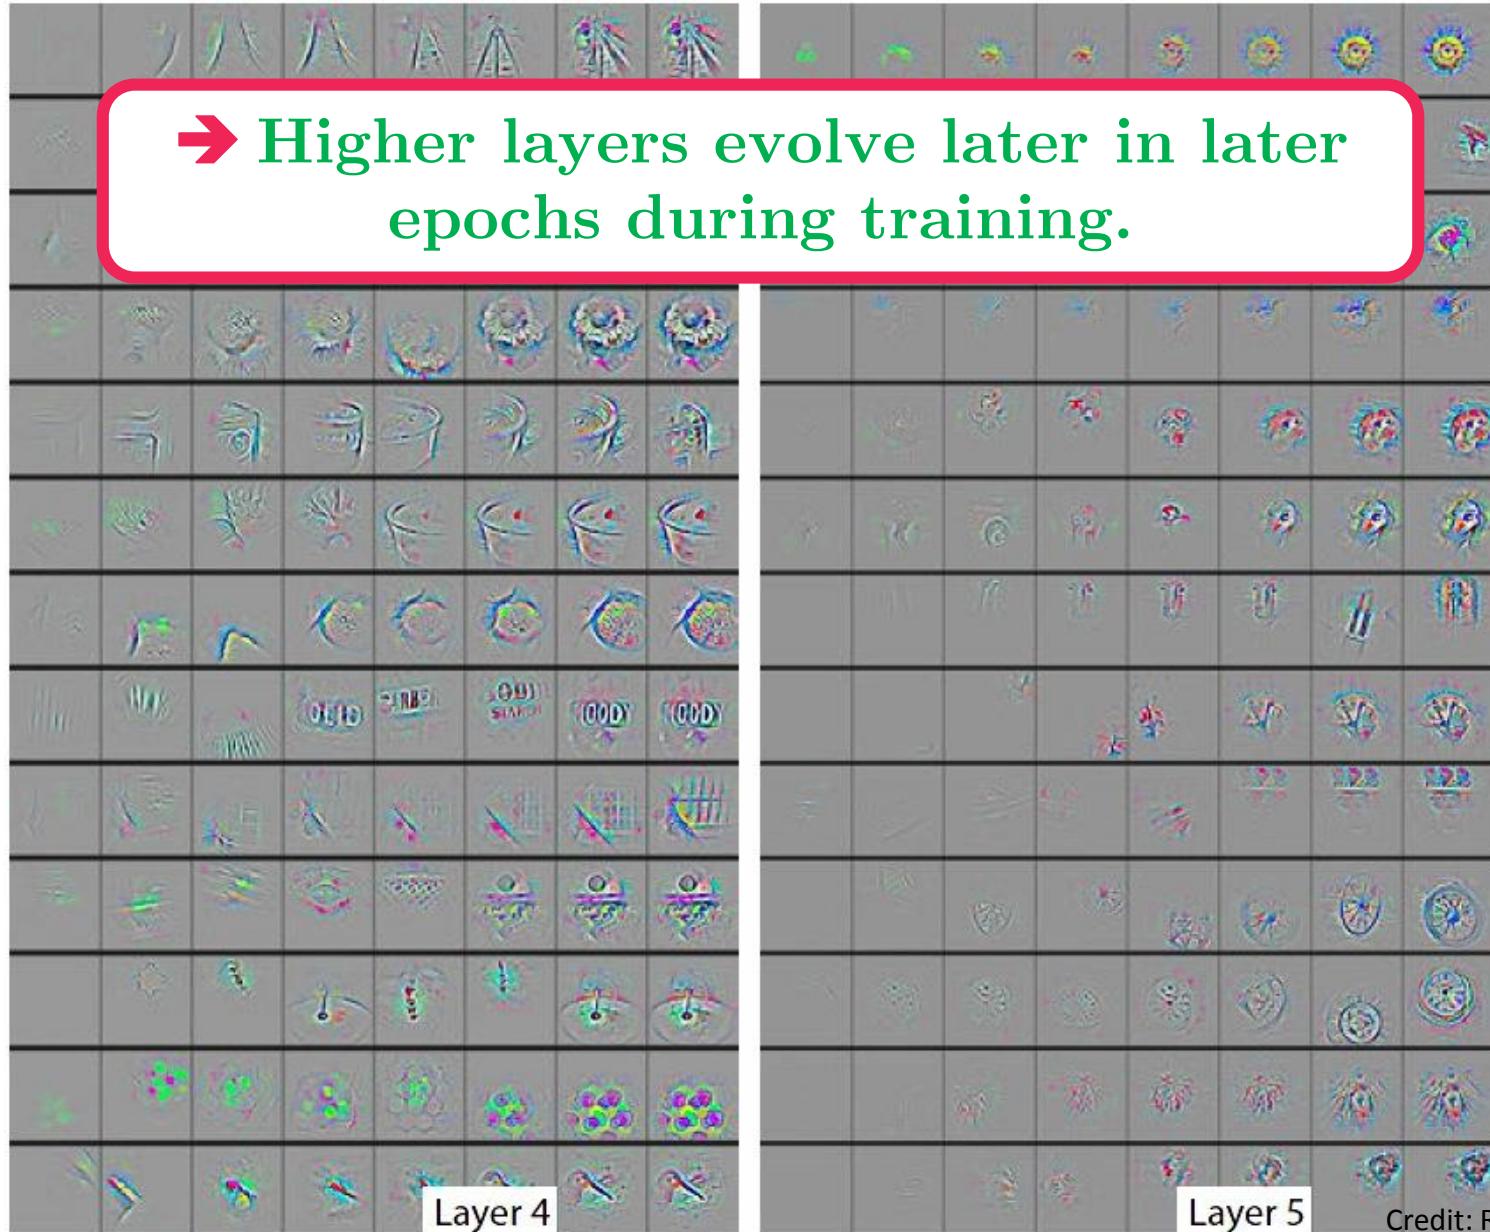


Evolution of Features During Training



Evolution of Features During Training

→ Higher layers evolve later in later epochs during training.

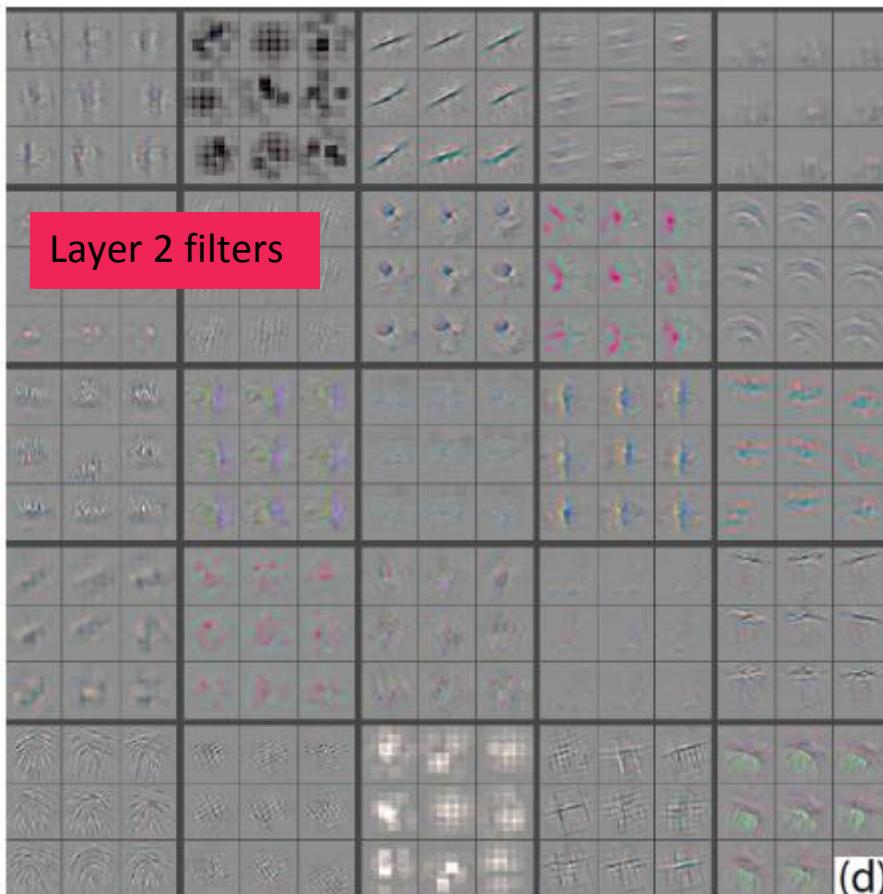


Credit: R. Fergus

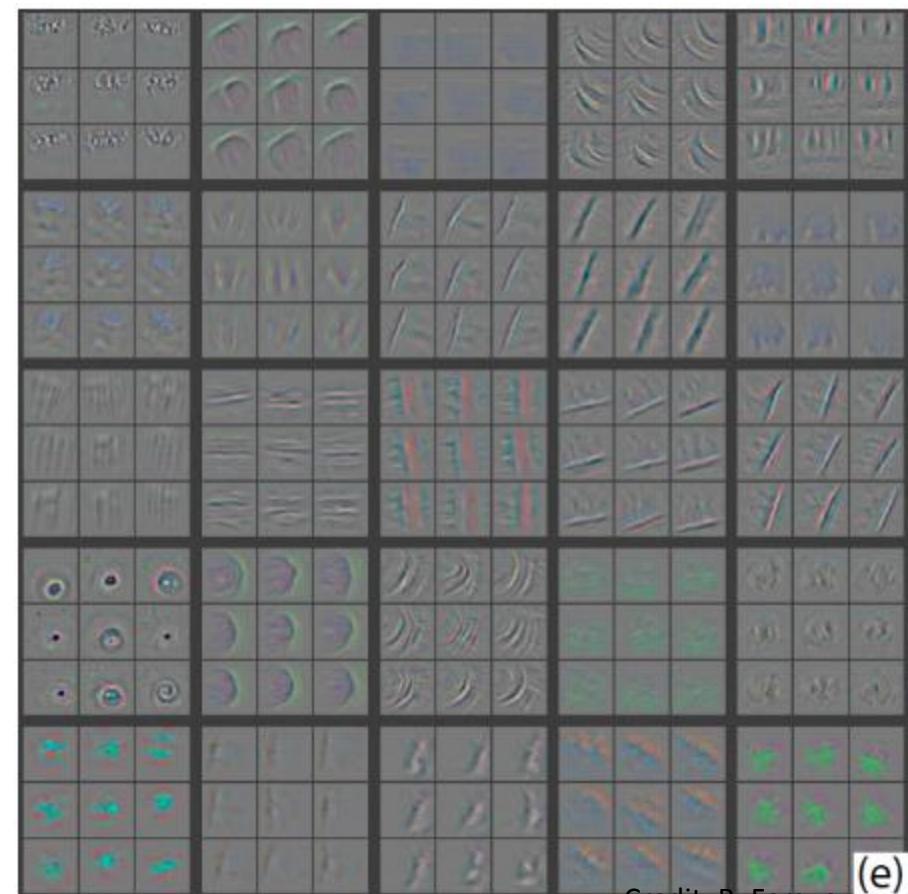
Visualizations can be used to improve model

- Visualization of Krizhevsky et al.'s architecture showed some problems with layers 1 and 2
- Alter architecture: smaller stride & filter size
 - Visualizations look better and Performance improves

Blocking artifacts



Smaller stride for convolution

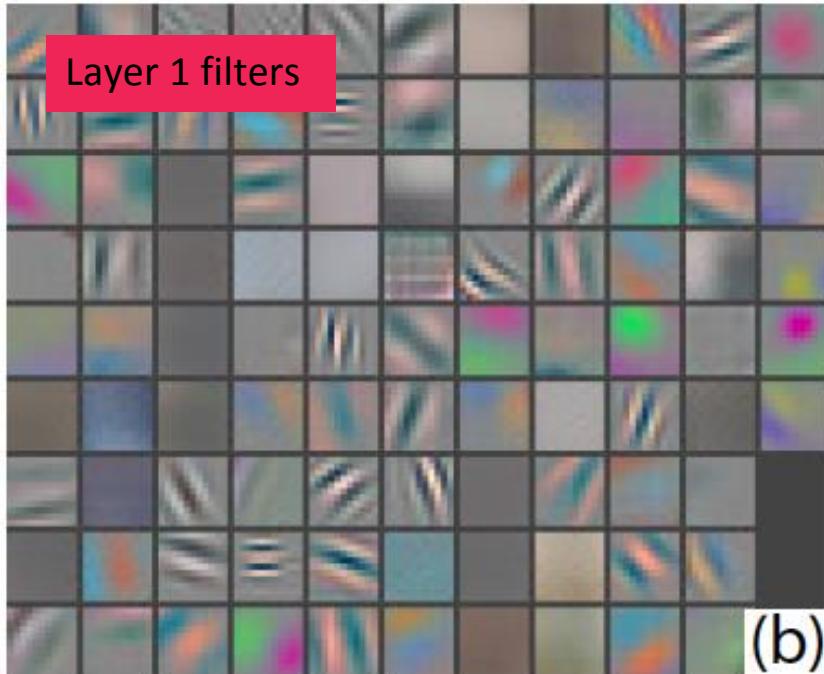


Credit: R. Fergus

Visualizations can be used to improve model

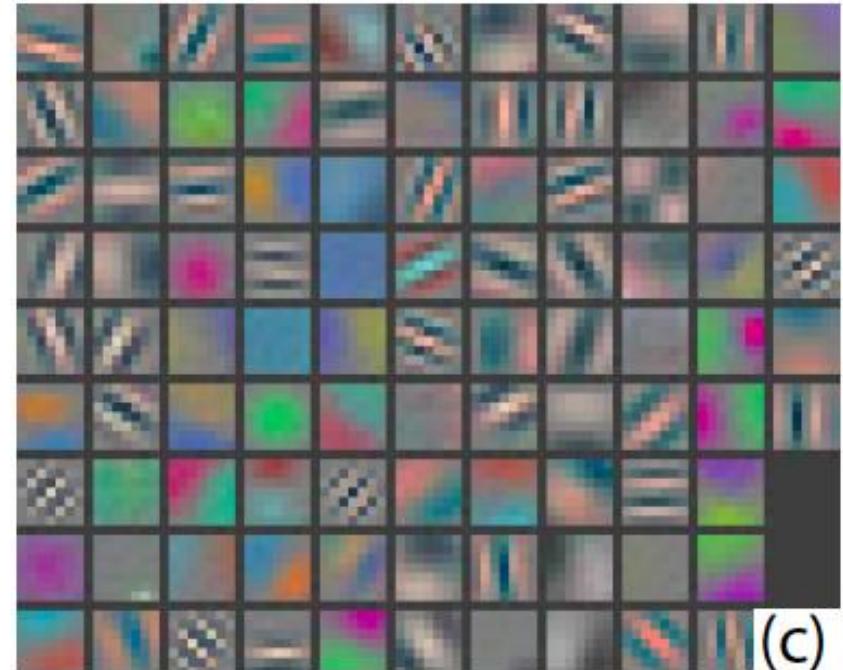
- Visualization of Krizhevsky et al.'s architecture showed some problems with layers 1 and 2
- Alter architecture: smaller stride & filter size
 - Visualizations look better and Performance improves

Too specific for low-level
+ dead filters



11x11 filters, stride 4

Restrict size (smaller)



7x7 filters, stride 2

Occlusion Experiment

- Mask parts of input with occluding square

- Monitor output
of classification
network

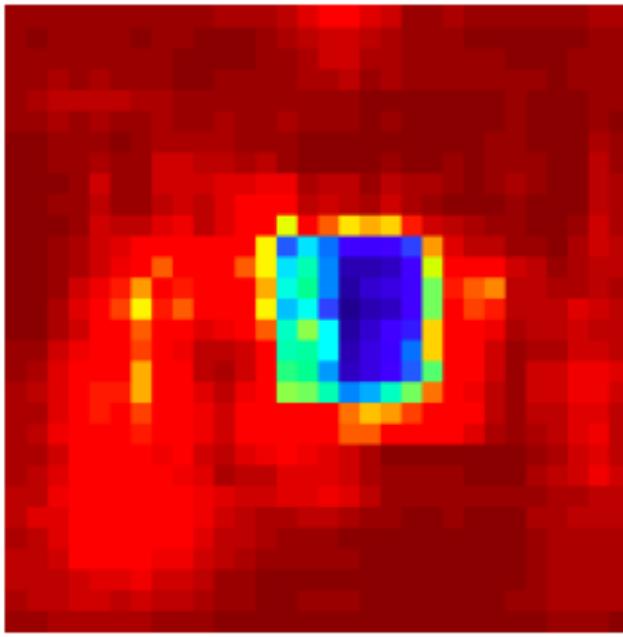


- Perhaps network using scene context?

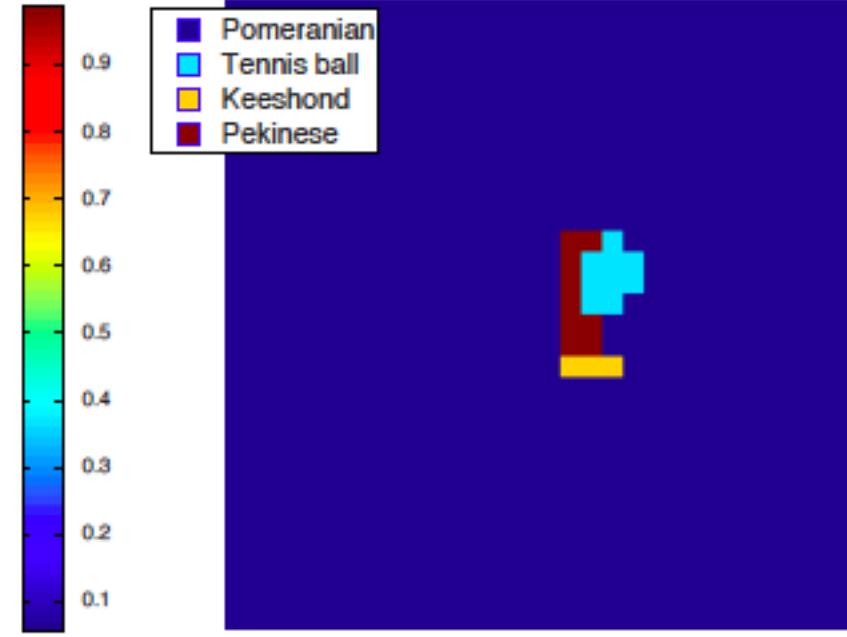
Input image



$p(\text{True class})$



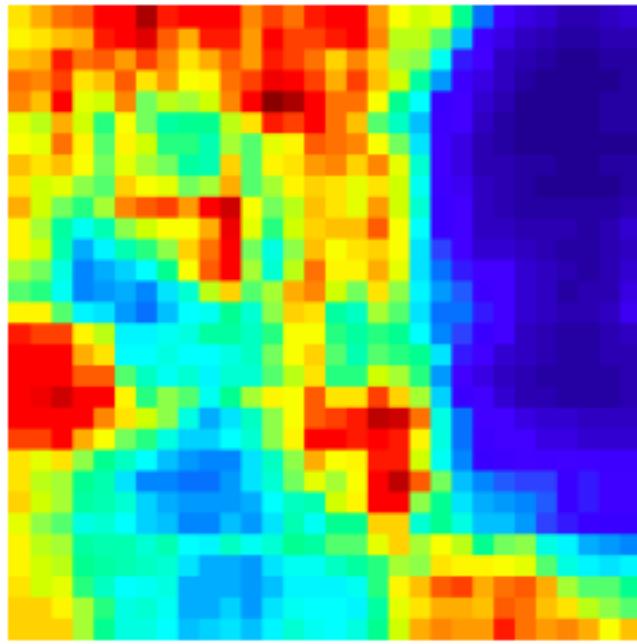
Most probable class



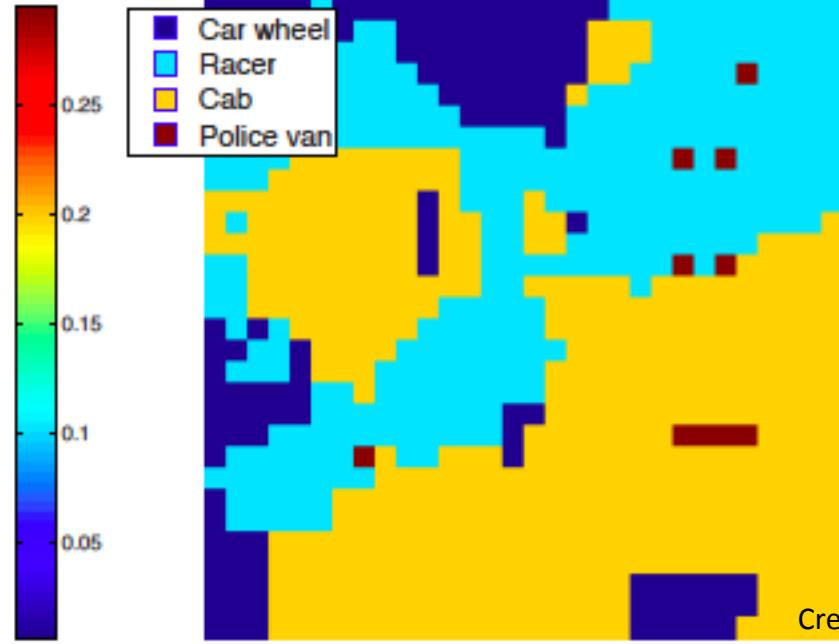
Input image



$p(\text{True class})$

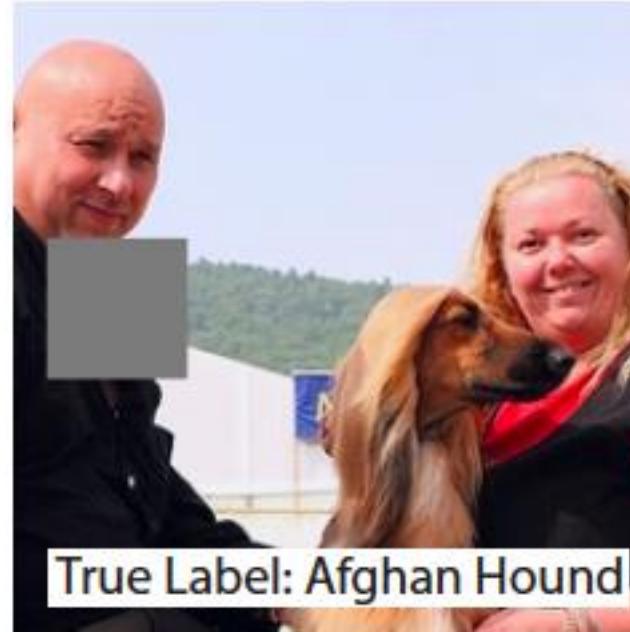


Most probable class

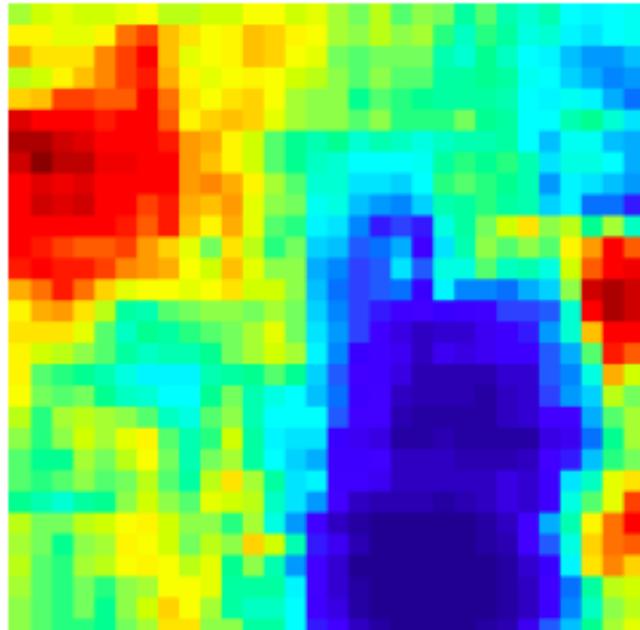


Credit: R. Fergus

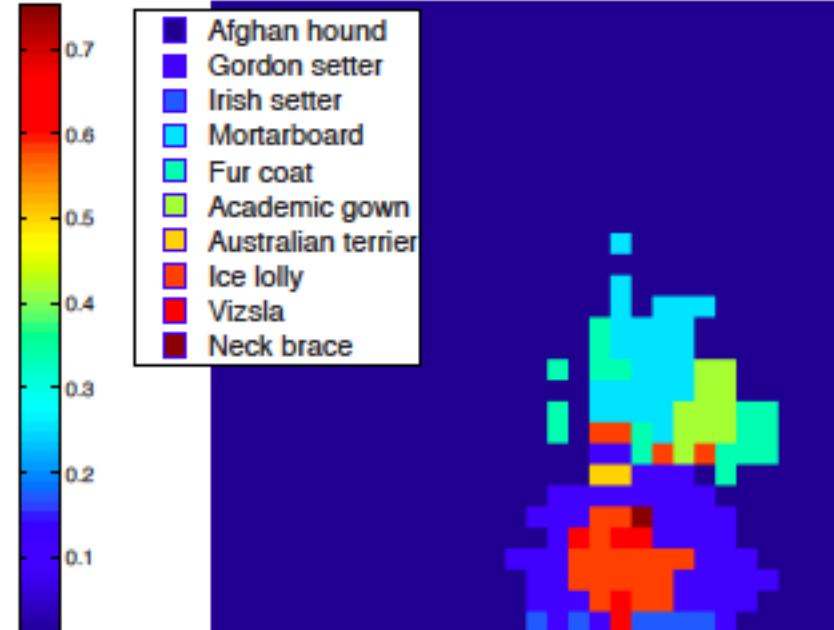
Input image



$p(\text{True class})$



Most probable class

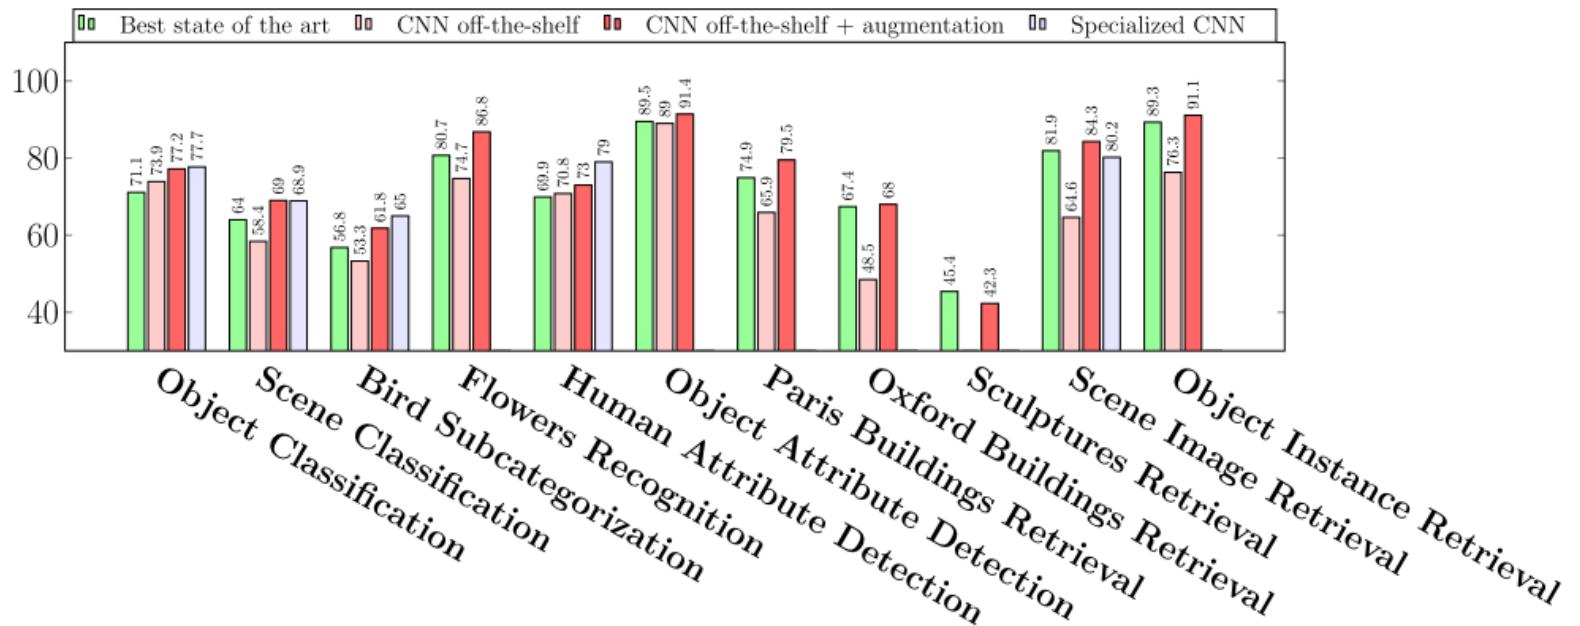


Credit: R. Fergus

Feature Generalization

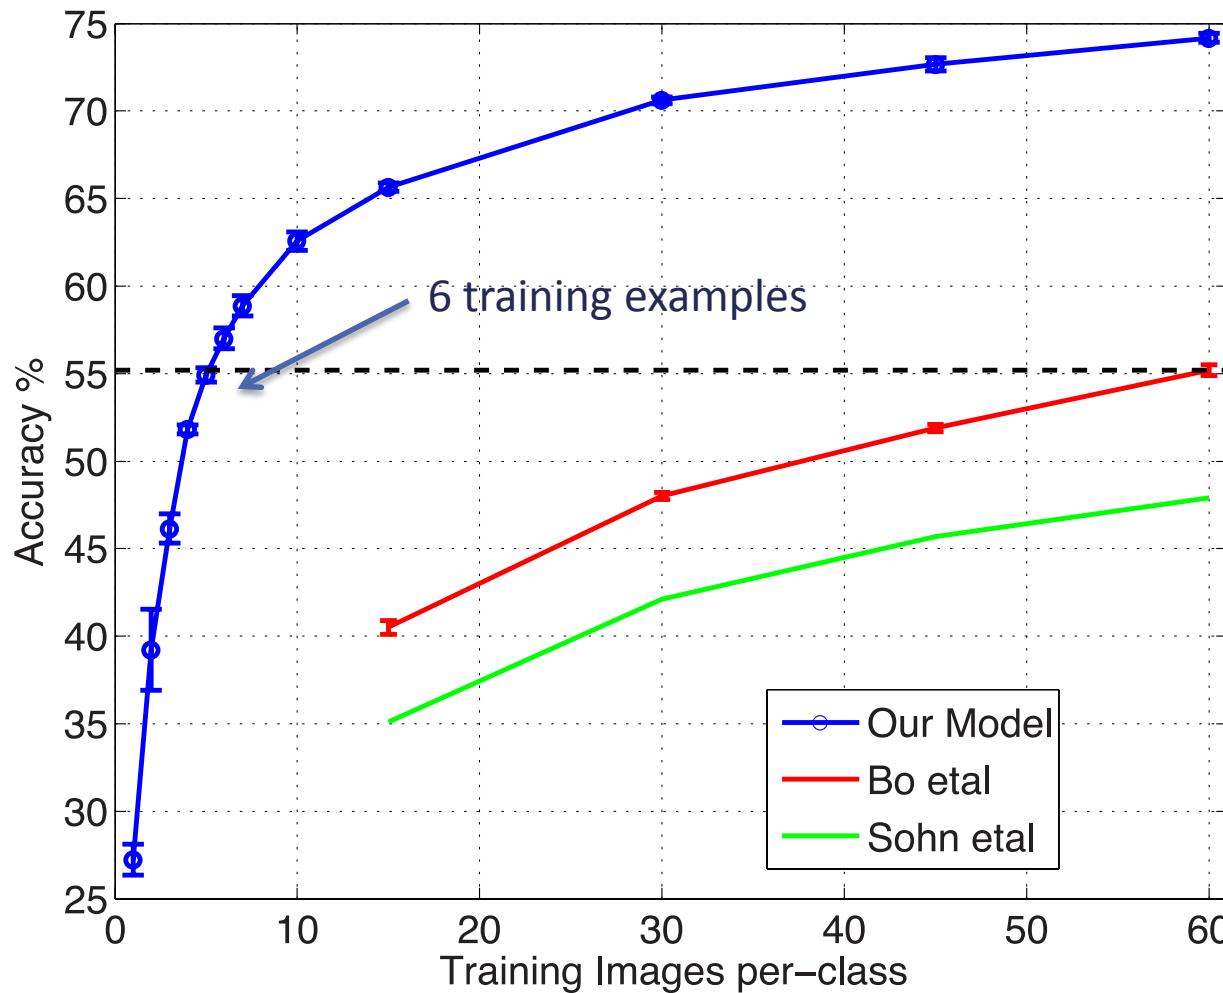
ImageNet pre-training

- **Labeled data is rare for detection:** leverage large classification labeled datasets for pre-training.
- **ImageNet Classification pretraining + fine-tuning on a different task** has been shown to work very well by many people.
 - [Razavian'14] took the off-the-shelf convnet **OverFeat + SVM classifier** on top and obtained many state-of-the-art or competitive results on 10+ datasets and visual tasks



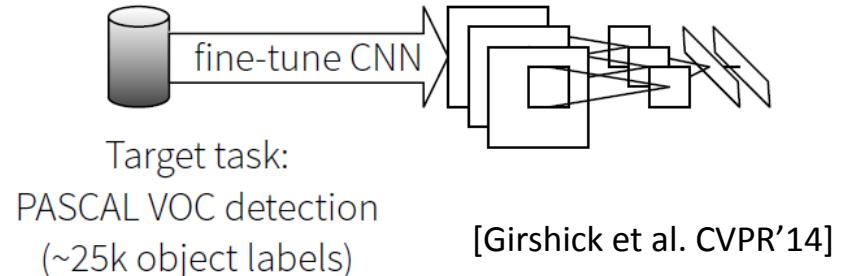
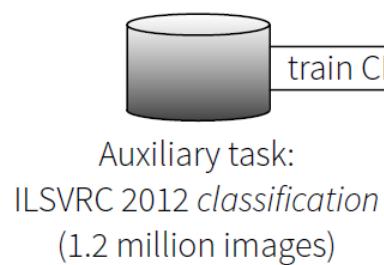
Classifier re-training on Caltech 256

State of the art accuracy with only 6 training samples/class



Feature sharing via transfer learning

- Pre-training allows to use big models for small datasets
 - For example: Pre-Train model on large ImageNet 2012 training set

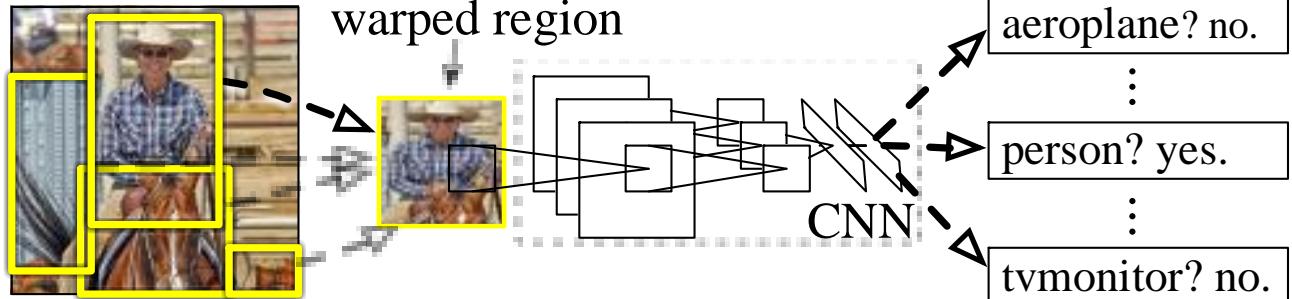


- Re-train on new dataset (fine tuning or transfer learning)
 - Either: Just the classifier-layer or the whole network
 - For fine-tuning pre-trained layers, the learning rate has to be lowered to avoid unlearning the pre-trained weights
 - For fine tuning new layers (e.g. the new classifier layer) the learning rate has to be higher
 - Better: Two stage fine-tuning
 - Stage 1: First only learn new layers with the learning rate of pre-trained layers set to zero
 - Stage 2: Use default learning rate to fine-tune everything (optimize all parameters jointly)
- Classify test set of new dataset

(Fine tuned) CNNs for detection on the Pascal dataset

- Combines bottom-up region proposals with rich features computed by a CNN

R-CNN: *Regions with CNN features*



1. Input image

2. Extract region proposals (~2k)

3. Compute CNN features

4. Classify regions

[Girshick et al. CVPR'14]

- Previous state-of-the art: 35.1% mean average precision
- scratch: Training on Pascal train+val data
- pre-train: Pre-training on ImageNet and just the classifier is trained on Pascal
- fine-tune: Two stage fine-tuning on Pascal

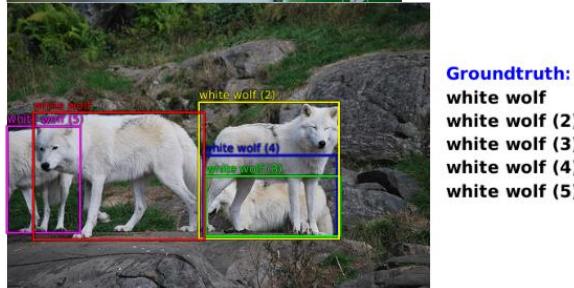
PASCAL-DET		
scratch	pre-train	fine-tune
40.7	45.5	54.1

CNNs have set a new state-of-the art for many tasks

- classification



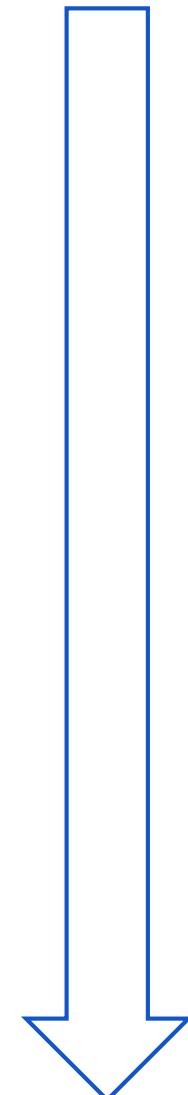
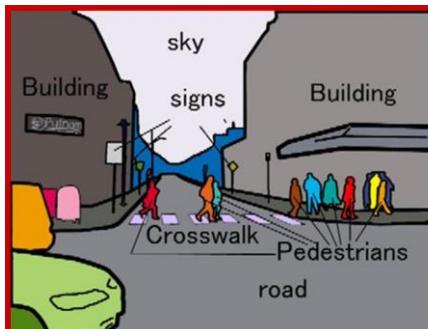
- localization



- detection



- segmentation



difficulty

Credit: P. Sermanet

... except for

Video Recognition

Goal: Scene Understanding

Car

Person

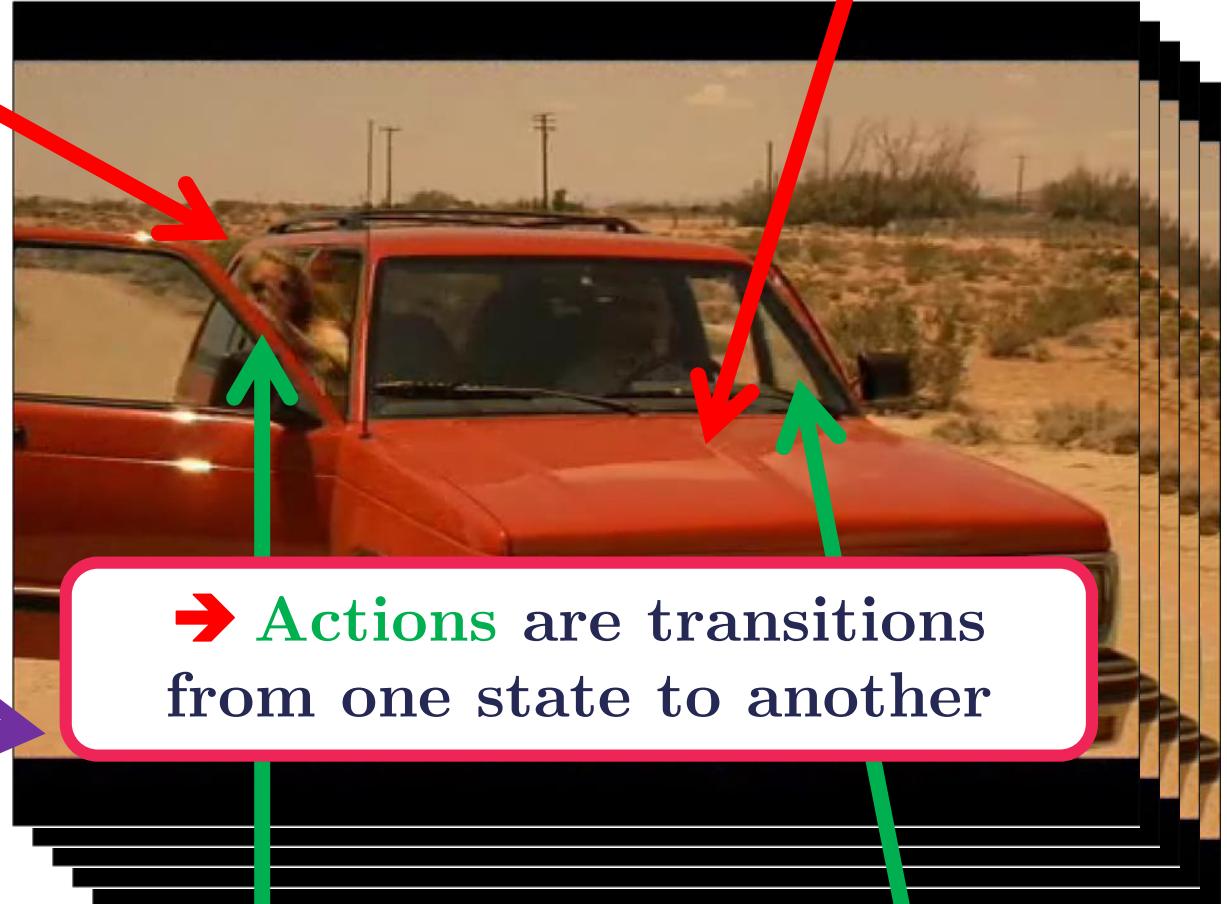
Scenes Objects

Actions

Activities

Desert

Leave a car in the desert



→ Actions are transitions
from one state to another

Temporal input sequence

Get out of car

Open door

Action classification

training samples



Large-scale Video Classification

Sports-1M dataset

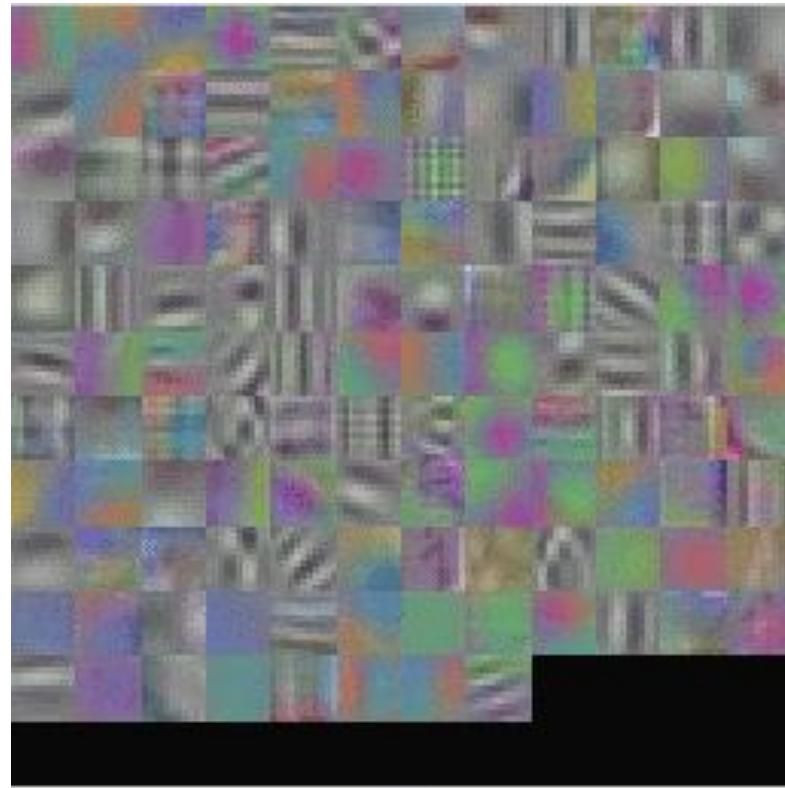
- 1 million YouTube videos in 487 classes of sports



Sports Video
Classification

Large-scale Video Classification

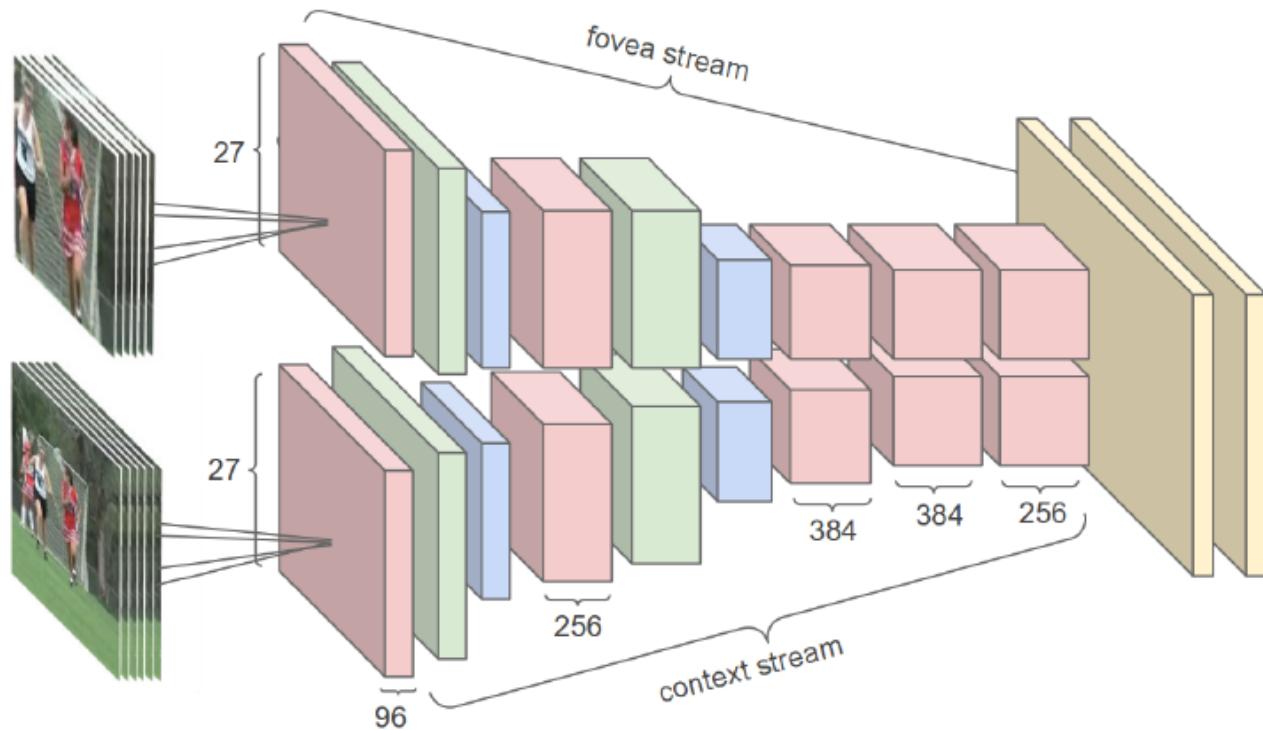
Learned features of the first Layer



Large-scale Video Classification

Multiresolution architecture

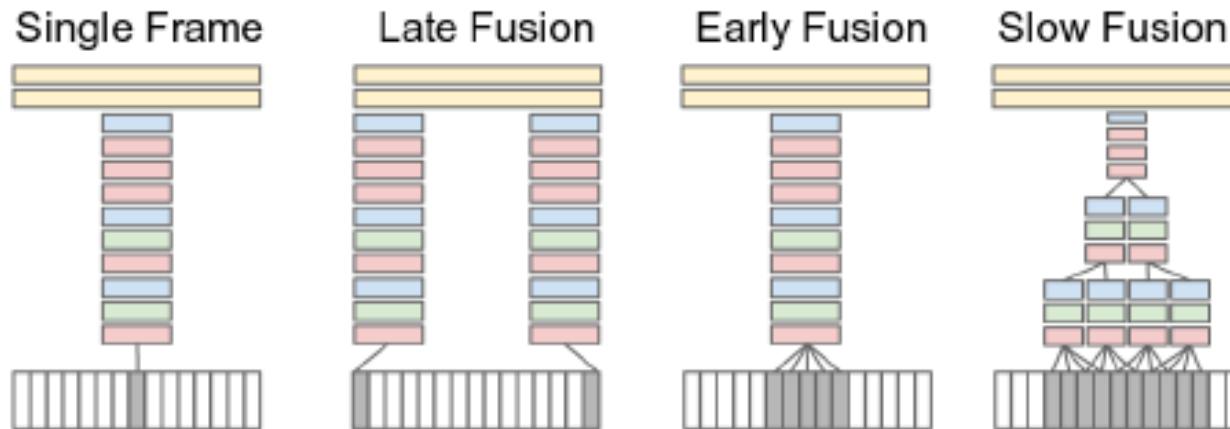
- Context + Fovea Stream



Karpathy et al., Large-scale Video Classification with Convolutional Neural Networks, CVPR 2014

Large-scale Video Classification

- Fusing information over temporal dimension through



Model	Clip Hit@1	Video Hit@1	Video Hit@5
Feature Histograms + Neural Net	-	55.3	-
Single-Frame	41.1	59.3	77.7
Single-Frame + Multires	42.4	60.0	78.5
Single-Frame Fovea Only	30.0	49.9	72.8
Single-Frame Context Only	38.1	56.0	77.2
Early Fusion	38.9	57.7	76.8
Late Fusion	40.7	59.3	78.7
Slow Fusion	41.9	60.9	80.2
CNN Average (Single+Early+Late+Slow)	41.4	63.9	82.4

Large-scale Video Classification

Transfer learning on UCF-101

- 13320 videos in 101 classes

Model	3-fold Accuracy
Soomro et al [22]	43.9%
Feature Histograms + Neural Net	59.0%
Train from scratch	41.3%
Fine-tune top layer	64.1%
Fine-tune top 3 layers	65.4%
Fine-tune all layers	62.2%

Table 3: Results on UCF-101 for various Transfer Learning approaches using the Slow Fusion network.

85.9% using Improved Dense Trajectories + Fisher Vectors [Wang et al. '13]

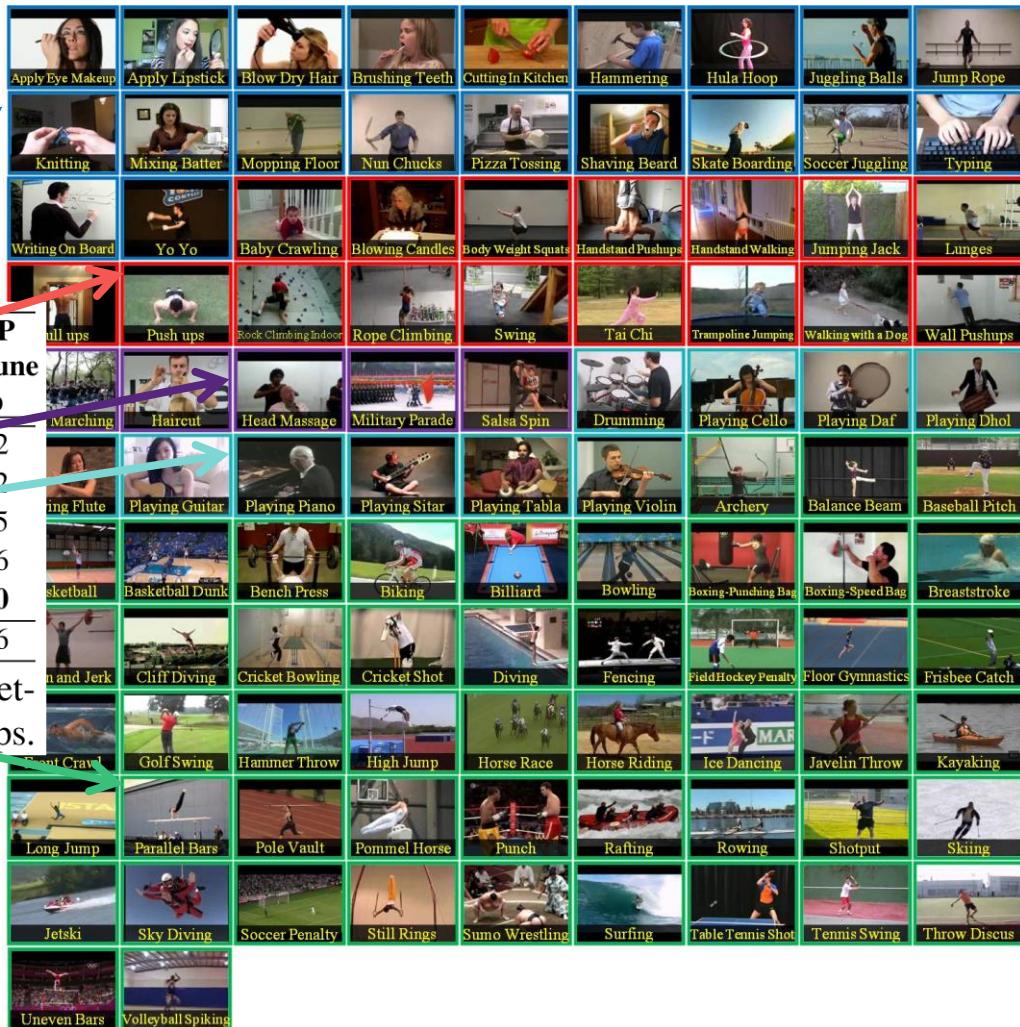
87.6% using Two-stream CNN [Simonyan and Zisserman '14]



Large-scale Video Classification

Transfer learning on UCF-101

- 13320 videos in 101 classes



Group	mAP from scratch	mAP fine-tune top 3	mAP fine-tune top
Human-Object Interaction	0.26	0.55	0.52
Body-Motion Only	0.32	0.57	0.52
Human-Human Interaction	0.40	0.68	0.65
Playing Musical Instruments	0.42	0.65	0.46
Sports	0.57	0.79	0.80
All groups	0.44	0.68	0.66

Table 4: Mean Average Precision of the Slow Fusion network on UCF-101 classes broken down by category groups.

State of the art in Action Recognition: Dense Trajectories + Fisher Vectors [Wang et al. '13]

- Low level primitive features are extracted densely at the first layer by tracking trajectories in a dense optical flow field

motion trajectory aligned pooling of primitives

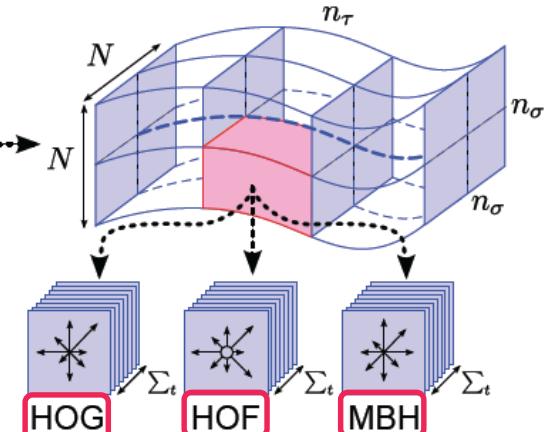
Dense sampling
in each spatial scale

Tracking in each spatial scale separately

Trajectory description

multiscale interest points

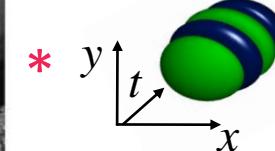
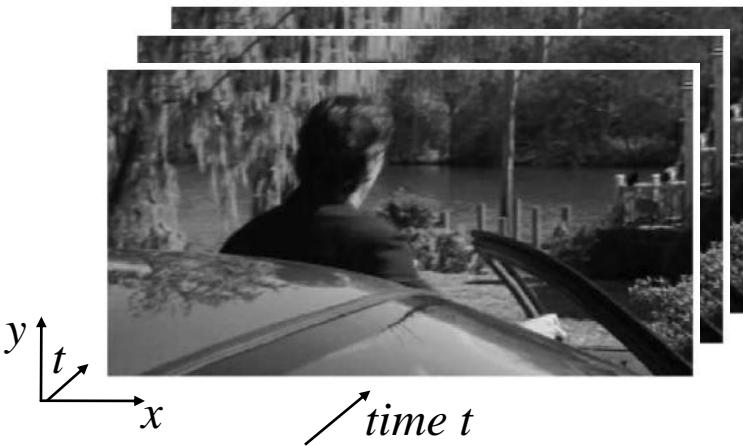
...creates feature descriptors



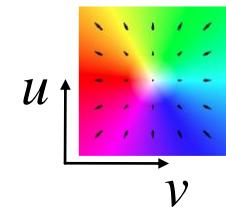
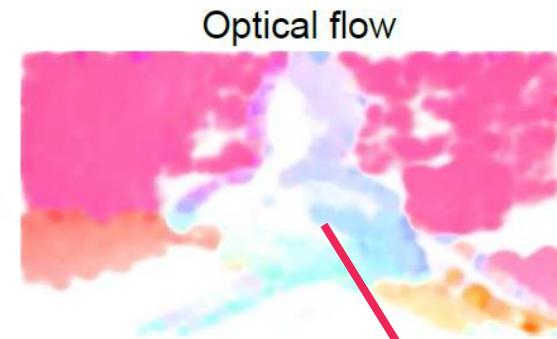
→ Optical flow also captures camera motion and parallax

within a convolutional framework

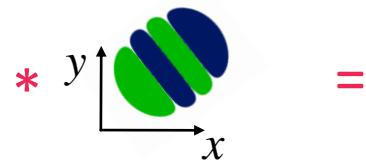
- Optical flow is the apparent motion of the brightness pattern between images
- Image gradients are the directional change of the intensity or colour in the image



“Temporal filtering
of the intensity”



Histograms of Optical Flow (HOF)



“Spatial filtering of
the intensity”

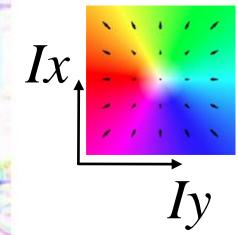
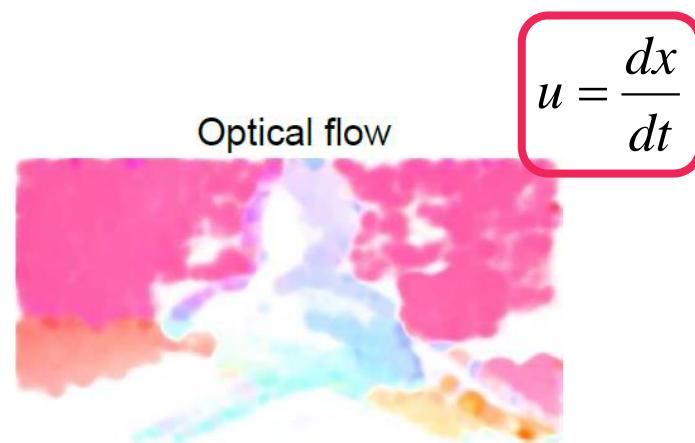


image $I(x, y, t)$

Histograms of Oriented Gradients (HOG)

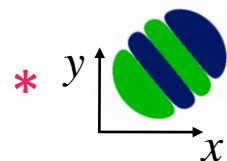
State of the art hand-crafted features within a convolutional framework

- Motion boundaries are the image gradients of the horizontal and vertical Optical flow components (i.e. u and v)



Optical flow

$$u = \frac{dx}{dt}$$

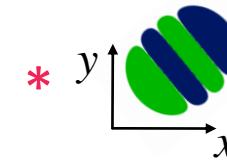


“Spatial filtering of
the horizontal flow”

Horizontal motion boundaries

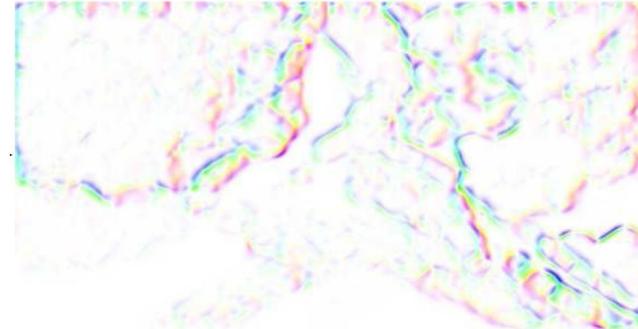


$$v = \frac{dy}{dt}$$

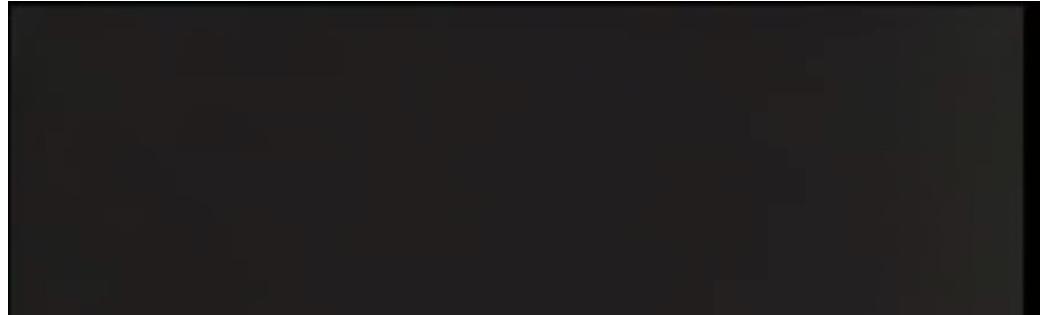


“Spatial filtering of
the vertical flow”

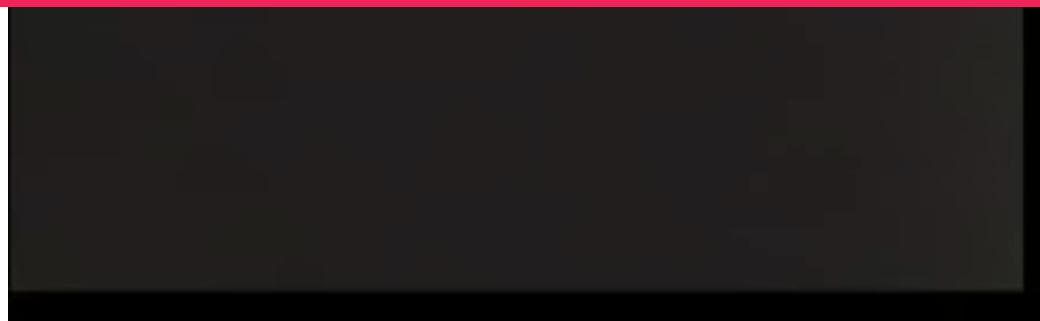
Vertical motion boundaries



Johansson: Perception of Biological Motion



→ Amazing what a human observer
can do without spatial information



Sources: Johansson, G. "Visual perception of biological motion and a model for its analysis." *Perception & Psychophysics*. 14(2):201-211. 1973.
Videos were made by JB Maas in 1971 (released via Houghton-Mifflin and now available on Youtube).

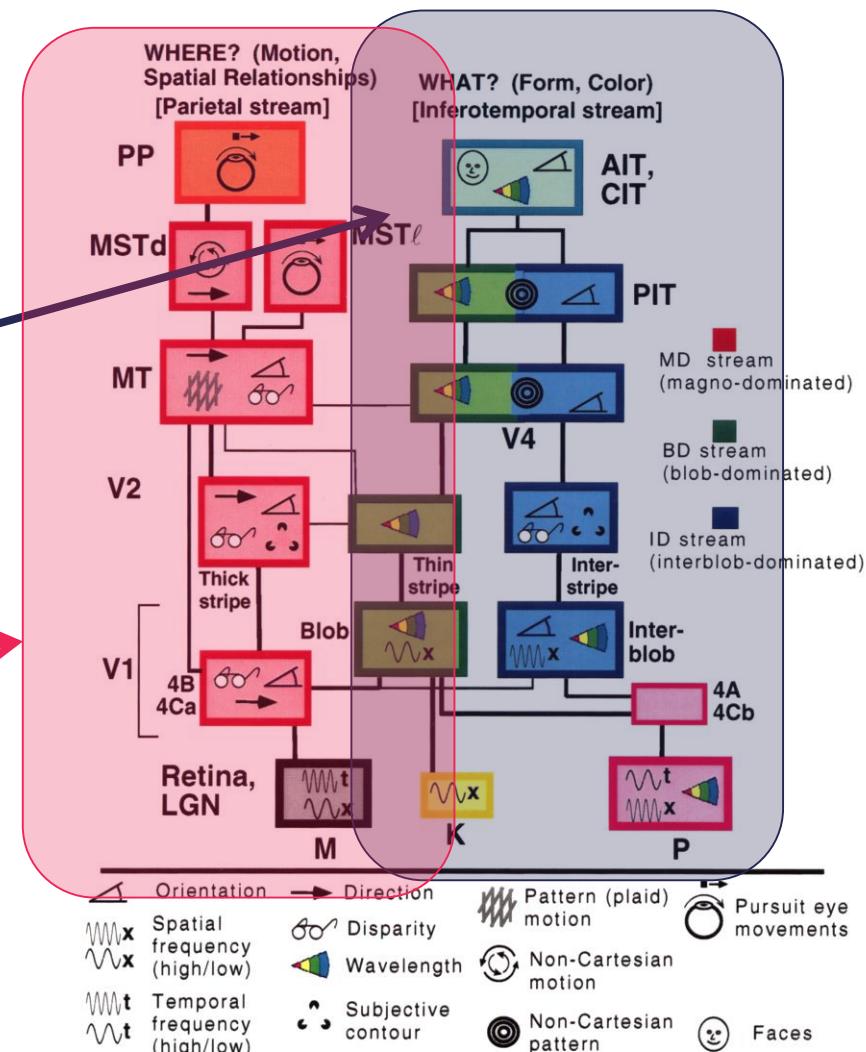
Motivation: Separate visual pathways for perception and action

The Human Visual Cortex has two **hierarchical** pathways

- Ventral stream performs object recognition
- Dorsal stream recognizes motion and locates objects

Spatial ConvNet?

Temporal ConvNet?



Two-Stream Convolutional Networks for Action Recognition in Videos

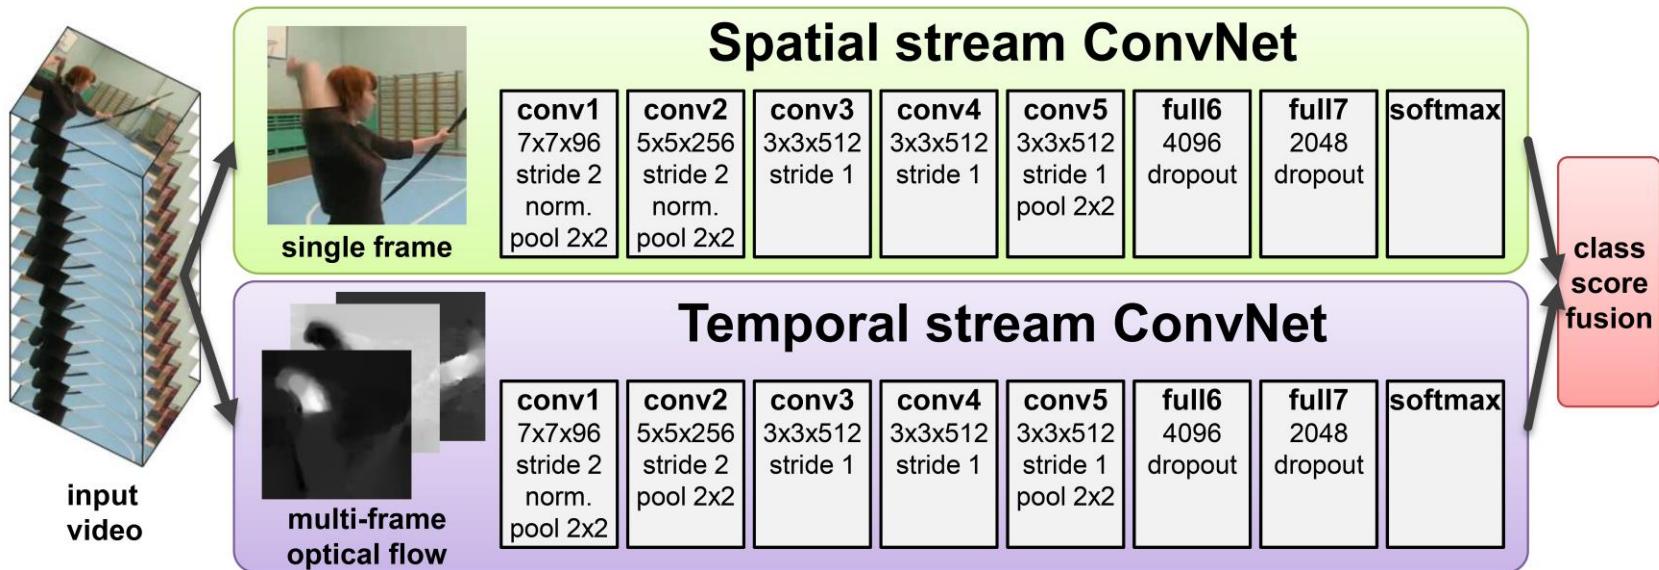


Figure 1: **Two-stream architecture for video classification.**

[Simonyan and Zisserman NIPS'14]

Individual processing of spatial and temporal information

- Using a separate ConvNet recognition stream for each
- Late fusion via softmax score averaging

Spatial stream ConvNet

CNN-F similar to Krizhevsky et al., NIPS 2012:

'ImageNet classification with deep convolutional networks'

conv1
64x11x11
stride 4

conv2
256x5x5
stride 1

conv3
256x3x3
stride 1

conv4
256x3x3

conv5
256x3x3

fc6
4096 d.o.

fc7
4096
drop-out

CNN-M similar to Zeiler and Fergus, CoRR 2013:

'Visualising and understanding convolutional networks'

conv1
96x7x7
stride 2

conv2
256x5x5
stride 2

conv3
512x3x3
stride 1

conv4
512x3x3

conv5
512x3x3

fc6
4096 d.o.

fc7
4096
drop-out

Same network (CNN-M) used for both streams

- Based on [Krizhevsky et al. NIPS'12]
- Better (\approx deeper) architectures exist now (see last lecture)
 - GoogLeNet
 - VGG Very Deep

[Chatfield et al. BMVC'14]



single frame

Spatial stream ConvNet

conv1
7x7x96
stride 2
norm.
pool 2x2

conv2
5x5x256
stride 2
norm.
pool 2x2

conv3
3x3x512
stride 1

conv4
3x3x512
stride 1

conv5
3x3x512
stride 1
pool 2x2

full6
4096
dropout

full7
2048
dropout

softmax

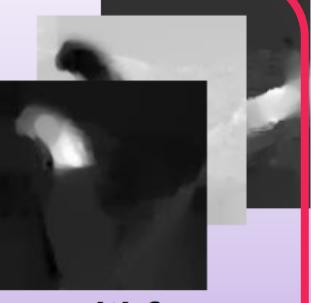
- Performs image classification on **single RGB frames**

Training:

- Supervised pre-training on ILSVRC (1.2M images in 1000 classes)
- Fine tuning of the softmax layer using the video frames

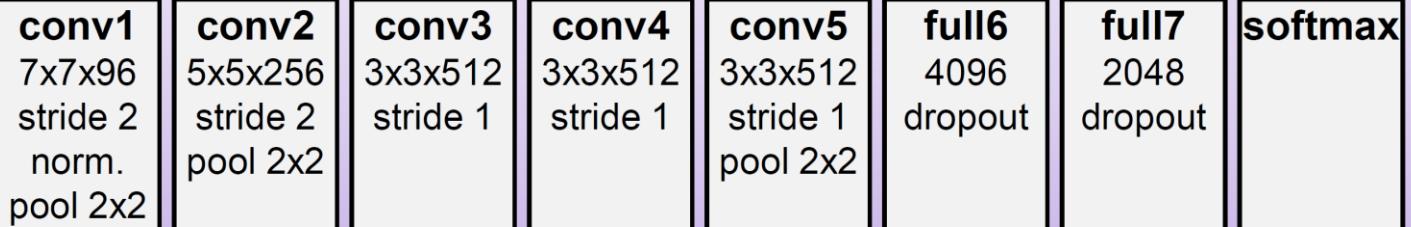
Testing

- ConvNet processes every 25th frame of a video
- Data augmentation: 10 ConvNet inputs for each frame (crops & flips)
- Results for all ConvNets are averaged



multi-frame
optical flow

Temporal stream ConvNet



Same model as in the spatial net except

- Optical flow over several frames acts as input

Displacement vector field between (a) and (b)



(a)

(b)



(c)

(d)

(e)

255

0

Figure 2: **Optical flow.** (a),(b): a pair of consecutive video frames with the area around a moving hand outlined with a cyan rectangle. (c): a close-up of dense optical flow in the outlined area; (d): horizontal component d^x of the displacement vector field (higher intensity corresponds to positive values, lower intensity to negative values). (e): vertical component d^y . Note how (d) and (e) highlight the moving hand and bow. The input to a ConvNet contains multiple flows (Sect. 3.1).

Horizontal and vertical flow is rescaled to [0, 255] for ConvNet input

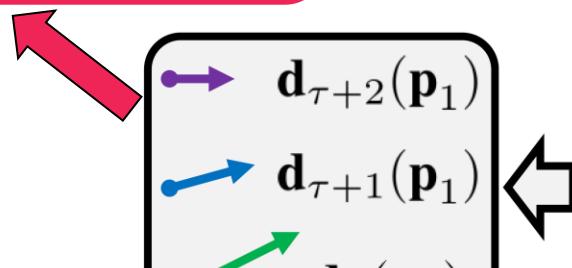
[Simonyan and Zisserman NIPS'14]

Optical flow stacking

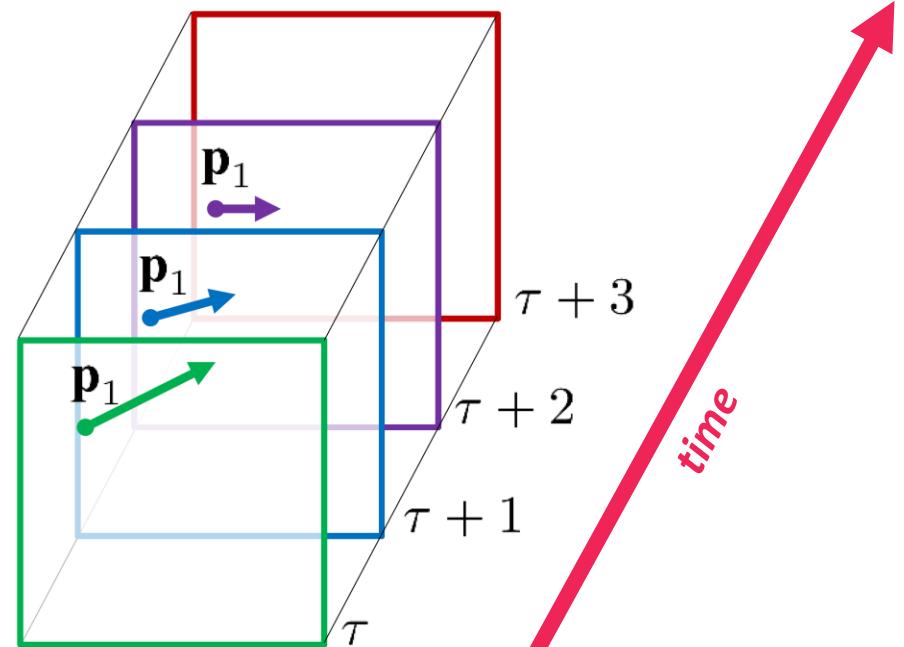
Stack horizontal and vertical displacement fields \mathbf{d}

- Optical flow, \mathbf{d} , over several frames, τ , acts as input I_τ to the network

Input I_τ at \mathbf{p}_1 represents motion at \mathbf{p}_1 across multiple frames



input volume channels
 $I_\tau(u, v, c)$ at point
 $\mathbf{p}_1 = (u, v)$

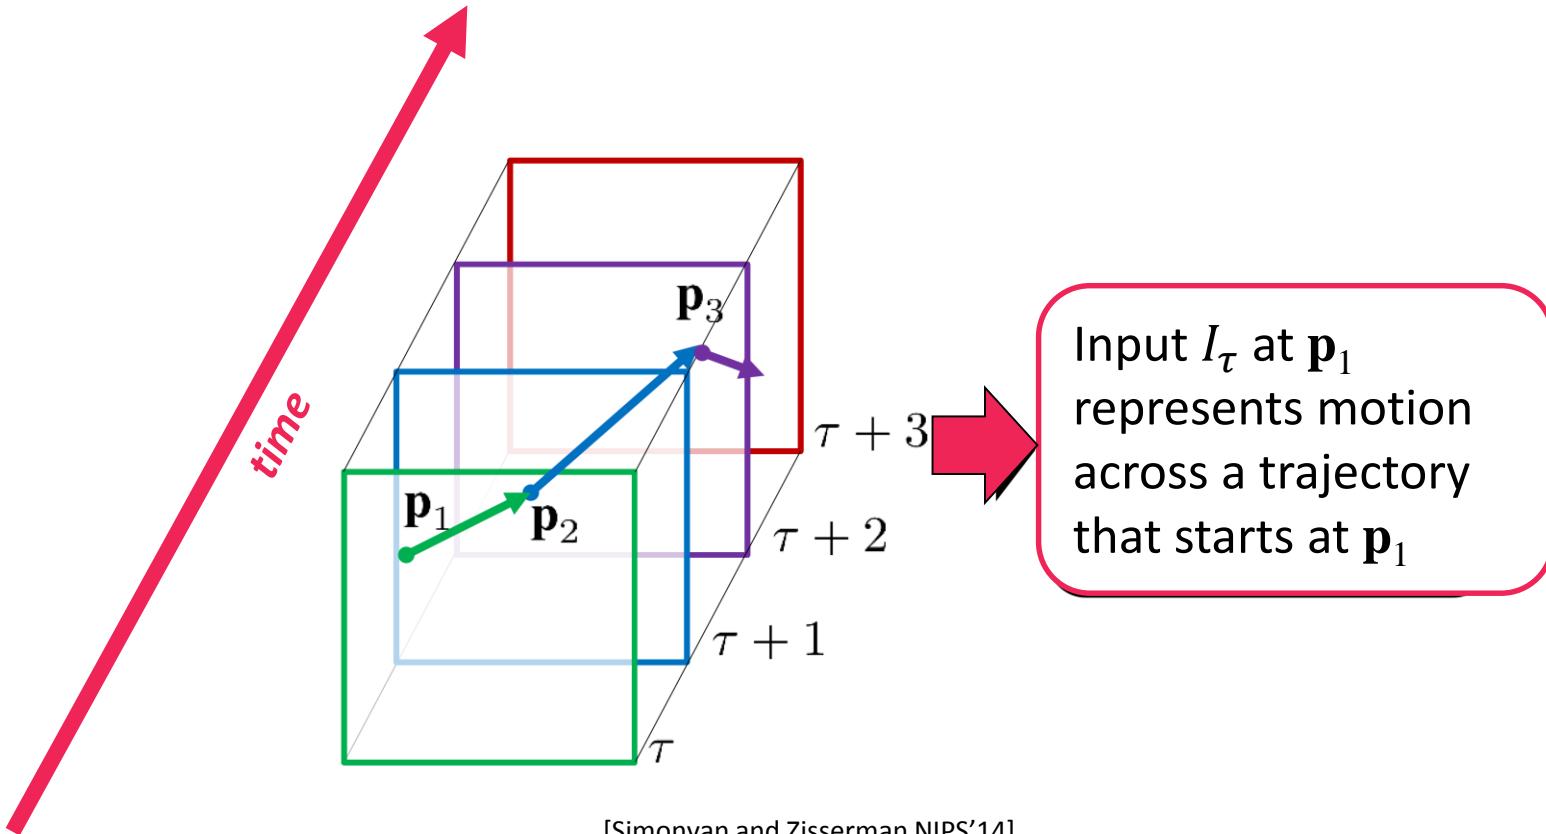


[Simonyan and Zisserman NIPS'14]

Trajectory stacking

Stack horizontal and vertical displacement fields \mathbf{d} along the trajectory

- Trajectory over several frames, τ , acts as input I_τ to the network



Empirical evaluation

Datasets

- UCF101: 101 classes, 13K videos, ~180 frames in a vid
- HMDB51: 51 classes, 6.8K videos
- Evaluation protocol: average classification accuracy over 3 train and test splits



Empirical evaluation: Overfit prevention

Action recognition datasets are rather small ($\approx 10K$ videos)

- Many images (=frames) but they are very similar

Spatial net overfit prevention:

- Supervised pre-training on large dataset (ILSVRC 1.2M images in 1000 classes)
- Fine tuning of the softmax layer using the video frames

Temporal net:

- Multitask Learning (train a model based on several loss functions for different tasks)
 - Each task has its own (softmax) loss
 - Total loss \approx sum over task losses
 - One Task = UCF101 classification, other task = HMDB51 classification
 - Datasets are not merged, however backprop operates on the sum of both losses

Empirical evaluation

Spatial net:

- Pre trained network is better than scratch
- Fine tuning the whole net is similar to re-train just the last layer
- Training from scratch is unpractical even with high dropout

Table 1: Individual ConvNets accuracy on UCF-101 (split 1).

(a) Spatial ConvNet.

Training setting	Dropout ratio	
	0.5	0.9
From scratch	42.5%	52.3%
Pre-trained + fine-tuning	70.8%	72.8%
Pre-trained + last layer	72.7%	59.9%

(b) Temporal ConvNet.

Input configuration	Mean subtraction	
	off	on
Single-frame optical flow ($L = 1$)	-	73.9%
Optical flow stacking (1) ($L = 5$)	-	80.4%
Optical flow stacking (1) ($L = 10$)	79.9%	81.0%
Trajectory stacking (2) ($L = 10$)	79.6%	80.2%
Optical flow stacking (1) ($L = 10$), bi-dir.	-	81.2%

Empirical evaluation

Temporal net:

- Flow or trajectory stacking improves significantly
($\approx 7\%$ improvement in accuracy)
- Mean subtraction brings only minor improvements

Table 1: Individual ConvNets accuracy on UCF-101 (split 1).

(a) Spatial ConvNet.

Training setting	Dropout ratio	
	0.5	0.9
From scratch	42.5%	52.3%
Pre-trained + fine-tuning	70.8%	72.8%
Pre-trained + last layer	72.7%	59.9%

(b) Temporal ConvNet.

Input configuration	Mean subtraction	
	off	on
Single-frame optical flow ($L = 1$)	-	73.9%
Optical flow stacking (1) ($L = 5$)	-	80.4%
Optical flow stacking (1) ($L = 10$)	79.9%	81.0%
Trajectory stacking (2) ($L = 10$)	79.6%	80.2%
Optical flow stacking (1) ($L = 10$), bi-dir.	-	81.2%

Empirical evaluation

Table 2: **Temporal ConvNet accuracy on HMDB-51 (split 1 with additional training data).**

Training setting	Accuracy
Training on HMDB-51 without additional data	46.6%
Fine-tuning a ConvNet, pre-trained on UCF-101	49.0%
Training on HMDB-51 with classes added from UCF-101	52.8%
Multi-task learning on HMDB-51 and UCF-101	55.4%

Temporal net, multi task learning:

- Additional data improves recognition
- Fine tuning a model that is trained only on a small dataset is challenging
 - Small learning rate → Net stays specialized on the original data
 - Large learning rate → Net overfits the new dataset
- Combining both datasets works better than fine tuning approach
- Multi-task learning works best
 - (But backpropagation operates on both datasets simultaneously)

Empirical evaluation: Comparison with the state of the art

Table 4: Mean accuracy (over three splits) on UCF-101 and HMDB-51.

Method	UCF-101	HMDB-51
Improved dense trajectories (IDT) [26, 27]	85.9%	57.2%
IDT with higher-dimensional encodings [20]	87.9%	61.1%
IDT with stacked Fisher encoding [21] (based on Deep Fisher Net [23])	-	66.8%
Spatio-temporal HMAX network [11, 16]	-	22.8%
“Slow fusion” spatio-temporal ConvNet [14]	65.4%	-
Spatial stream ConvNet	73.0%	40.5%
Temporal stream ConvNet	83.7%	54.6%
Two-stream model (fusion by averaging)	86.9%	58.0%
Two-stream model (fusion by SVM)	88.0%	59.4%

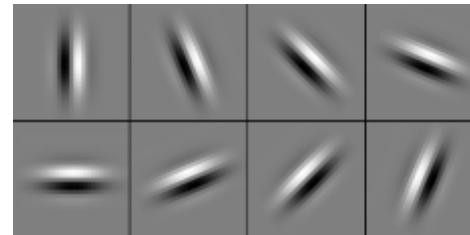
- Spatial / temporal stream network is better than spatiotemporal processing
- Hand-crafted features are still better
- Datasets too small?

Compare: Hand-Crafted Descriptors

Image
Pixels
(HOG
/SIFT...)

Image
Motion
(HOF /
MBH ...)

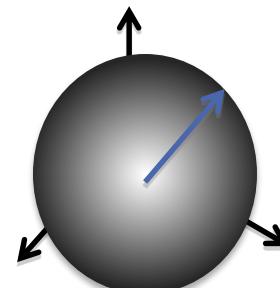
Apply
Gabor filters



→ Is the network able to generalize
hand-crafted representations?



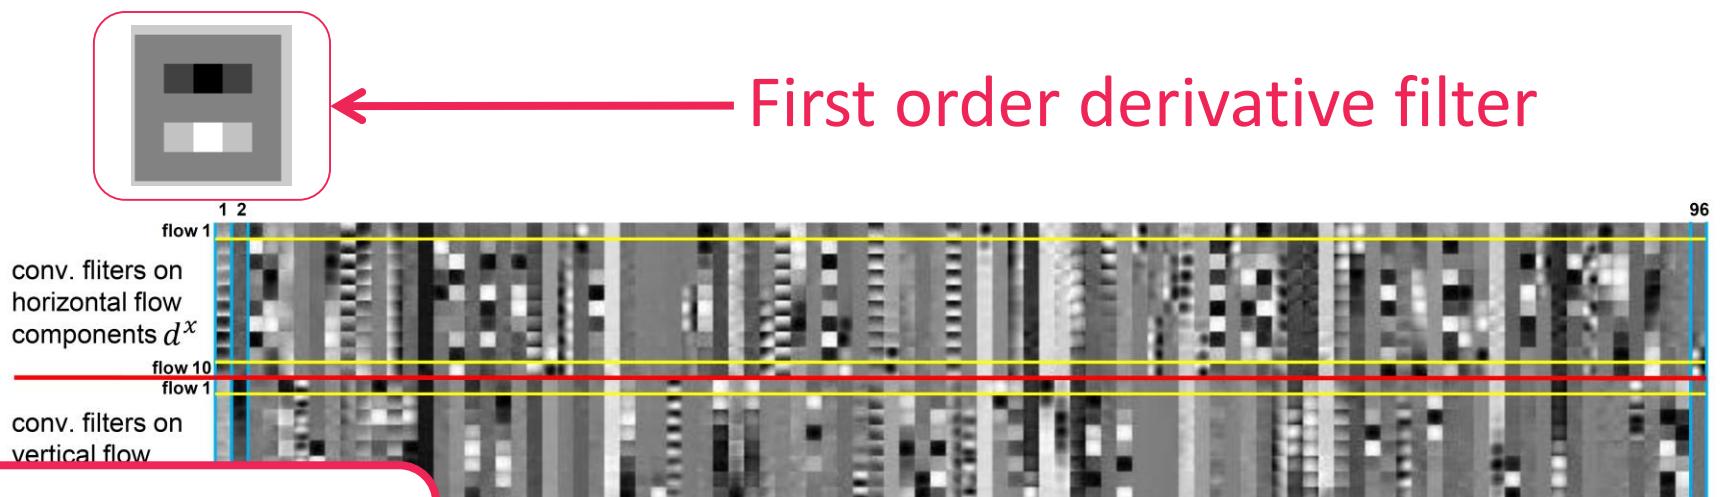
Normalize to unit
length



Feature
Vector

Relation to convolutional networks

- Trajectory over several frames acts as input to the network
- HOG (Histograms of Oriented Gradients) \approx single layer in the spatial network
- HOF (Histograms of Oriented Gradients) \approx single layer in the temporal network
- MBH (Histograms of Oriented Gradients) \approx single layer in the temporal network



\approx motion
change in time

temporal
derivative

spatial
derivative

\approx motion change in
space \approx MBH

[Simonvan and Zisserman NIPS'14]

Figure 4: **First-layer convolutional filters** lead to a feature map that is split into 96 columns and 20 rows: each column corresponds to a spatial input channel.

sation
input

Summary

- Convolutional Networks (ConvNets) for Image Classification
 - Overall architecture defines operations in each layer → Krizhevsky, A., Sutskever, I. and Hinton, G. E., *ImageNet Classification with Deep Convolutional Neural Networks*, NIPS 2012
 - Visualizations improve results on ImageNet → M. Zeiler & R. Fergus, *Visualizing and Understanding Convolutional Networks*, ECCV, 2014
 - Fine-tuning on other datasets helps
- Representations for Video Classification
 - Hand-designed features are still competitive → Wang et al., *Action Recognition by Dense Trajectories*, CVPR 2011.
 - Straightforward application of spatiotemporal ConvNets performs worse → Karpathy et al., *Large-scale Video Classification with Convolutional Neural Networks*, CVPR 2014
 - Two-stream ConvNets are able to generalize hand-crafted representations → K. Simonyan & A. Zisserman, *Two-Stream Convolutional Networks for Action Recognition in Videos*, NIPS 2014